



# Kubernetes Arabic Course

Ahm

# Kubernetes From Scratch - Hands-on

01 - Introduction

Ahmed ElBakry



# Course Contents

## 01 Introduction

- Course Introduction
- What is Kubernetes
- Kubernetes Features

## 02 Kubernetes Architecture

- Cluster Architecture
- Master Components
- Worker Components

## 03 Setup Lab Cluster

- Setup Single Node Cluster (Minikube)
- Setup Kind Cluster
- Setup Internal Docker Registry

## 04 Basic Concepts

- Namespaces
- Resource Quota
- Labels and Selectors
- Nodes



# Course Contents

## 05 Pods

- Pods Explained
- Pods demo
- Multi-container pod

## 06 Deployments

- Deployment Explained
- Deployment Demo
- Rolling Update

## 07 Services

- Service Explained
- Service Types
- Service Demo

## 08 Ingress

- Ingress Explained
- Deploy Nginx Ingress Controller
- Ingress Demo



# Course Contents

## 09 Deploy Your First App

- Containerize The Application
- Deploy the Application
- Volumes
- Secrets - ConfigMaps

## 10 Deploying With Helm

- Kubernetes Package Manager (Helm)
- Deploy Your First Helm Chart
- Helm Templates

## 11 Monitoring

- Deploy Prometheus
- Deploy Alert Manager
- Explore Prometheus & Alert Manager

## 12 Summary

- Wrap Up
- What's Next

# About The Instructor



## AHMED ELBAKRY

- Sr Cloud Architect
- + 12 Years IT Experience
- RHCE, VCP, Google Professional Cloud Architect
- Kubernetes, Terraform, Python, GCP, AWS, etc.

 /ahmed.elbakry.39

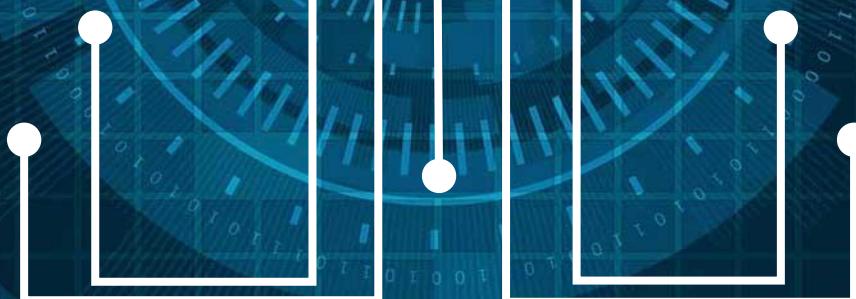
 @ahmed43068401

 /in/ahmed-elbakry-56261134





# THANK YOU

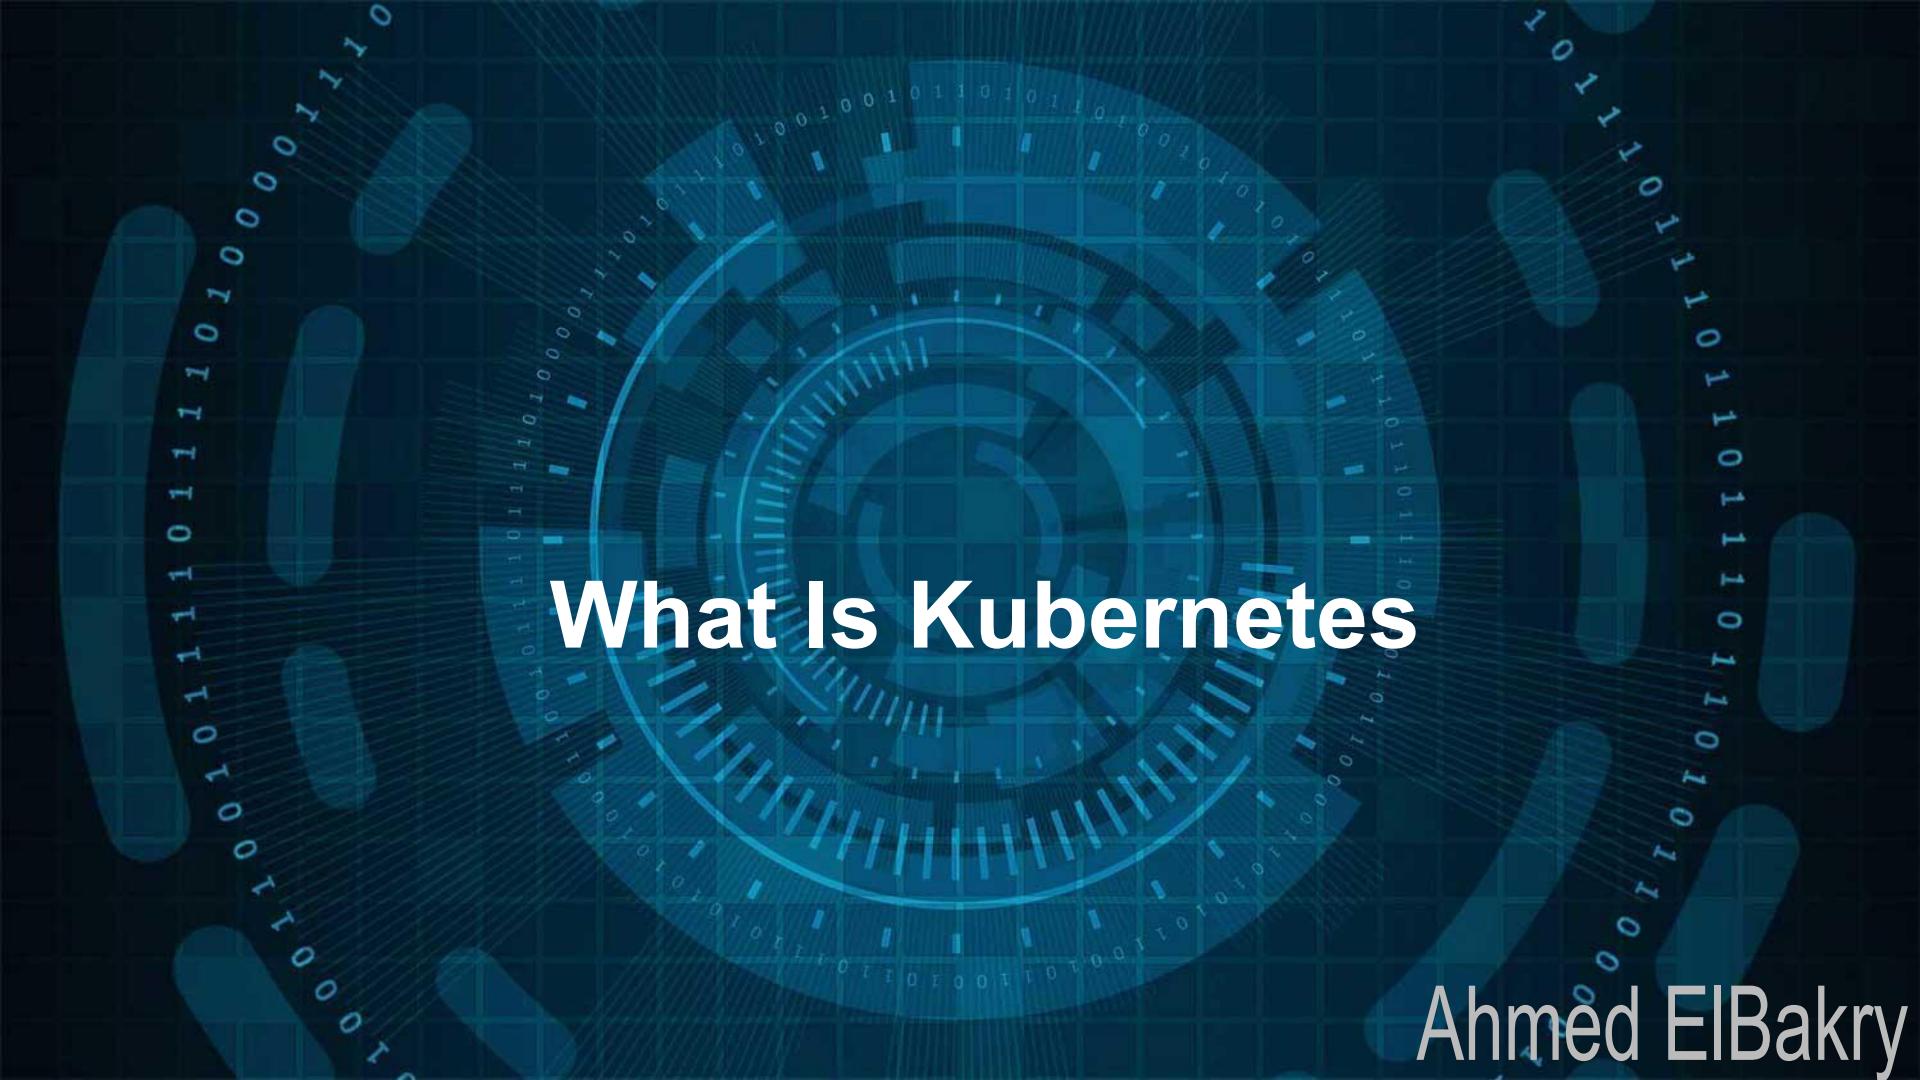


Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

02 - What Is Kubernetes

Ahmed ElBakry

The background of the slide features a dark blue abstract design. It includes several large, semi-transparent white circles of varying sizes. Overlaid on these circles are numerous binary digits (0s and 1s) arranged in curved paths, suggesting data flow or computation. A fine grid of light blue lines covers the entire background, adding to the technical and futuristic feel.

# What Is Kubernetes

Ahmed ElBakry

# What is Kubernetes

## Container Orchestration

It's an open source container orchestration platform, to automate and manage containerized applications.

## Cloud Agnostic

you can run it anywhere on bare metal or in any cloud provider infrastructure.

## Developed at Google

It was developed at Google and released as open source in 2014, So it combines the best practices of running containerized workloads at google for 15 years.



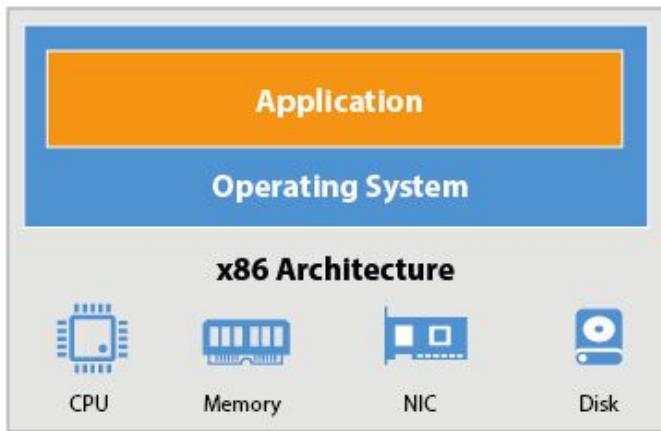
## Zero Downtime Deployment

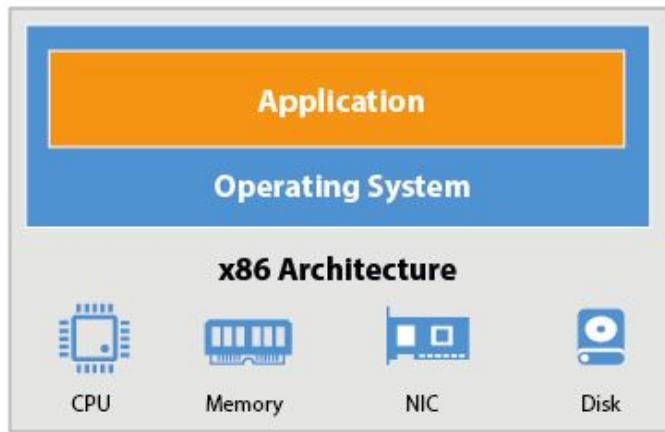
It became the industry standard for deploying containers in production

# Container Orchestration

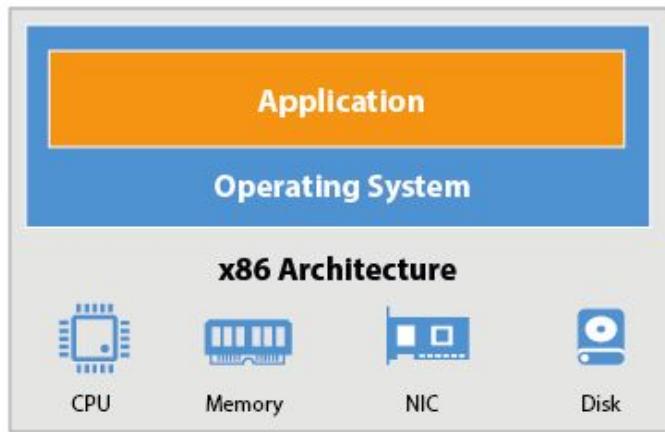
**It's an open source container orchestration platform, to automate and manage containerized applications.**

Ahmed ElBakry

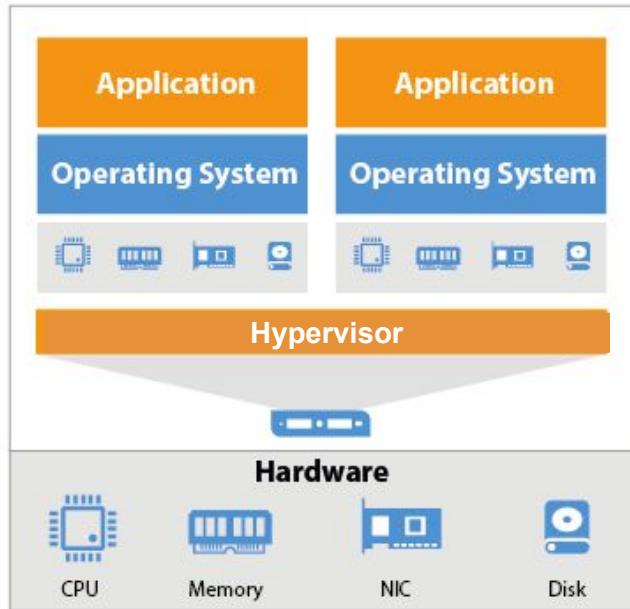




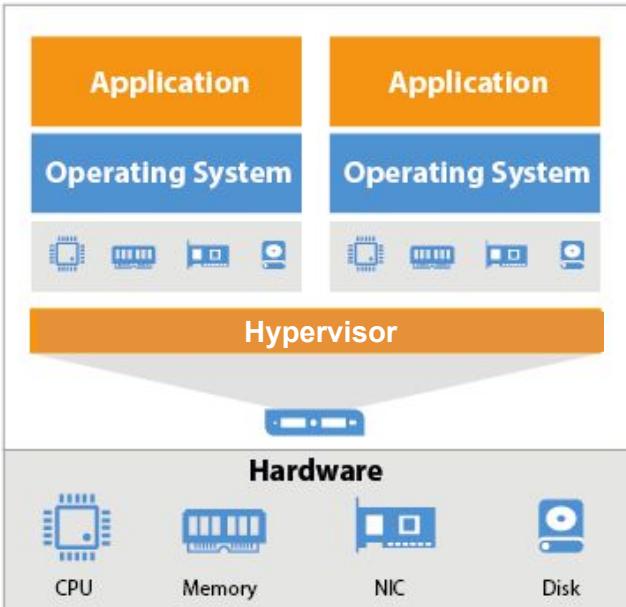
Ahmed ElBakry



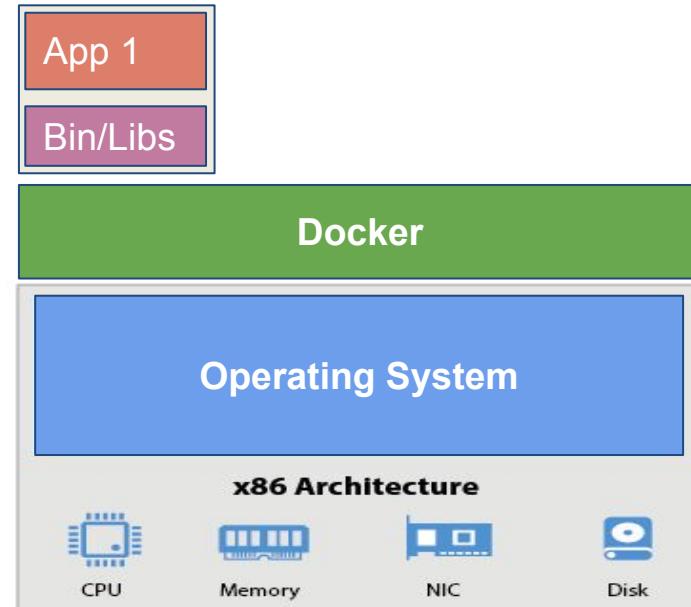
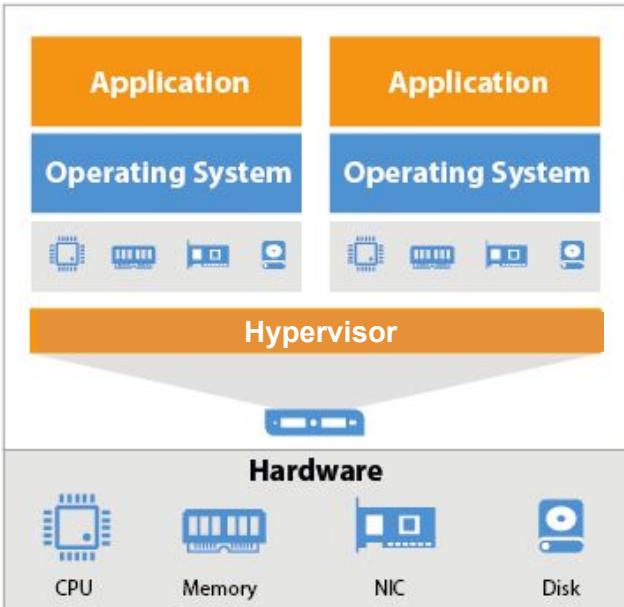
Ahmed ElBakry



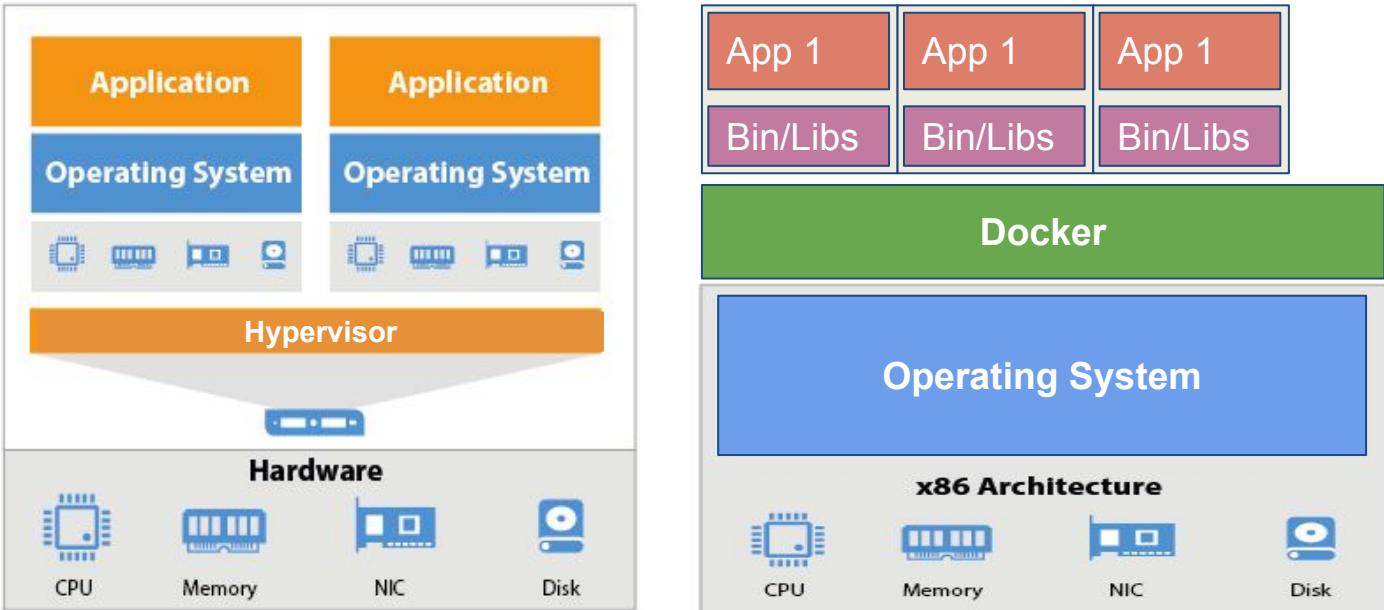
Ahmed ElBakry

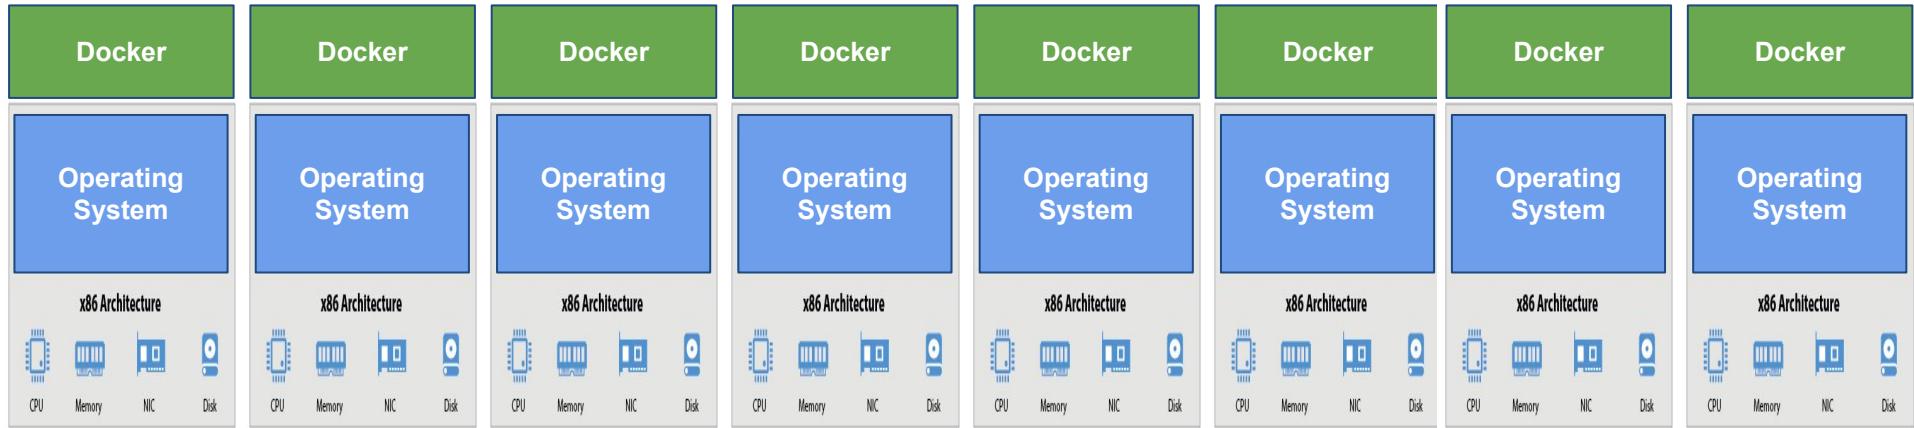


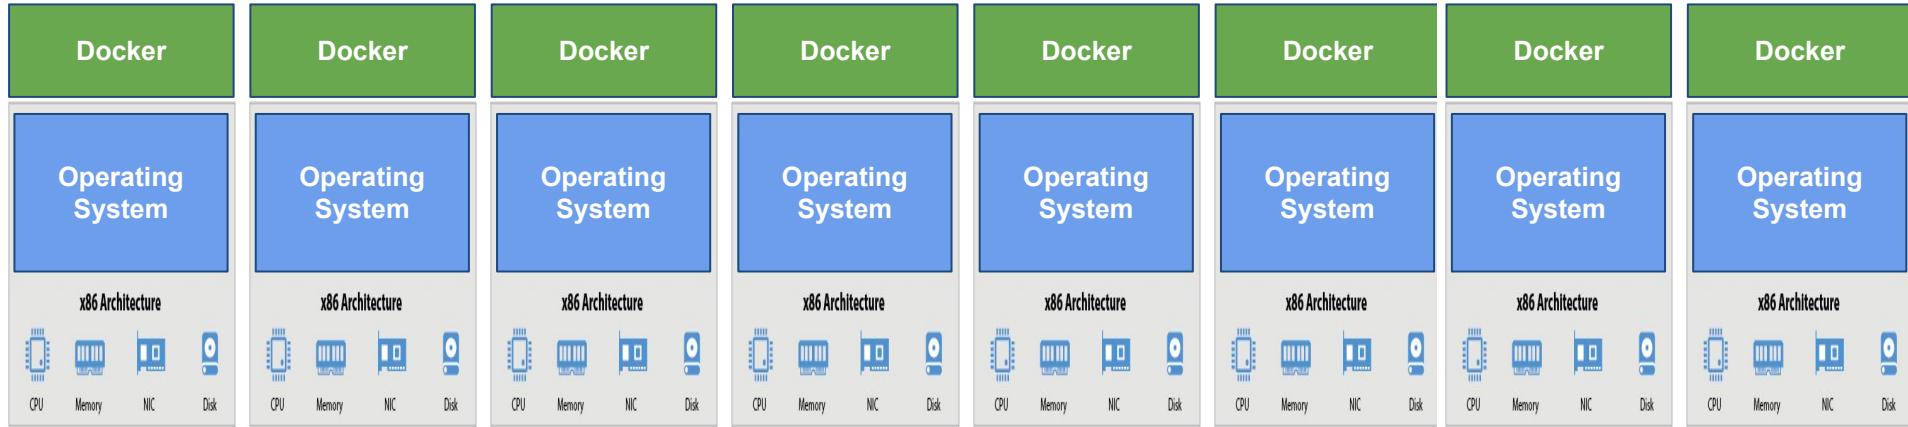
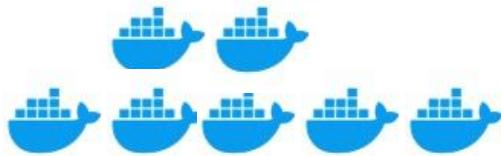
Animated ElBakry

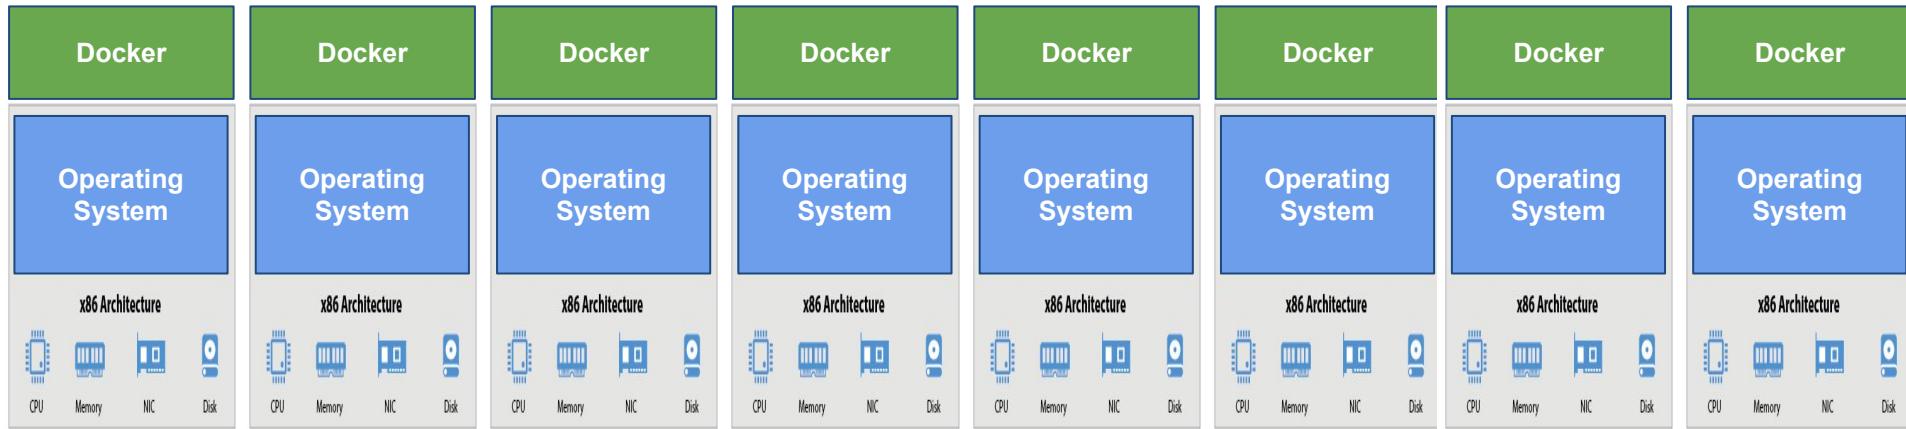
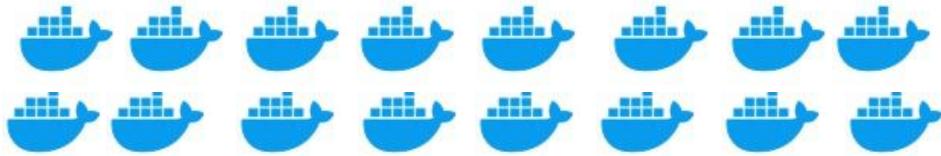
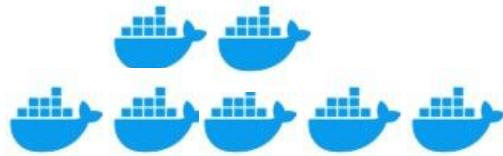


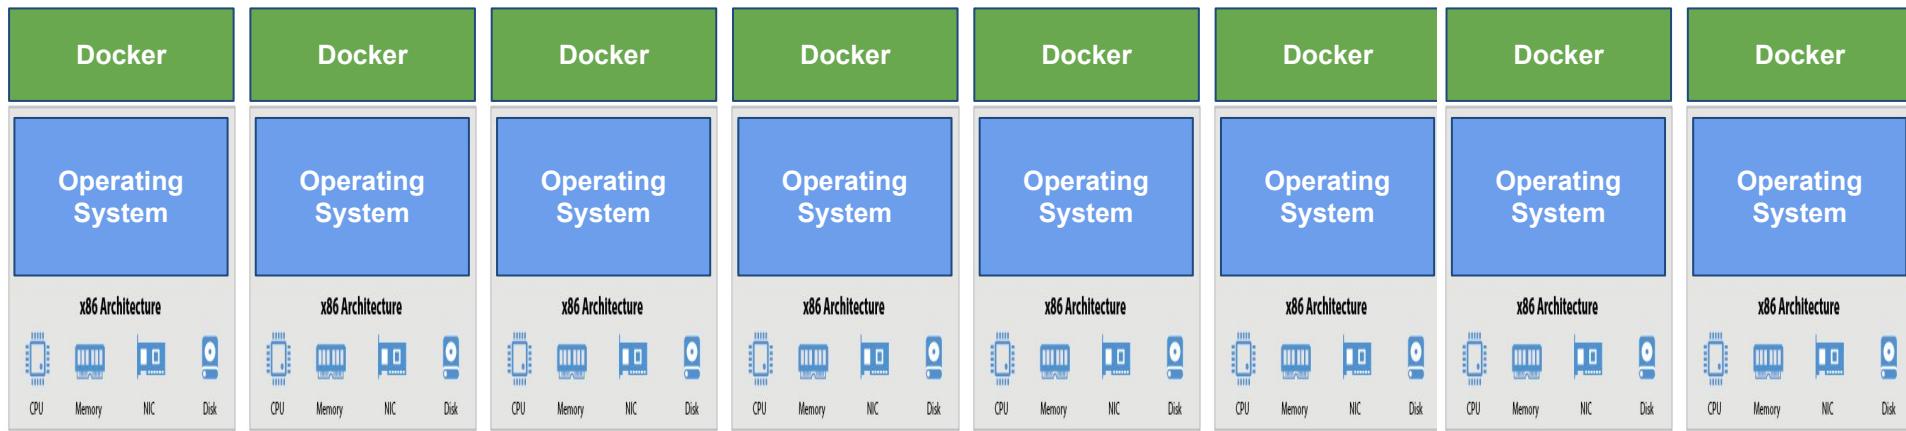
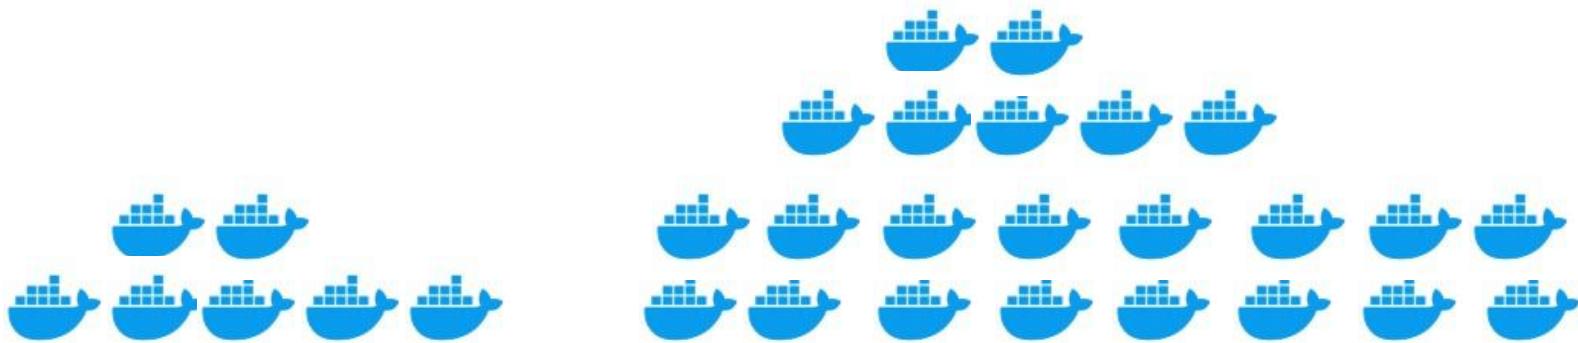
Ali Mohamed ElBakry

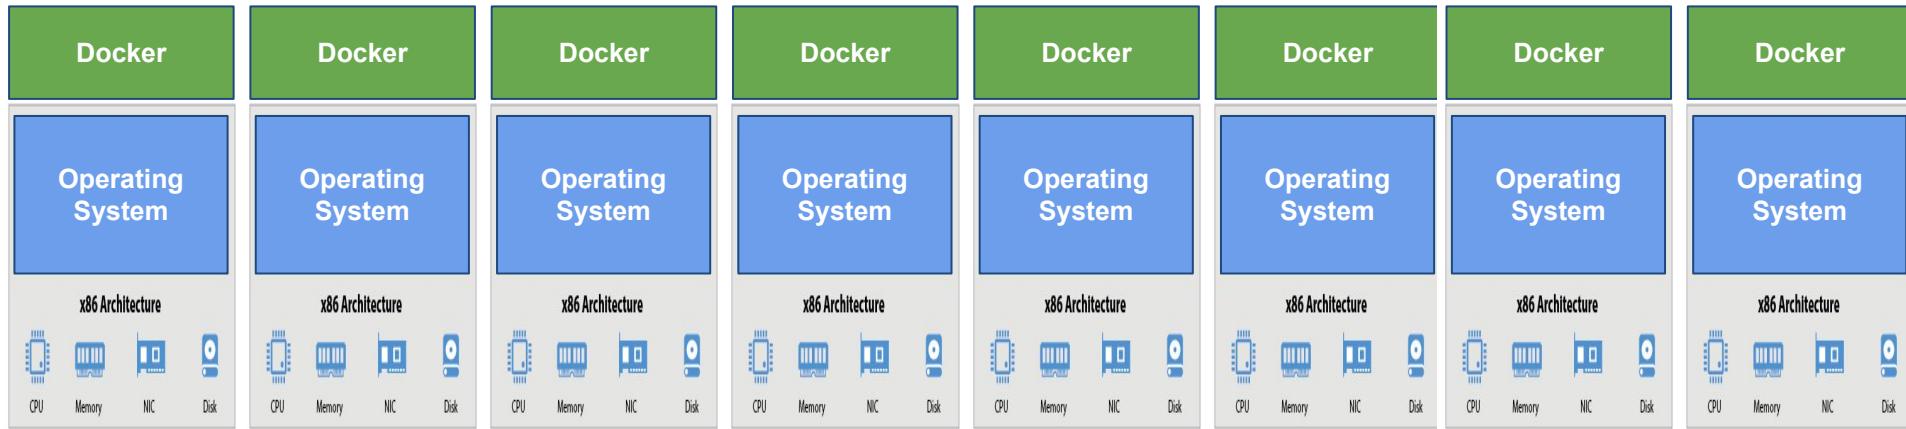
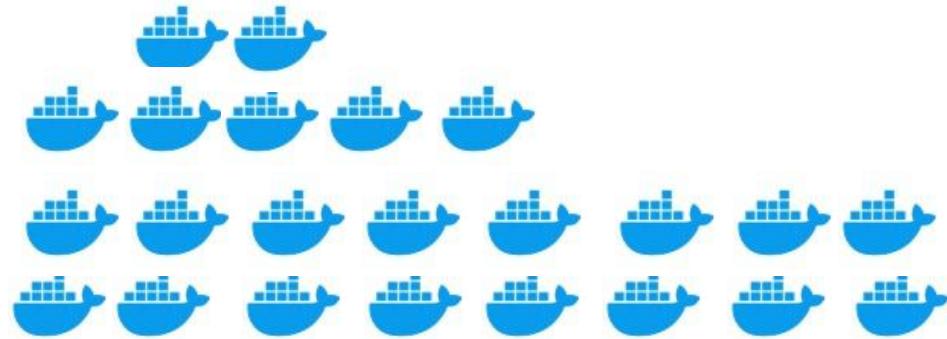
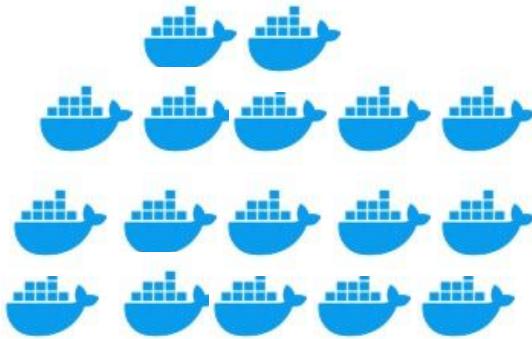


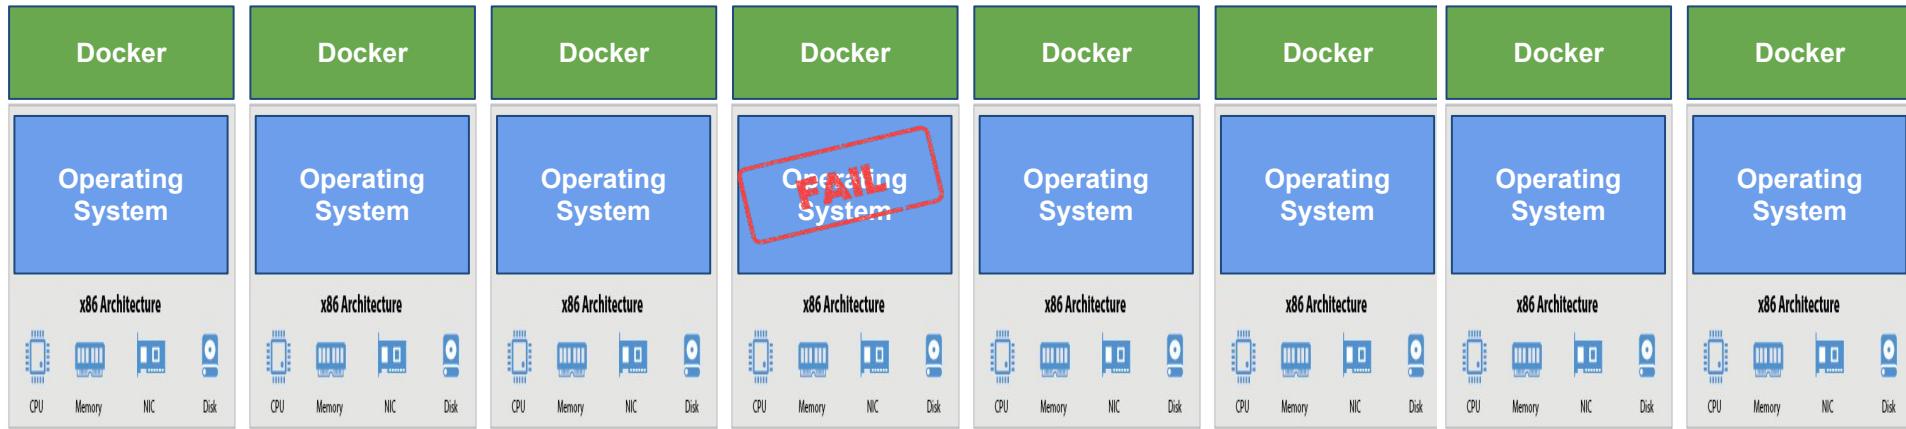
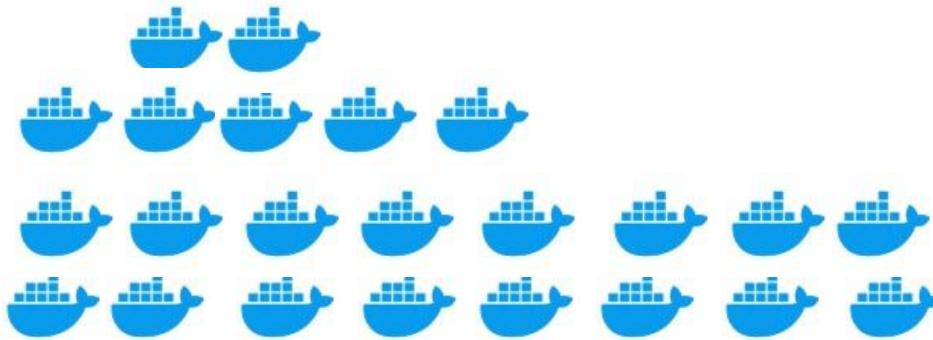
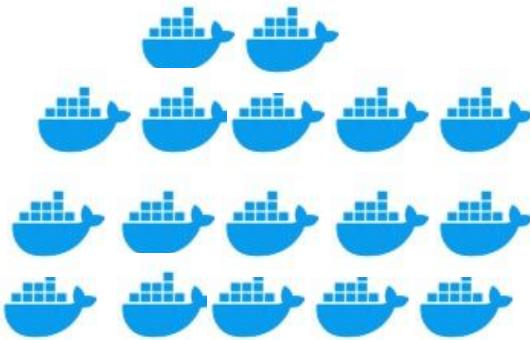


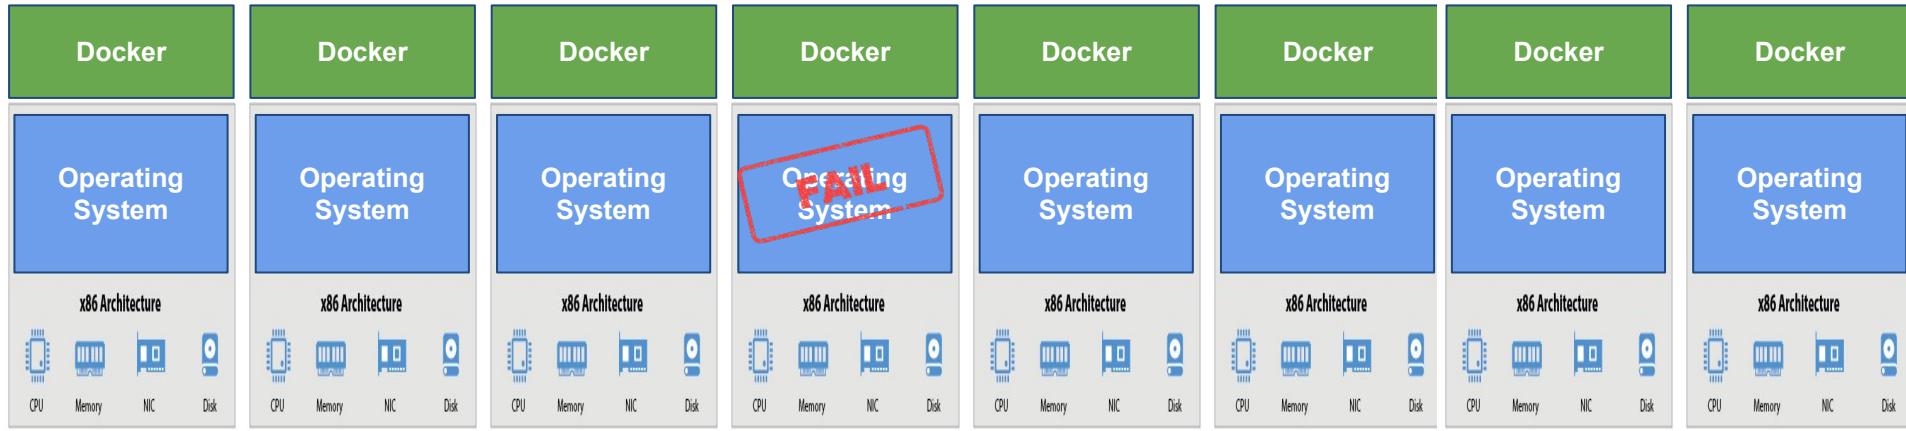
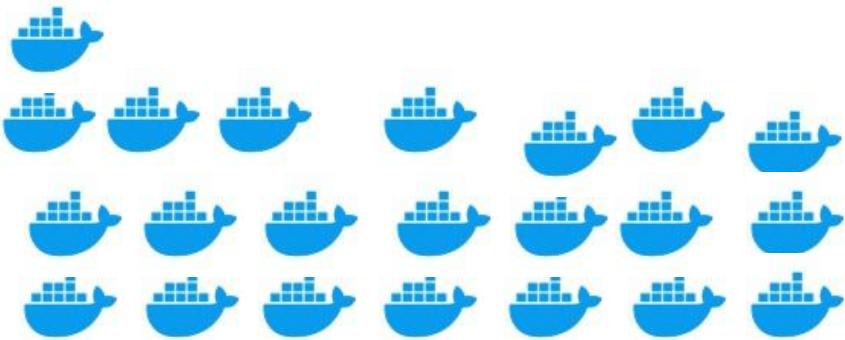
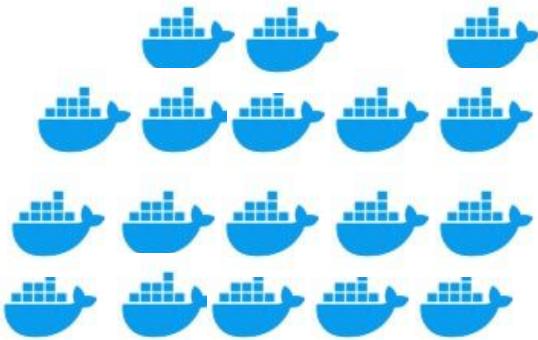


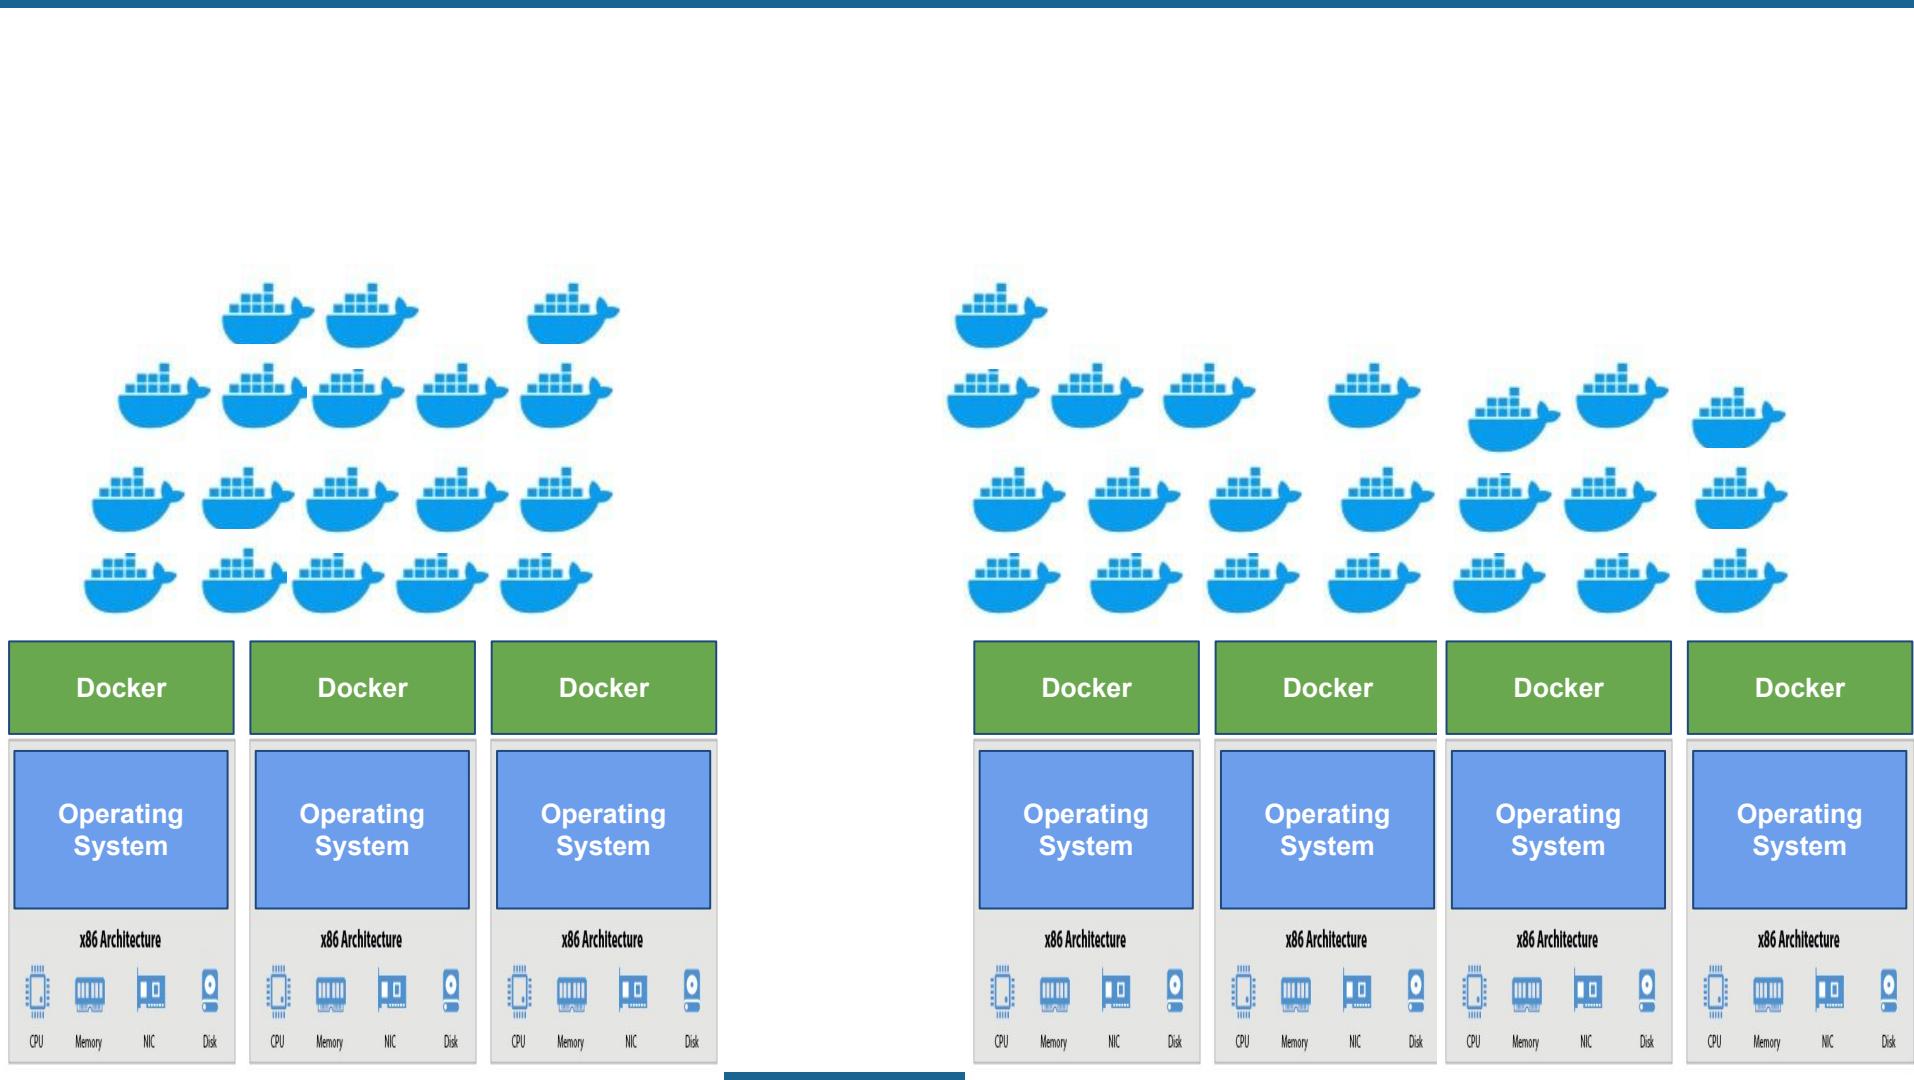


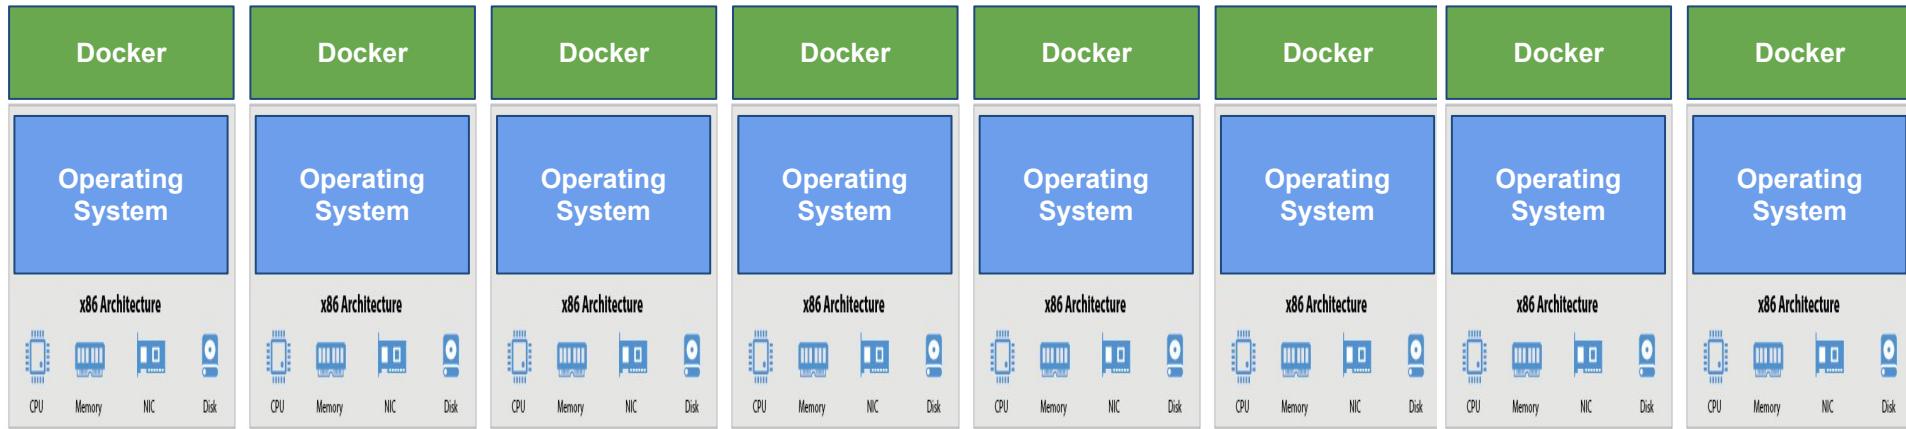
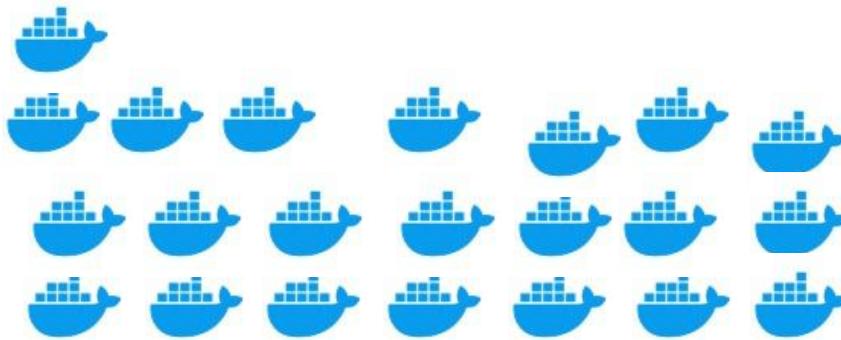
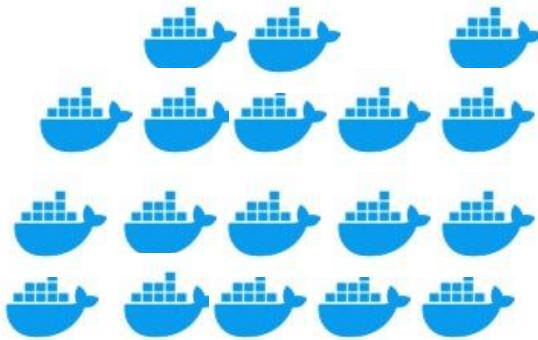


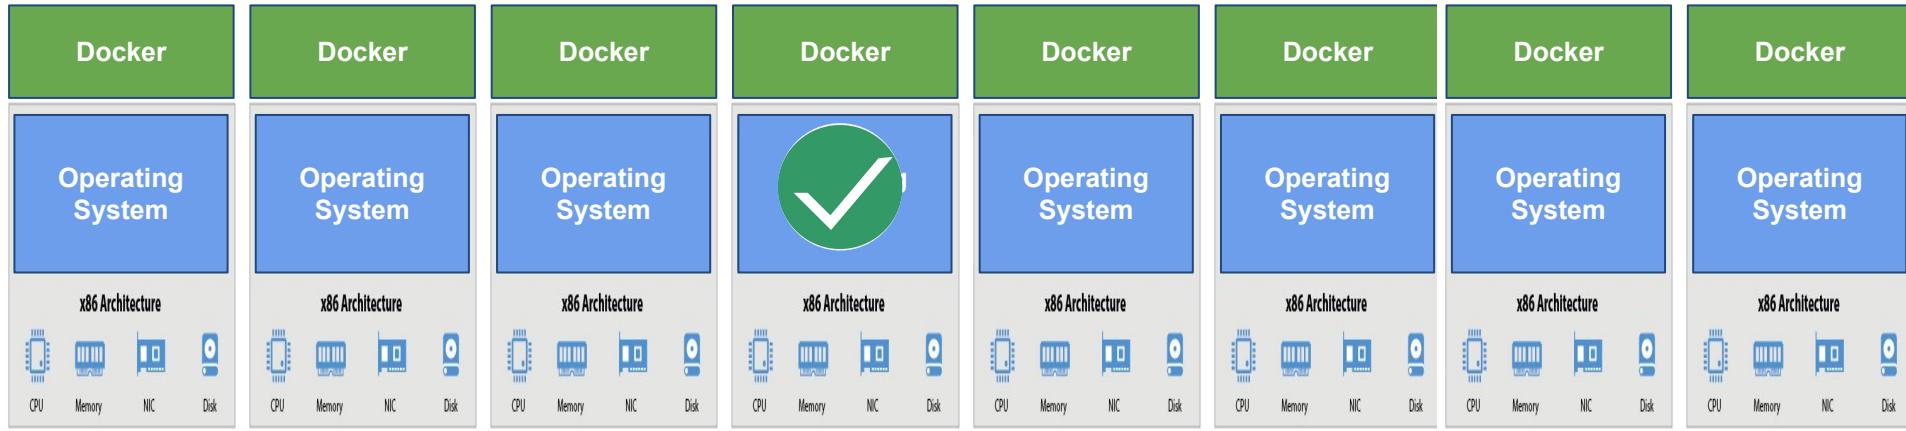
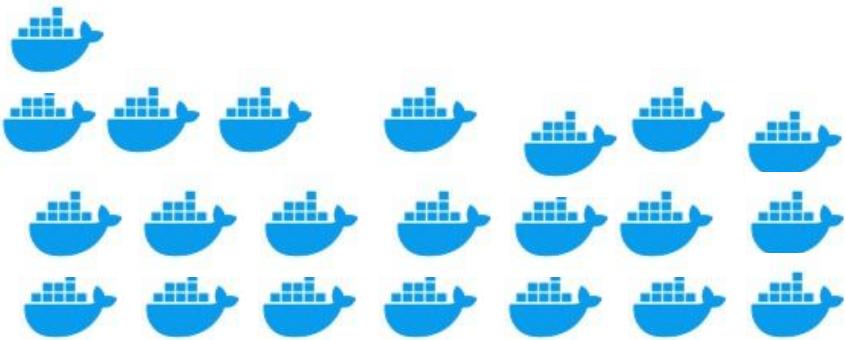
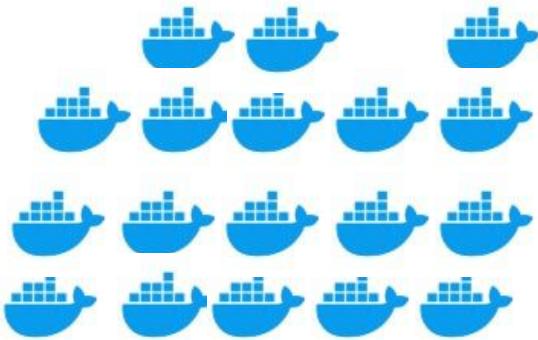


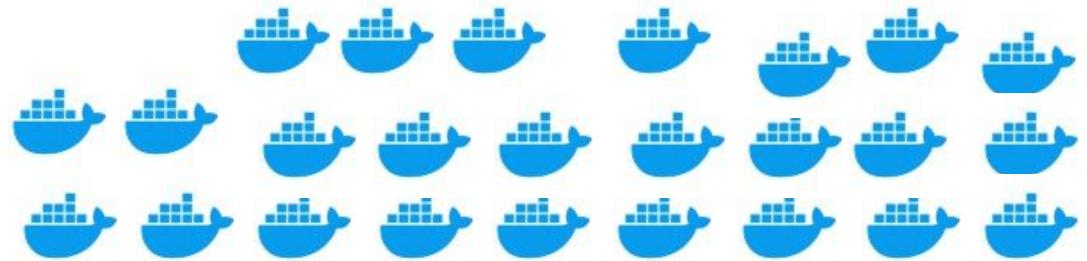
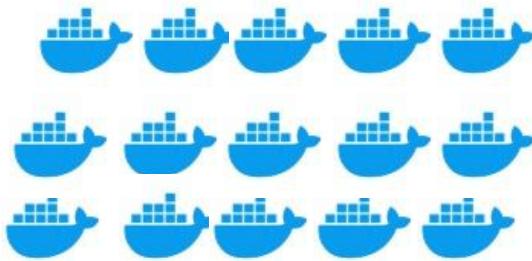












| Docker           |
|------------------|------------------|------------------|------------------|------------------|------------------|------------------|------------------|
| Operating System |
| x86 Architecture |
| CPU              | Memory           | NIC              | Disk             | CPU              | Memory           | NIC              | Disk             |

# Final Note

## 01 What does Kubernetes Provide for the Infrastructure

### Automated Infrastructure

- Compute
- Networking
- Storage
- Firewall
- Monitoring
- Alerting

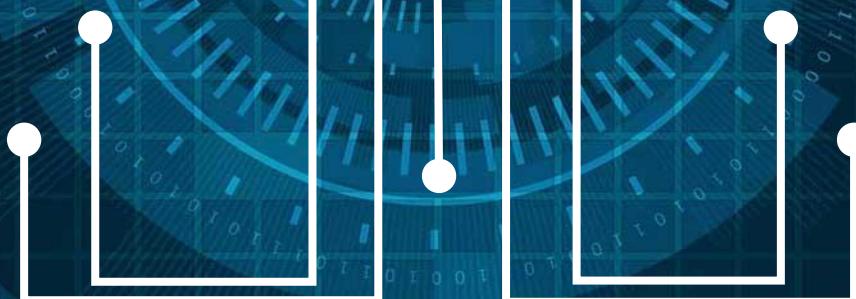
## 02 What does Kubernetes Provide for the Developer

### Automated Deployment

- Deployment
- Self Healing
- Load Balancing
- Service Discovery
- Monitoring
- Logging
- Alerting



# THANK YOU



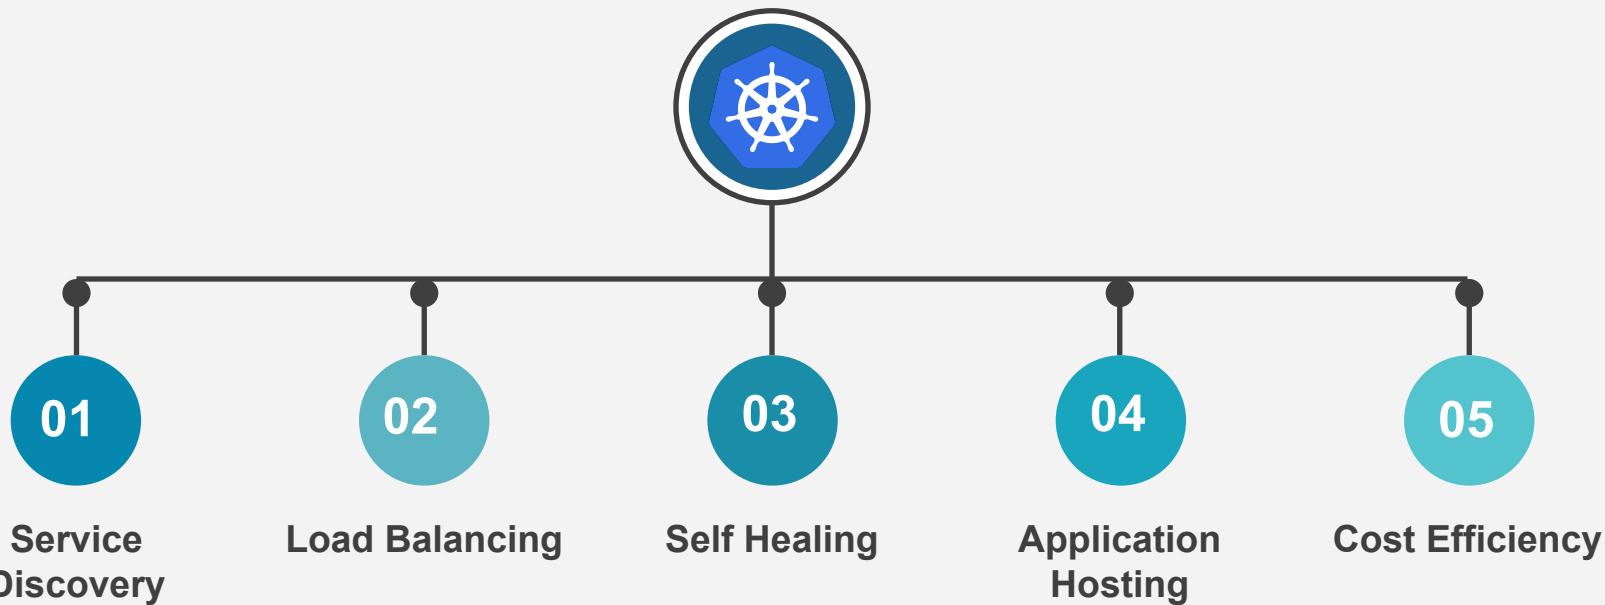
Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

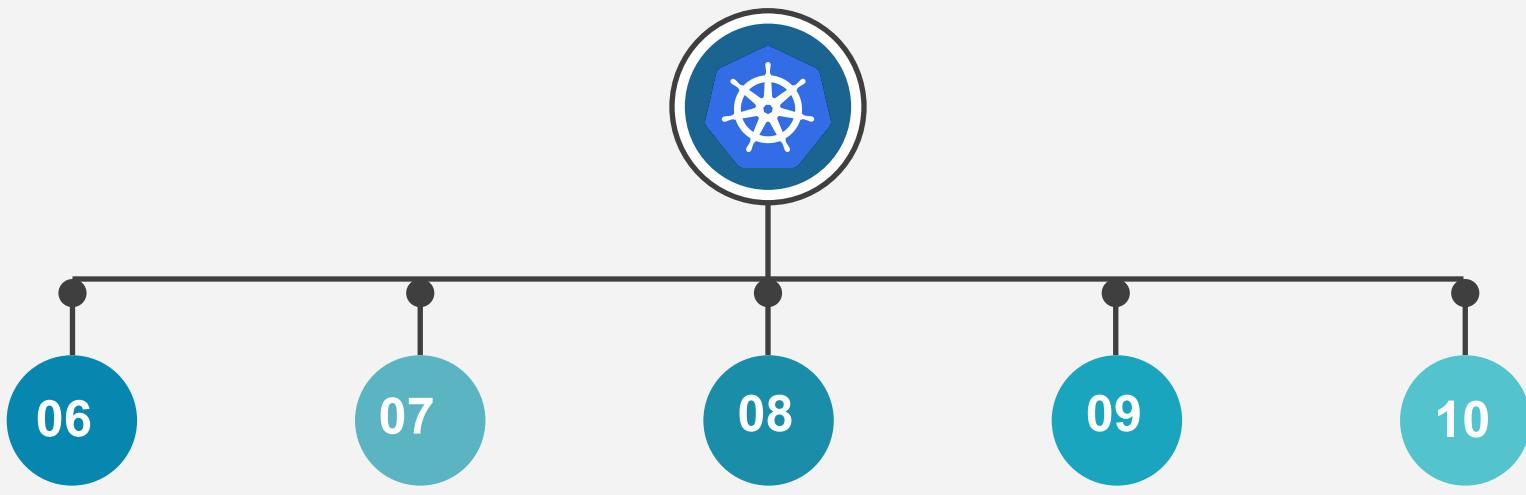
03 - Kubernetes Features

Ahmed ElBakry

# Kubernetes Features



# Kubernetes Features



06 Faster Deployment  
07 Fine-grained Access Control  
08 Batch Jobs  
09 Autoscaling  
10 Declarative Configuration

# Kubernetes Features

01

## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

03

## Self Healing

Kubernetes ensures that the Desired state and Actual state are always in sync by monitoring the container healthcheck and restarting the containers which fail this healthcheck.

04

## Application Hosting

Kubernetes hosts containerized applications. So the developers doesn't care whether kubernetes is running on bare metal or on the cloud

# Kubernetes Features

05

## Faster Deployment

Kubernetes allows you to deliver a self-service Platform-as-a-Service (PaaS) that creates a hardware layer abstraction for development teams.

06

## Fine-grained Access Control

Kubernetes RBAC allows the cluster maintainers to configure fine-grained access control to Kubernetes resources

07

## Batch Jobs

Kubernetes supports running automated jobs with the CronJob object. They are similar to the Unix/Linux Cron tasks.

08

## Autoscaling

One of the cool features of Kubernetes is that it can monitor your workload and scale it up or down based on the CPU utilization or memory consumption

# Kubernetes Features

09

## Cost Efficiency

Kubernetes optimizes infrastructure cost by efficiently managing the infrastructure resources  
Reduction of administration time, etc.

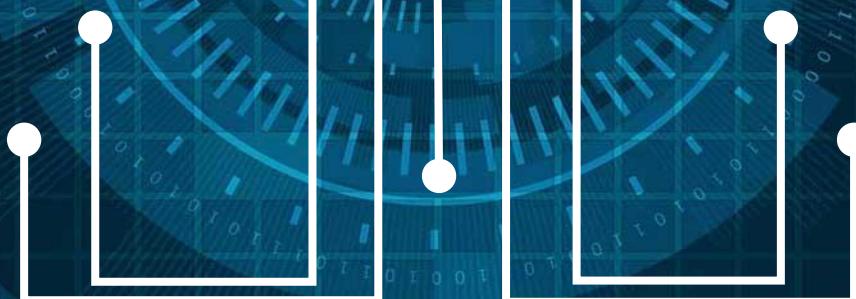
10

## Declarative Configuration

We provide the API server with the manifest files (YAML or JSON) describing how we want the Kubernetes cluster to look.



# THANK YOU



Ahmed ElBakry



## **Module 2**

### **Kubernetes Architecture**

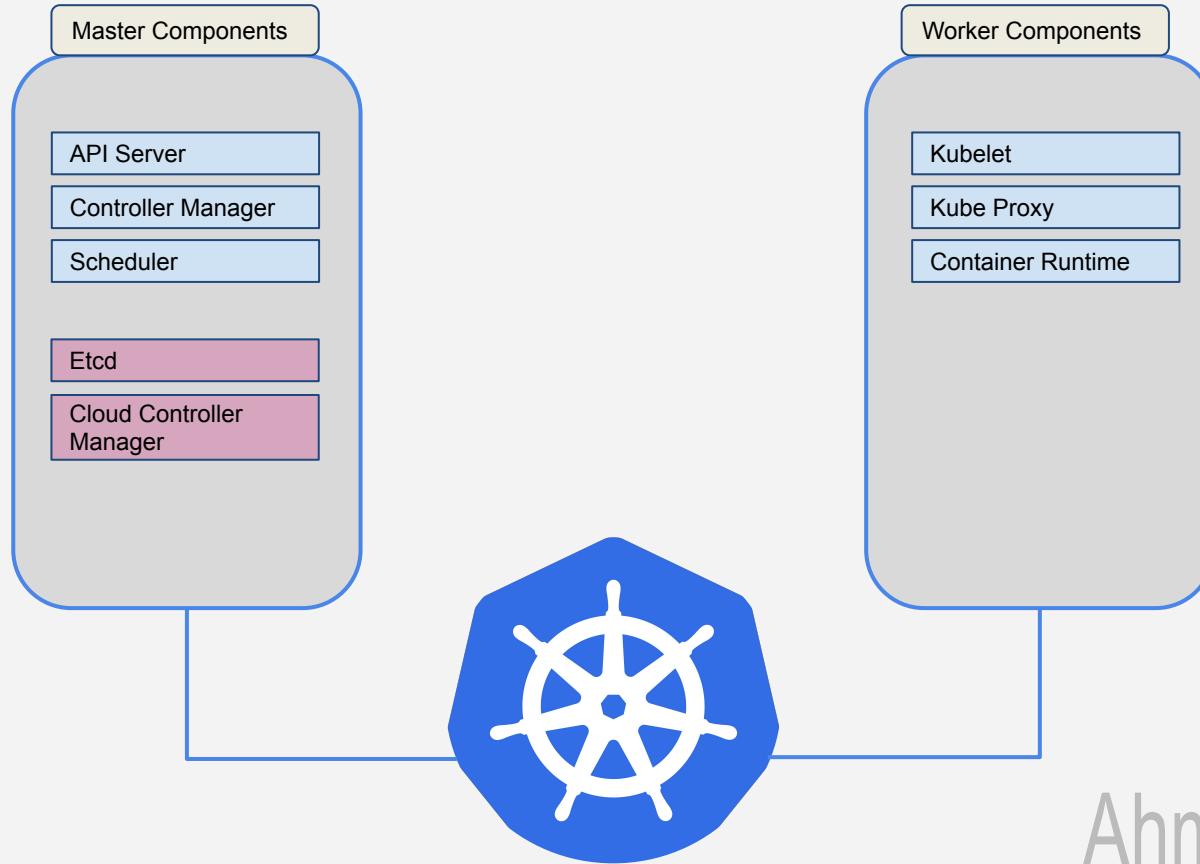
Ahm

# Kubernetes From Scratch - Hands-on

04 - Cluster Architecture

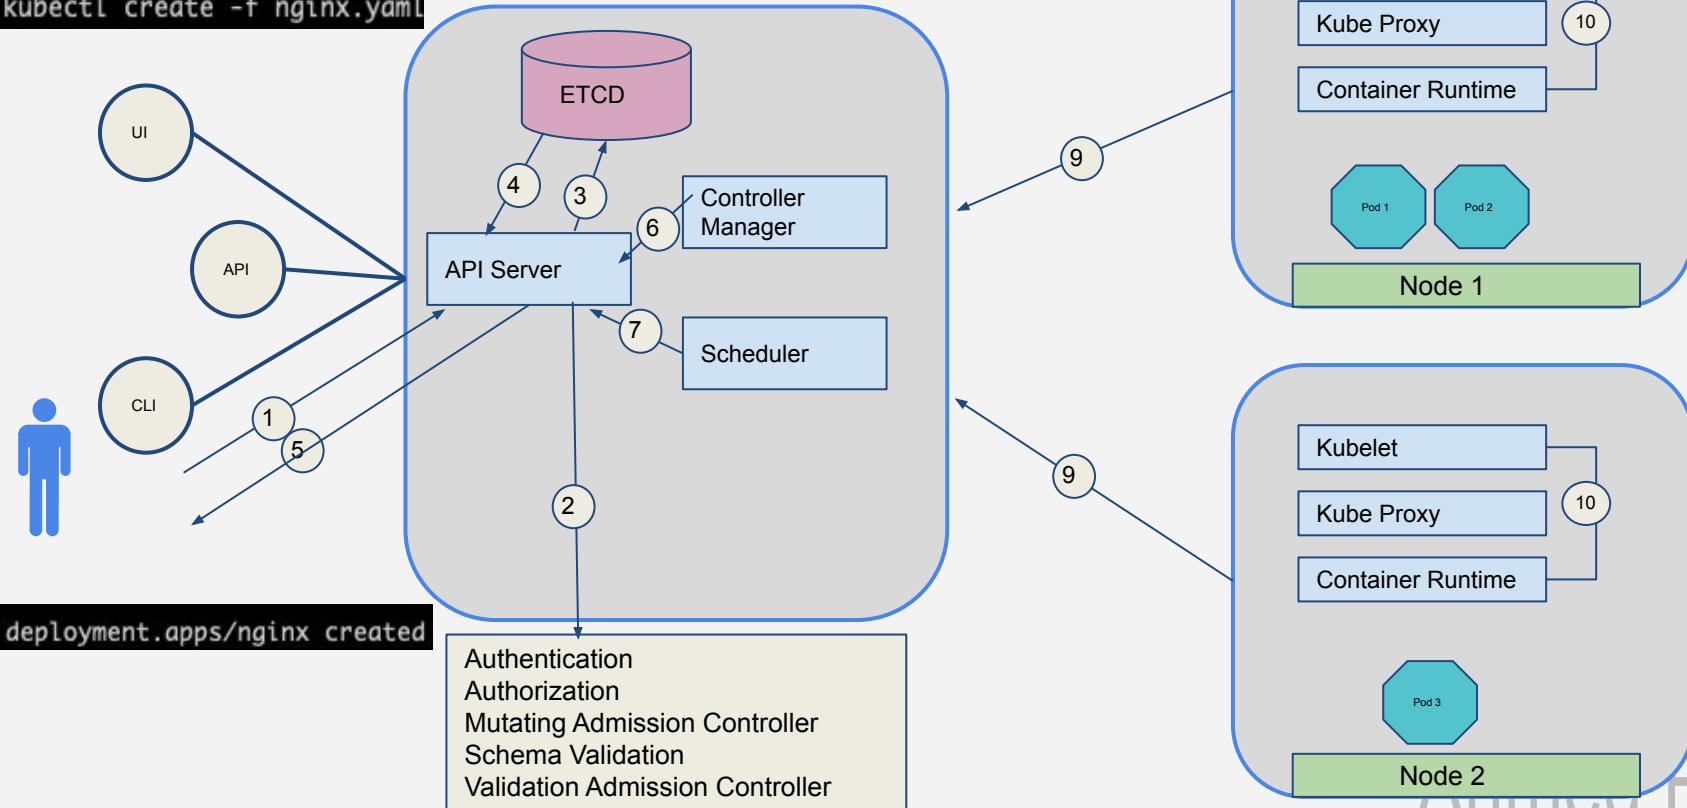
Ahmed ElBakry

# Cluster Architecture



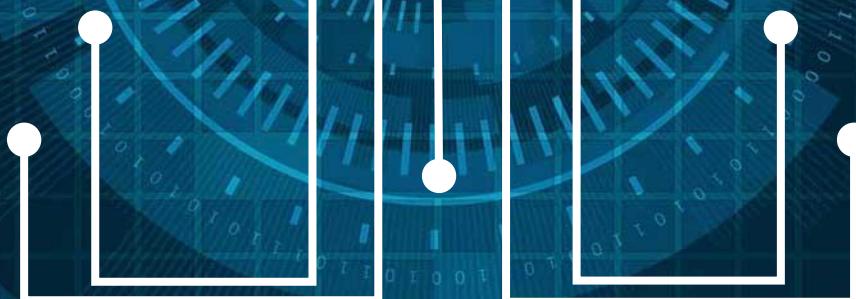
# Cluster Architecture

```
kubectl create -f nginx.yaml
```



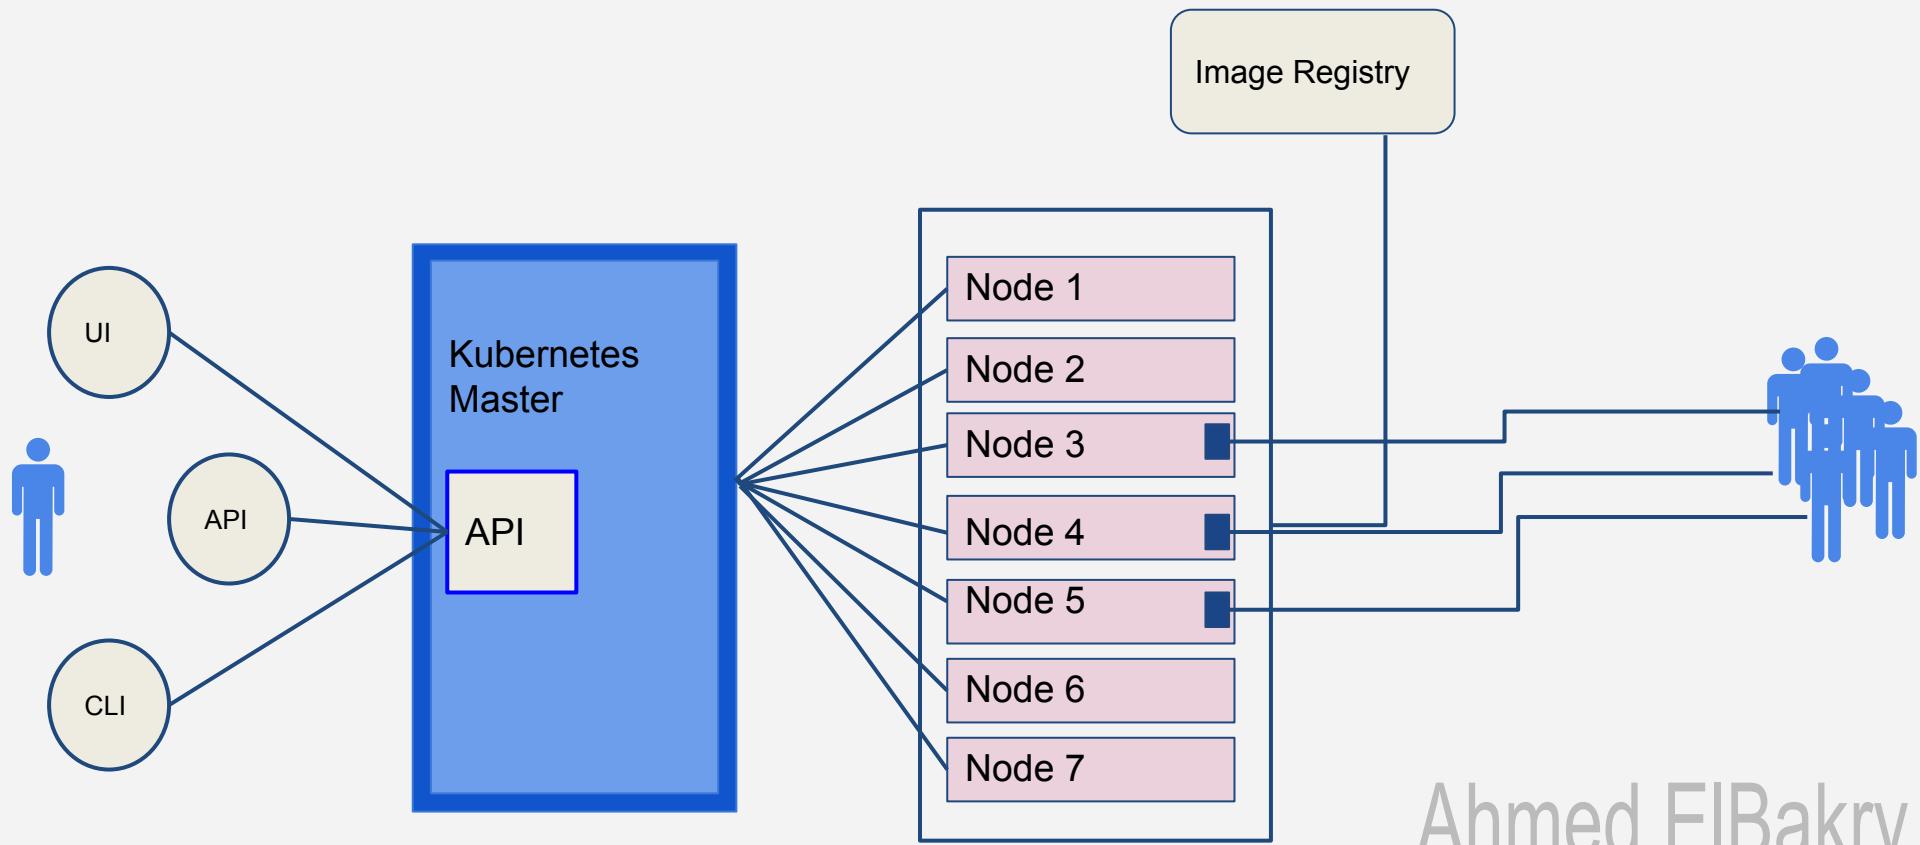


# THANK YOU



Ahmed ElBakry

# Cluster Architecture

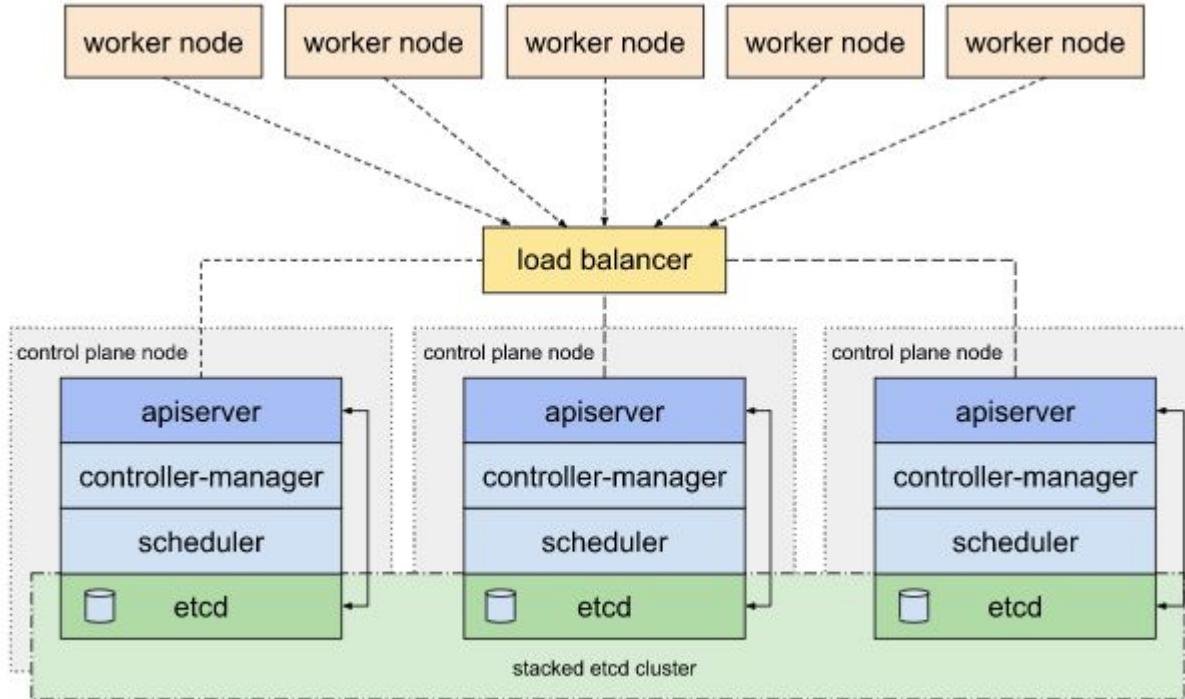


# Kubernetes From Scratch - Hands-on

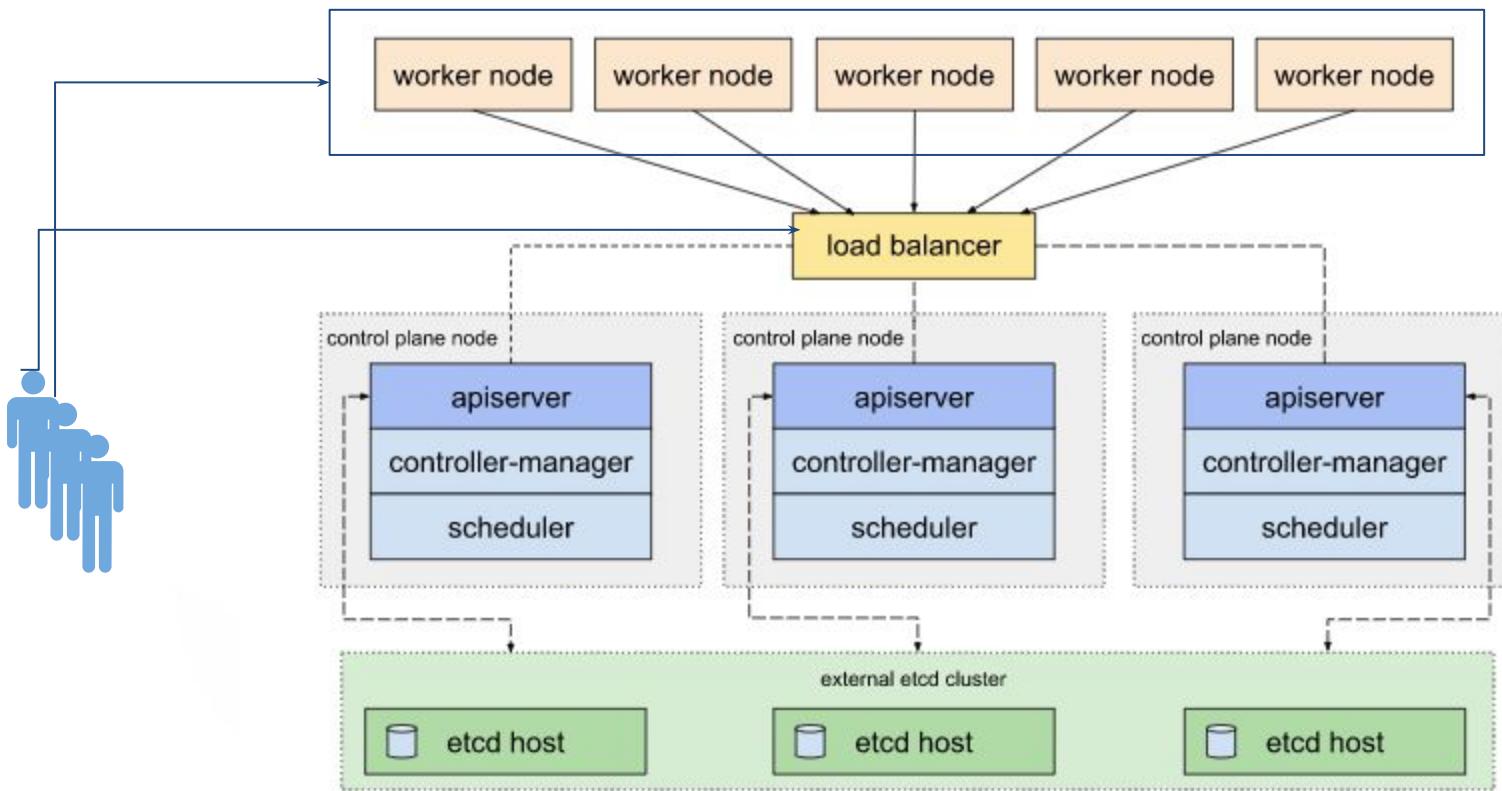
05 - Master Components

Ahmed ElBakry

# Master Components

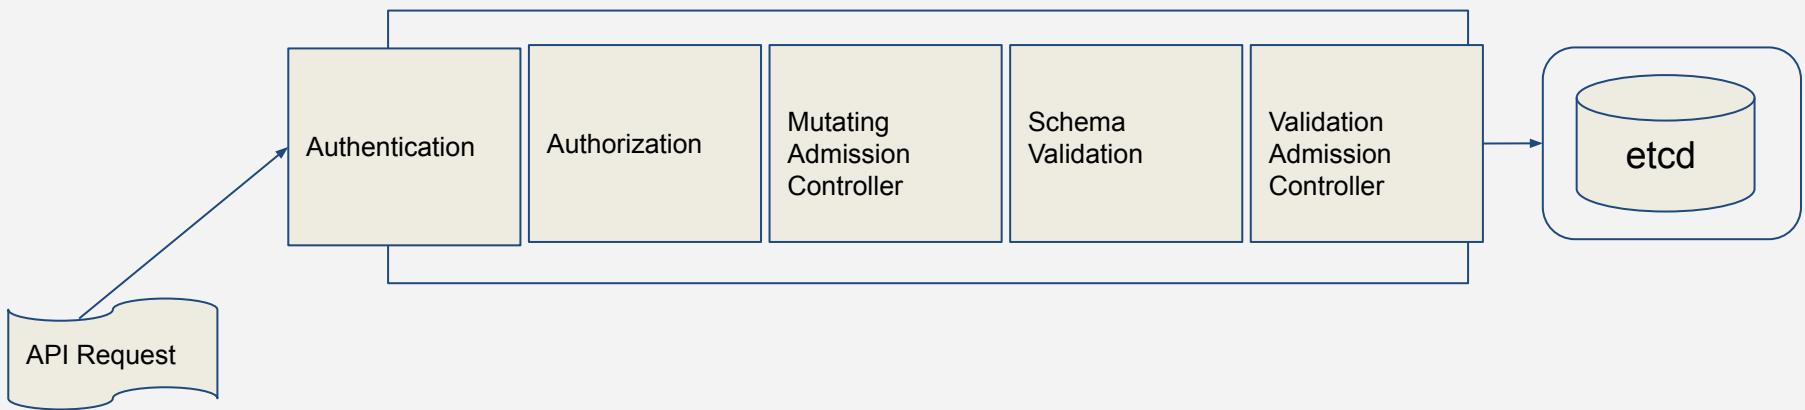


# Master Components



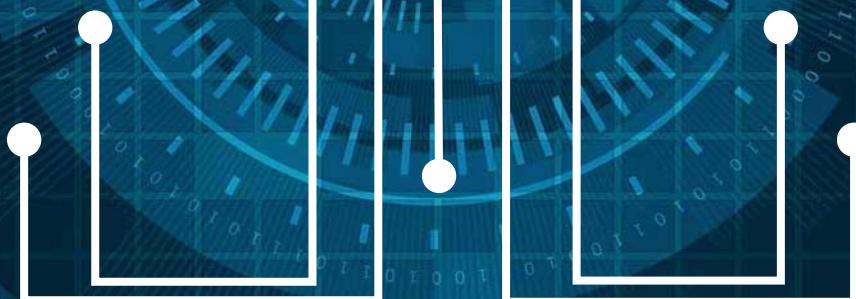
# Master Components

## Authentication, Authorization, and Admission Control





# THANK YOU



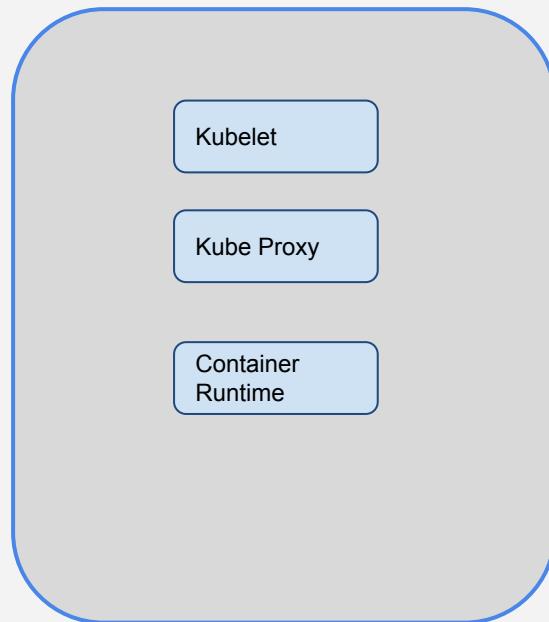
Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

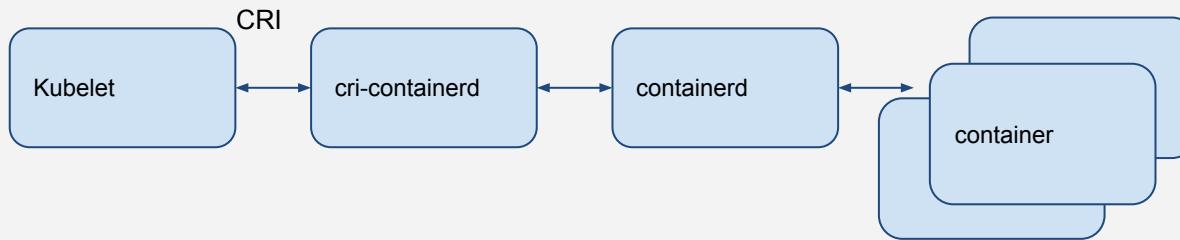
06 - Worker Components

Ahmed ElBakry

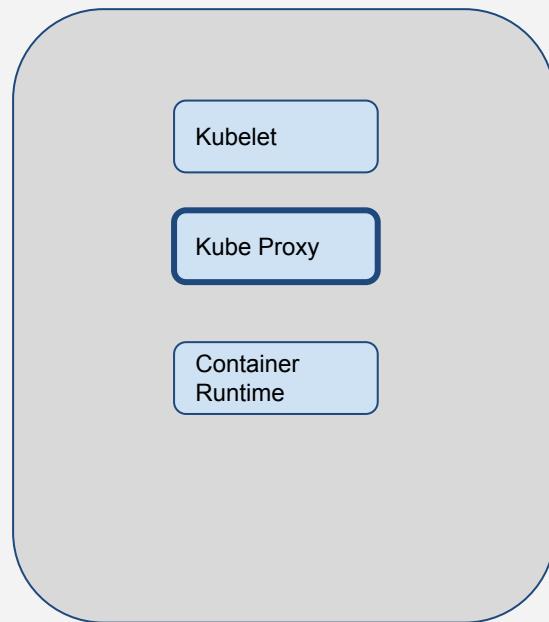
# Worker Components



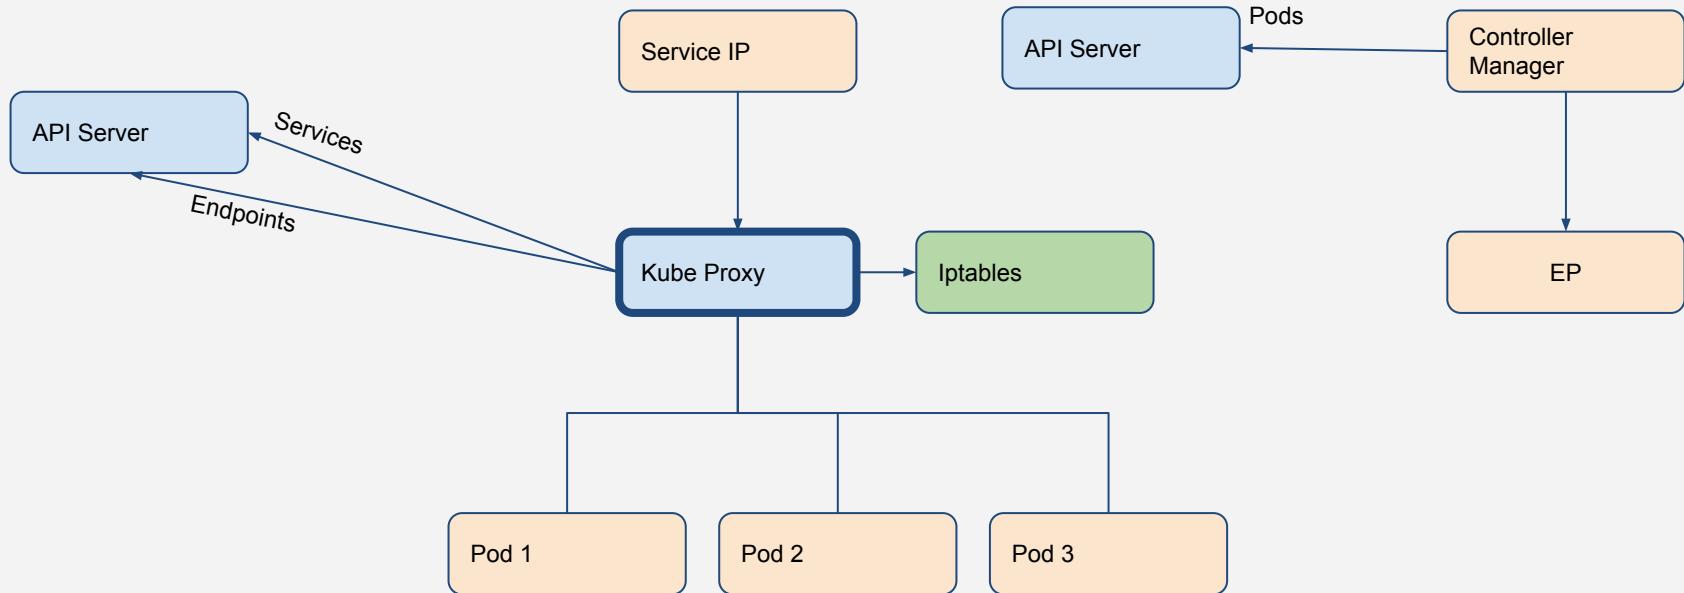
# Worker Components



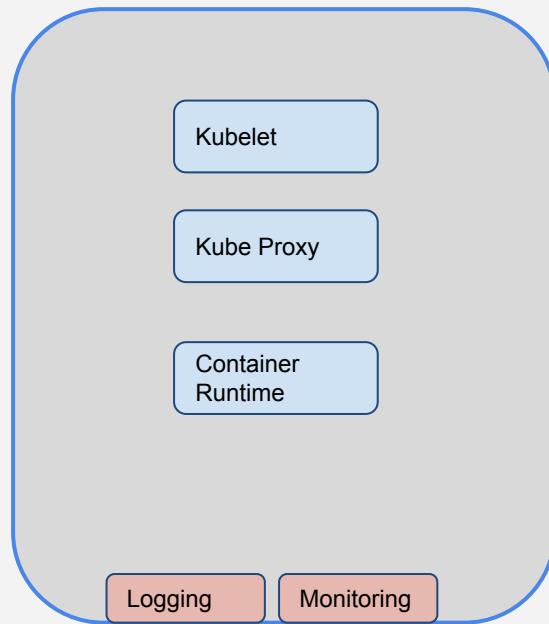
# Worker Components



# Worker Components

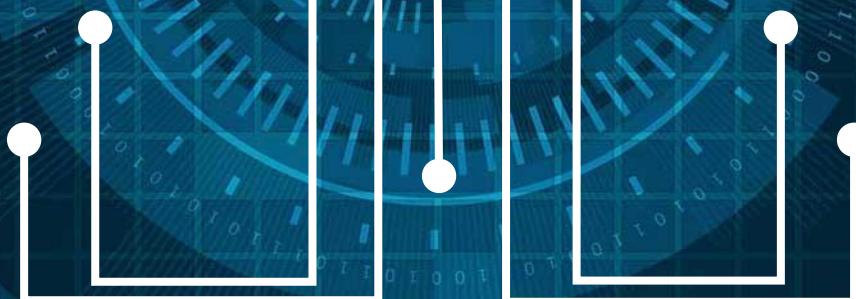


# Worker Components





# THANK YOU



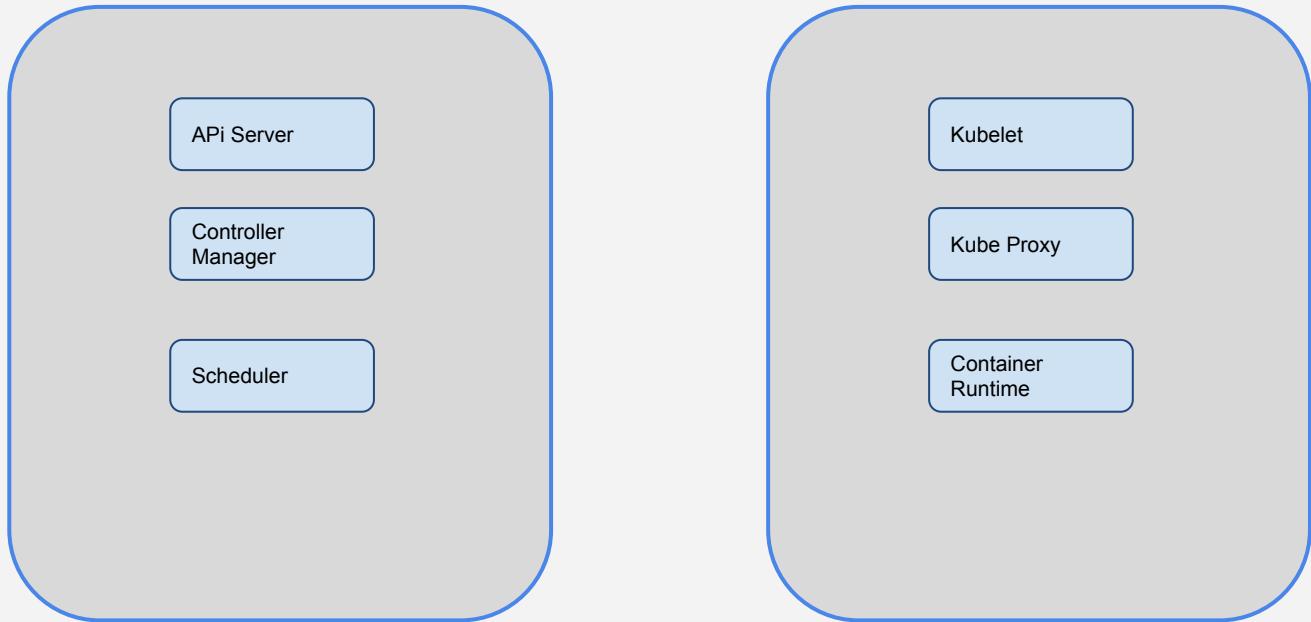
Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

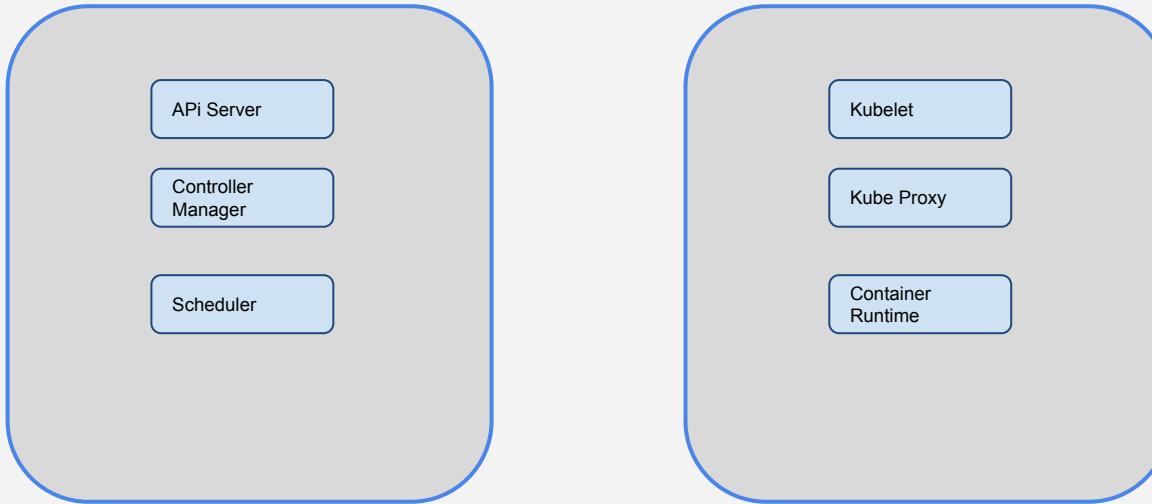
07 - Cluster Architecture Demo

Ahmed ElBakry

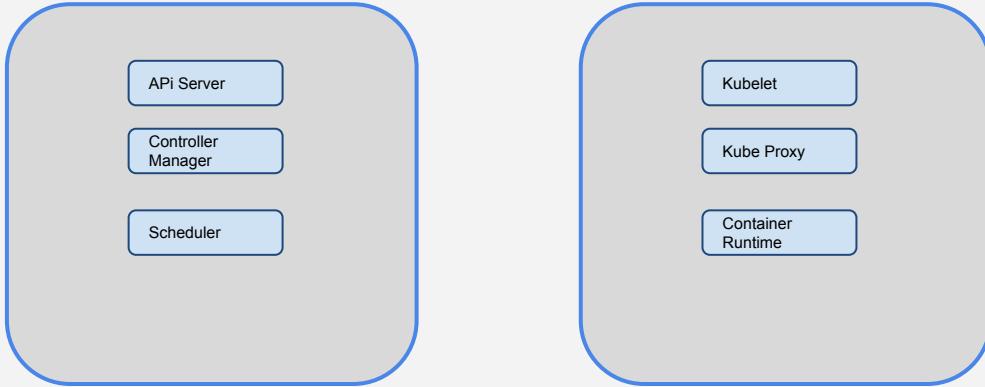
# Kubernetes Architecture Demo



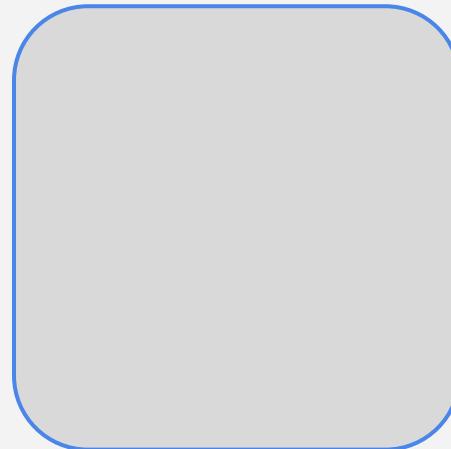
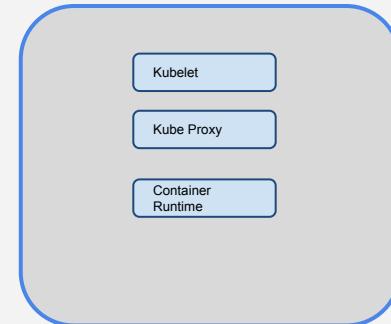
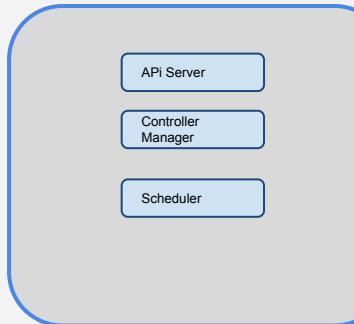
# Kubernetes Architecture Demo



# Kubernetes Architecture Demo

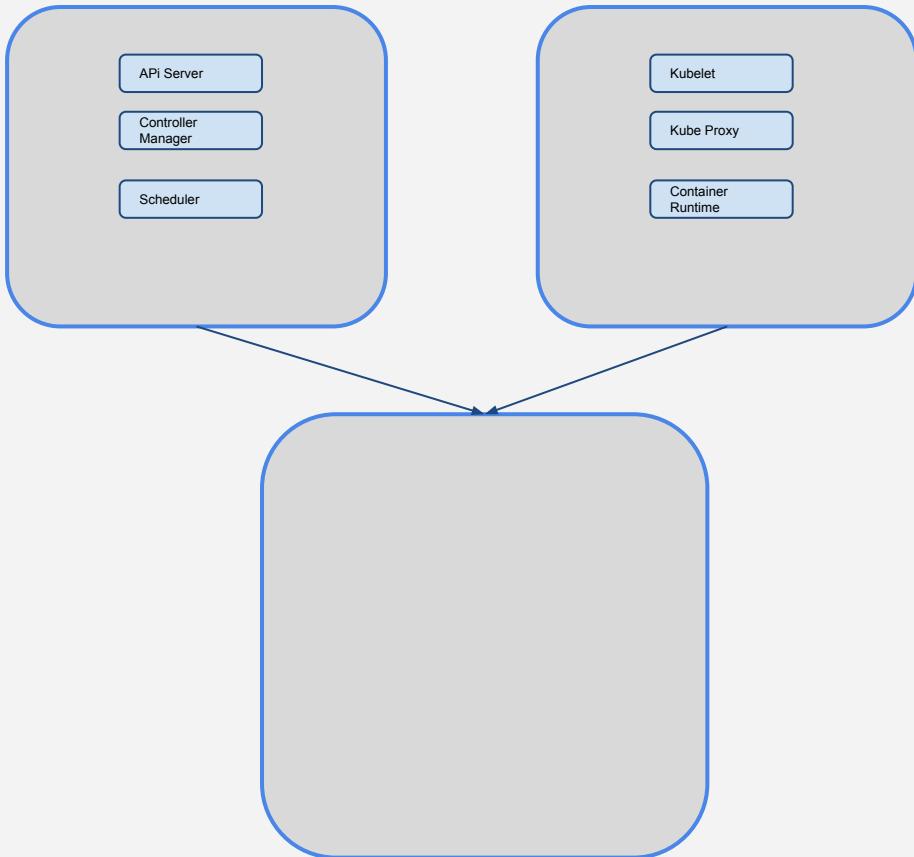


# Kubernetes Architecture Demo



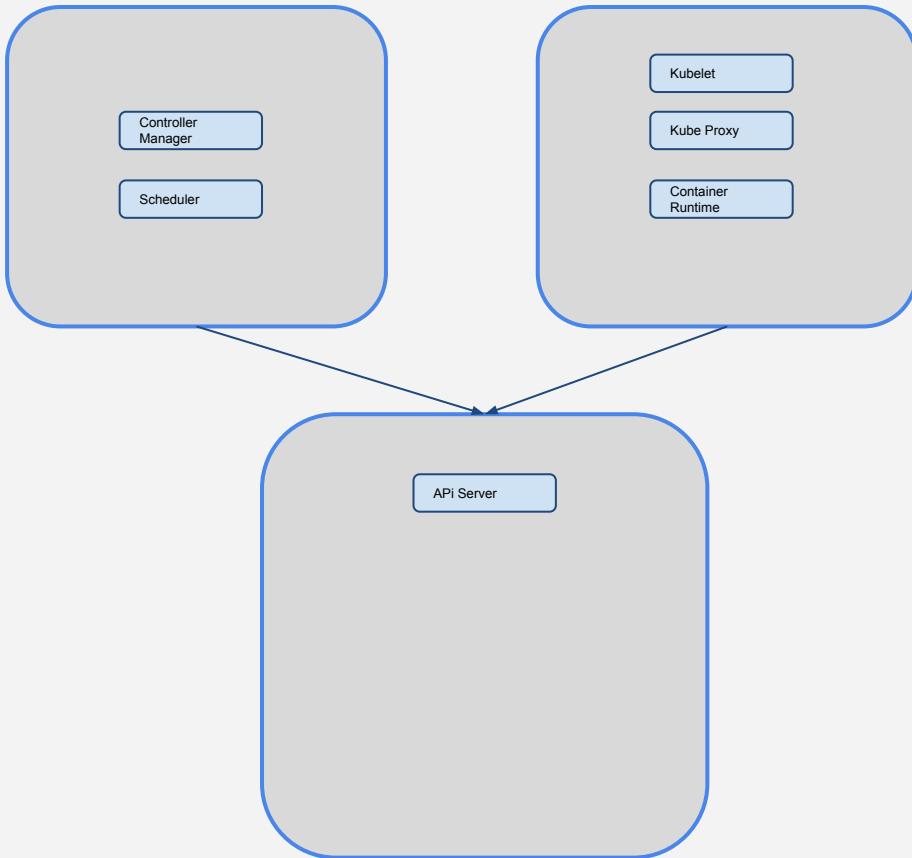
Ahmed ElBakry

# Kubernetes Architecture Demo



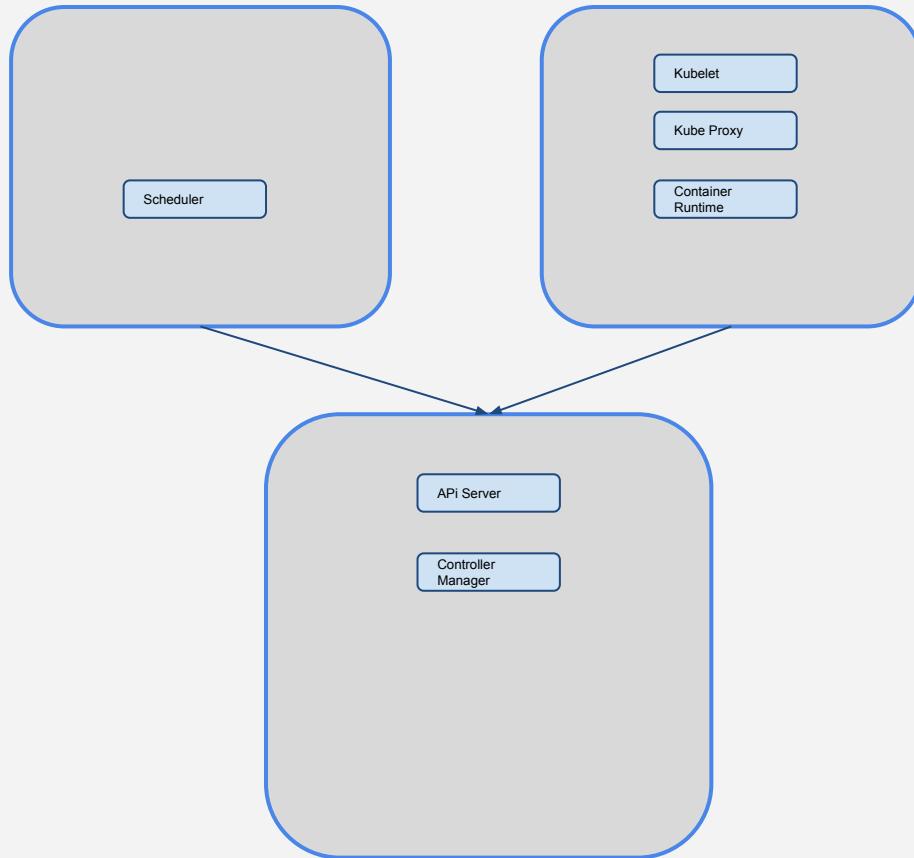
Ahmed ElBakry

# Kubernetes Architecture Demo



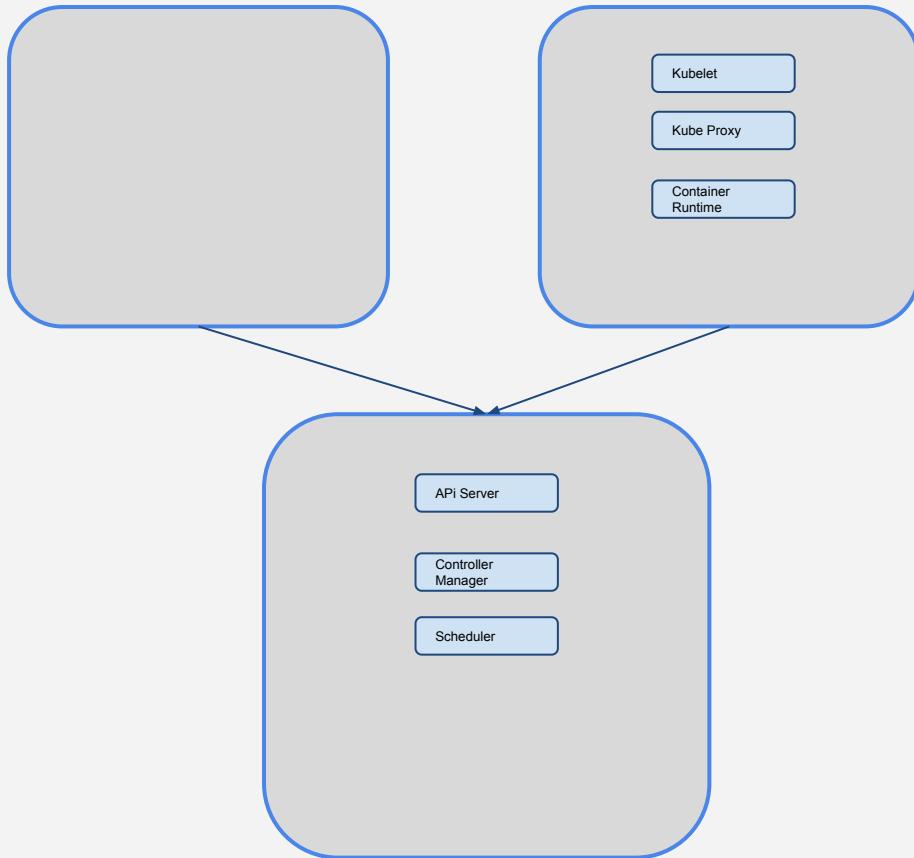
Ahmed ElBakry

# Kubernetes Architecture Demo



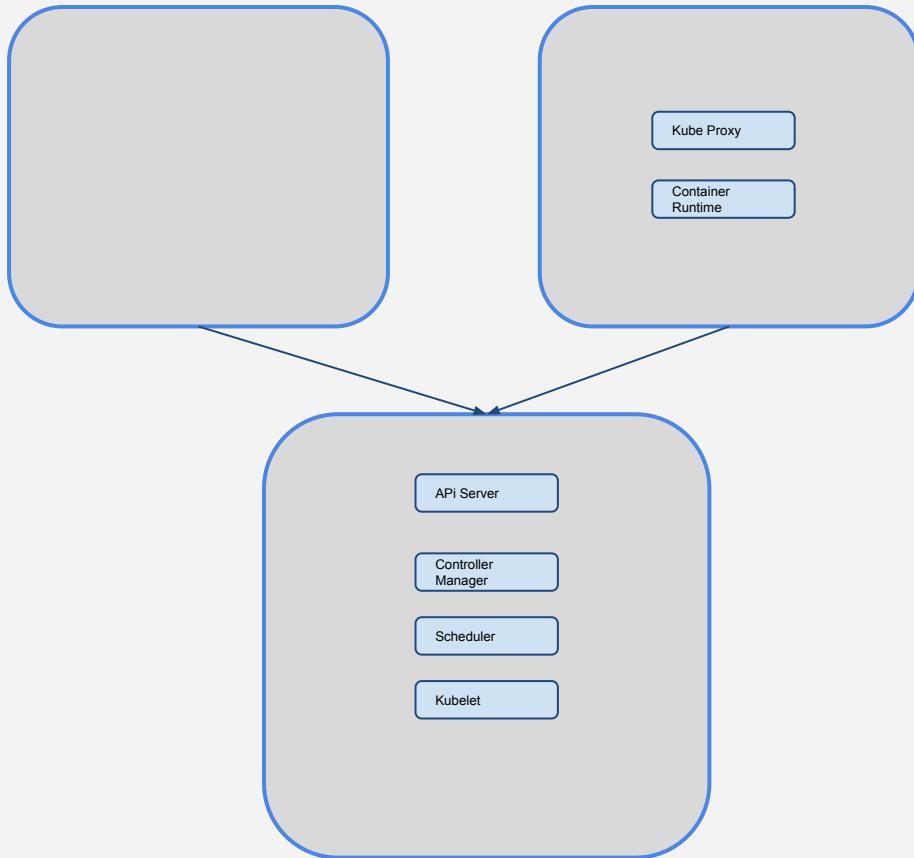
Ahmed ElBakry

# Kubernetes Architecture Demo



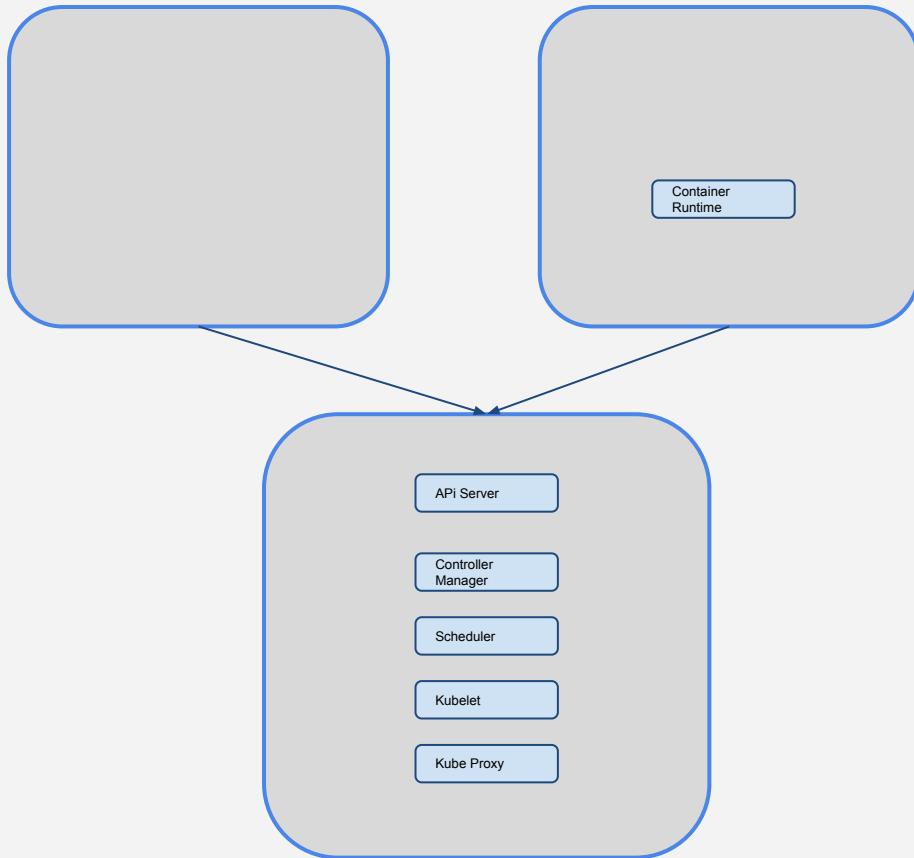
Ahmed ElBakry

# Kubernetes Architecture Demo



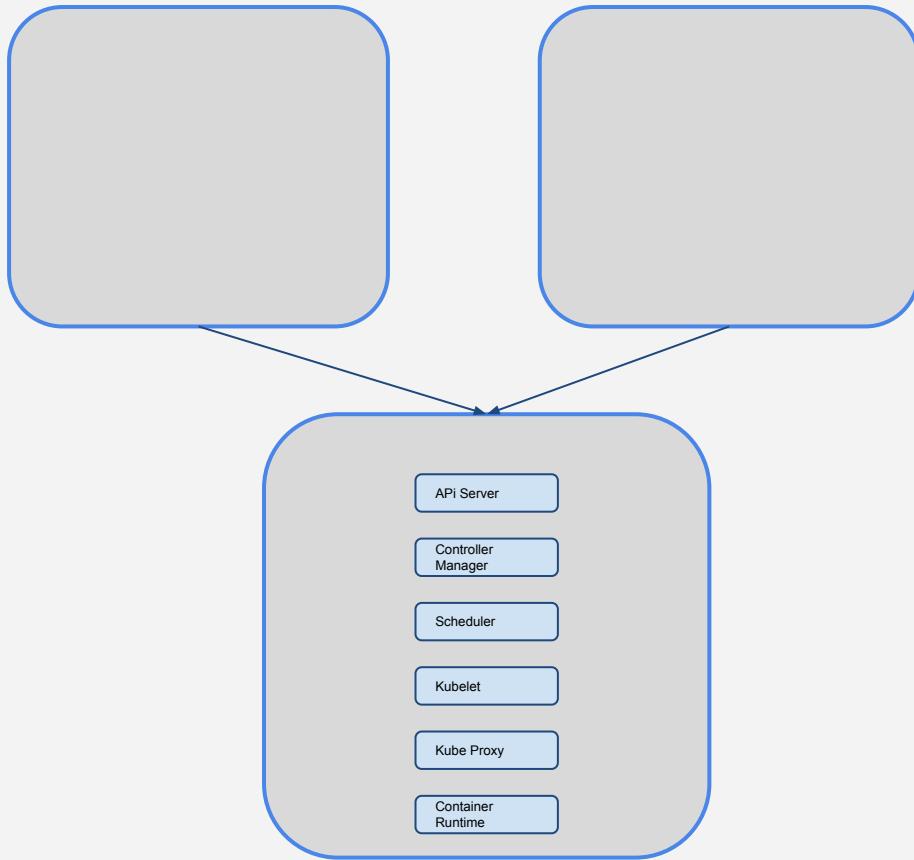
Ahmed ElBakry

# Kubernetes Architecture Demo



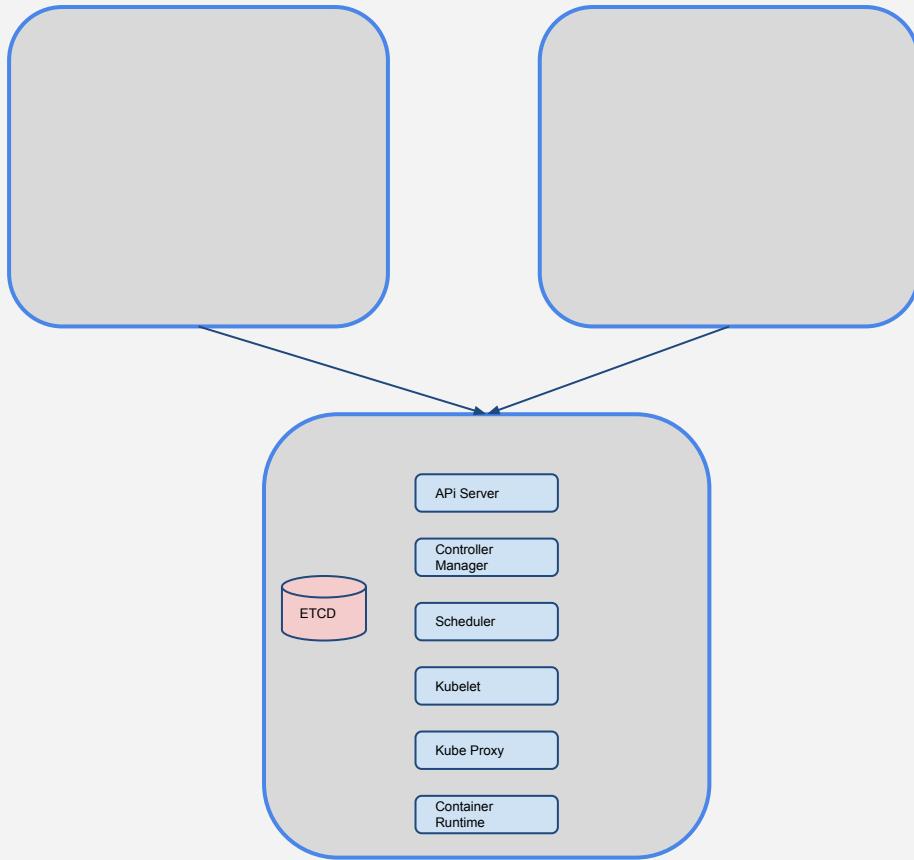
Ahmed ElBakry

# Kubernetes Architecture Demo



Ahmed ElBakry

# Kubernetes Architecture Demo



Ahmed ElBakry



# **Module 3**

## **Setup Lab Cluster**

Ahm

# Setup Lab Cluster

## In this module



Setup Single  
Node Cluster  
Using Minikube



Setup Multi  
Node Cluster  
Using Kind



PRIVATE DOCKER CONTAINER IMAGES

Setup Internal/  
Private Docker  
Registry

Ahmed ElBakry

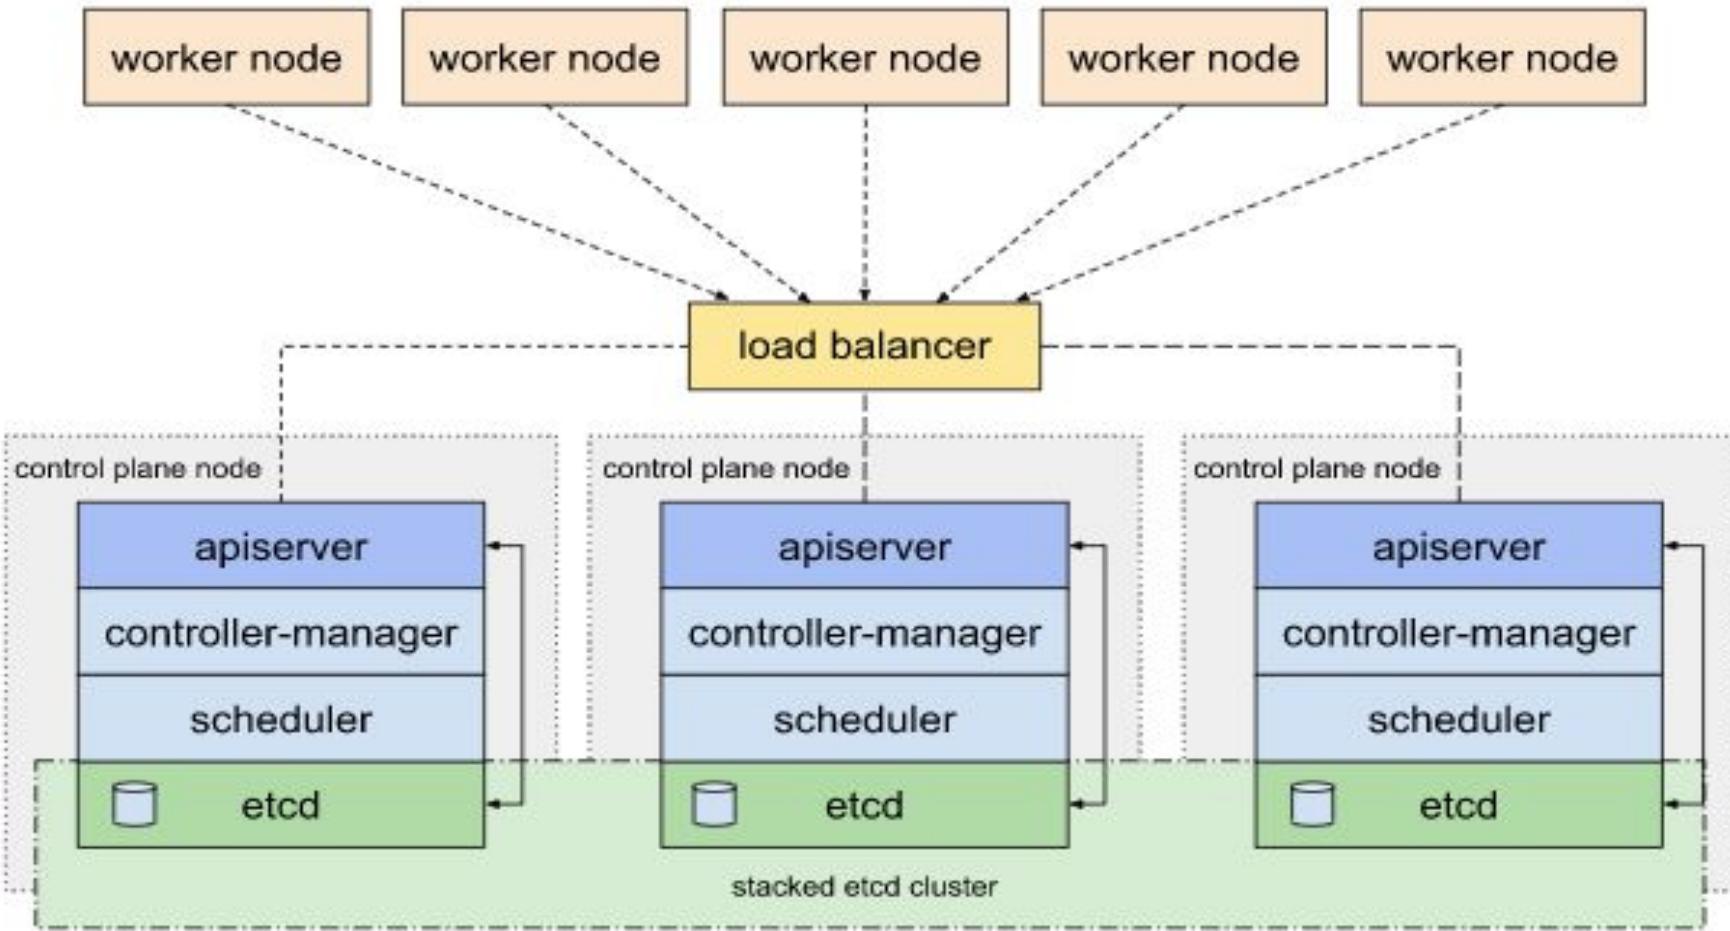
# Setup Single Node Cluster

## In this video



- What is Minikube
- Install Hypervisor (VirtualBox)
- Install Minikube
- Setup Kubectl

Ahmed ElBakry



# What is Minikube

- Minikube is a tool that makes it easy to run Kubernetes locally
- Minikube runs a single-node Kubernetes cluster inside a Virtual Machine (VM)
- Supports the following features



- DNS
  - NodePorts
  - ConfigMaps and Secrets
  - Dashboards
  - Container Runtime: Docker, CRI-O, and containerd
  - Enabling CNI (Container Network Interface)
  - Ingress
- Requirements:
- 2 CPU or more
  - 2GB of free memory
  - 20GB of free disk space
  - Internet connection
  - Container or virtual machine manager, such as: Docker, Hyperkit, Hyper-V, KVM, Parallels, Podman, VirtualBox, or VMWare

## Minikube

### Master processes

- ❖ api-server
- ❖ controller-manager
- ❖ scheduler
- ❖ etcd

### Worker processes

- ❖ kube-proxy
- ❖ kubelet
- ❖ container runtime

Ahmed ElBakry

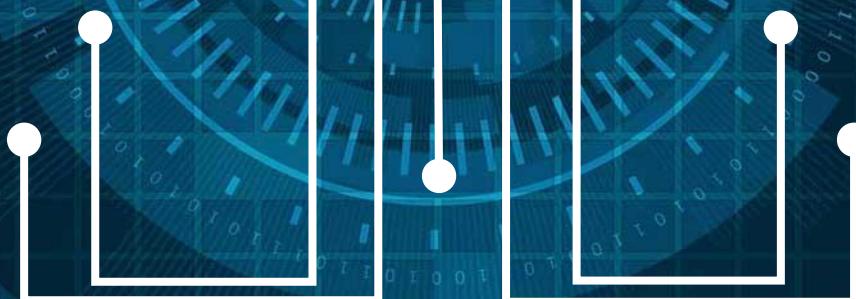


**DEMO**

Amr ElBakry



# THANK YOU



Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

09 - Setup Multi Node Cluster (Kind)

Ahmed ElBakry

# Setup Multi Node Cluster

## In this video



- What is Kind
- Install Docker
- Install Kind
- Explore The Cluster
- Check The kubeconfig File

# What is Kind

- kind is a tool for running local Kubernetes clusters using Docker container “nodes”.
- kind supports Linux, macOS and Windows
- Run multiple clusters on different versions.
- Can be used for
  - Testing Kubernetes upgrades
  - Testing components with Kubernetes new versions
  - Testing with CI/CD



بالعربي DevOps-Kubernetes-Docker



Ahmed ElBakry

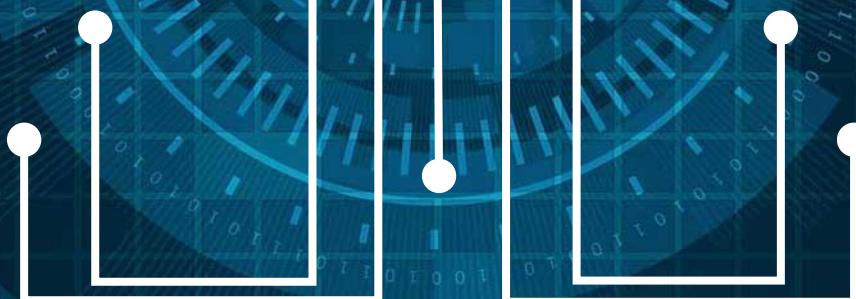


**DEMO**

Amr ElBakry



# THANK YOU



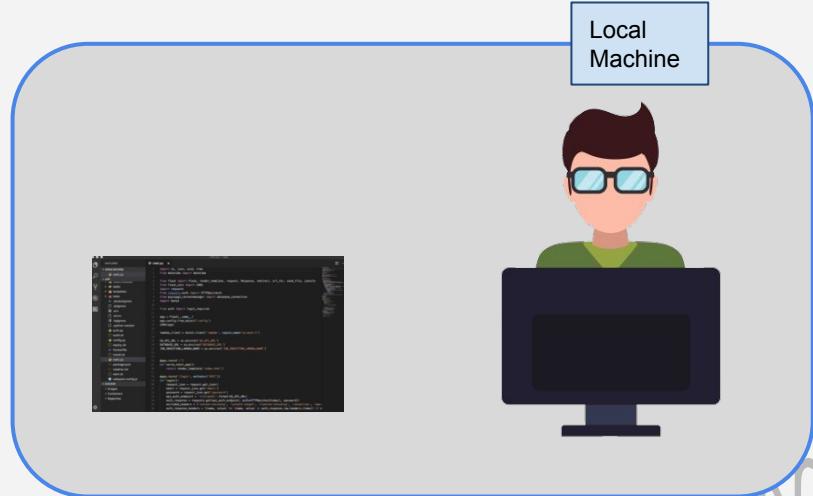
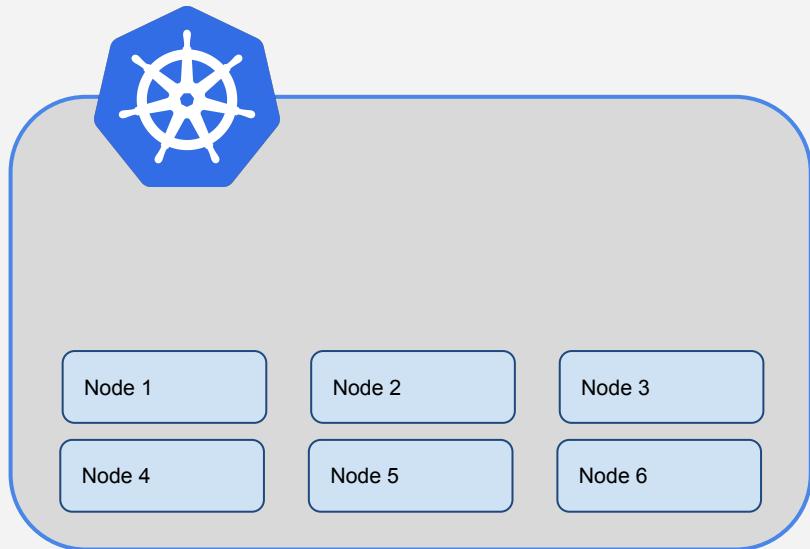
Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

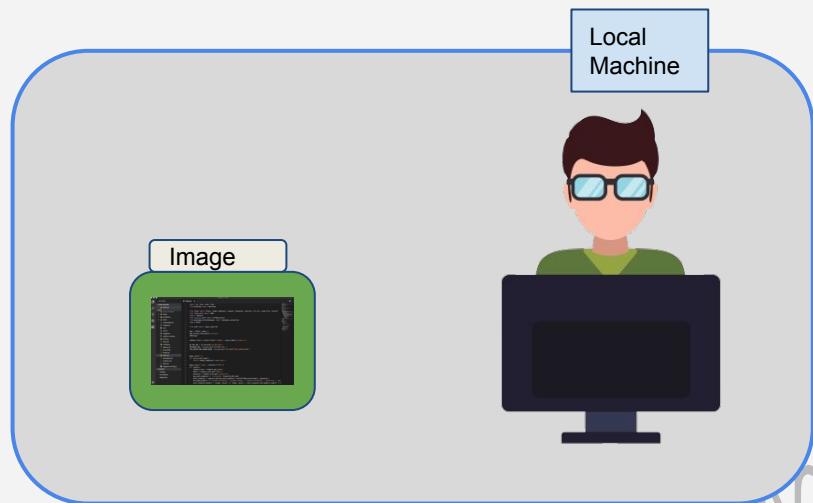
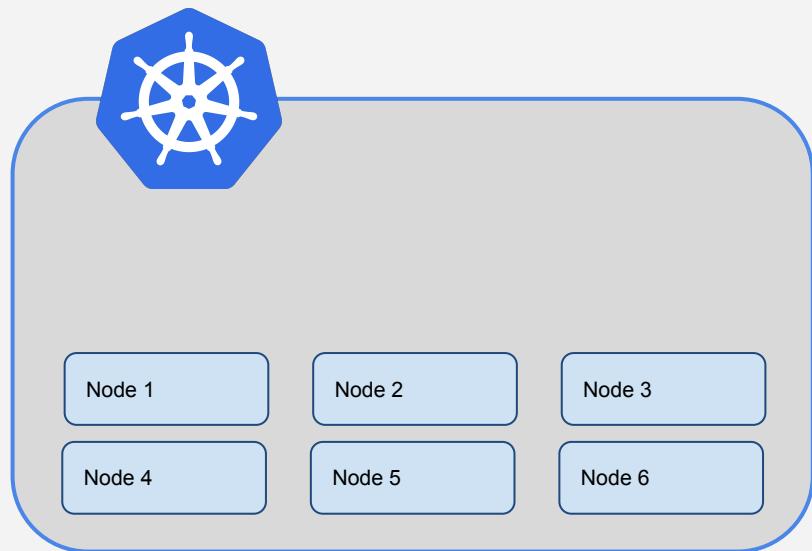
10 - Setup Internal Docker Registry

Ahmed ElBakry

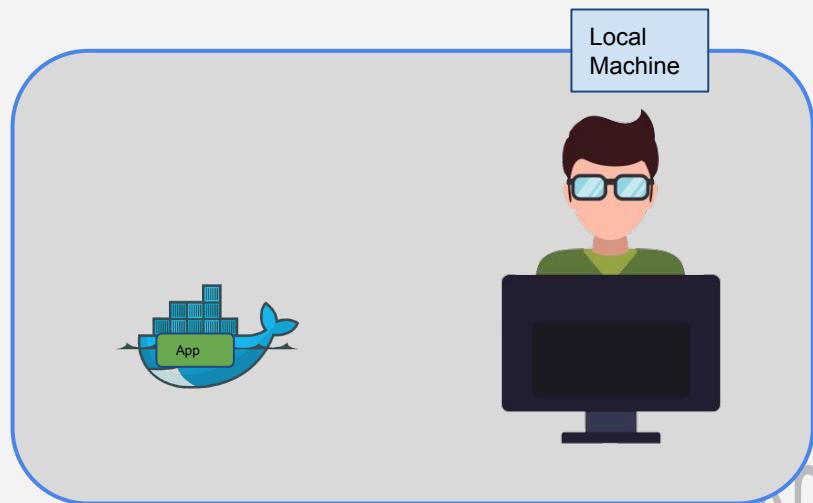
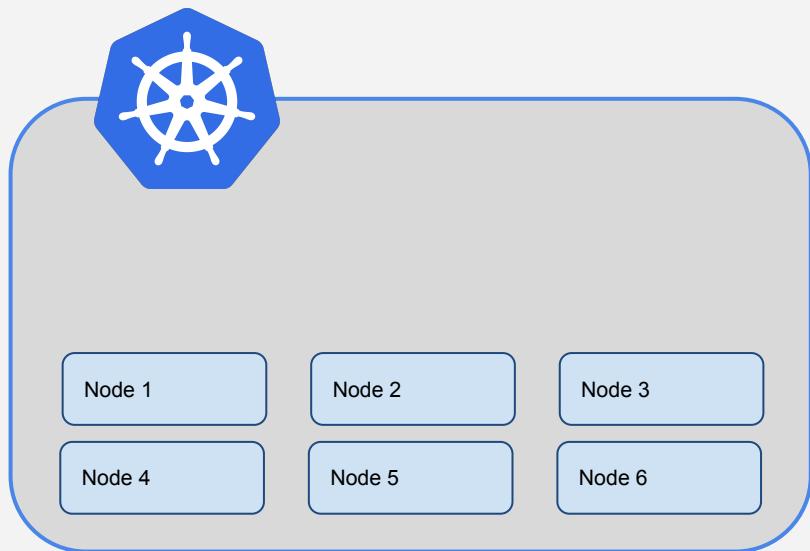
# Setup Internal Docker Registry



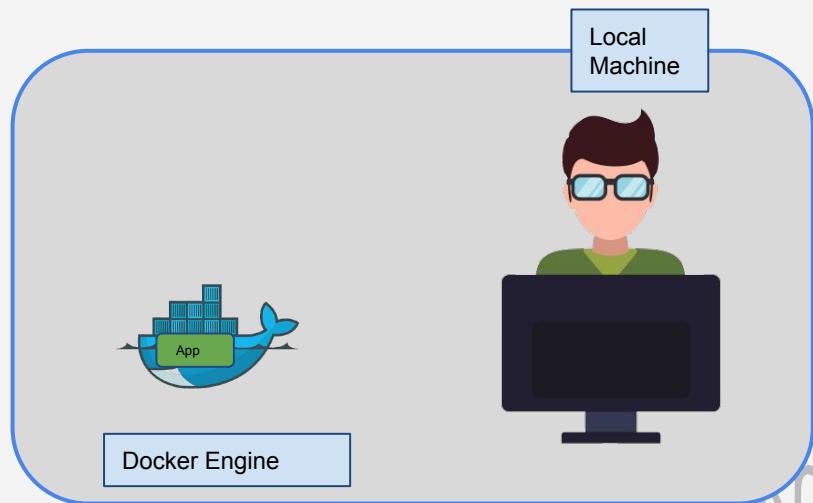
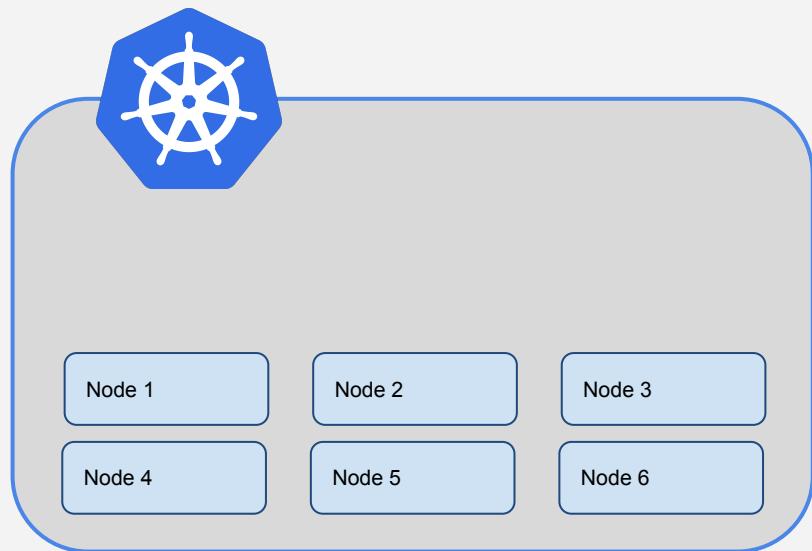
# Setup Internal Docker Registry



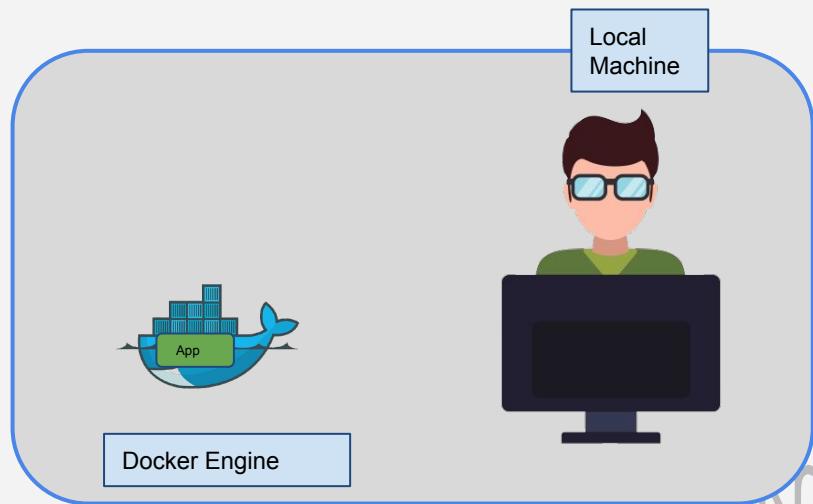
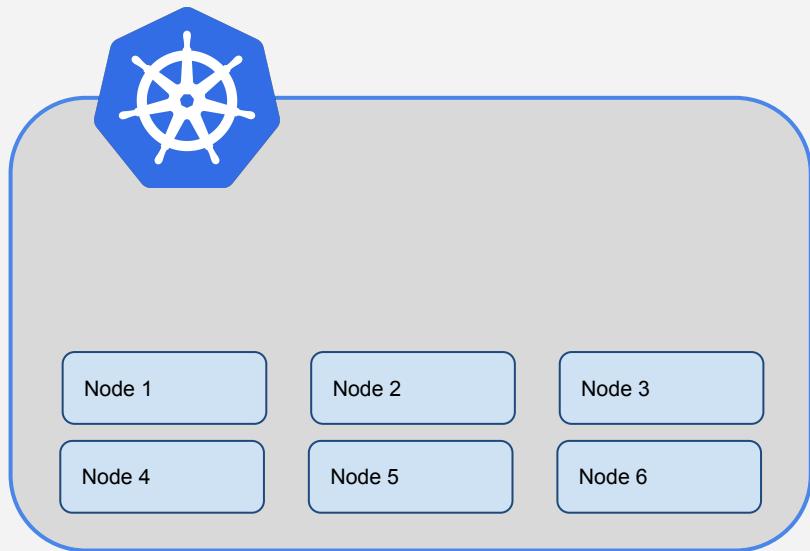
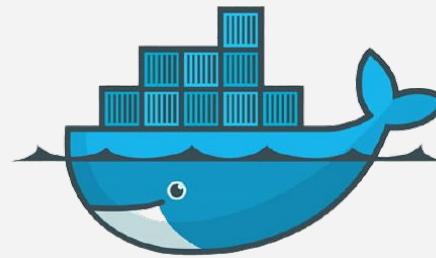
# Setup Internal Docker Registry



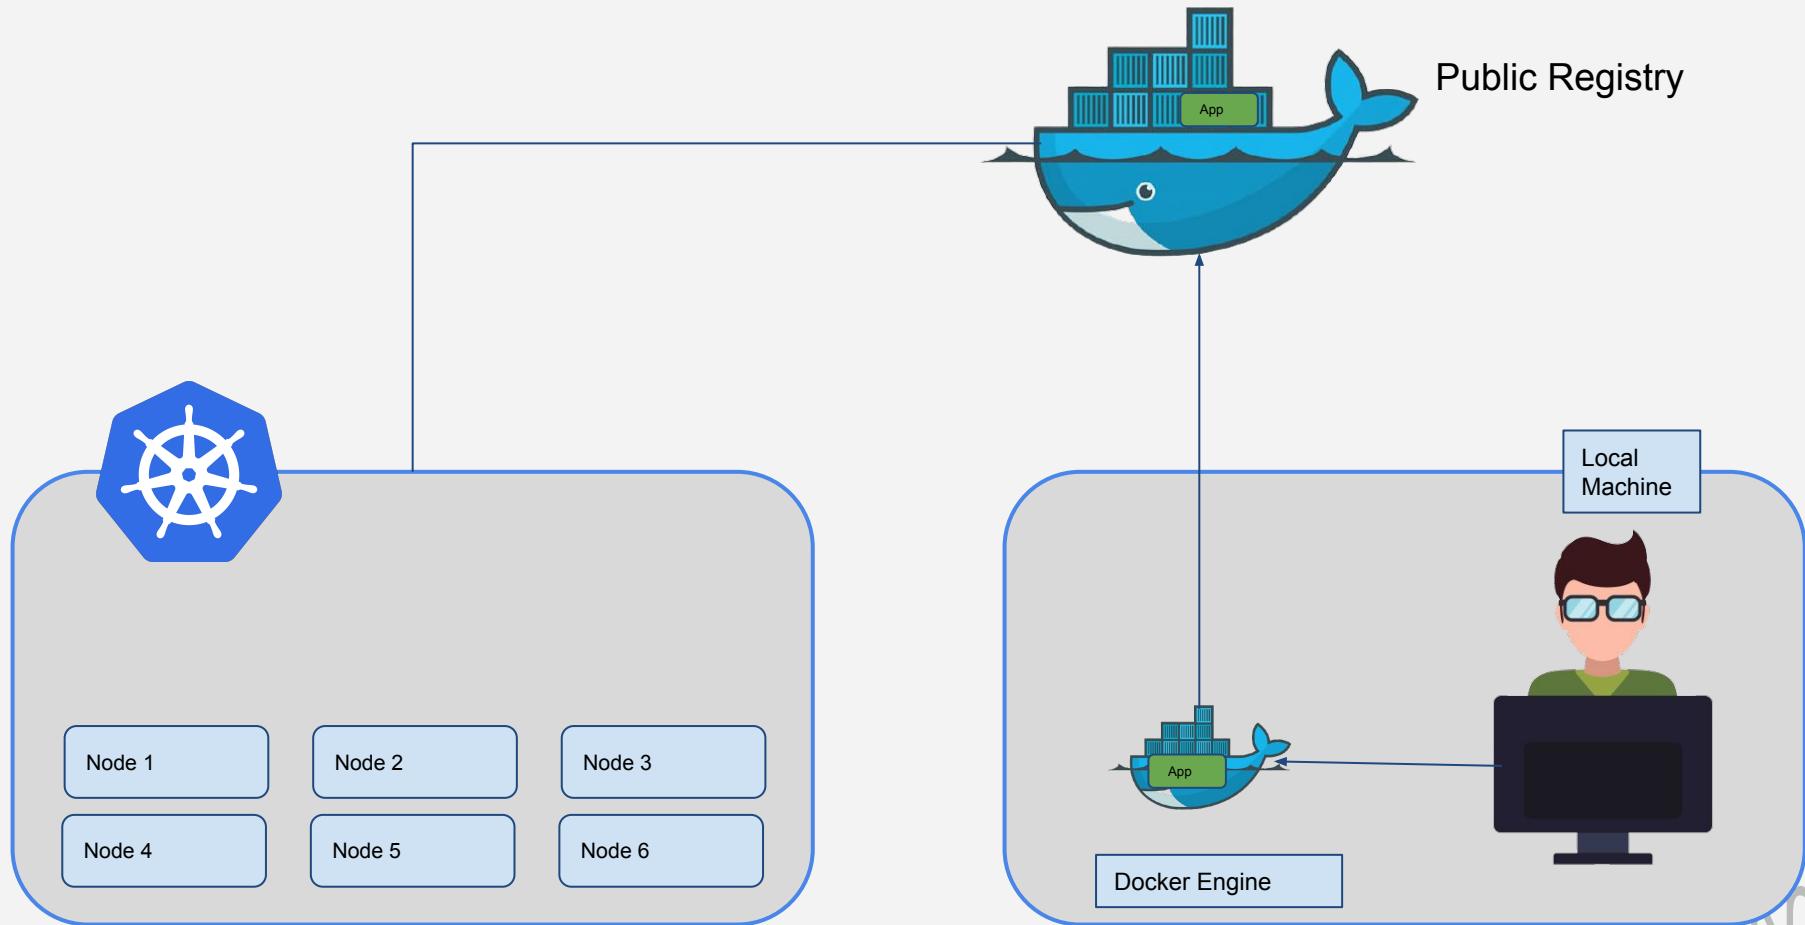
# Setup Internal Docker Registry



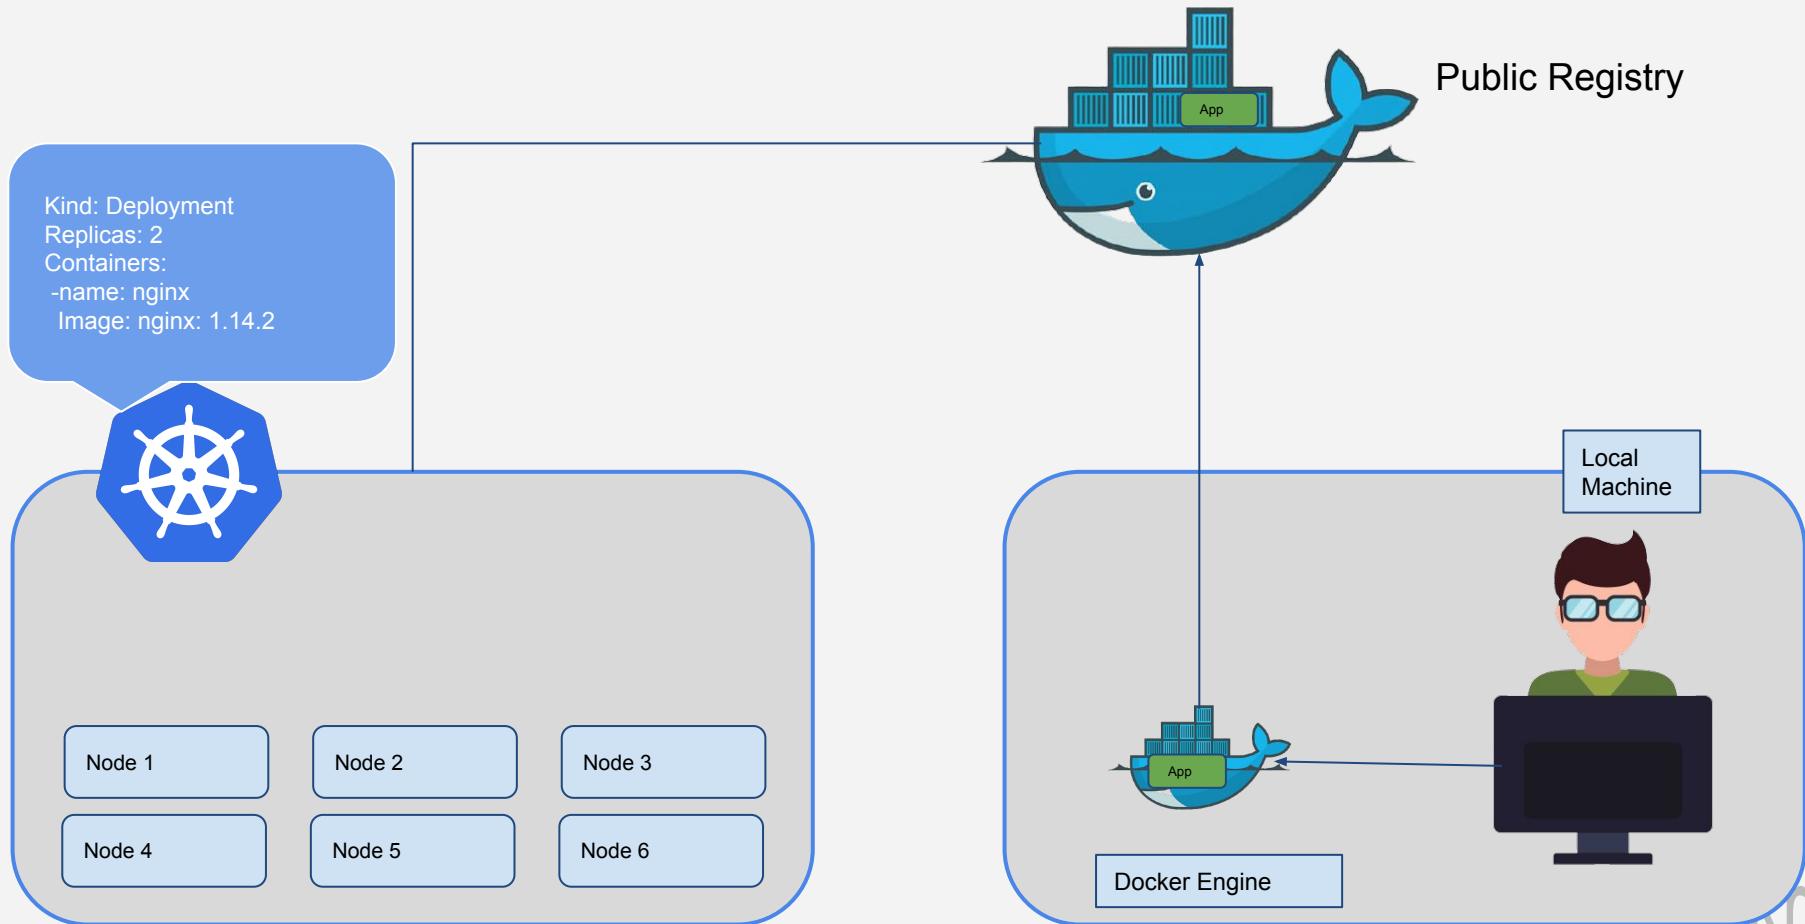
# Setup Internal Docker Registry



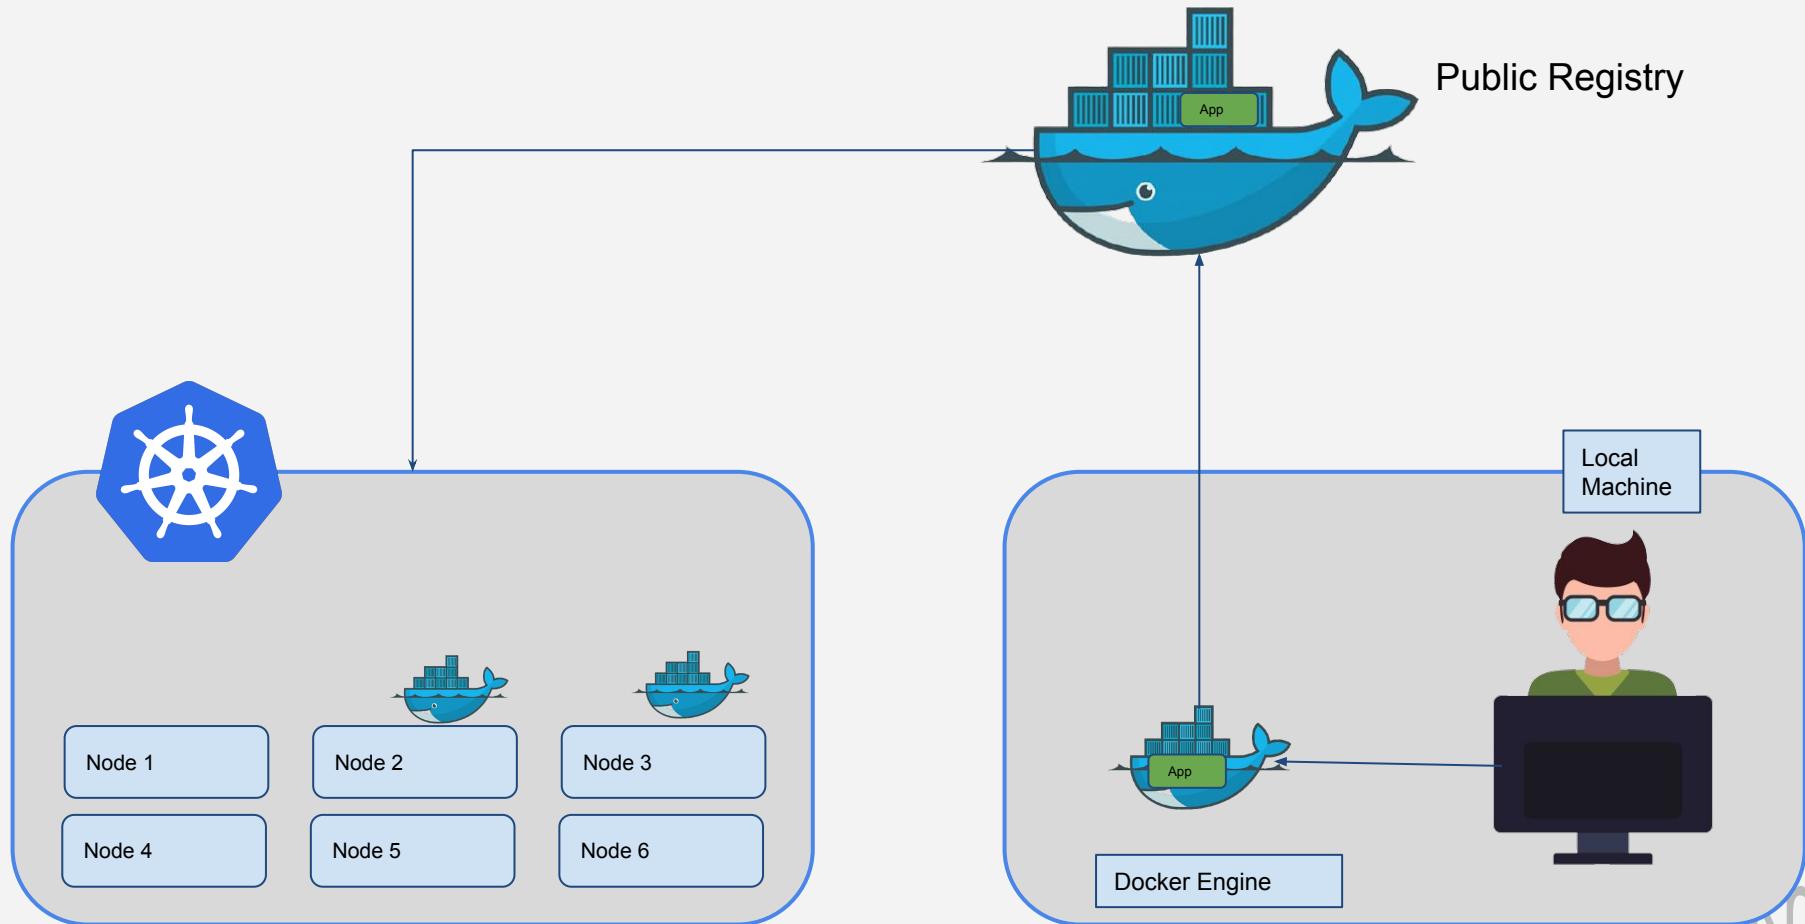
# Setup Internal Docker Registry



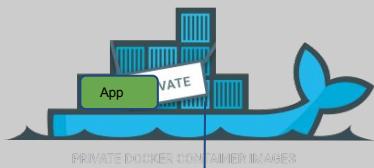
# Setup Internal Docker Registry



# Setup Internal Docker Registry



# Setup Internal Docker Registry



Node 1

Node 4

Node 2

Node 5

Node 3

Node 6

Docker Engine

Local  
Machine



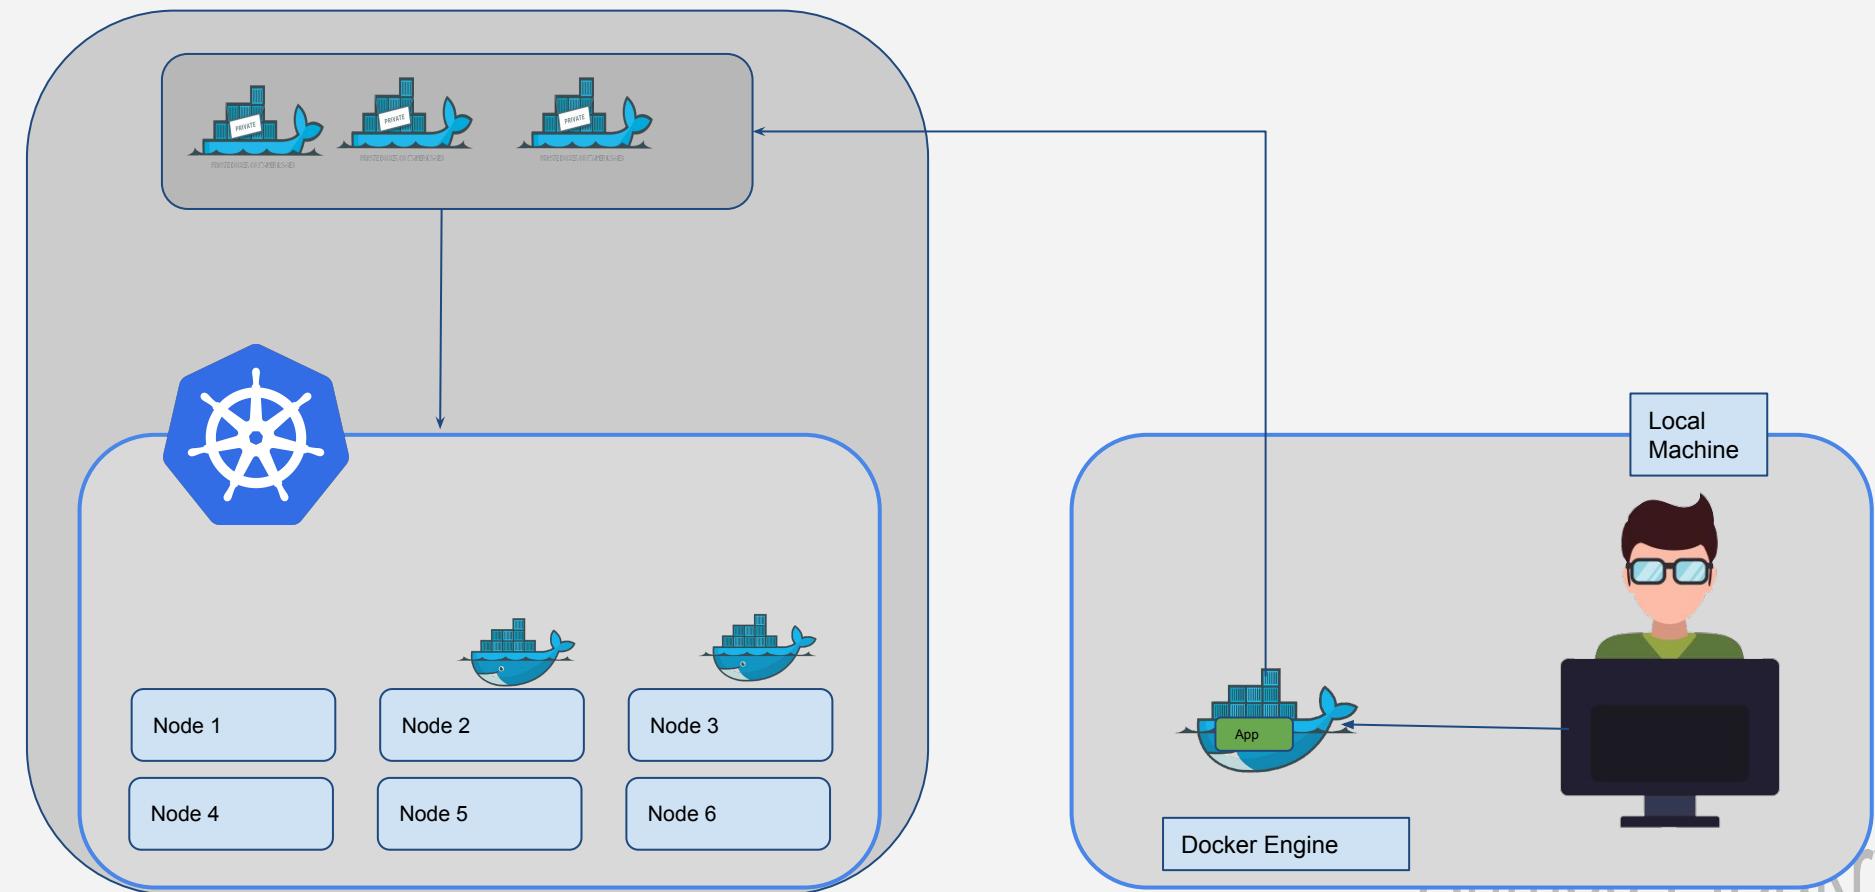
# Setup Internal Docker Registry

## Demo

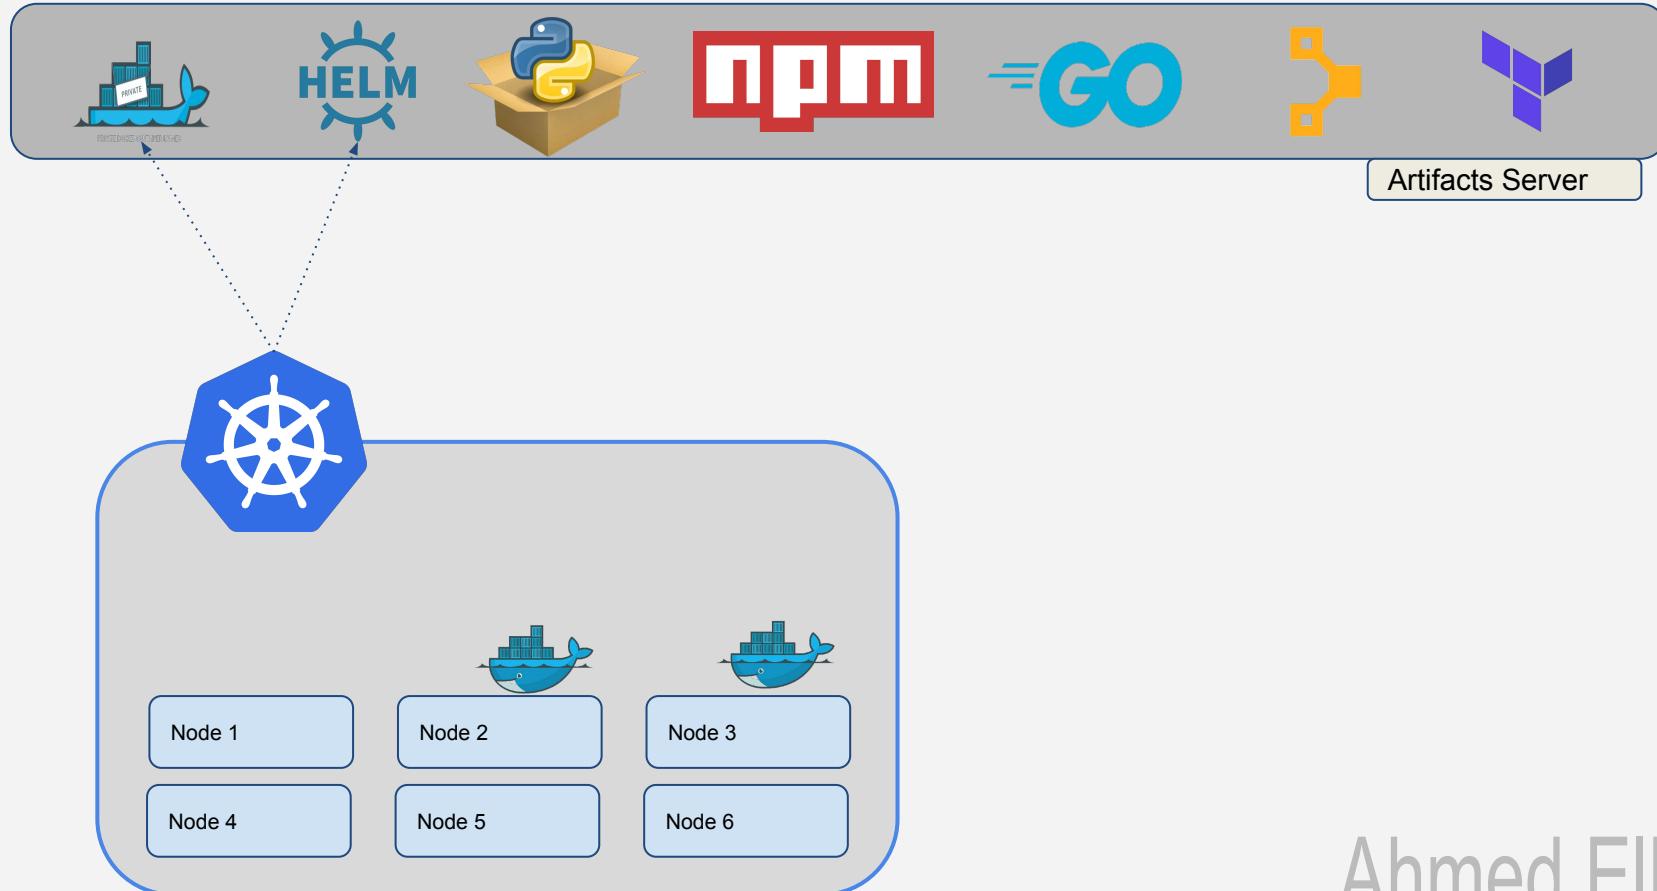


- Run Local Application
- Containerize The Application
- Push The Application to Public Docker Registry
- Deploy The Application to Kubernetes
- Deploy Internal/Private Docker Registry
- Push The Application to Internal Registry
- Deploy The Application to Kubernetes

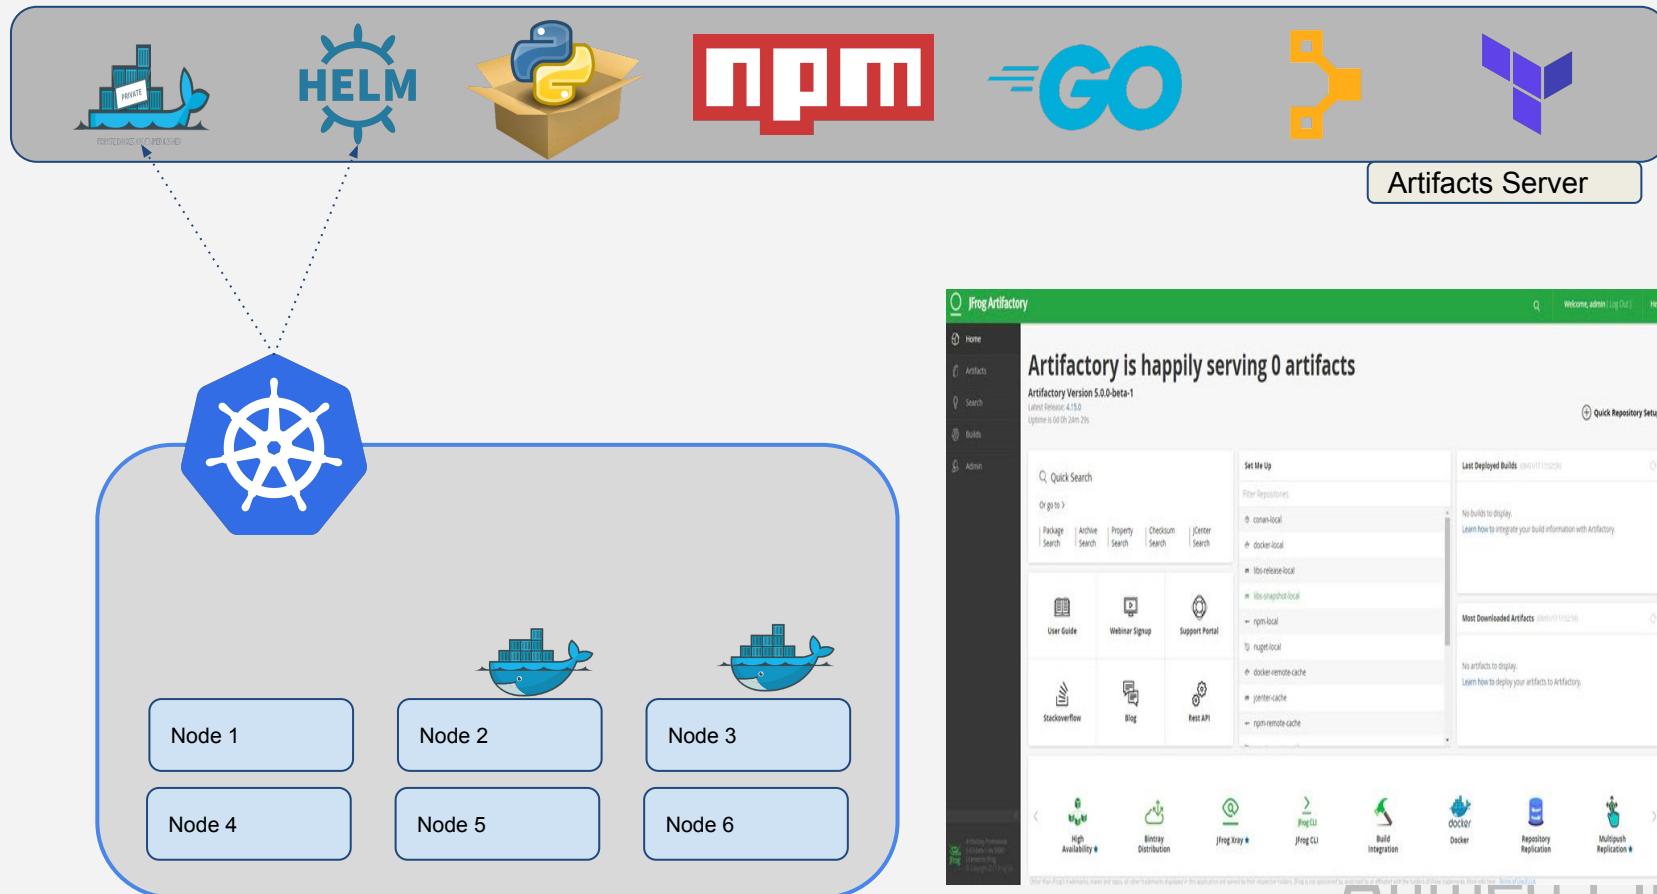
# Setup Internal Docker Registry



# Setup Internal Docker Registry

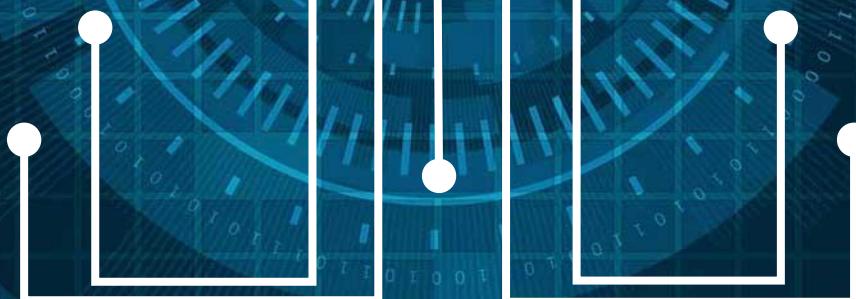


# Setup Internal Docker Registry





# THANK YOU



Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

12 - Namespaces

Ahmed ElBakry



# **Module 4**

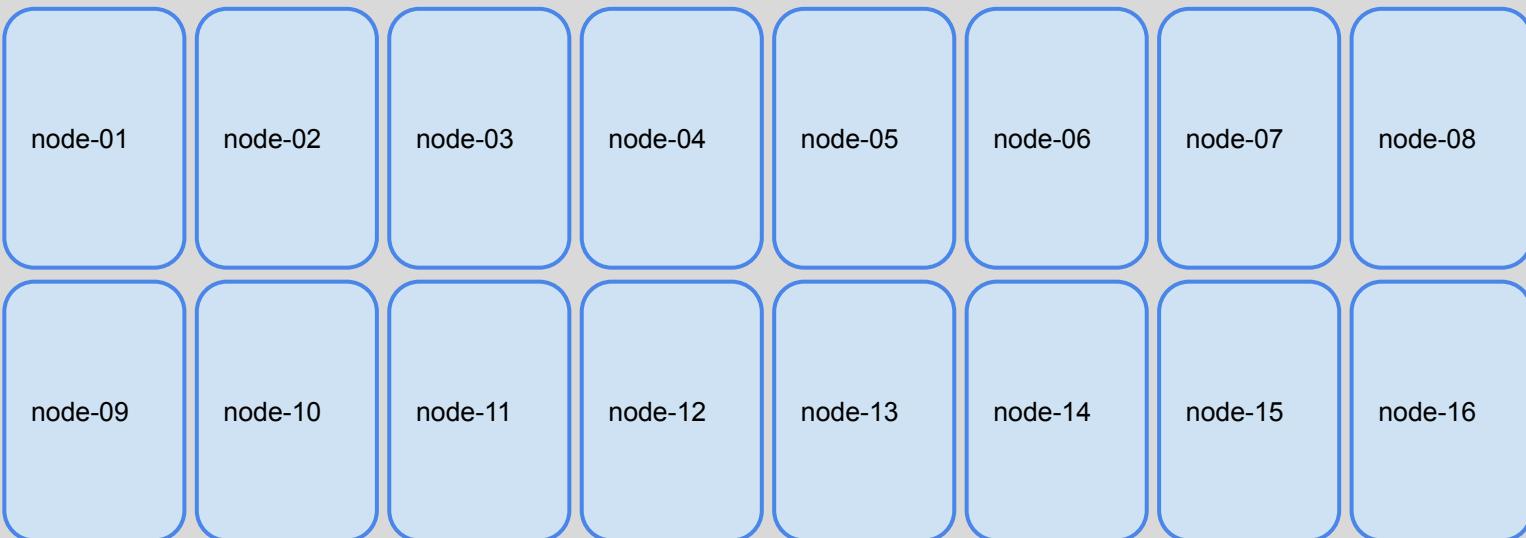
## **Basic Concepts**

Ahm

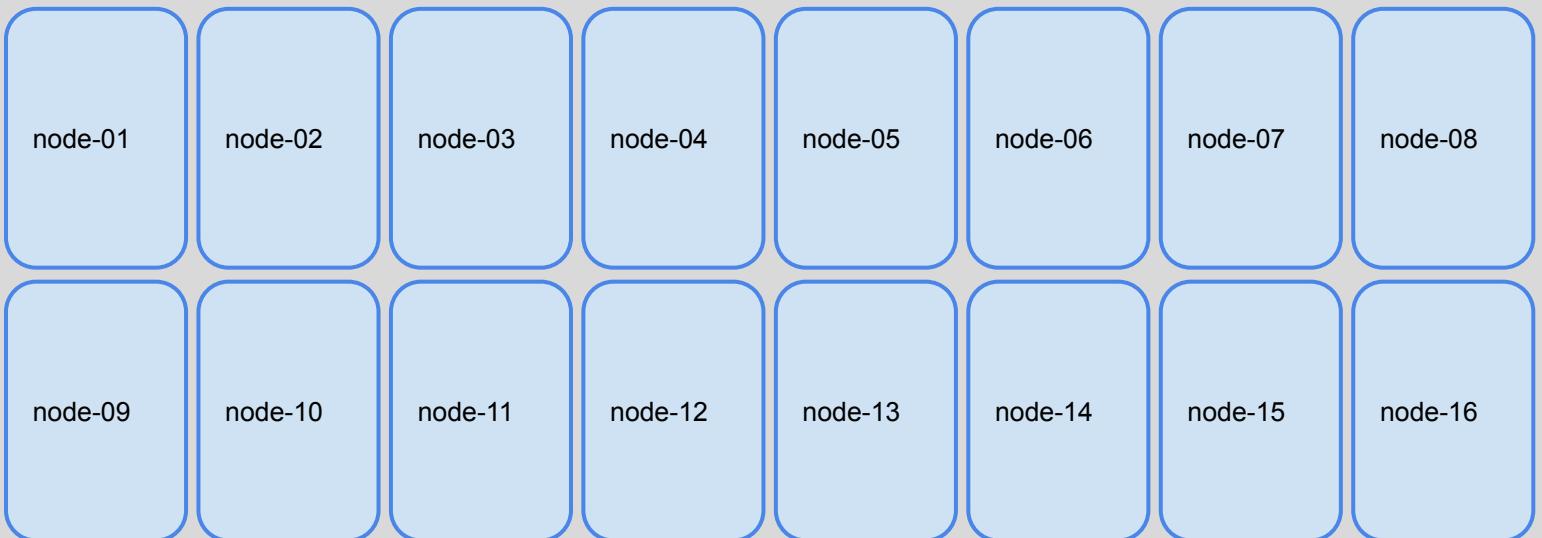
# Namespaces



# Namespaces



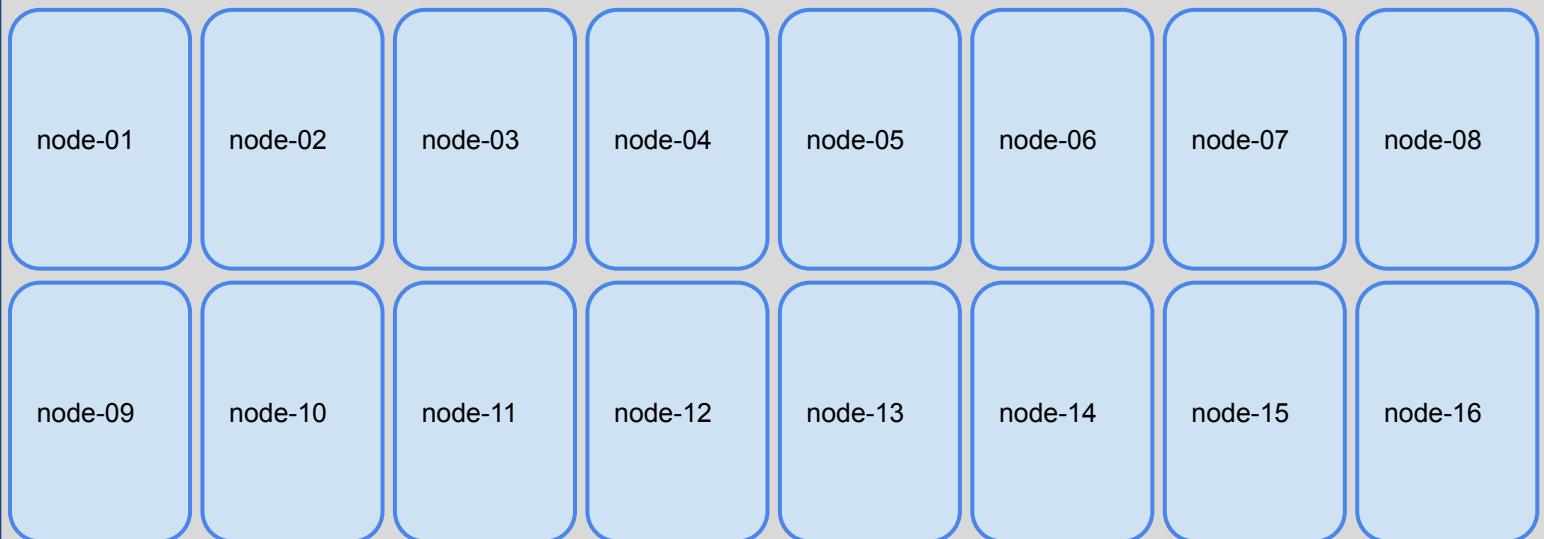
# Namespaces



# Namespaces



**Available Space**  
CPU: 160  
Memory: 320 GB



# Namespaces



Frontend

Backend

Database

# Namespaces

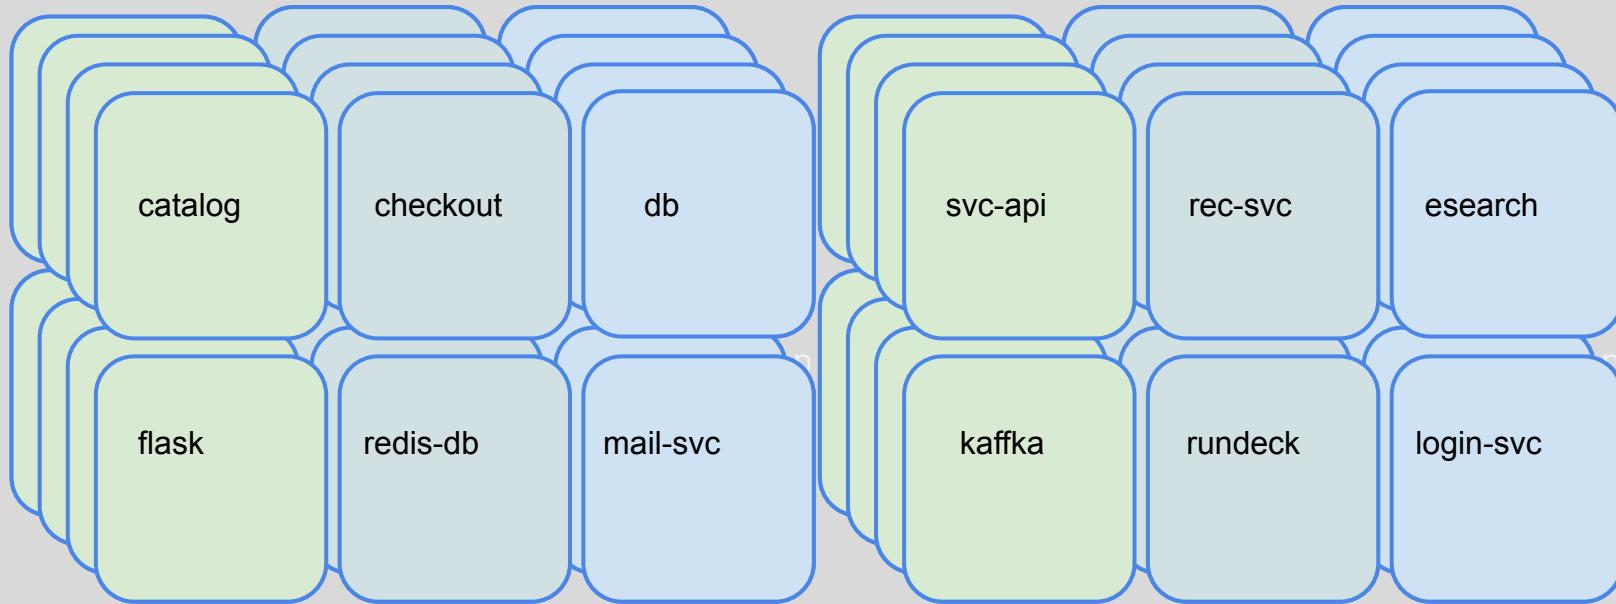


Production

Development

Test

# Namespaces



# Namespaces

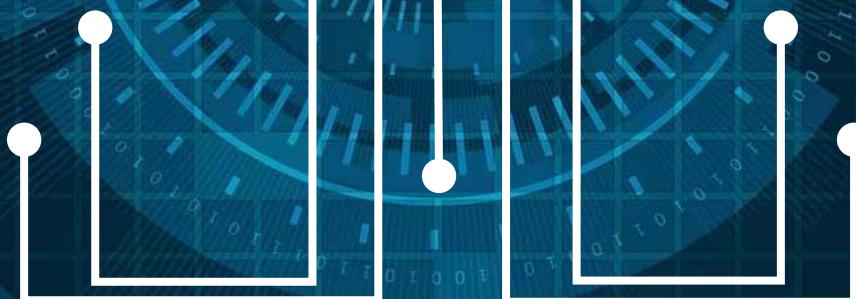
## Demo



- ❑ List namespaces
- ❑ Create namespaces
- ❑ Switch namespaces.
- ❑ Create objects in different namespaces
- ❑ Namespaced resource vs non-namespaced resource
- ❑ Access services in different namespaces.
- ❑ Limit and isolate access with network policy



# THANK YOU



Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

13 - Resources, Resource Quota, and  
LimitRanges

Ahmed ElBakry

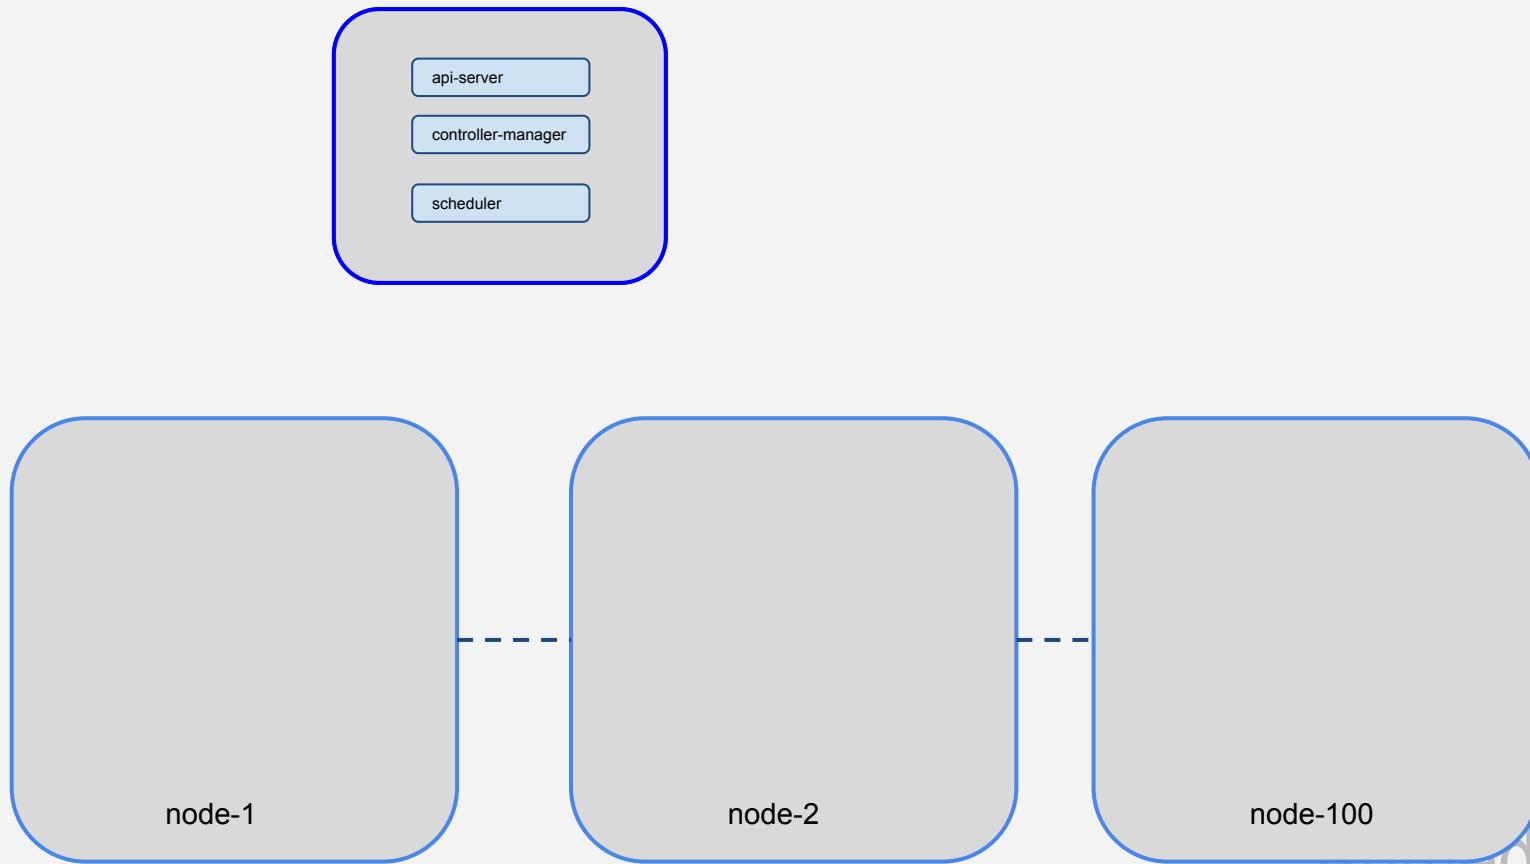
# Resource Quota

## In this video

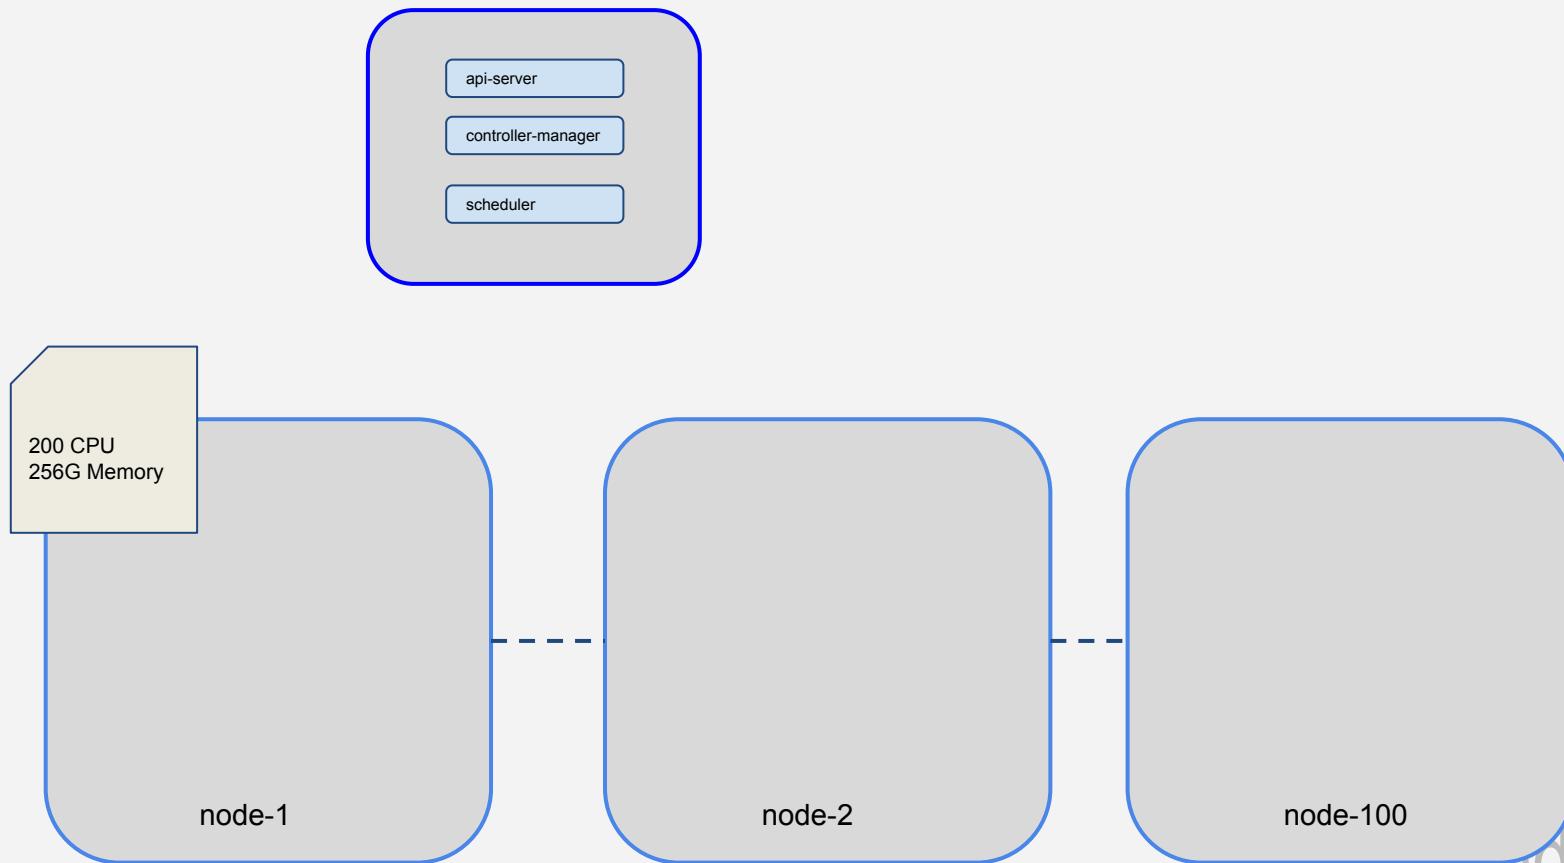
- Resources
- Limit Range
- Resource Quota
- Demo

Ahmed ElBakry

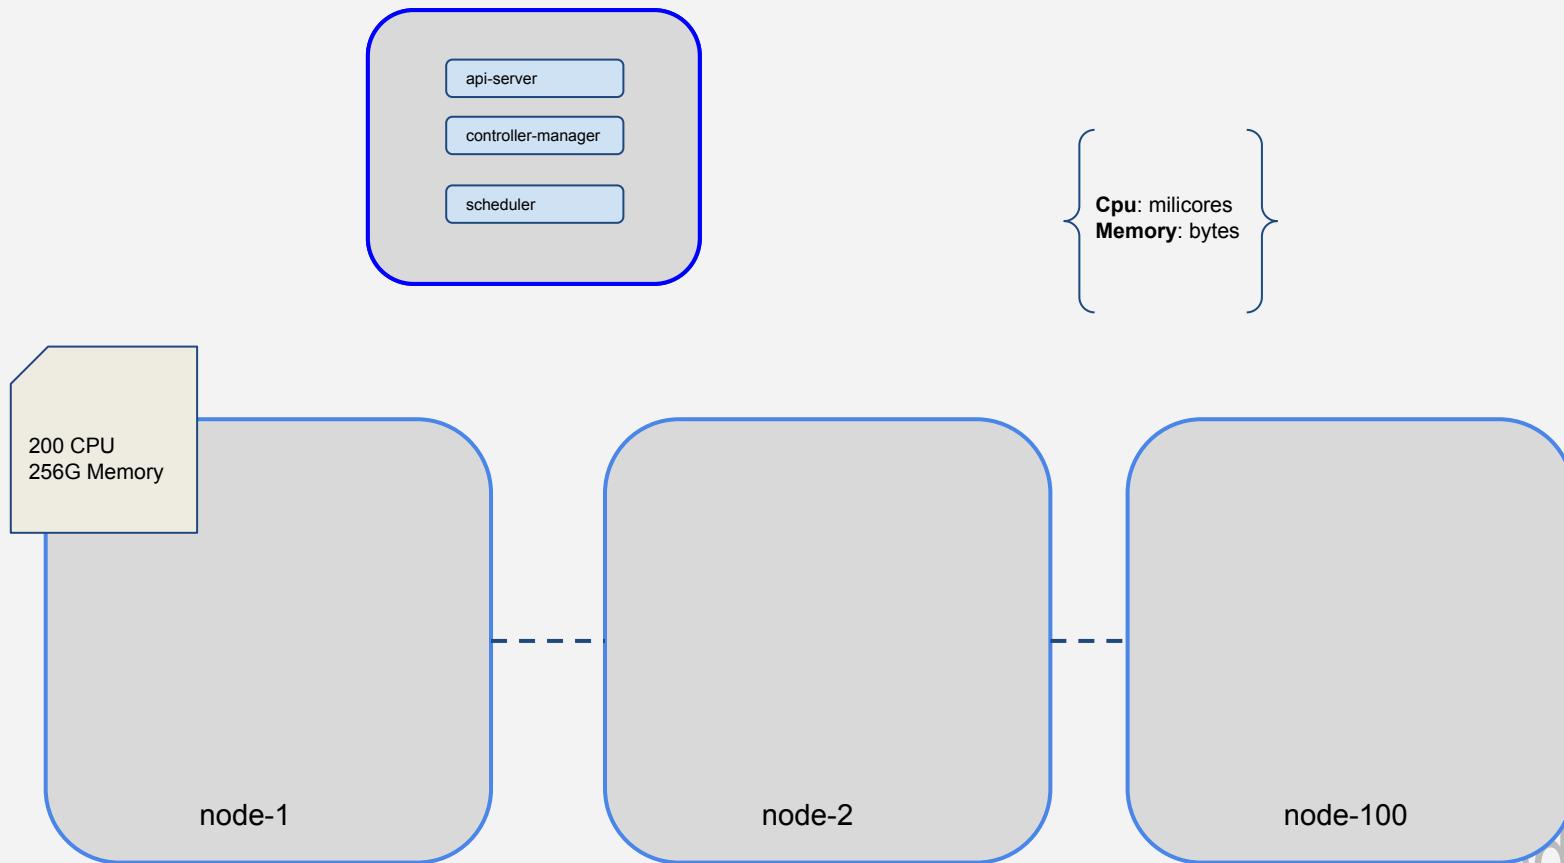
# Resources



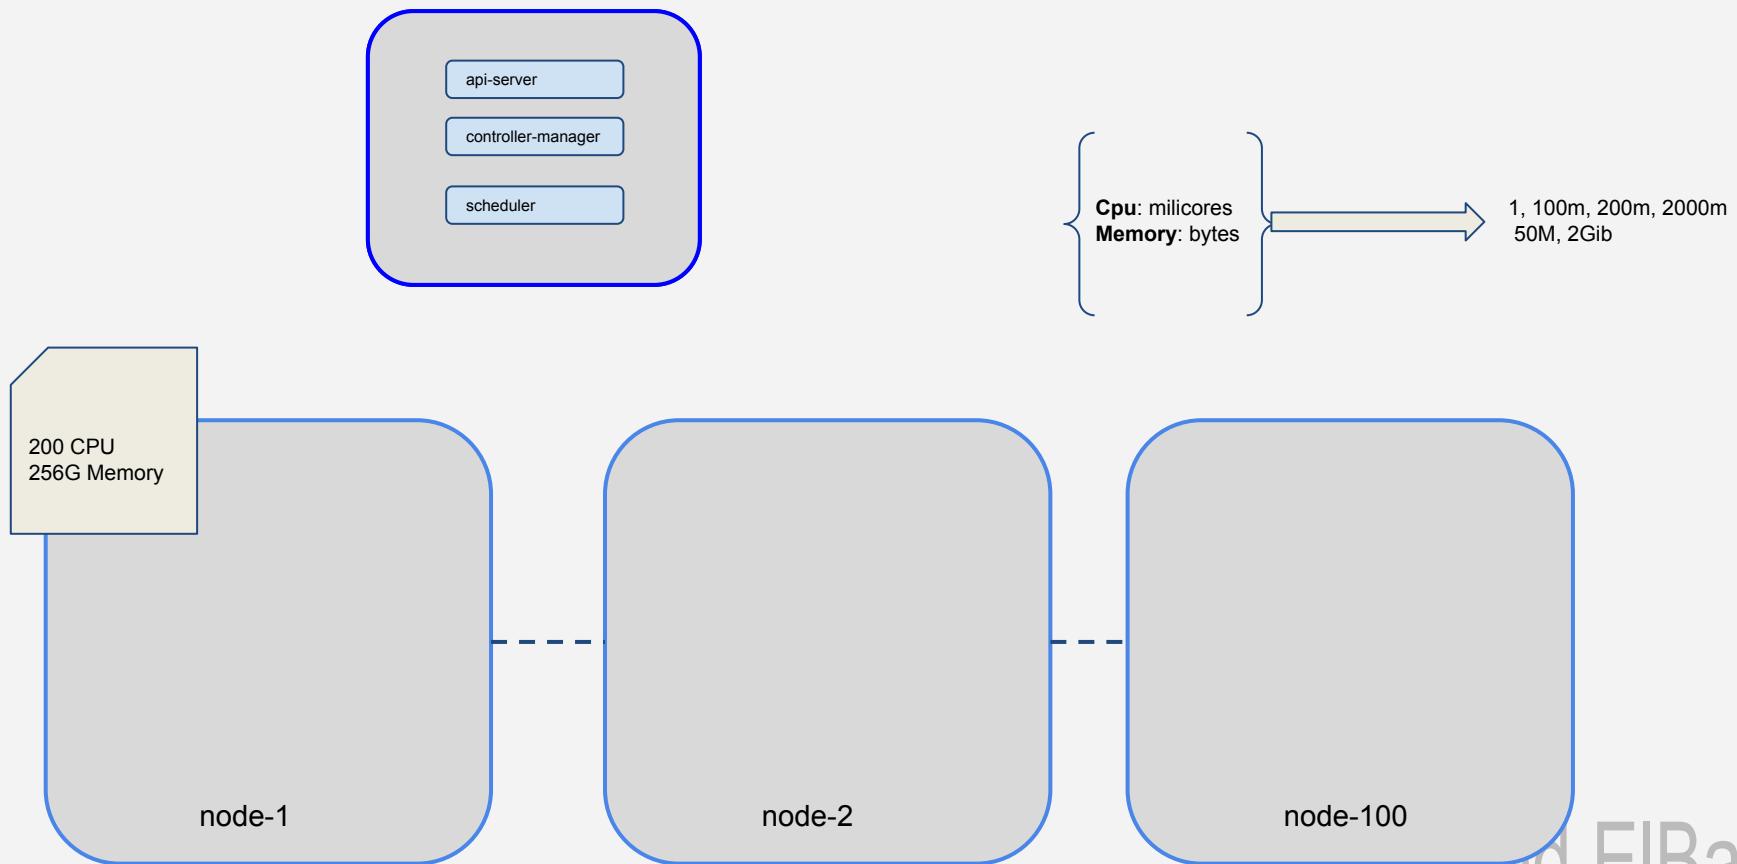
# Resources



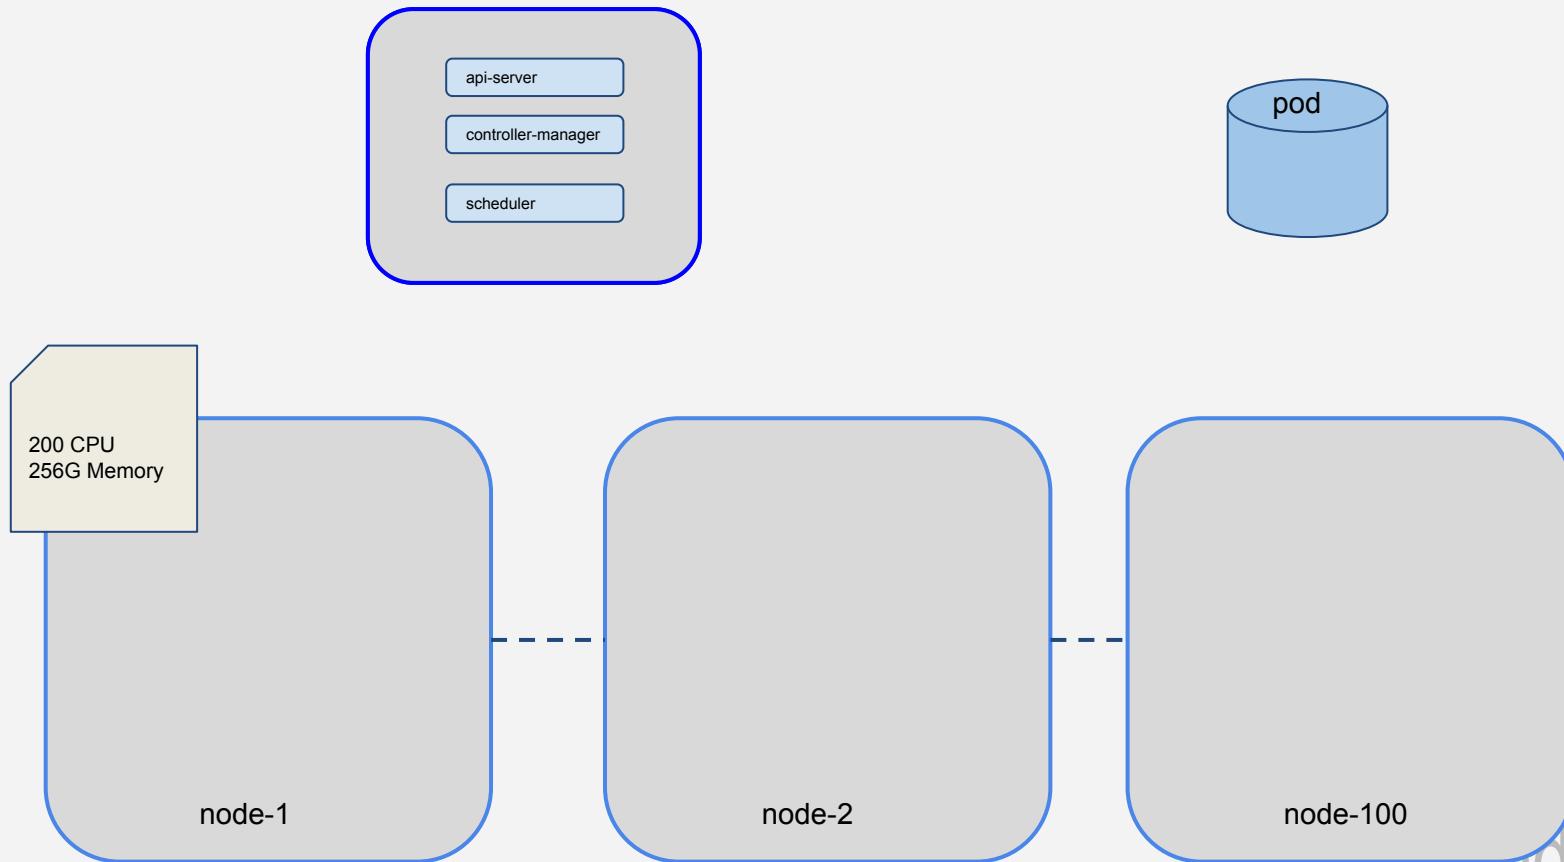
# Resources



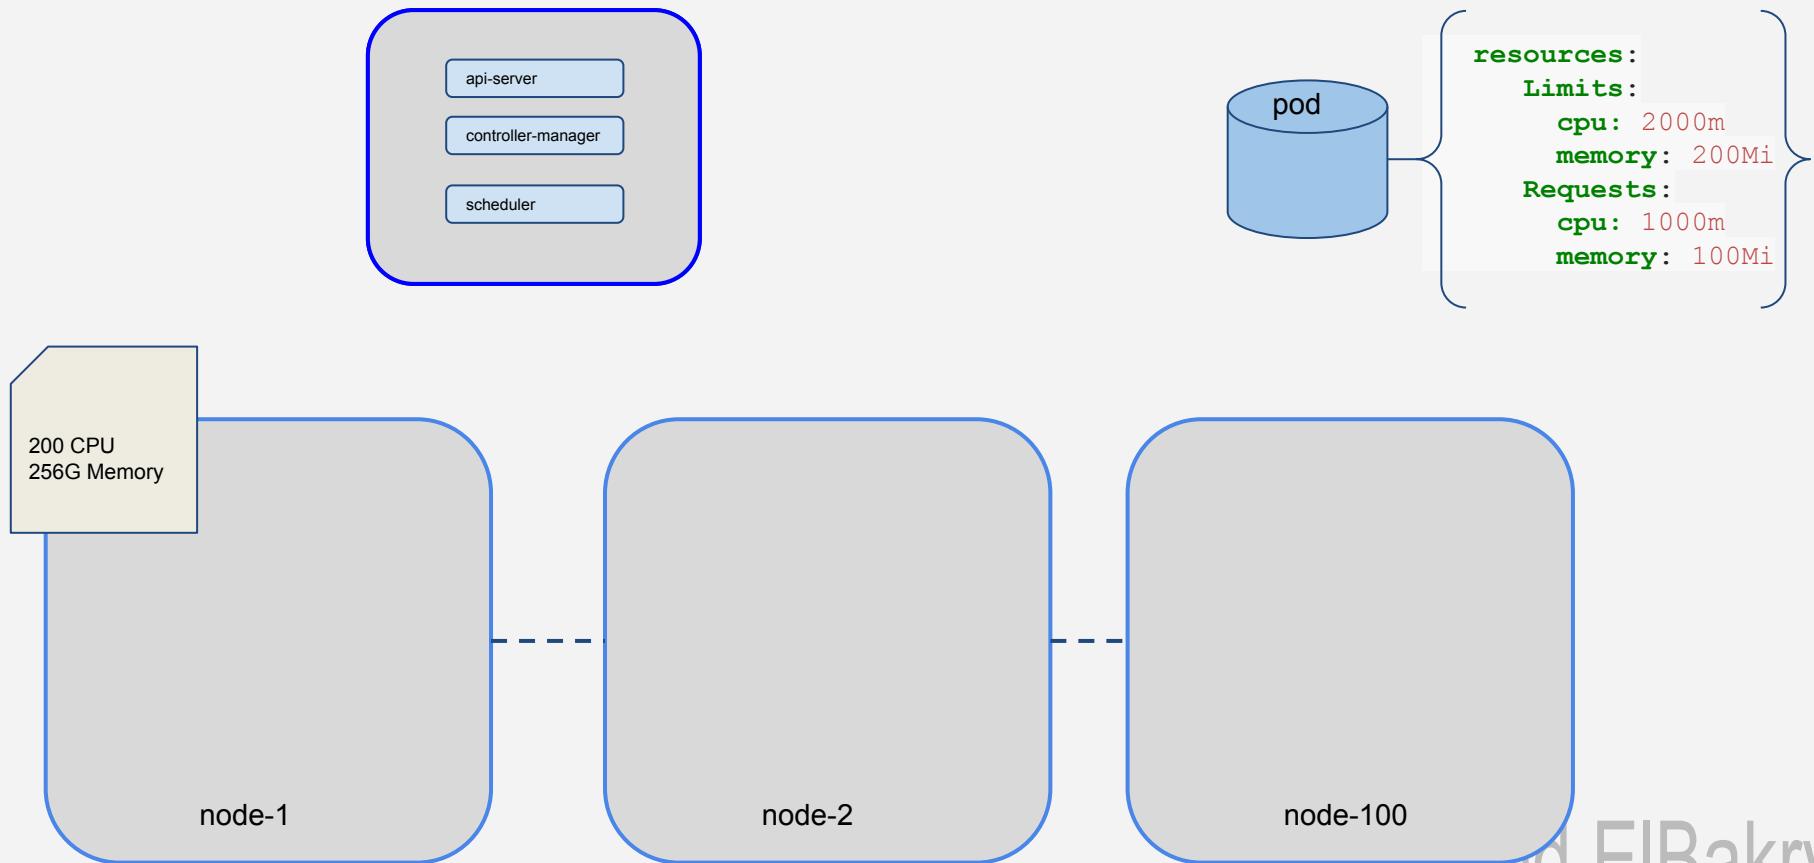
# Resources



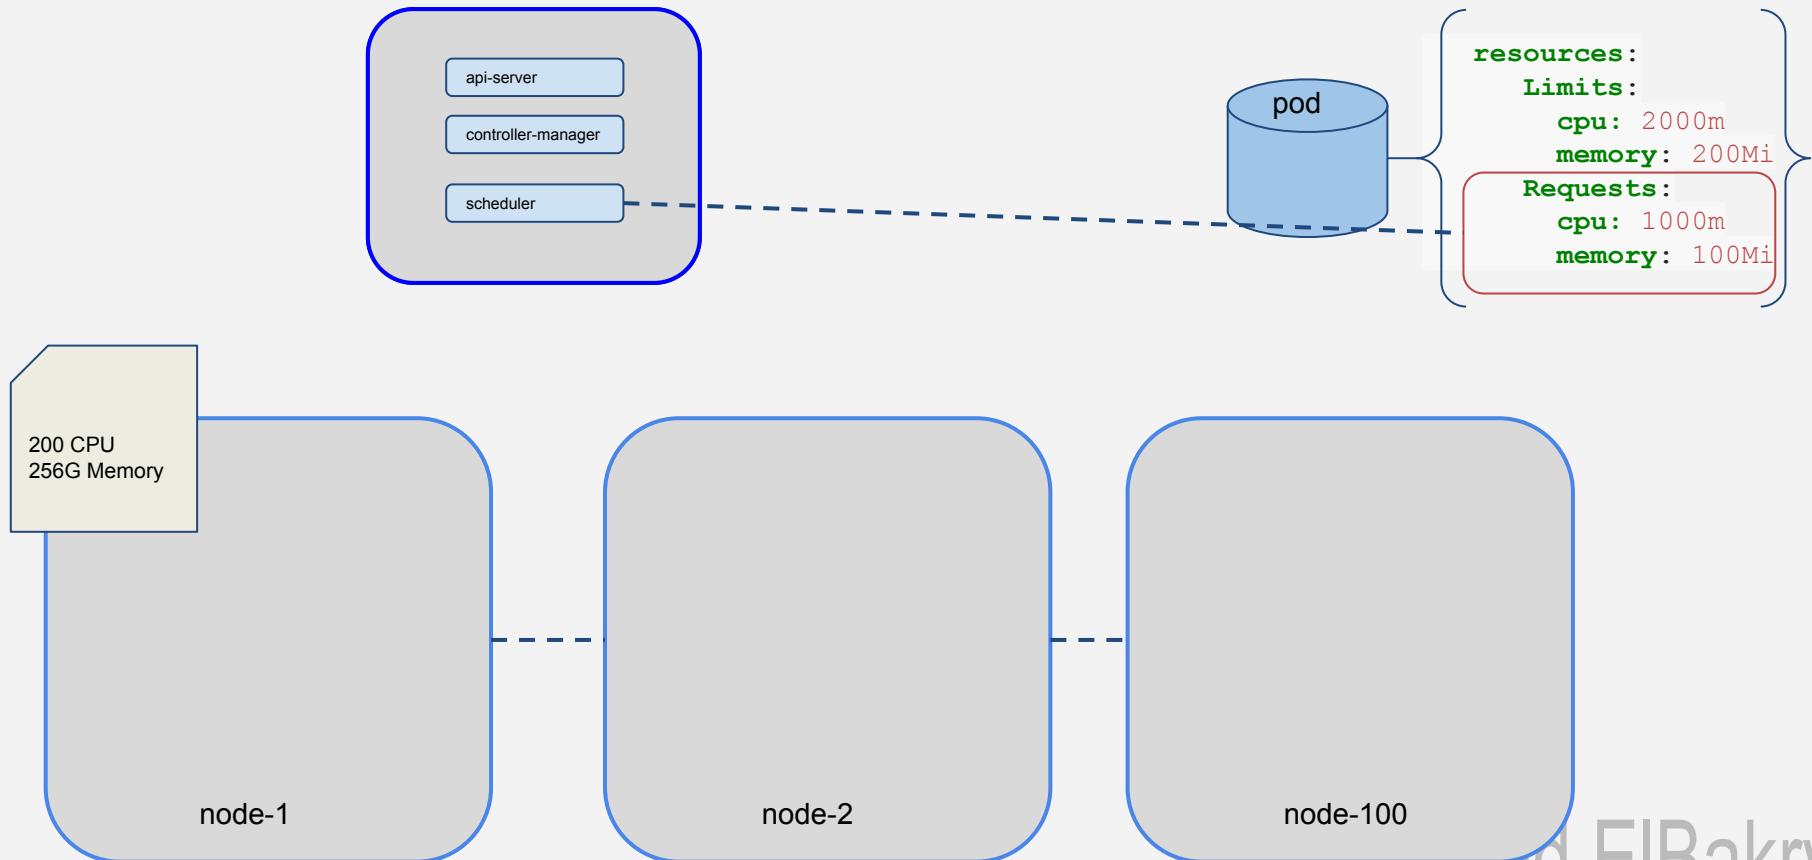
# Resources



# Resources



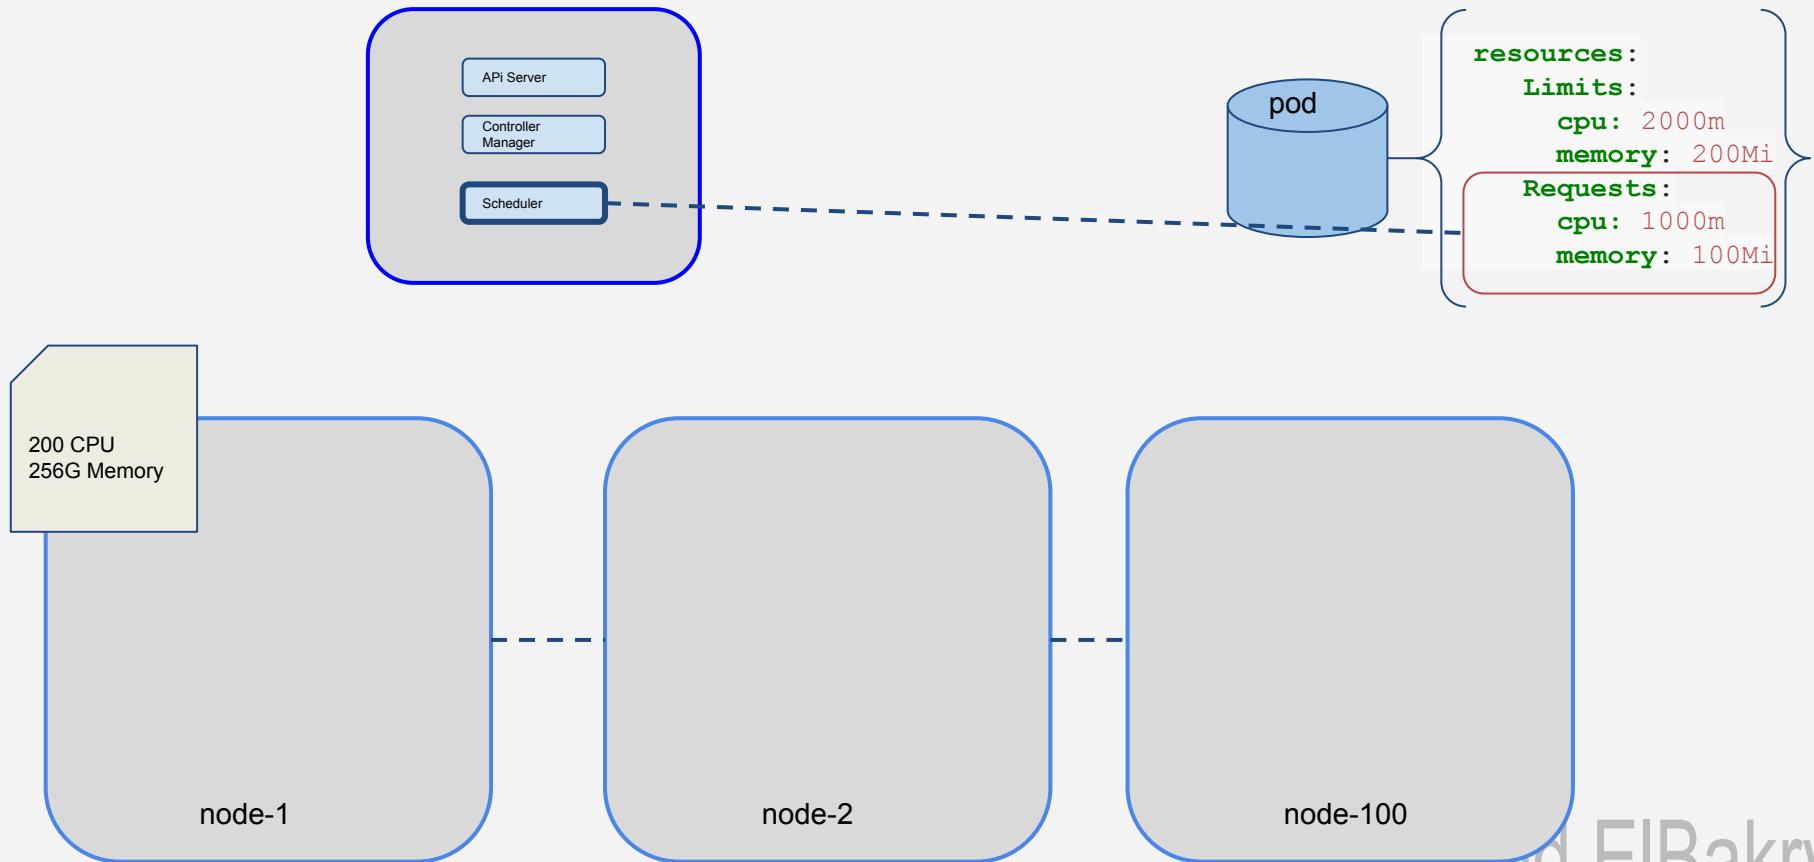
# Resources



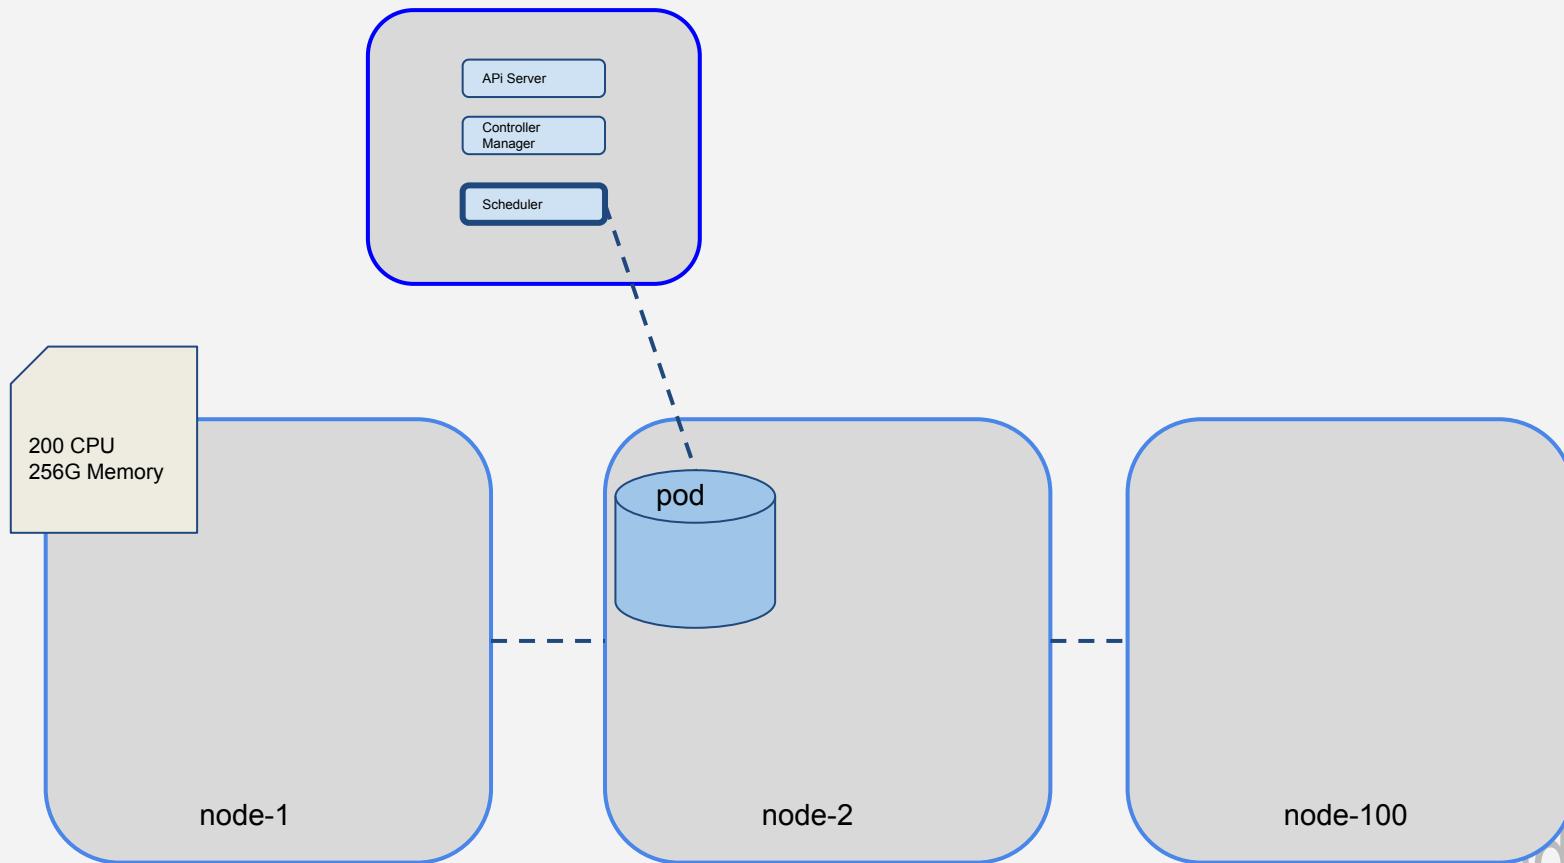
# Quality of Service for Pods

Guaranteed	Burstable	BestEffort
<pre>apiVersion: v1 kind: Pod metadata:   name: nginx   namespace: nginx spec:   containers:     - name: nginx       image: nginx       resources:         limits:           memory: "500Mi"           cpu: "1000m"         requests:           memory: "500Mi"           cpu: "1000m"</pre>	<pre>apiVersion: v1 kind: Pod metadata:   name: nginx   namespace: nginx spec:   containers:     - name: nginx       image: nginx       resources:         limits:           memory: "500Mi"           cpu: "1000m"         requests:           memory: "200Mi"           cpu: "500m"</pre>	<pre>apiVersion: v1 kind: Pod metadata:   name: nginx   namespace: nginx spec:   containers:     - name: nginx       image: nginx</pre>

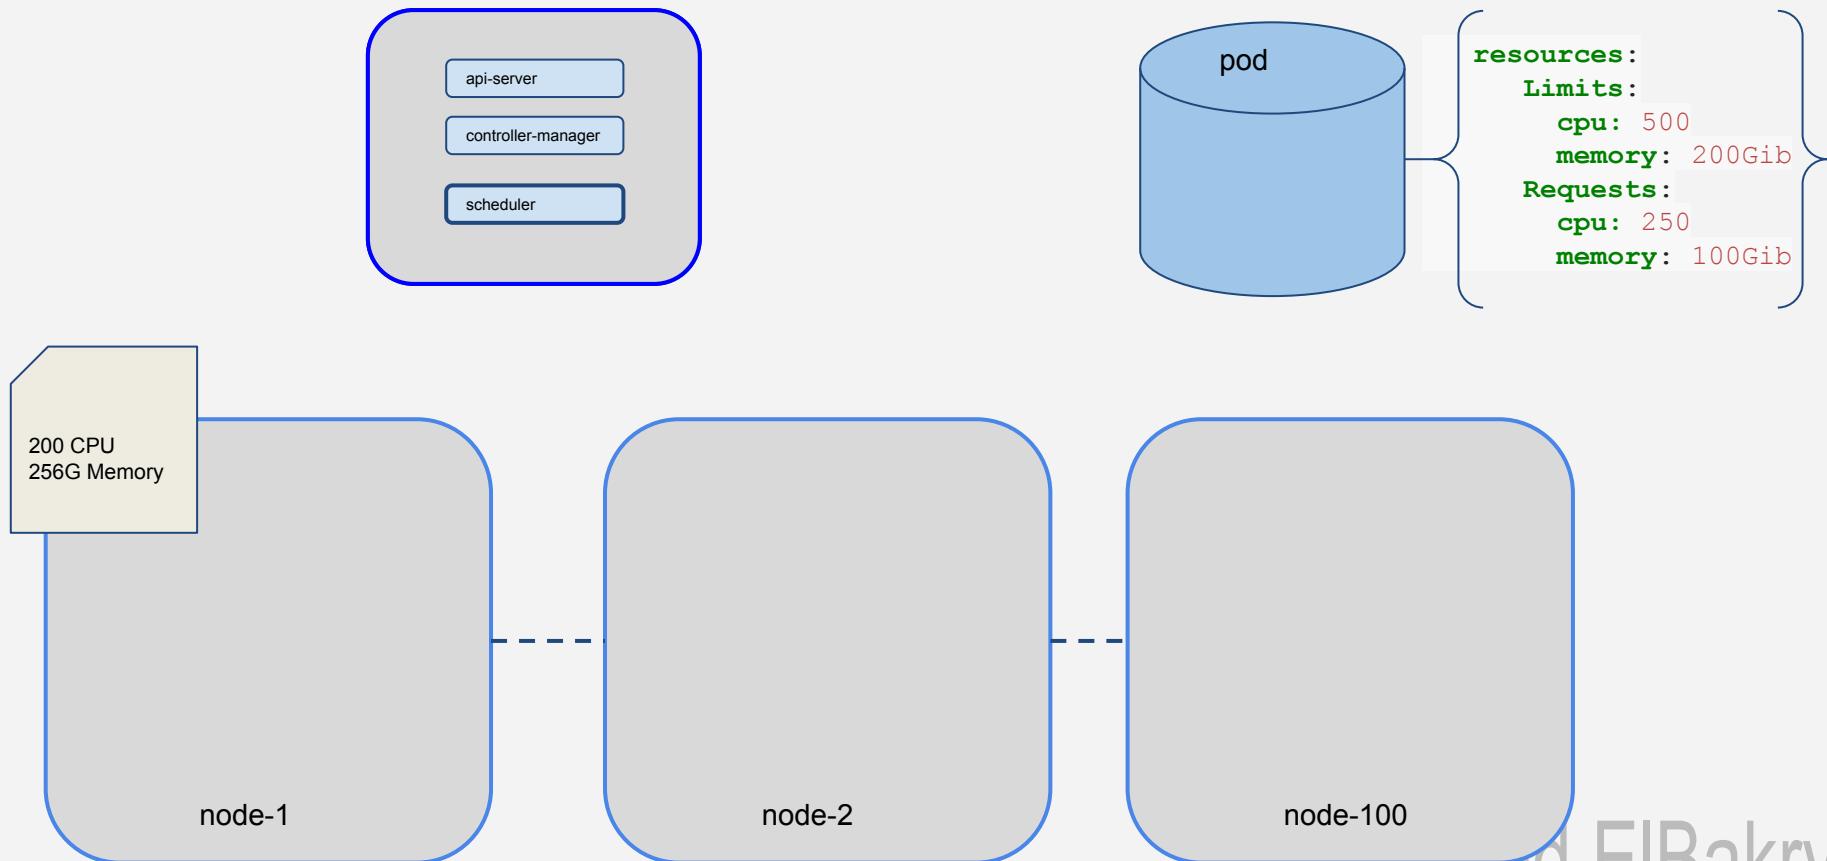
# Resources



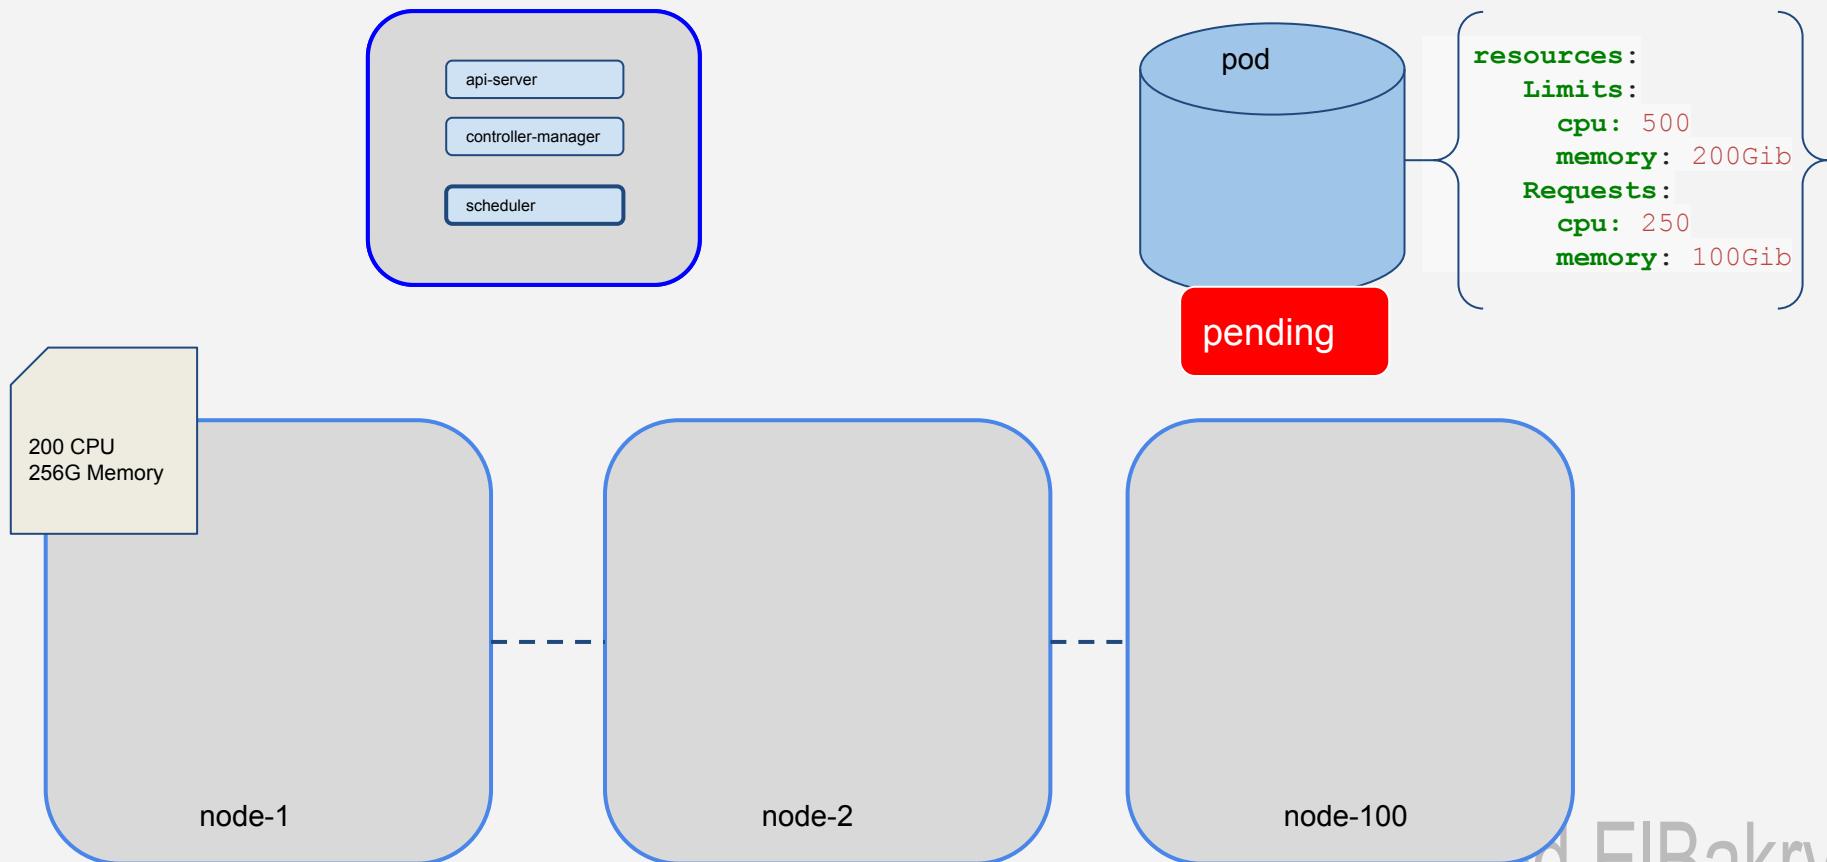
# Resources



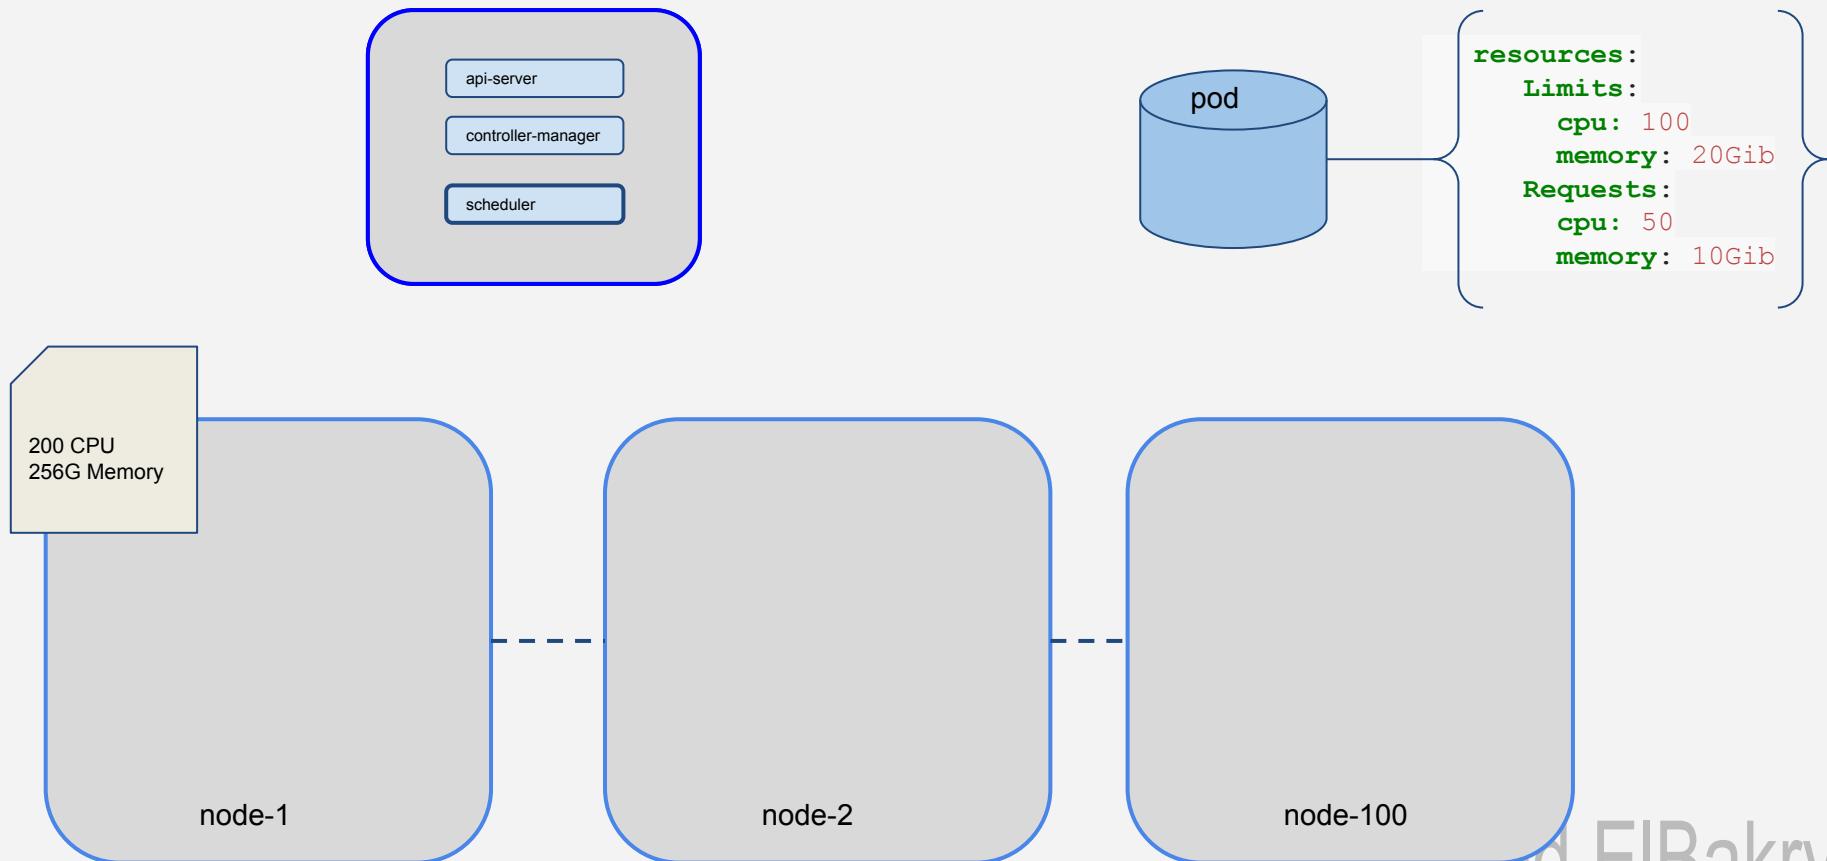
# Resources



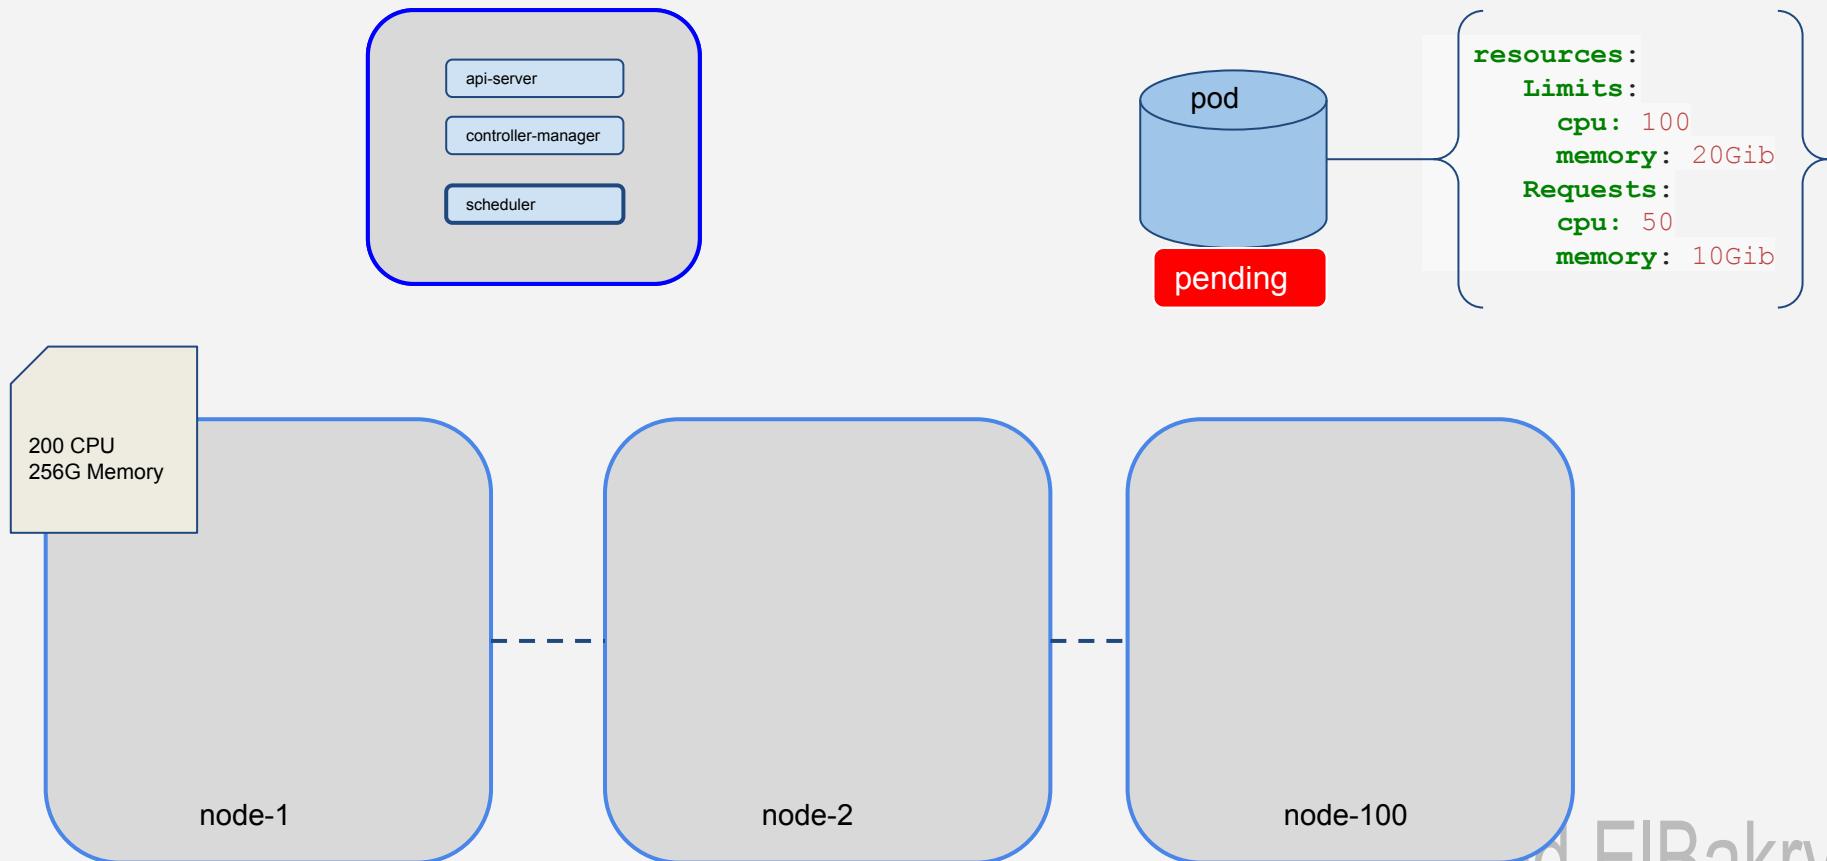
# Resources



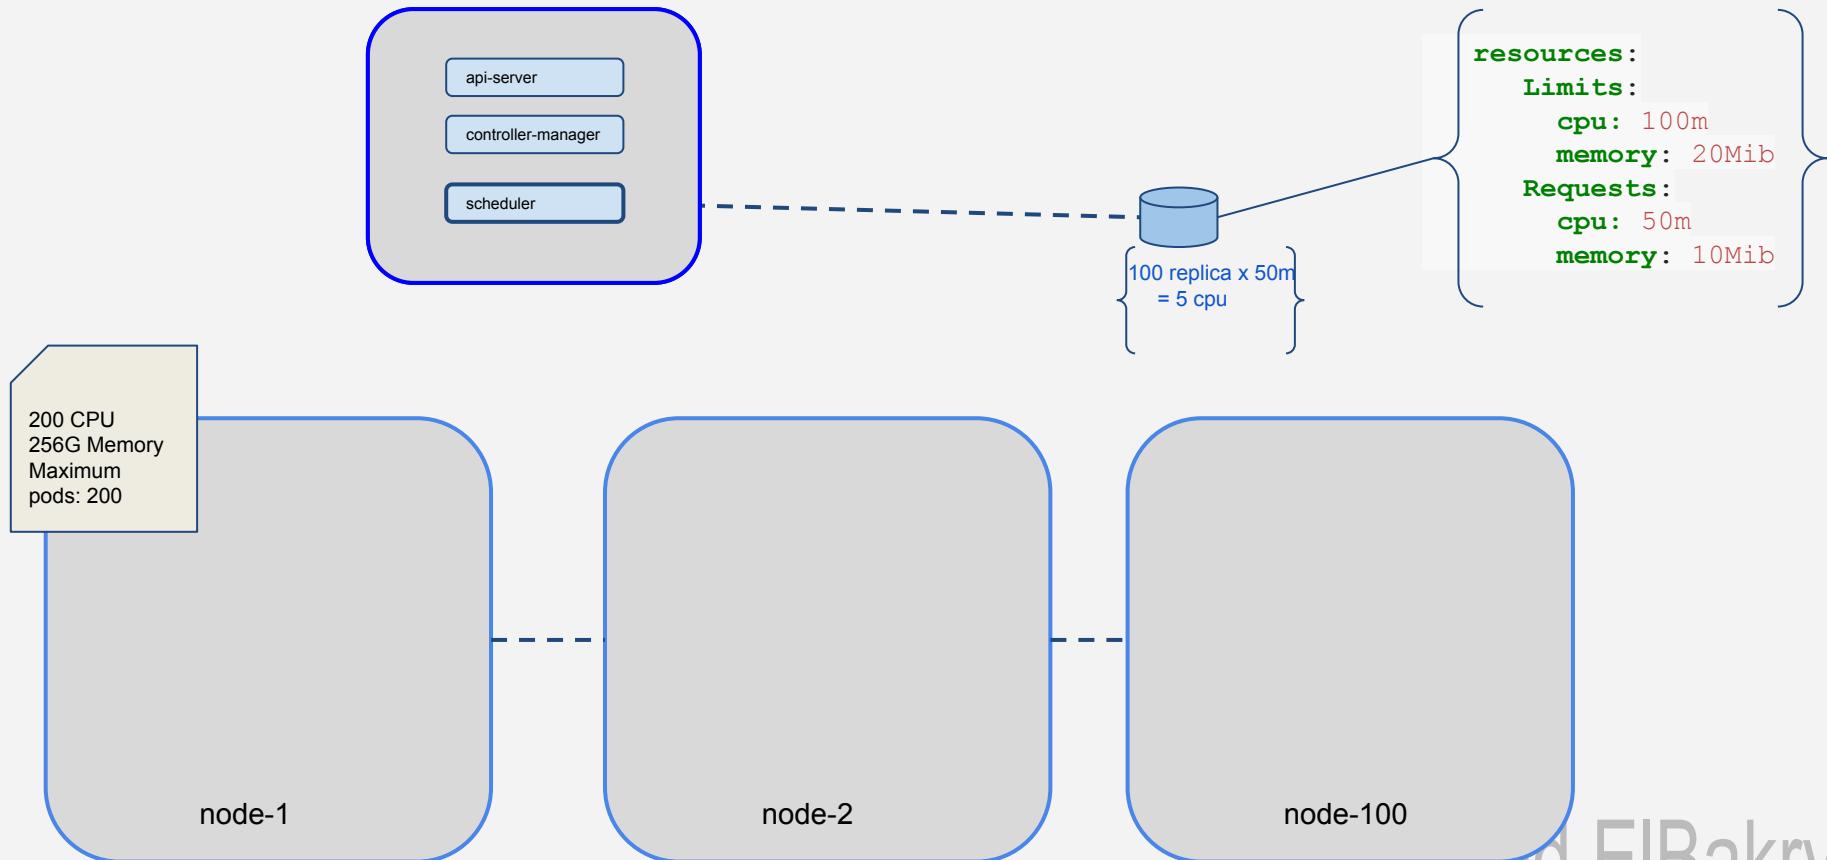
# Resources



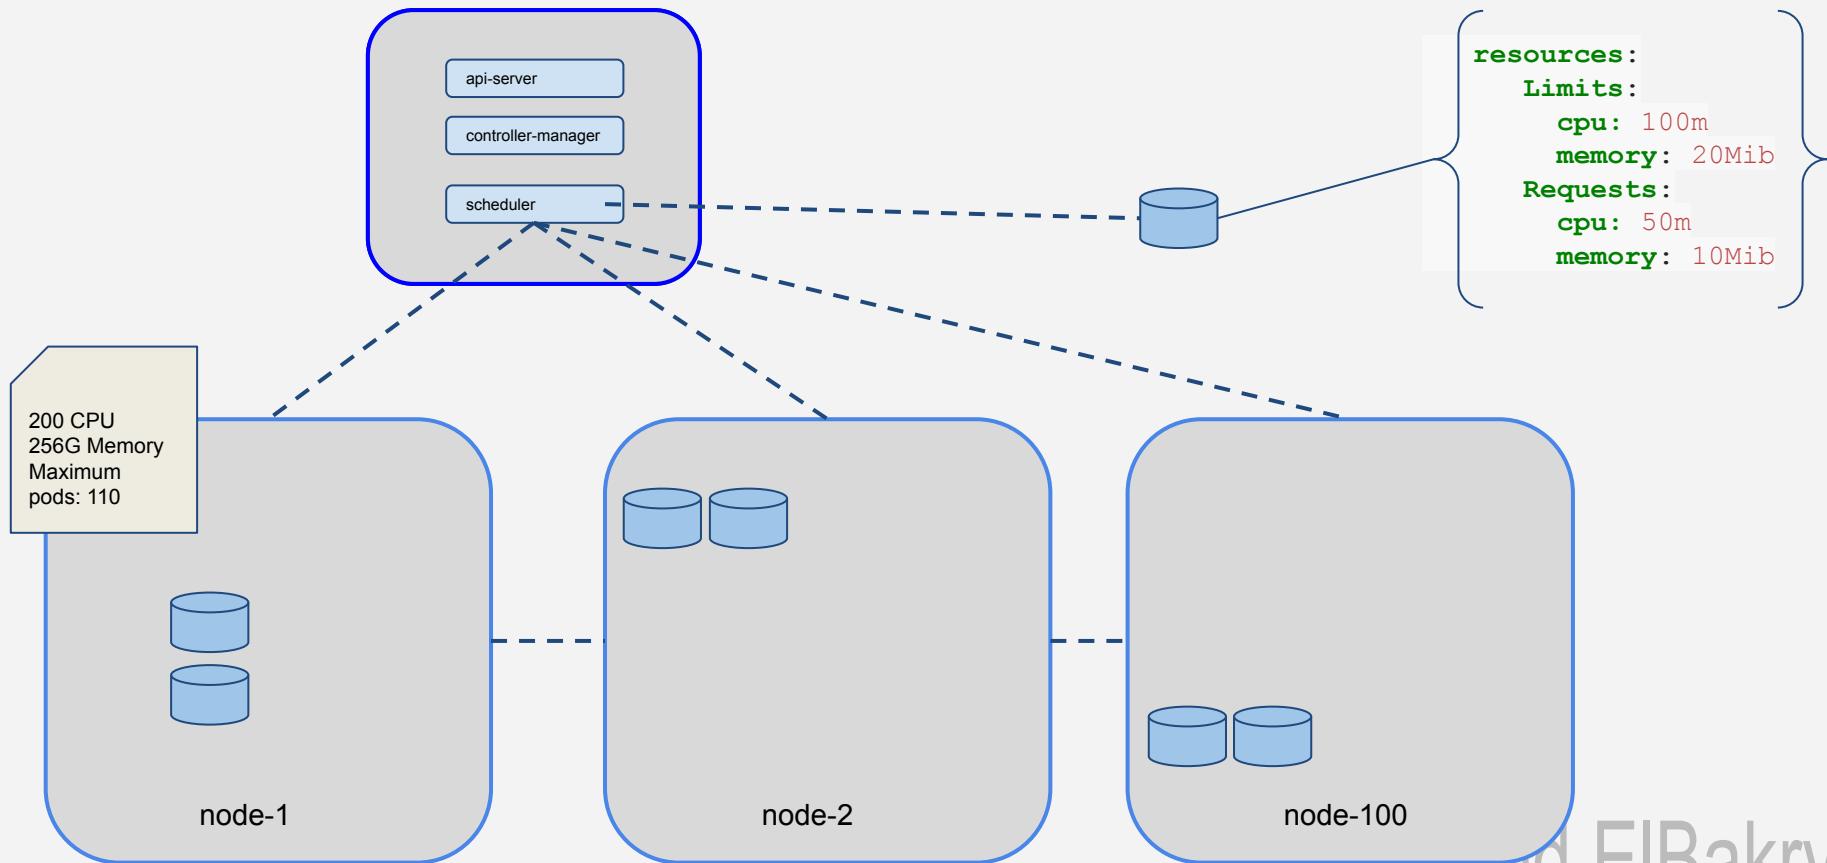
# Resources



# Resources



# Resources



# Limit Ranges

```
apiVersion: v1
kind: LimitRange
metadata:
  name: limitrange-demo
spec:
  limits:
    - max:
        memory: 1Gi
      min:
        memory: 500Mi
      type: Container
```



Pod creation will fail if:

- Memory is less than 500Mi
- Memory is more than 1Gi

# Limit Ranges

```
apiVersion: v1
kind: LimitRange
metadata:
  name: limitrange-demo
spec:
  limits:
  - max:
      memory: 1Gi
    min:
      memory: 500Mi
    type: Container
```

```
apiVersion: v1
kind: LimitRange
metadata:
  name: limitrange-demo-2
spec:
  limits:
  - type: Container
    max:
      cpu: 2
      memory: 500Mi
    Min:
      cpu: 100m
      memory: 50Mi
    defaultRequest:
      cpu: 200m
      memory: 100Mi
```

Pod creation will fail if:

- Memory is less than 50Mi or more than 500Mi
- CPU is less than 100m or more than 2 cpus.

# Resource Quota



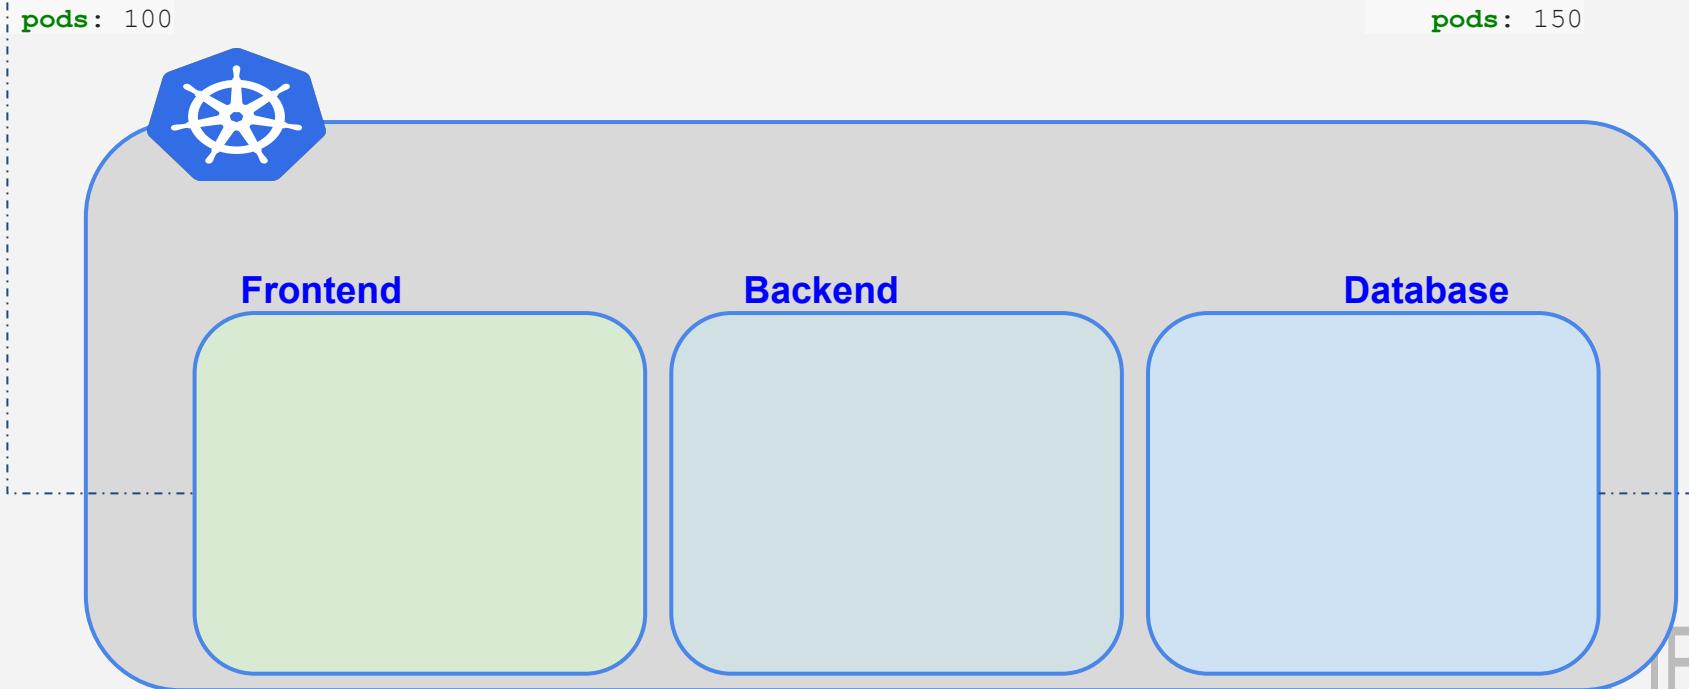
Frontend

Backend

Database

# Resource Quota

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: frontend-rq
spec:
  hard:
    requests.cpu: "50"
    requests.memory: 100Gi
    limits.cpu: "100"
    limits.memory: 200Gi
  pods: 100
```



```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: database-rq
spec:
  hard:
    requests.cpu: "100"
    requests.memory: 500Gi
    limits.cpu: "200"
    limits.memory: 800Gi
  pods: 150
```

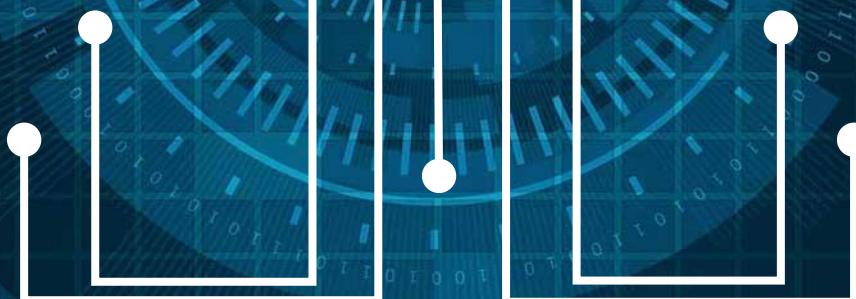


**DEMO**

Amr ElBakry



# THANK YOU



Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

14 - Resources Demo

Ahmed ElBakry

# Resource Quota

## Demo



- ❑ Check node allocatable space
- ❑ Create pods with different resources
  - ❑ Resources not available in the nodes
  - ❑ Very tiny resources
  - ❑ No resources
- ❑ Create LimitRanges
- ❑ Create Resource Quota
- ❑

# Kubernetes From Scratch - Hands-on

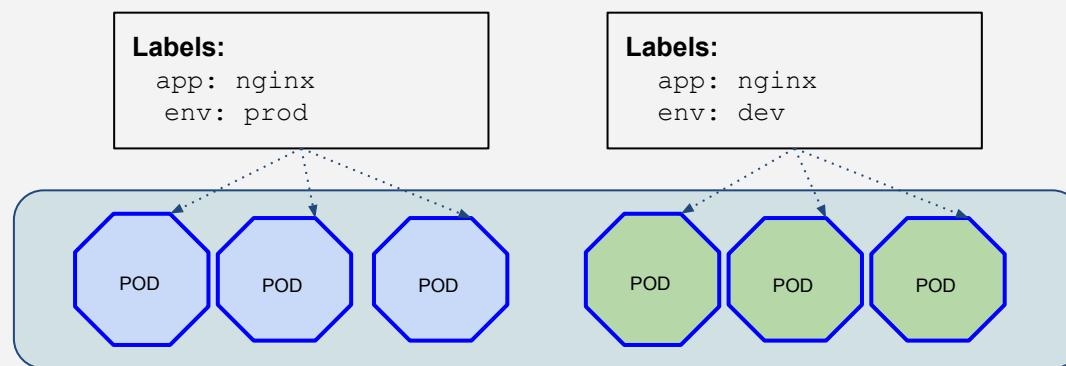
14 - Labels and Selectors

Ahmed ElBakry

# Labels and Selectors

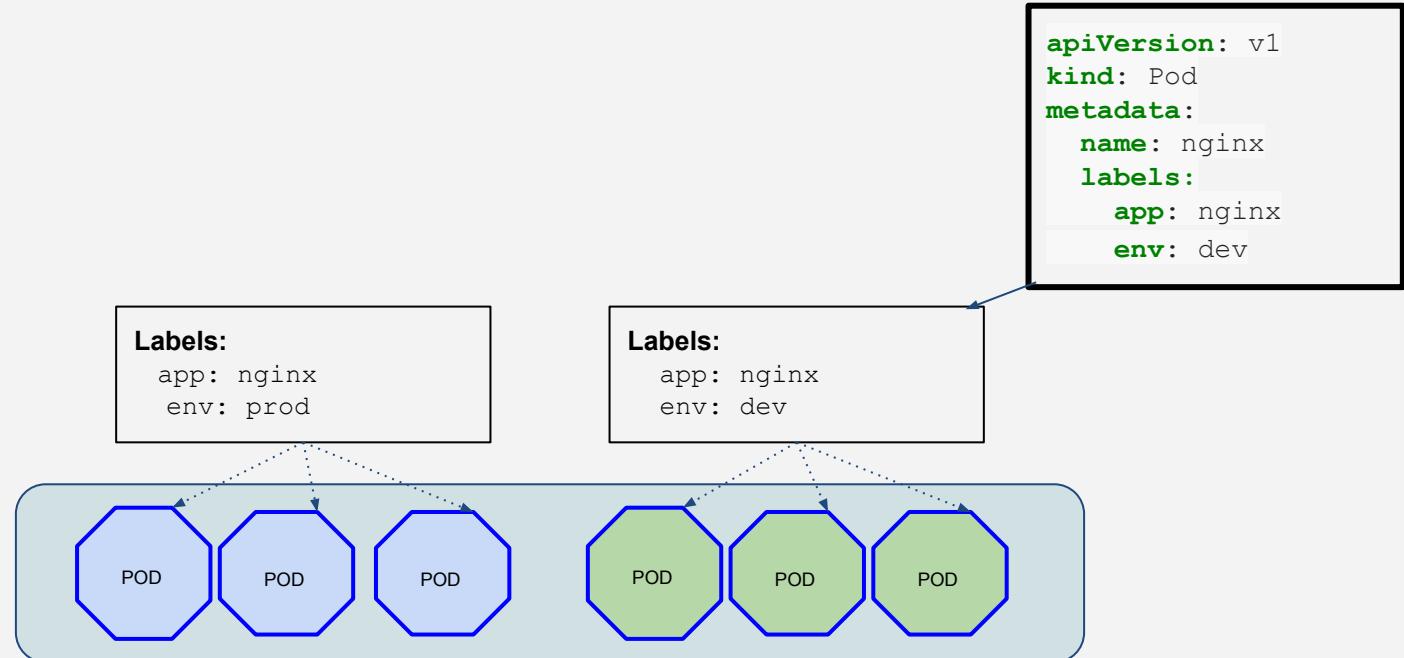
- ❑ Labels are key/value pairs used to identify objects
- ❑ Labels can be attached to almost any Kubernetes Object
- ❑ Labels are mutable, they can be modified at anytime.
- ❑ Selectors are a way to select objects based on their labels
- ❑ Types of Selectors
  - ❑ Equality-based selectors
    - ❑ environment = production
    - ❑ tier != frontend
  - ❑ Set-based selectors
    - ❑ environment in (production, qa)
    - ❑ tier notin (frontend, backend)

# Labels and Selectors



- Labels and Selectors

# Labels and Selectors



# Labels and Selectors

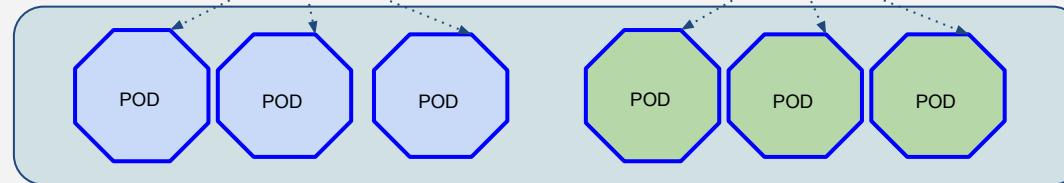
```
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  selector:
    app: nginx
```

**Labels:**

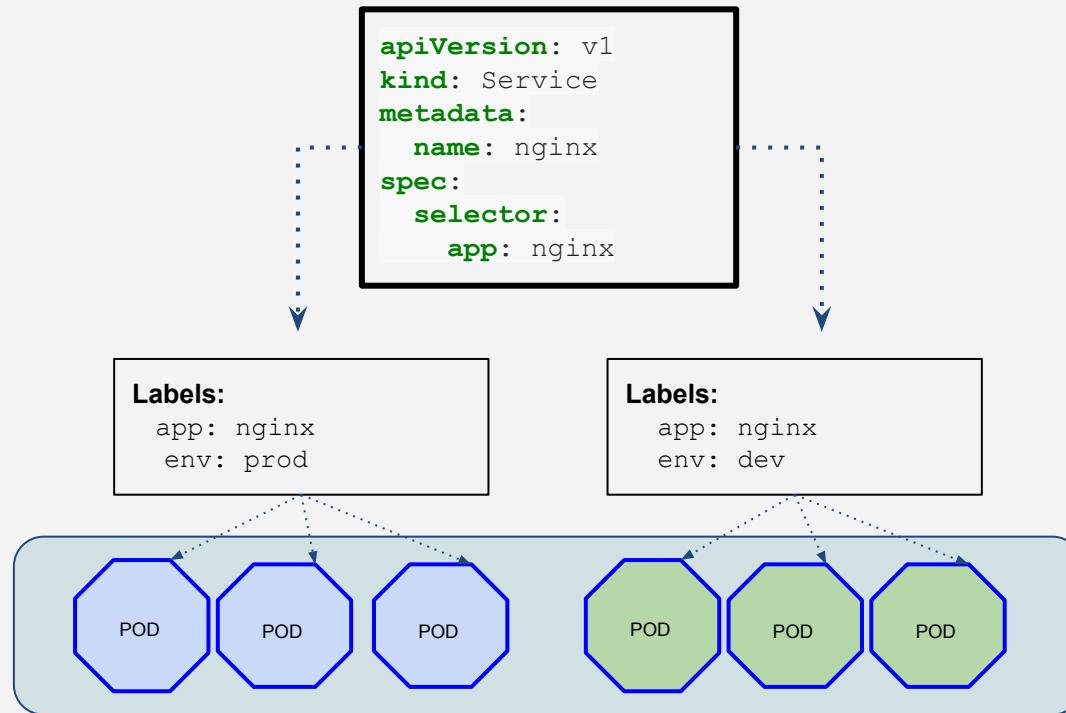
```
app: nginx
env: prod
```

**Labels:**

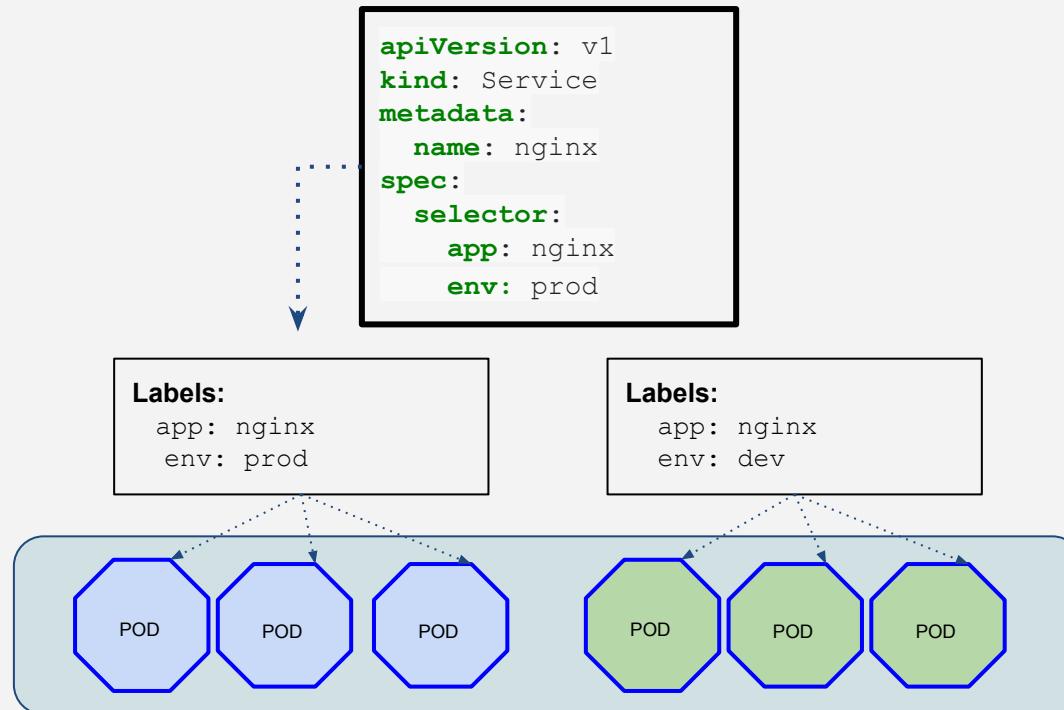
```
app: nginx
env: dev
```



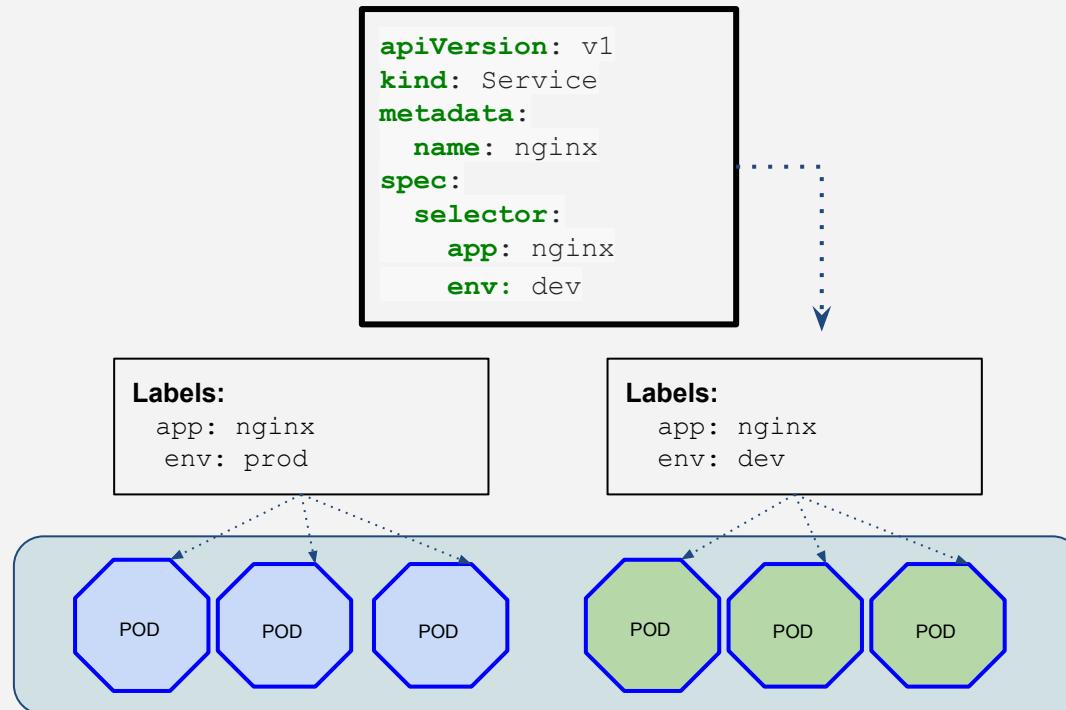
# Labels and Selectors



# Labels and Selectors

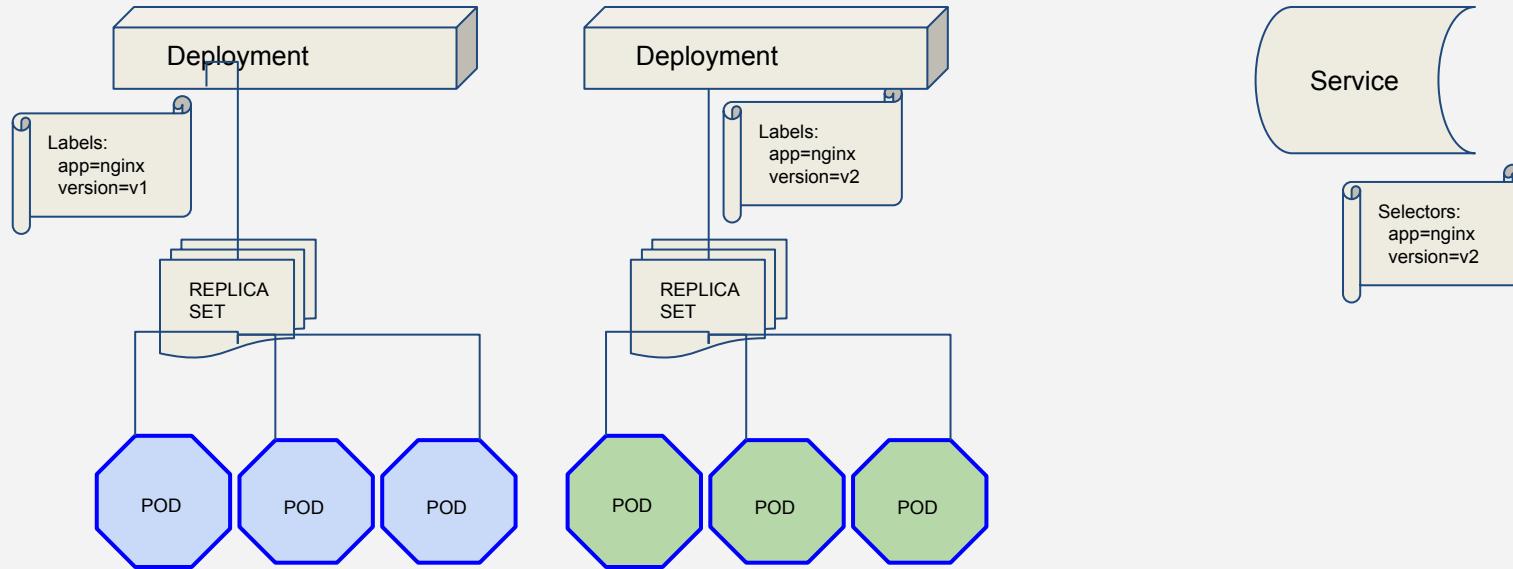


# Labels and Selectors



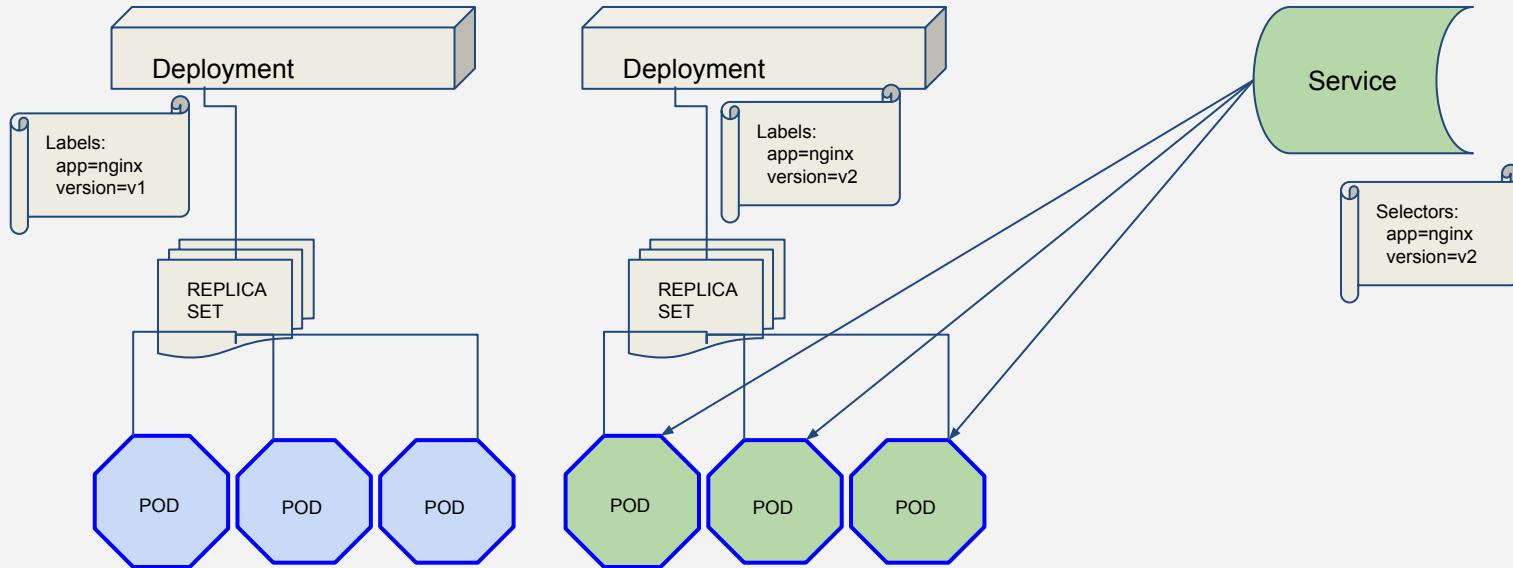
- Labels and Selectors

# Labels and Selectors



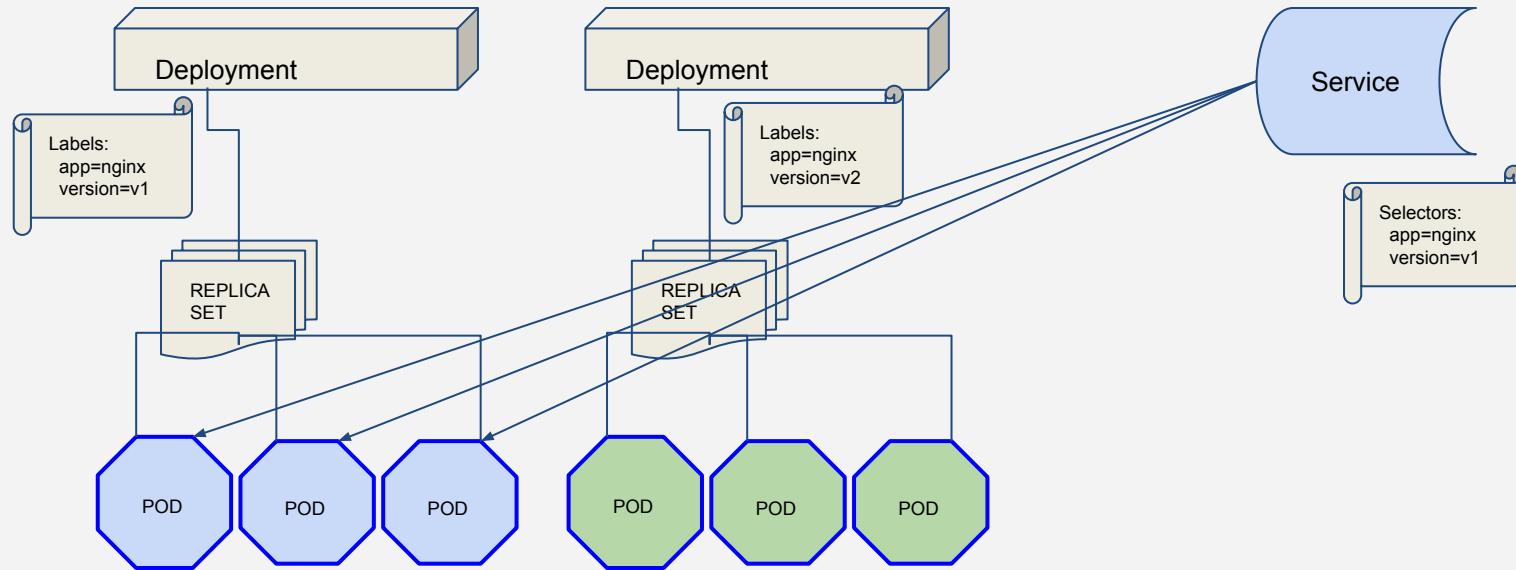
- Labels and Selectors

# Labels and Selectors

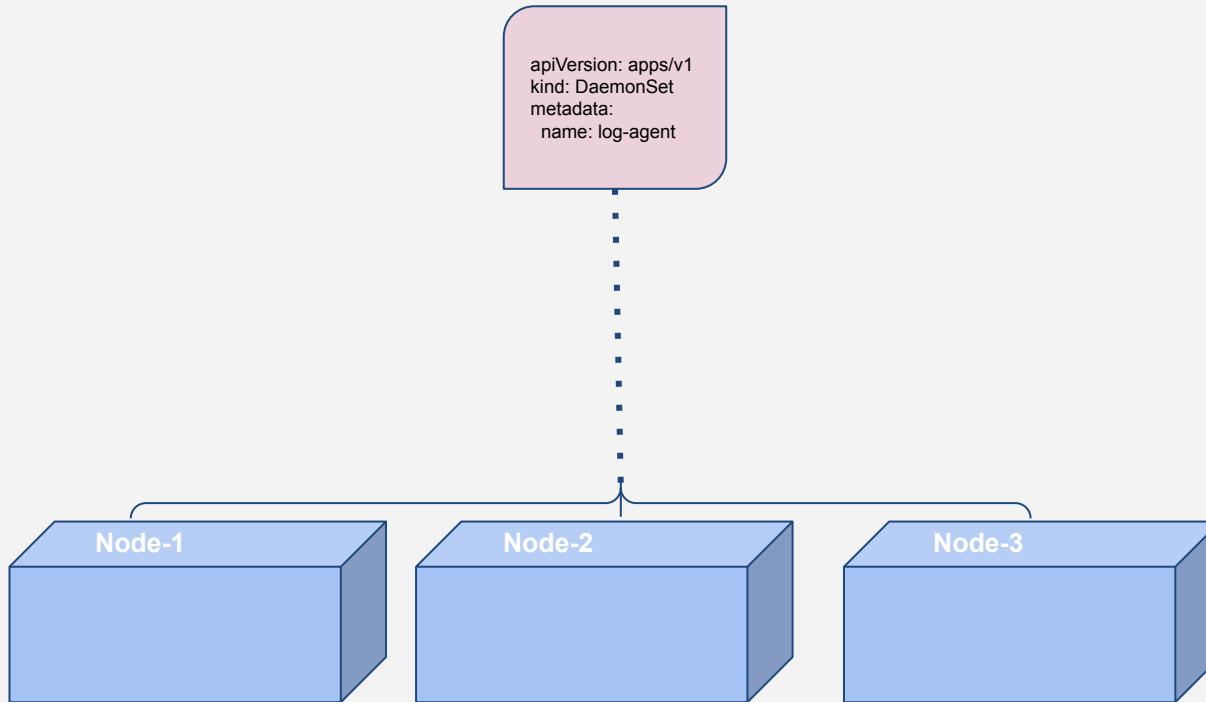


- Labels and Selectors

# Labels and Selectors

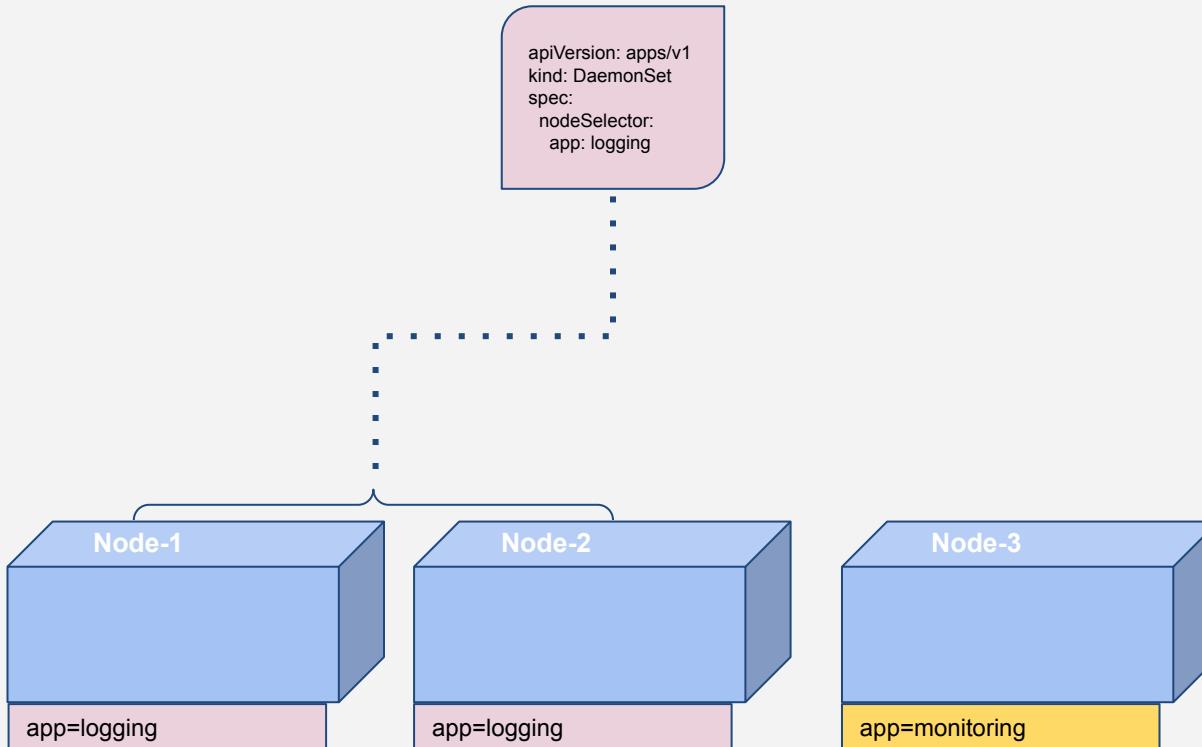


# Labels and Selectors



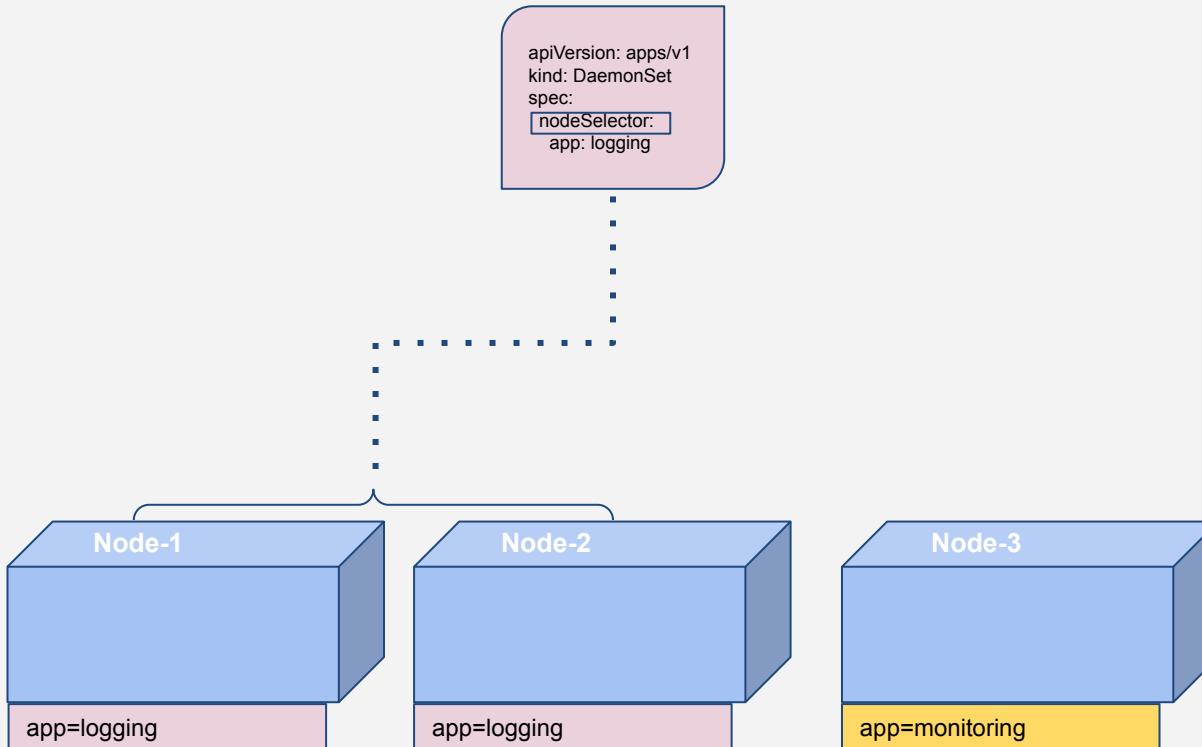
- Labels and Selectors

# Labels and Selectors

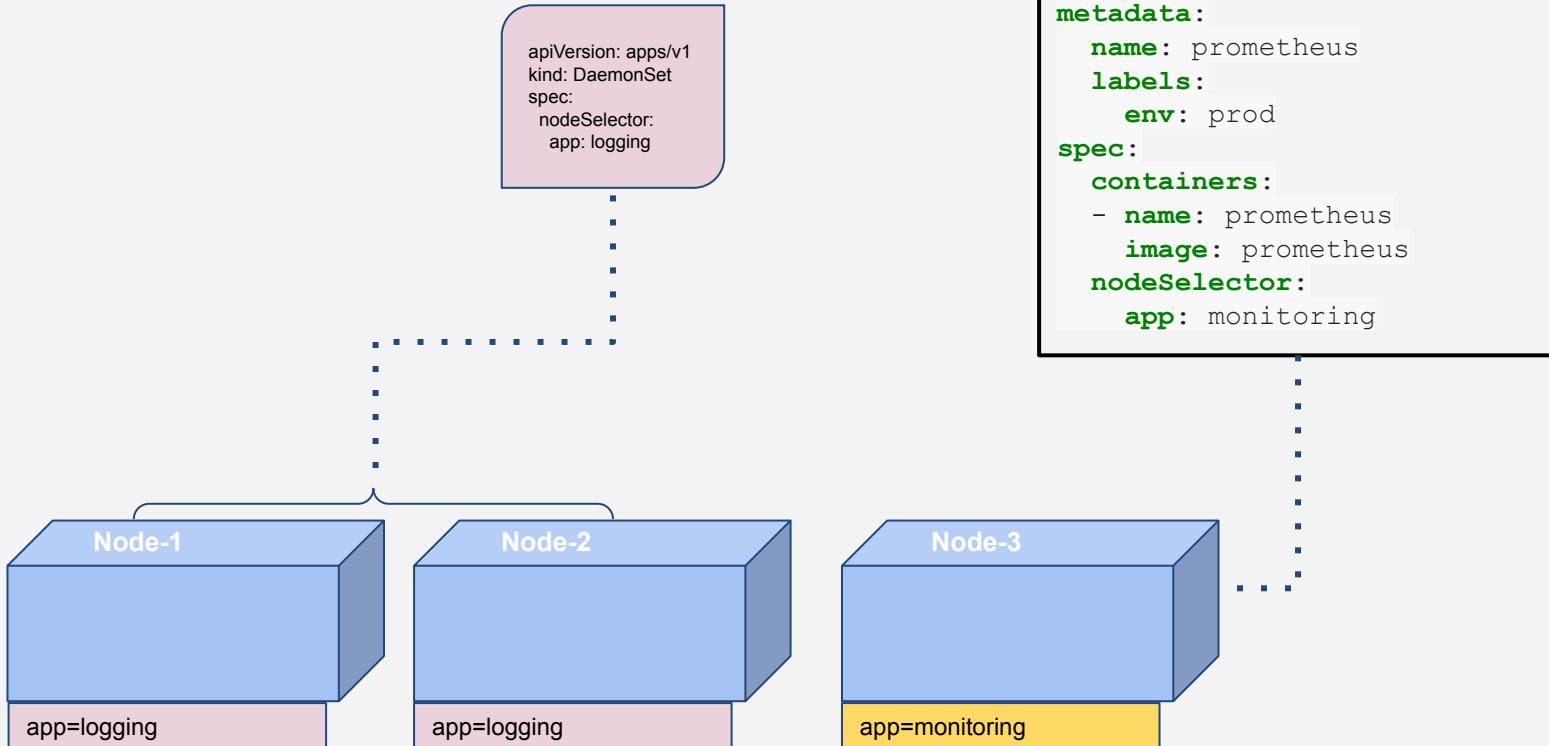


- Labels and Selectors

# Labels and Selectors



# Labels and Selectors



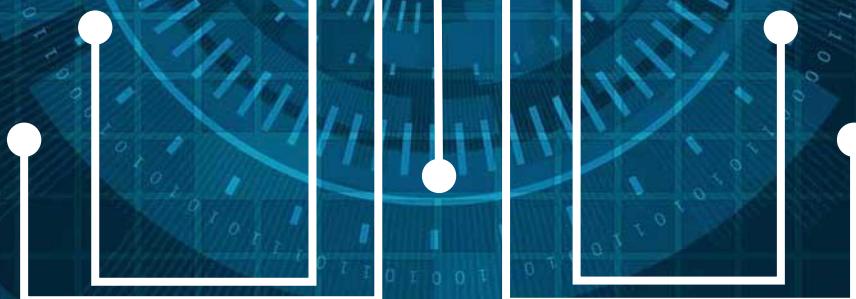


**DEMO**

Amr ElBakry



# THANK YOU



Ahmed ElBakry

# Kubernetes From Scratch - Hands-on

15 - Nodes

Ahmed ElBakry

- Labels and Selectors

# Nodes

# Nodes

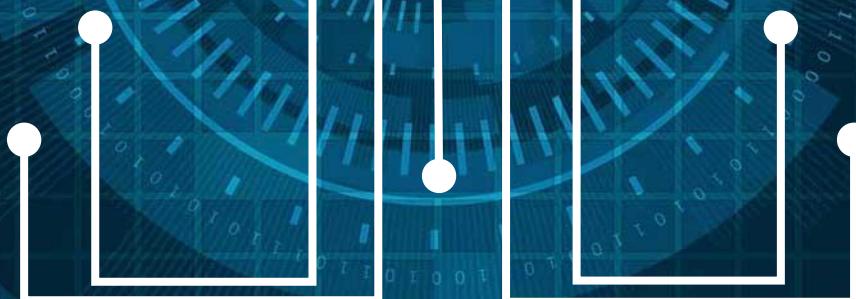
## Demo



- ❑ Capacity and Allocatable
- ❑ Check service running on nodes
- ❑ Check running pods on the nodes
- ❑ Taints and Tolerations
- ❑ Check node status
- ❑ Taint the node
- ❑ Drain the node
- ❑
- ❑



# THANK YOU



Ahmed ElBakry



# **Module 5**

## **Pods**

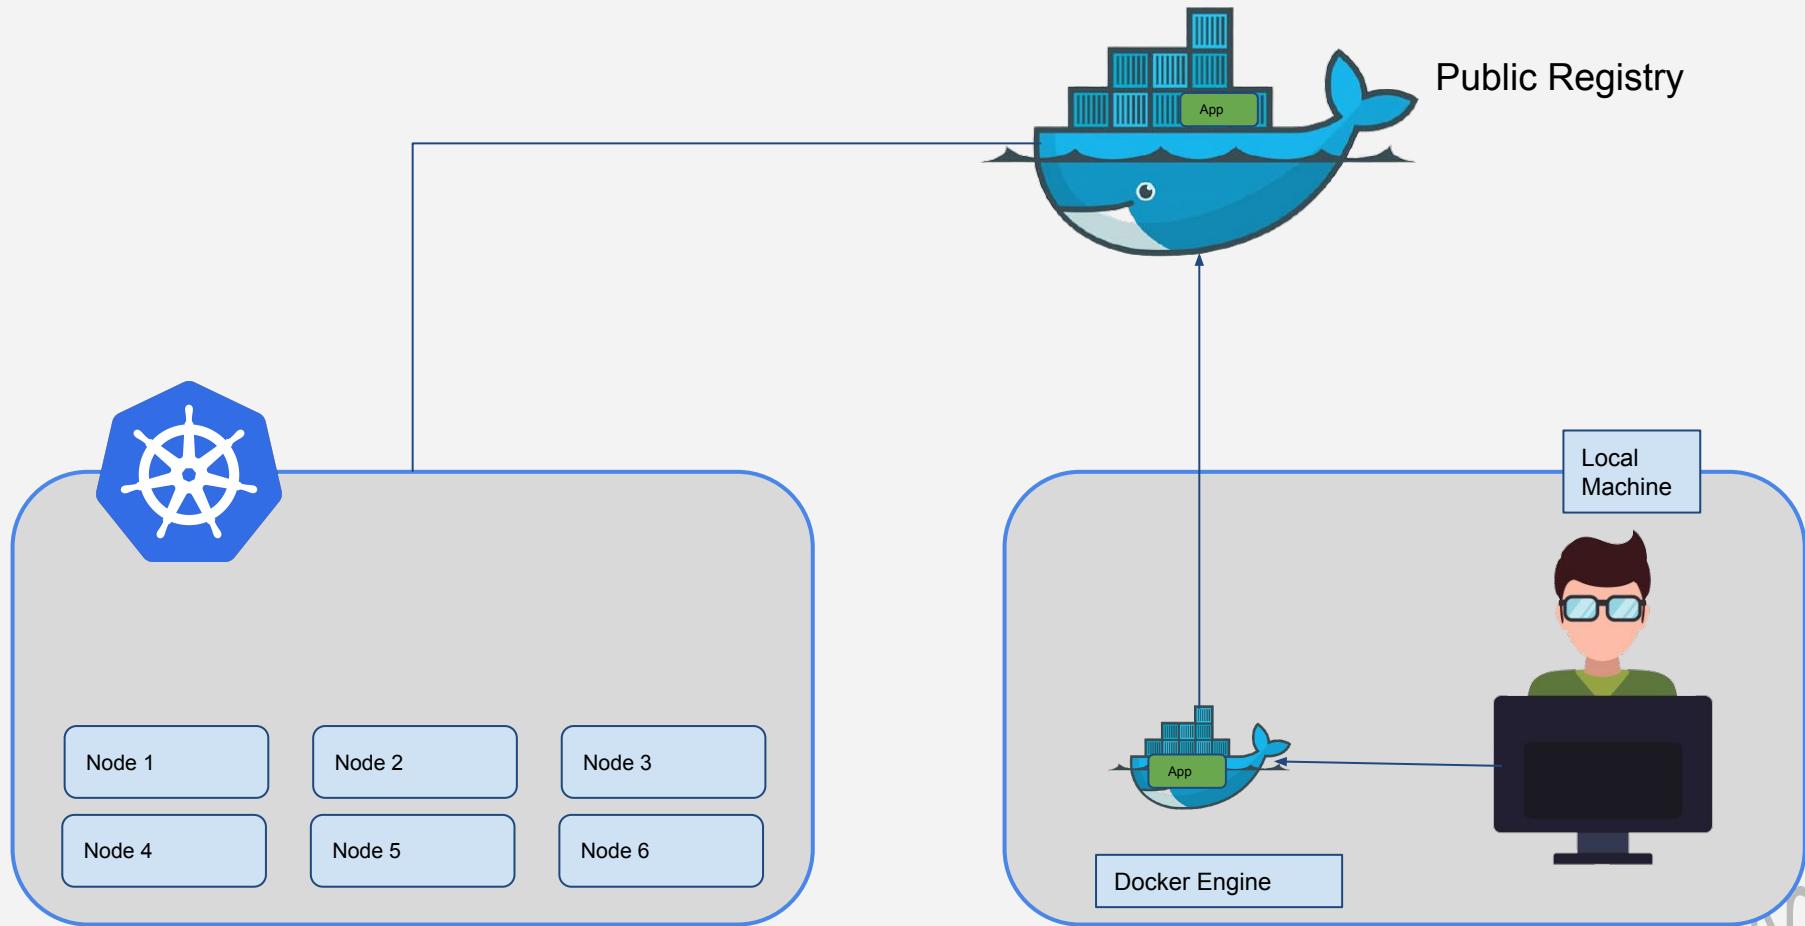
Ahm

# Pods

- Pods are the smallest deployable units in Kubernetes.
- A Pod can be one or more containers.

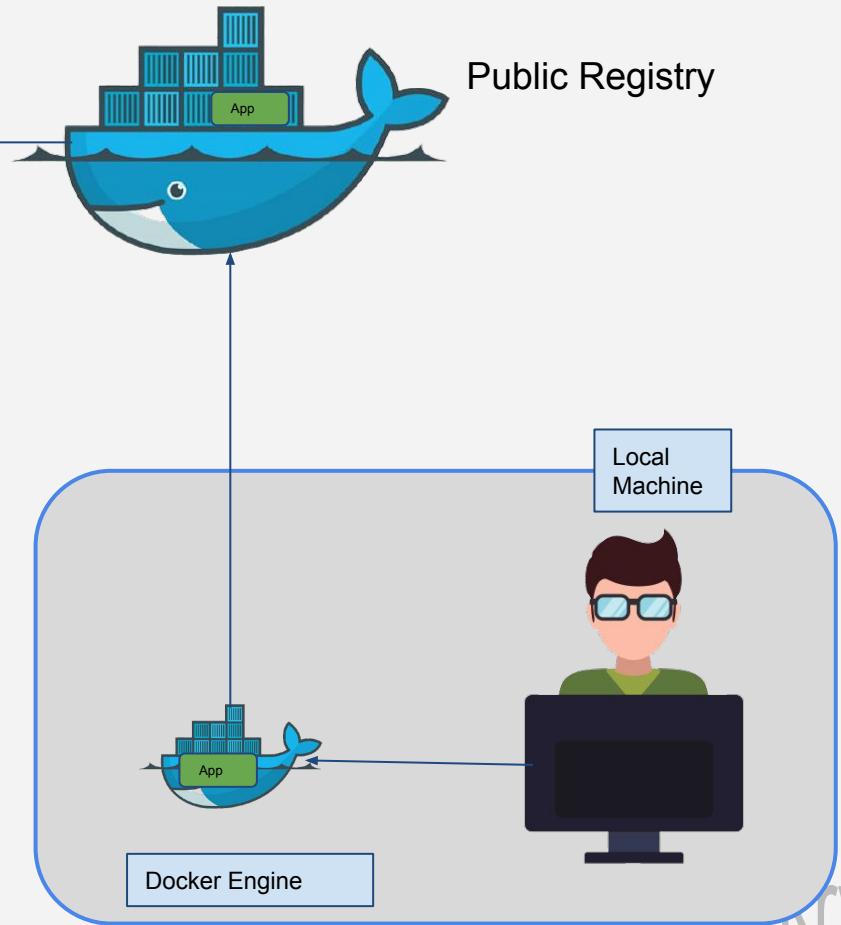
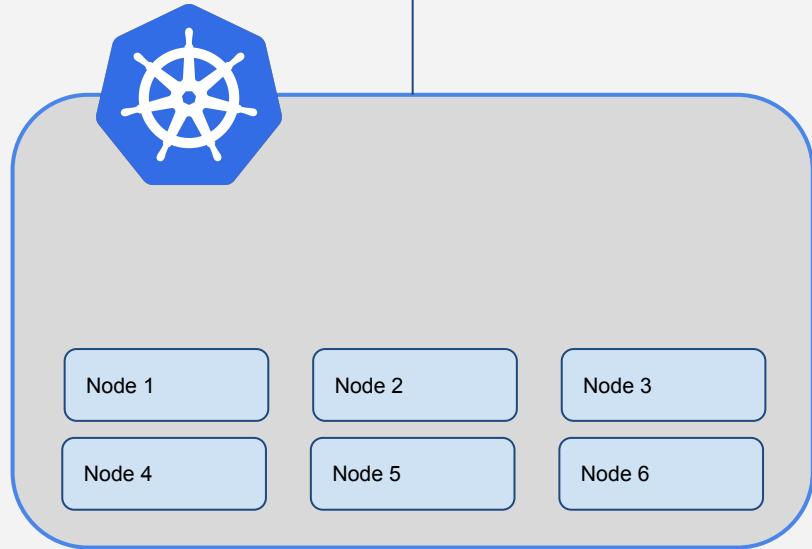


# Pods



# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

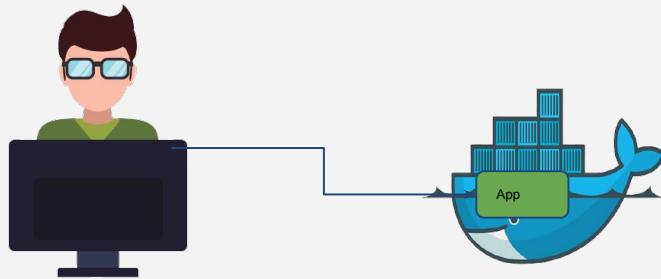


# Pods

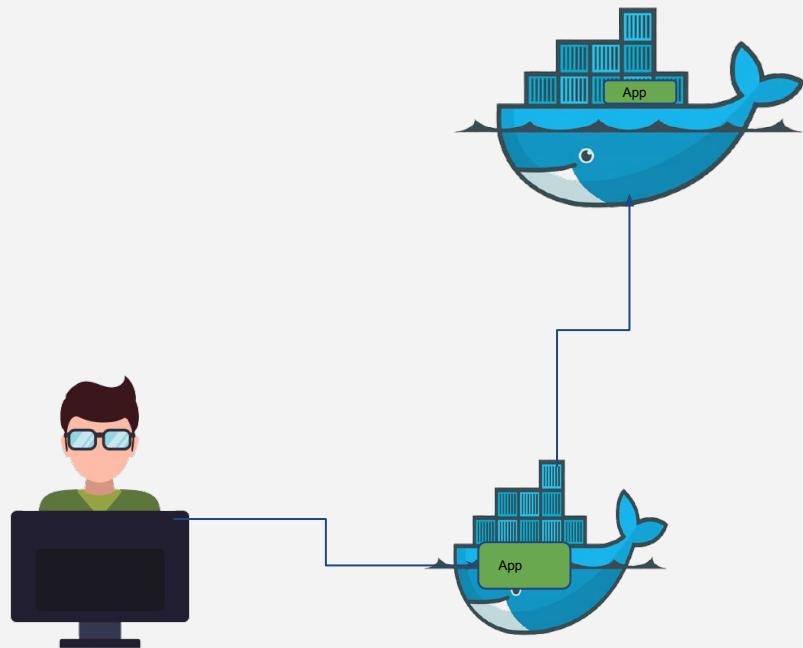


Ahmed ElBakry

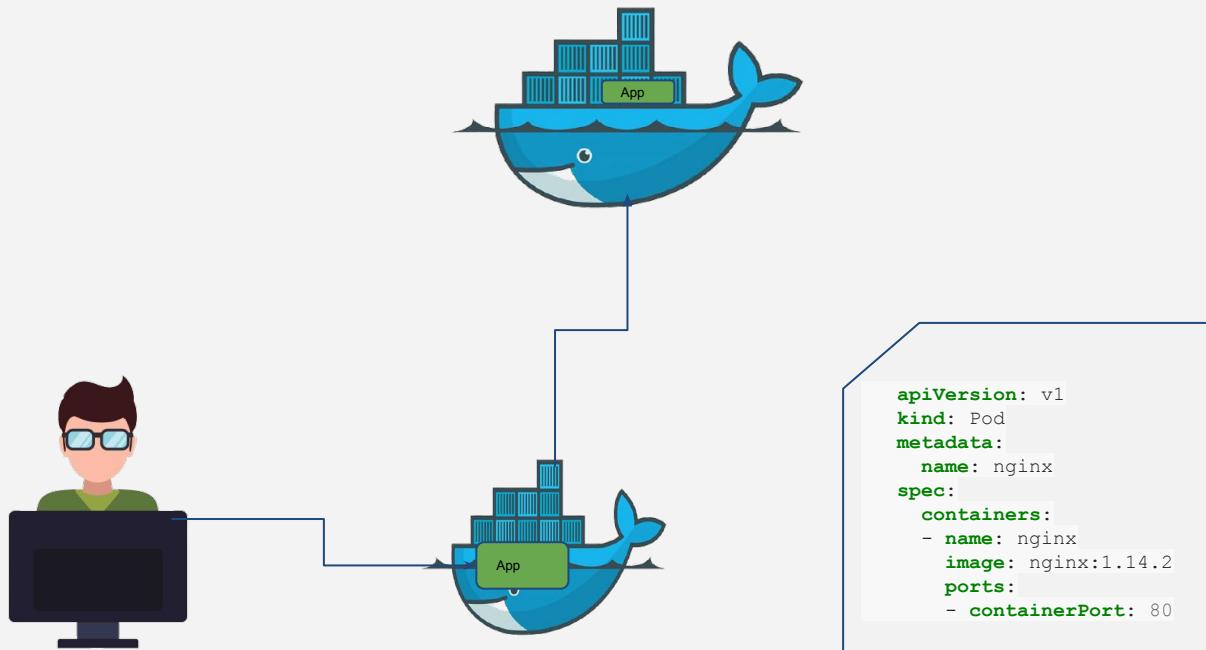
# Pods



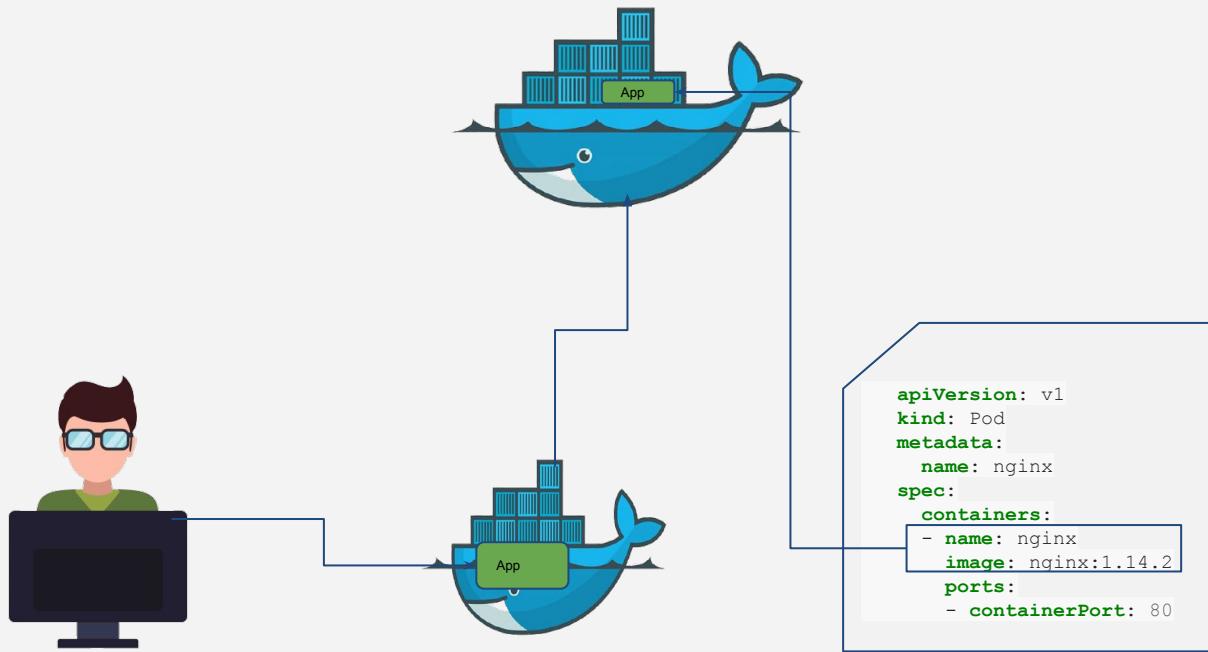
# Pods



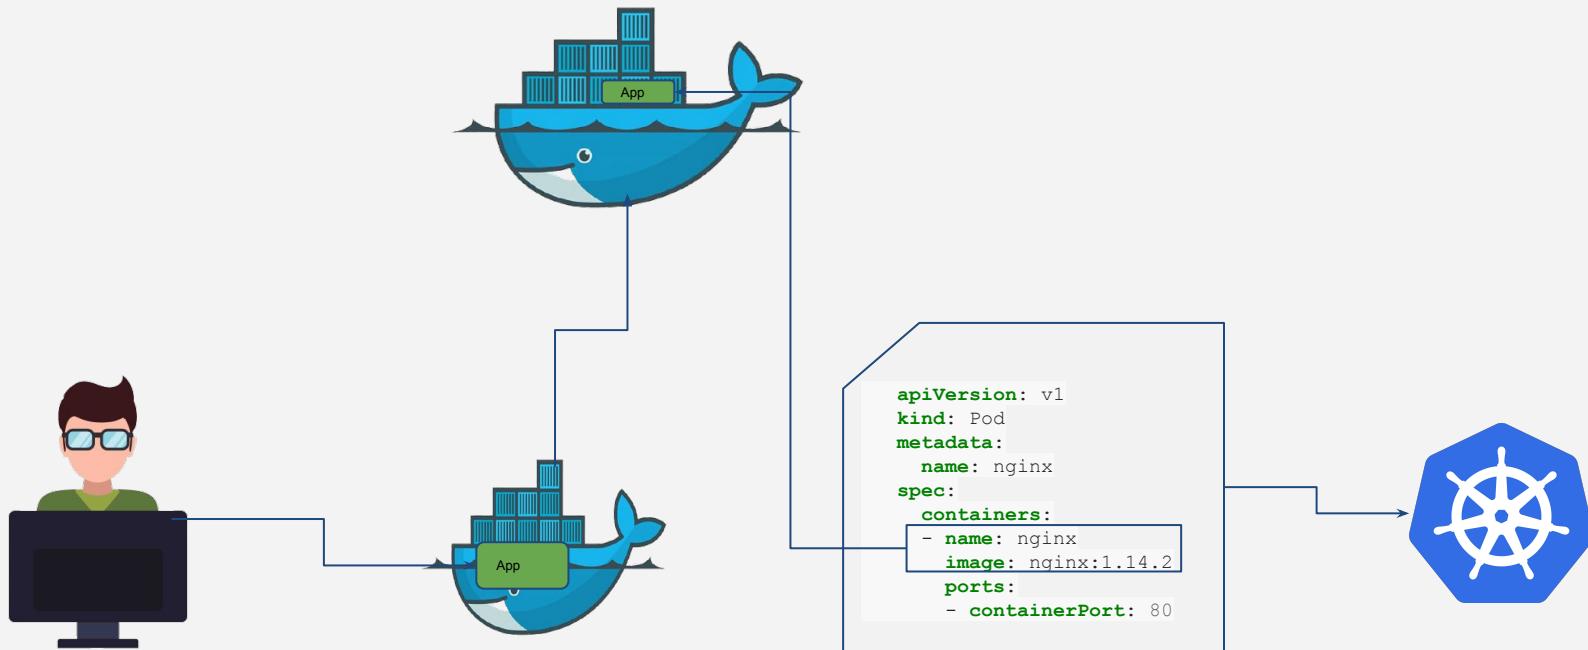
# Pods



# Pods

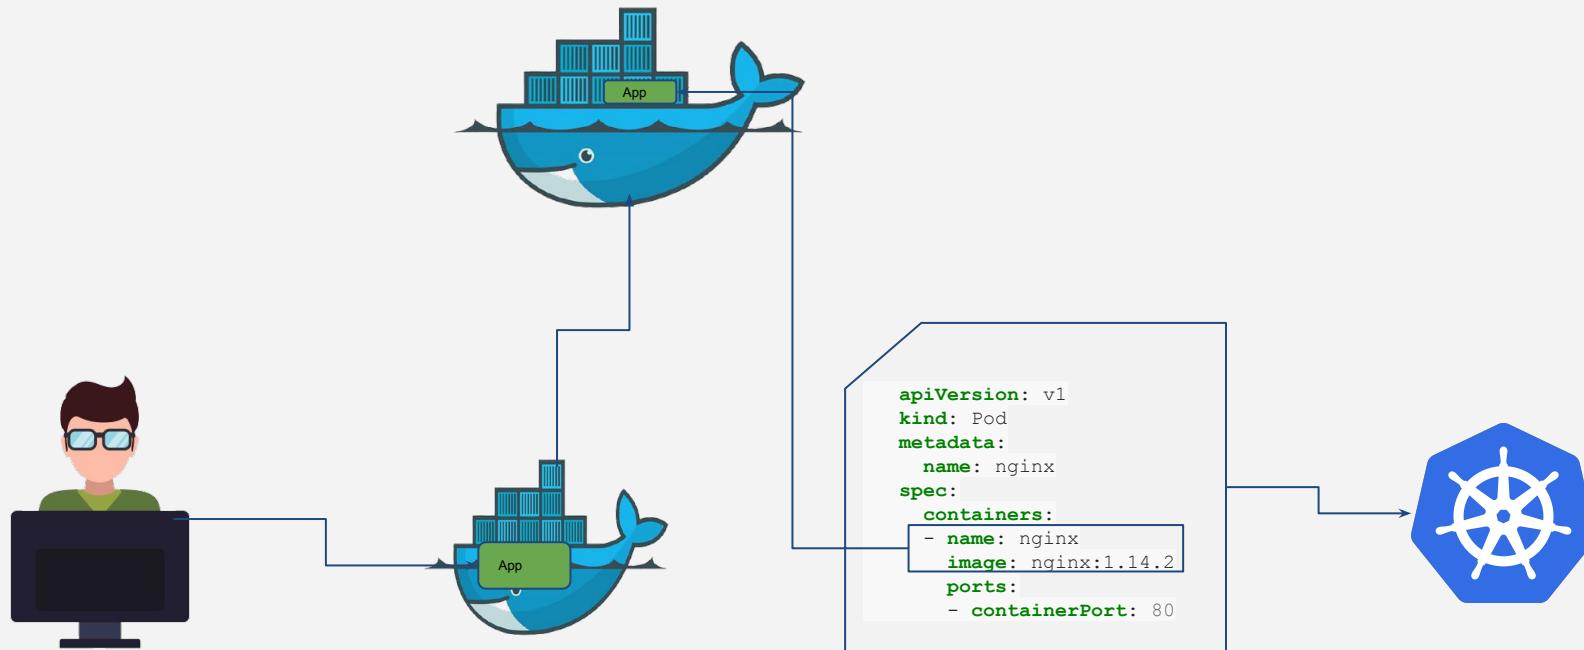


# Pods



Ahmed ElBakry

# Pods



Ahmed ElBakry

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

# Pods

```
(* |kind-c11:default)-> kubectl explain pod
```

KIND: Pod  
VERSION: v1

## DESCRIPTION:

Pod is a collection of containers that can run on a host. This resource is created by clients and scheduled onto hosts.

## FIELDS:

apiVersion <string>

APIVersion defines the versioned schema of this representation of an object. Servers should convert recognized schemas to the latest internal value, and may reject unrecognized values. More info:  
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources>

kind <string>

Kind is a string value representing the REST resource this object represents. Servers may infer this from the endpoint the client submits requests to. Cannot be updated. In CamelCase. More info:  
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds>

metadata <Object>

Standard object's metadata. More info:  
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#metadata>

spec <Object>

Specification of the desired behavior of the pod. More info:  
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status>

status <Object>

Most recently observed status of the pod. This data may not be up to date.  
Populated by the system. Read-only. More info:  
<https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status>

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

# Pods

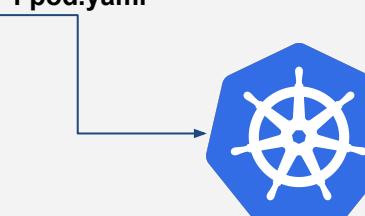
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

(⌘ |kind-c11:default)-> kubectl create -f pod.yaml

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

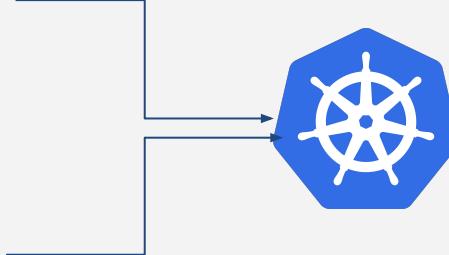
(⎈ |kind-c11:default)-> [kubectl create -f pod.yaml](#)



# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

(\* |kind-c11:default)-> kubectl create -f pod.yaml



(\* |kind-c11:default)-> kubectl get pod

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

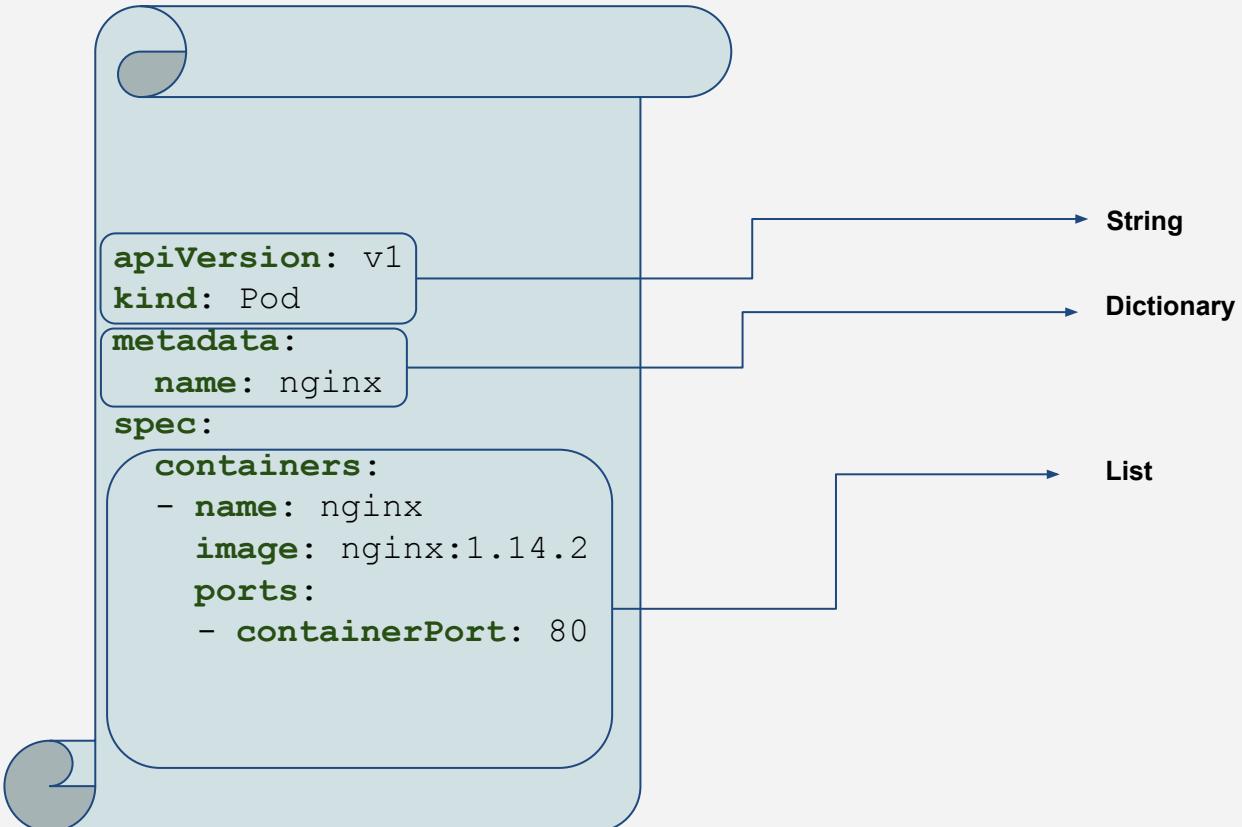
String

# Pods

The diagram illustrates the structure of a Kubernetes Pod specification. The main structure is a light blue rounded rectangle. Inside, there are several key fields:

- apiVersion: v1**: An annotation pointing to a **String**.
- kind: Pod**: An annotation pointing to a **Dictionary**.
- metadata:** A nested section containing:
  - name: nginx**: An annotation pointing to a **String**.
- spec:** A nested section containing:
  - containers:** A list containing:
    - name: nginx**
    - image: nginx:1.14.2**
  - ports:** A list containing:
    - containerPort: 80**

# Pods



# Pods

```
# This is the pod
# Definition yaml file

apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

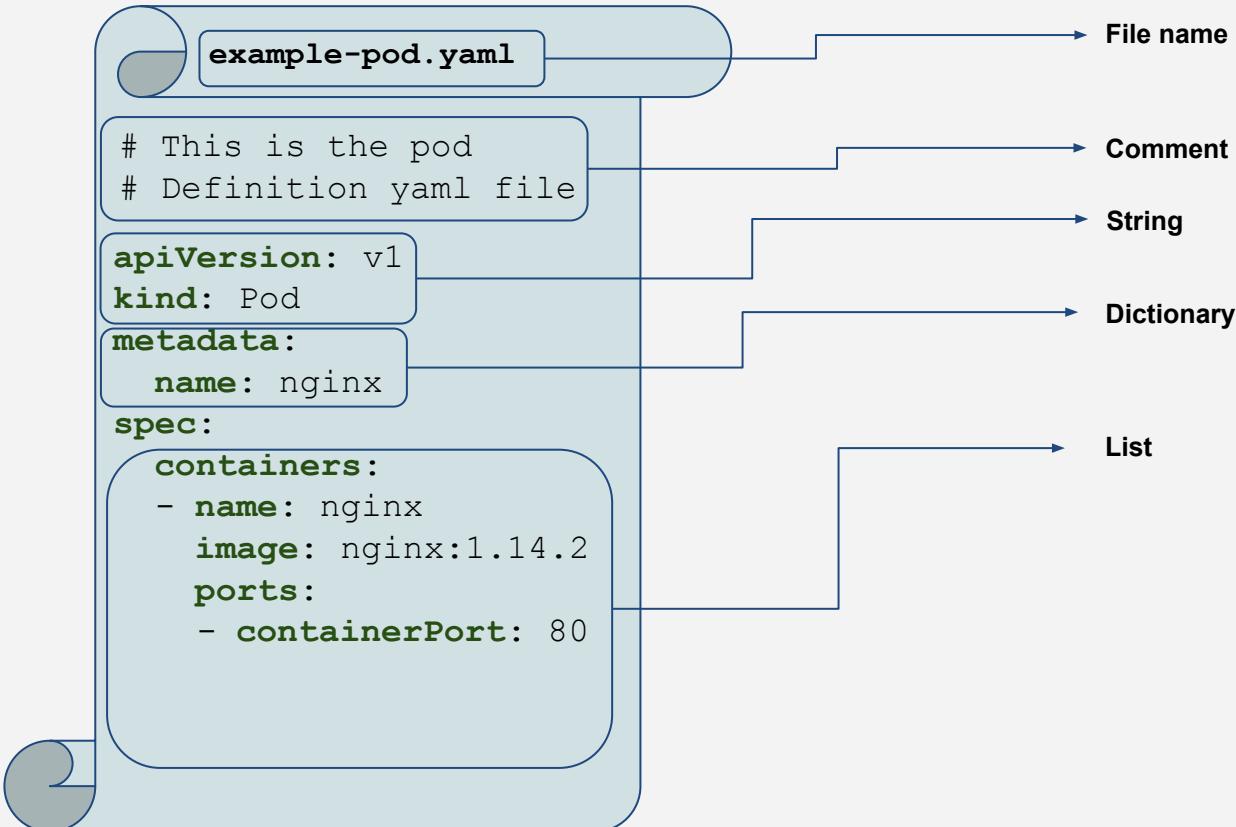
Comment

String

Dictionary

List

# Pods



# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
    - name: busybox
      image: busybox:1.20
      volumeMounts:
        - name: config-vol
          mountPath: /config
      command:
        - /bin/sh
        - -ec
        - |
          if [ -f /persistent/log ]; then
            echo -n "Found old state starting at "; head -n1 /persistent/log
          else
            echo -n "Starting with a fresh state"
          fi
          while sleep 5; do date | tee -a /persistent/log; done
  volumes:
    - name: config-vol
      persistentVolumeClaim:
        claimName: my-lamp-site-data
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: kubernetes.io/hostname
                operator: Exists
containers:
  - name: busybox
    image: busybox:1.20
    volumeMounts:
      - name: config-vol
        mountPath: /config
    command:
      - /bin/sh
      - -ec
      - |
        if [ -f /persistent/log ]; then
          echo -n "Found old state starting at "; head -n1 /persistent/log
        else
          echo -n "Starting with a fresh state"
        fi
        while sleep 5; do date | tee -a /persistent/log; done
volumes:
  - name: config-vol
    persistentVolumeClaim:
      claimName: my-lamp-site-data
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
    - name: busybox
      image: busybox:1.20
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "500m"
      volumeMounts:
        - name: config-vol
          mountPath: /config
      command:
        - /bin/sh
        - -ec
        - |
          if [ -f /persistent/log ]; then
            echo -n "Found old state starting at "; head -n1 /persistent/log
          else
            echo -n "Starting with a fresh state"
          fi
          while sleep 5; do date | tee -a /persistent/log; done
  volumes:
    - name: config-vol
      persistentVolumeClaim:
        claimName: my-lamp-site-data
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
    - name: busybox
      image: busybox:1.20
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "500m"
      volumeMounts:
        - name: config-vol
          mountPath: /config
      command:
        - /bin/sh
        - -ec
        - |
          if [ -f /persistent/log ]; then
            echo -n "Found old state starting at "; head -n1 /persistent/log
          else
            echo -n "Starting with a fresh state"
          fi
          while sleep 5; do date | tee -a /persistent/log; done
  volumes:
  - name: config-vol
    persistentVolumeClaim:
      claimName: my-lamp-site-data
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
    ports:
      - containerPort: 80
```

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    ports:
    - containerPort: 80
```

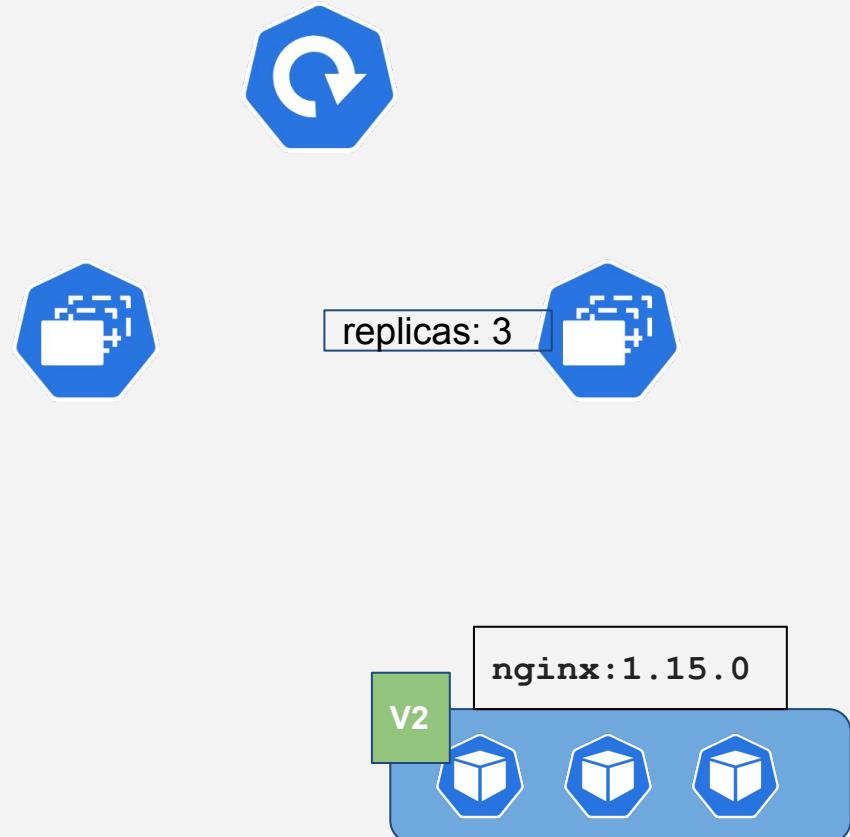
(※ |kind-c1:default)-> kubectl create -f /tmp/pod.yaml  
The Pod "nginx" is invalid: spec.containers[0].image: Required value

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



Ahmed ElBakry

# Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - image: nginx:1.14.2
      name: nginx
      ports:
        - containerPort: 80
```

# Pods

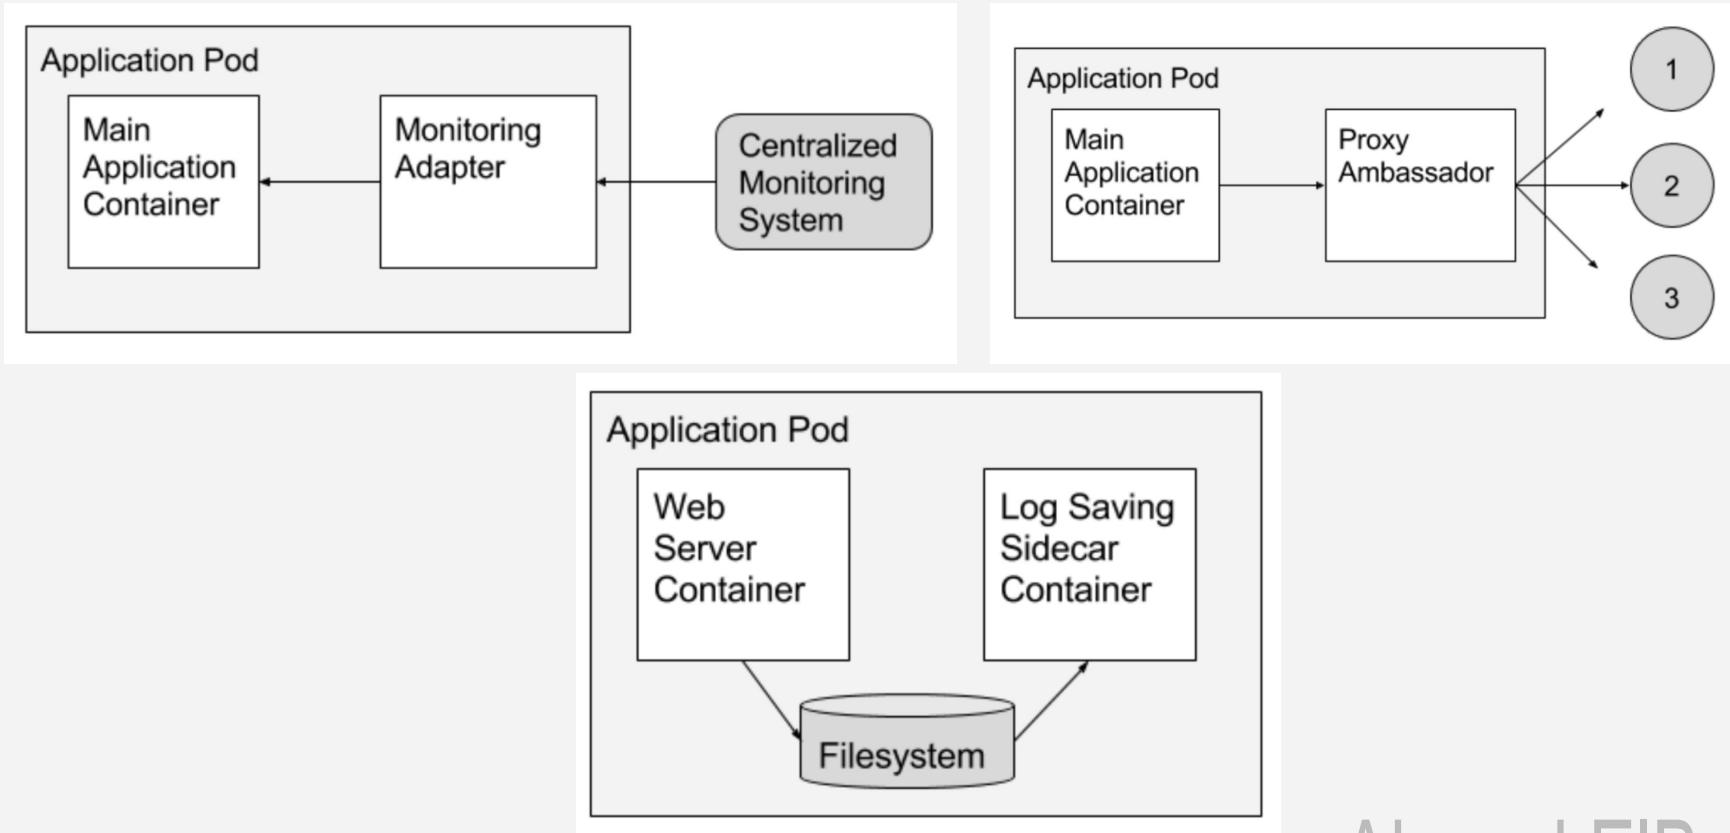
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
    - name: php-fpm
      image: php-fpm:1.14.2
      resources:
        limits:
          memory: "200Mi"
        requests:
          memory: "100Mi"
```

# Pods

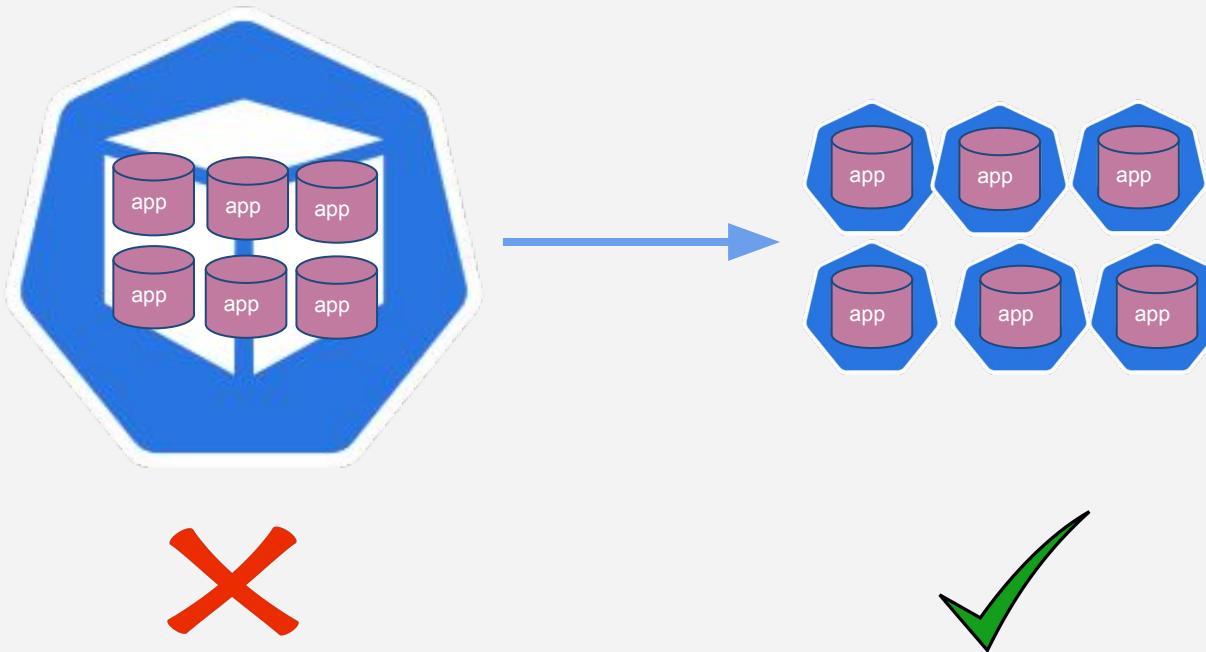
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
    - name: php-fpm
      image: php-fpm:1.12
      resources:
        limits:
          memory: "200Mi"
        requests:
          memory: "100Mi"
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: php-fpm
      image: php-fpm:1.12
      resources:
        limits:
          memory: "200Mi"
        requests:
          memory: "100Mi"
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

# Multi Container Pod

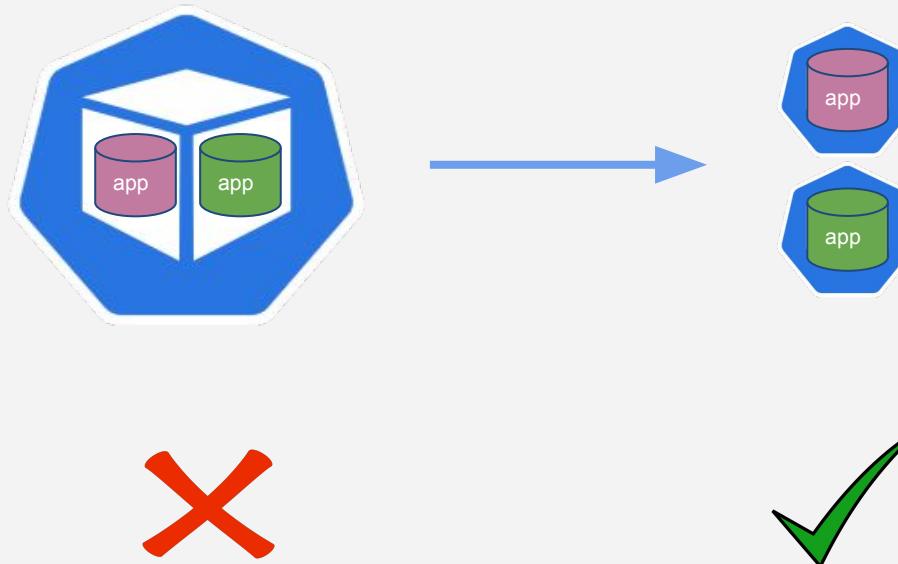


# Multi Container Pod



Ahmed ElBakry

# Multi Container Pod



Ahmed ElBakry



# Module 6

## Deployments

Ahm

# Deployments

## What is Kubernetes Deployment

### Managing The Lifecycle of The Application:

- Create the Pods
- Rollout new version of the application
- Rollback to previous version
- Scale the application
- Self healing

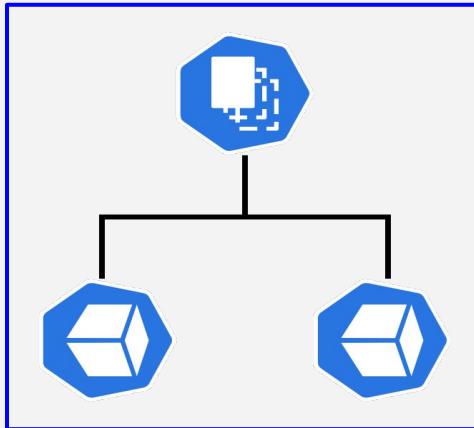
# Deployments



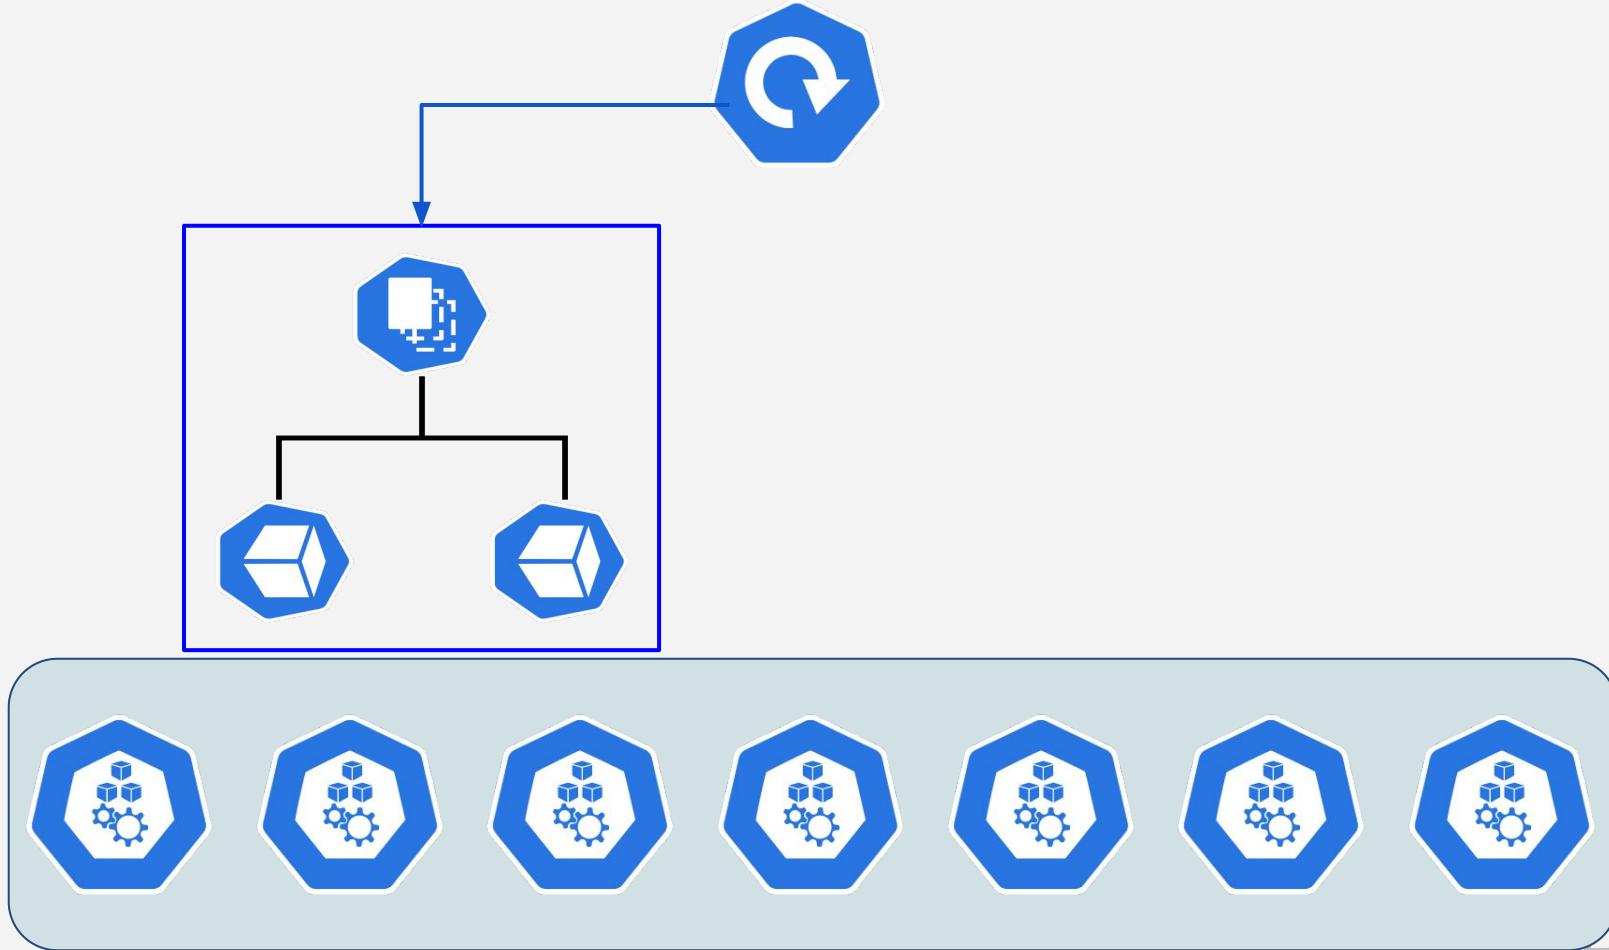
# Deployments



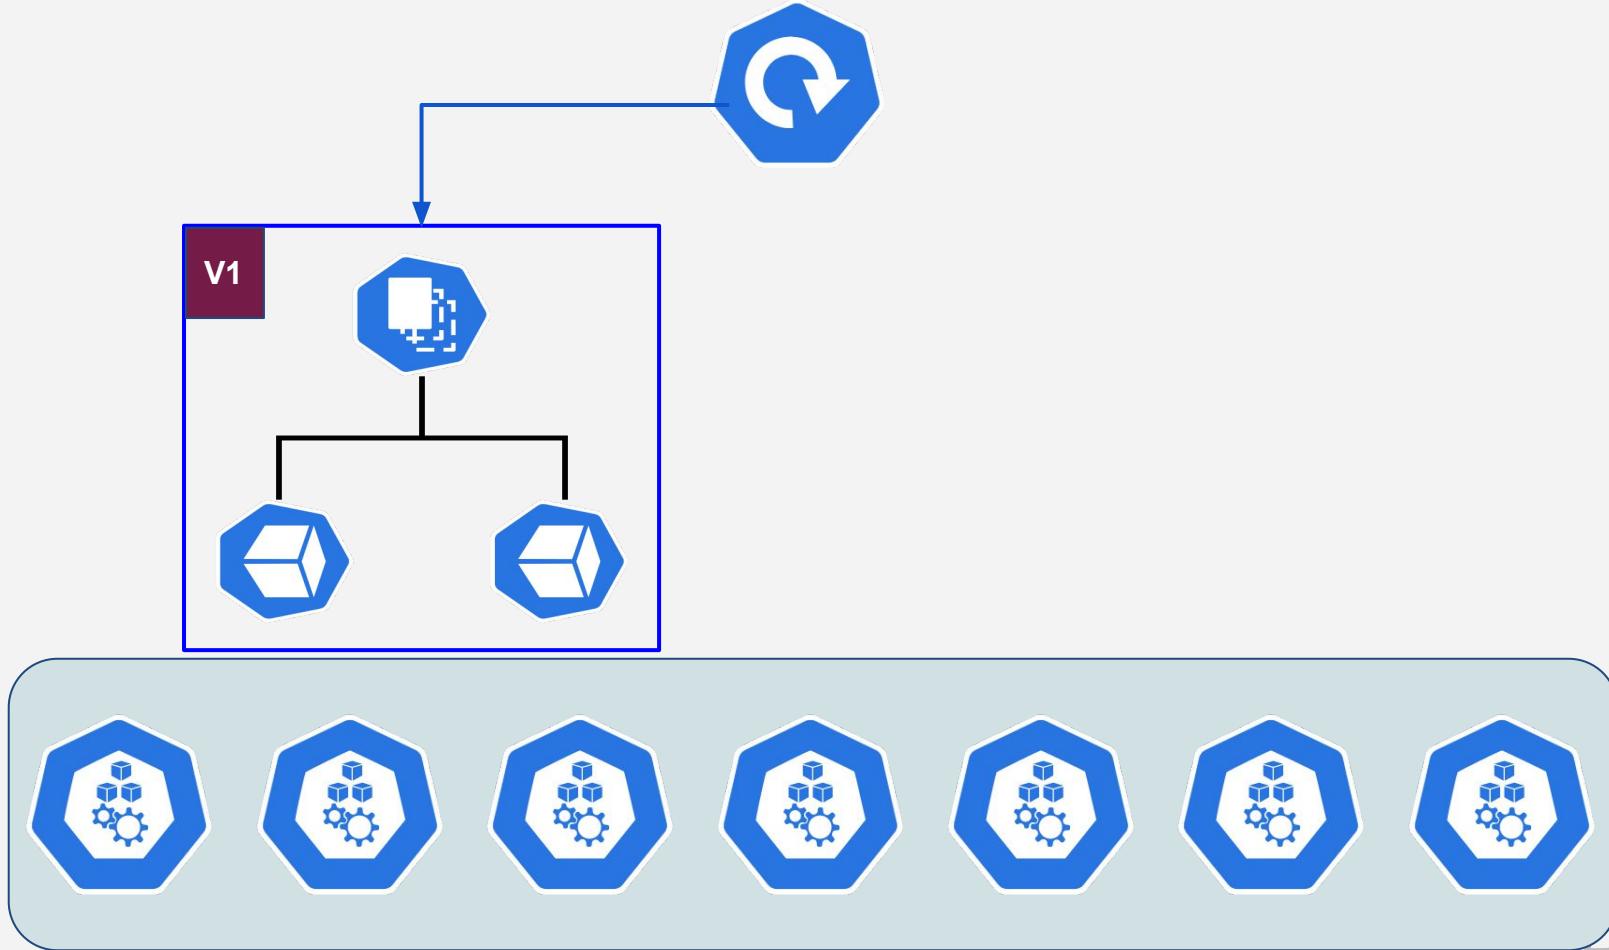
# Deployments



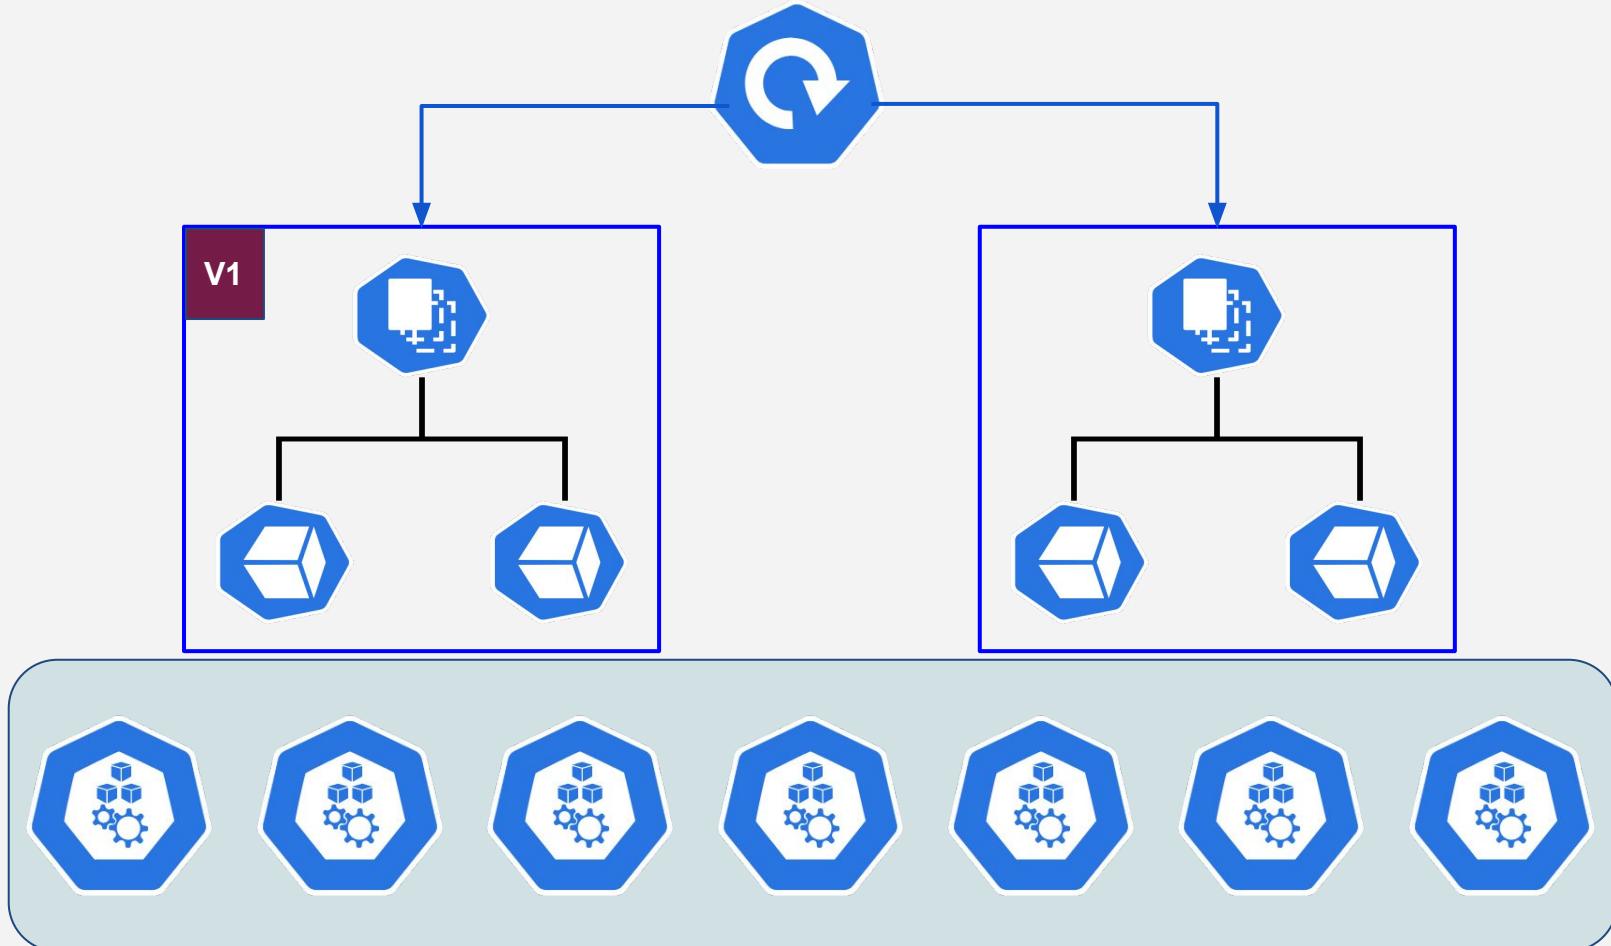
# Deployments



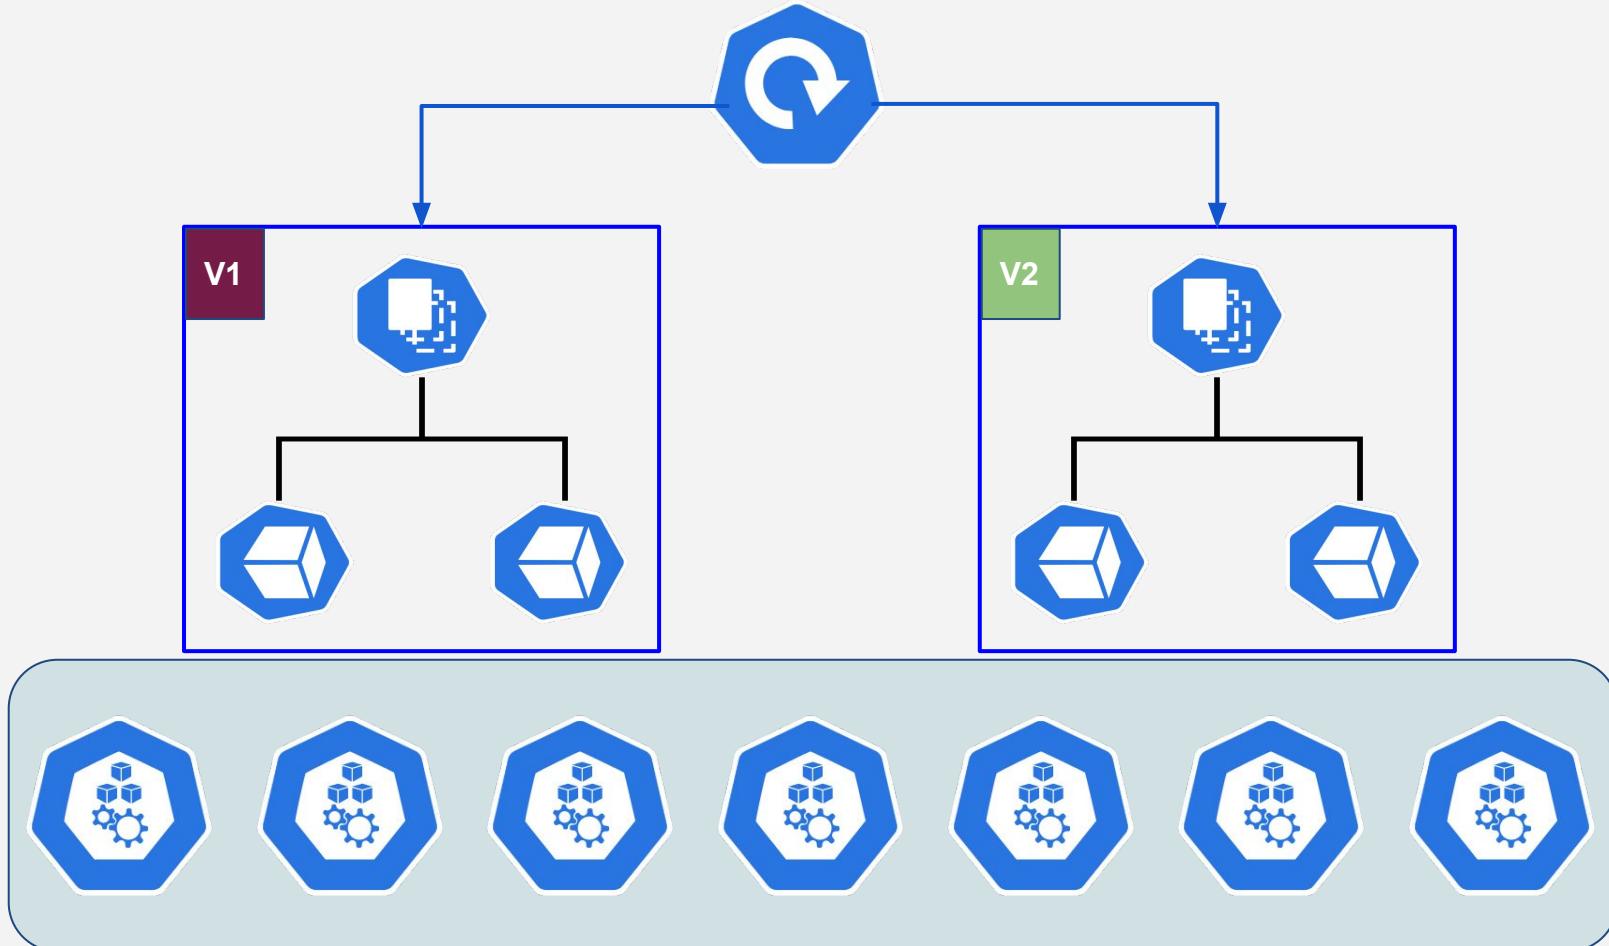
# Deployments



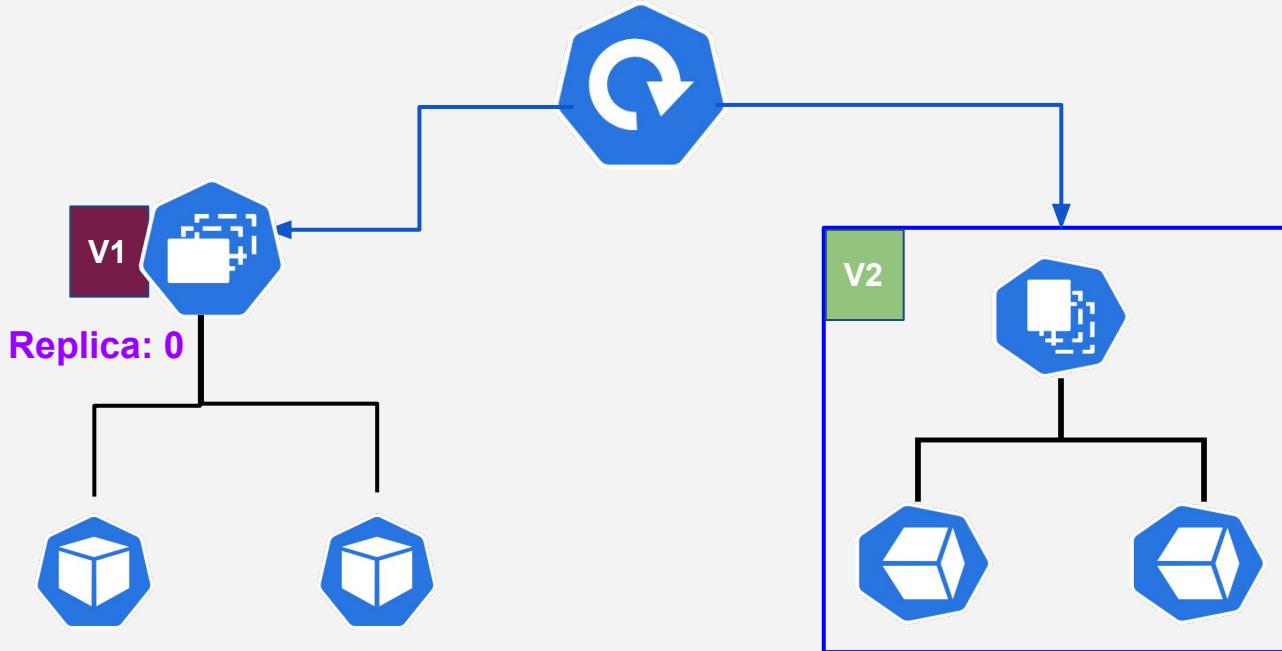
# Deployments



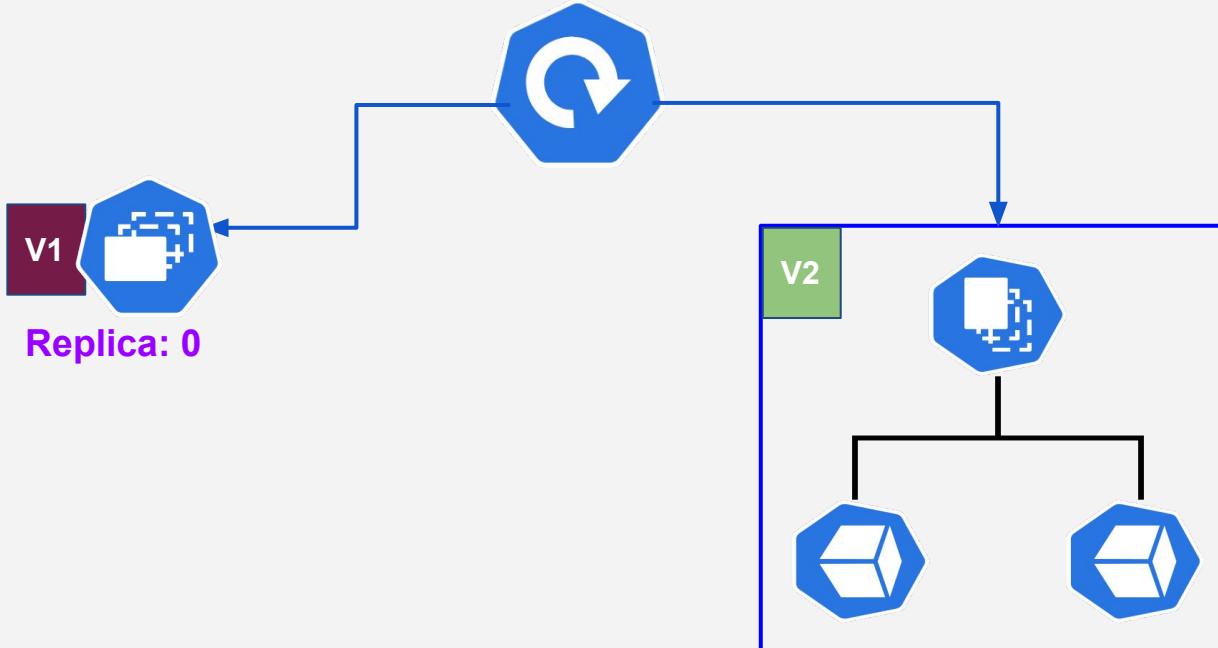
# Deployments



# Deployments



# Deployments



# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

# Deployments



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



replicas: 3



nginx:1.14.2

Ahmed ElBakry

# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```

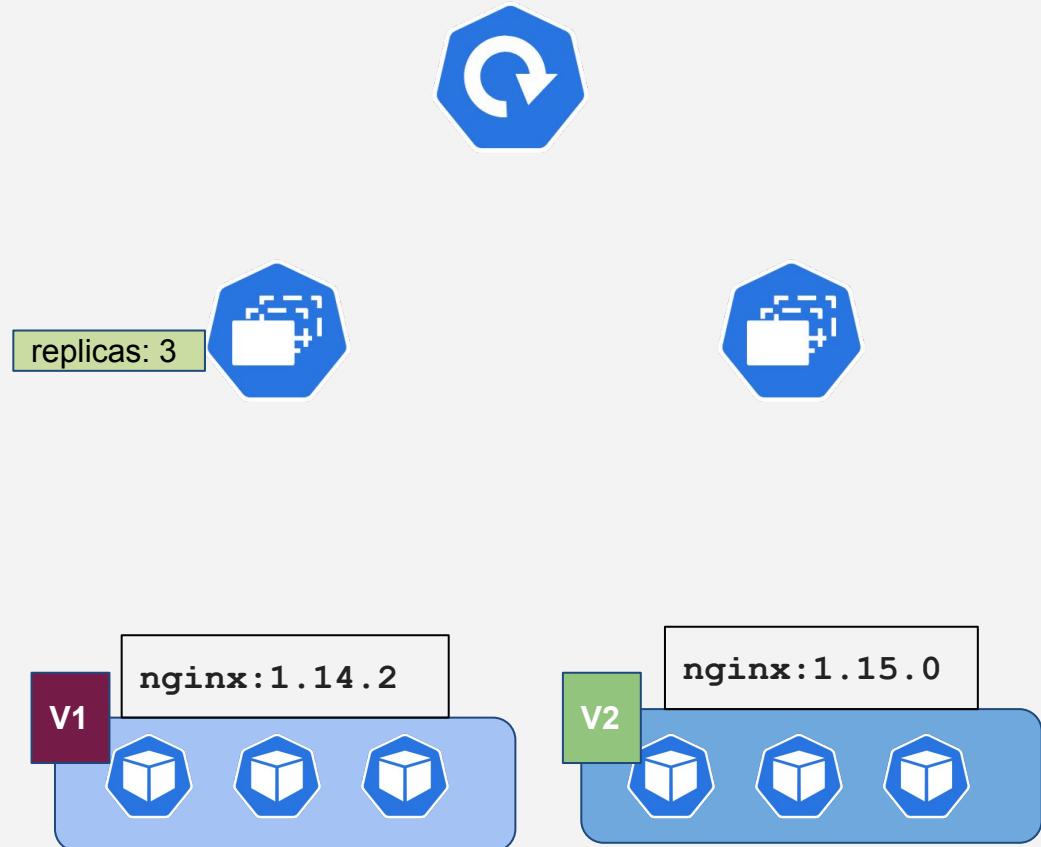


replicas: 3



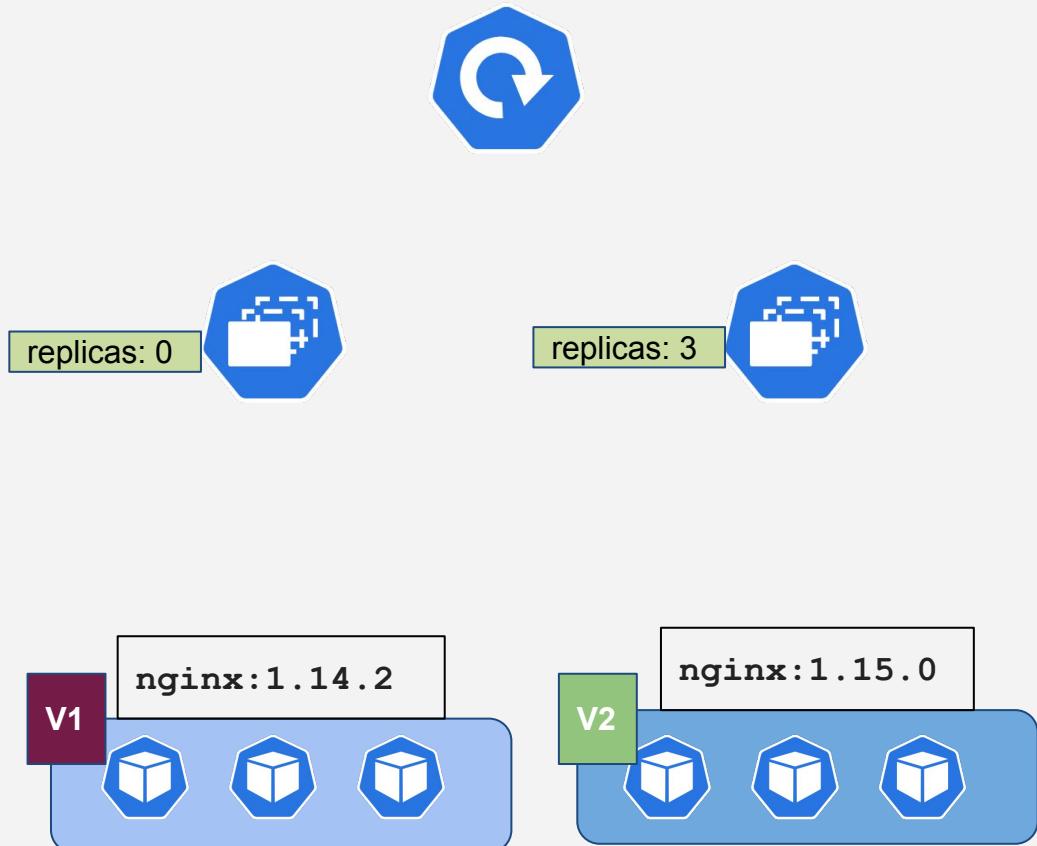
# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



replicas: 3



replicas: 3



Ahmed ElBakry

# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.15.0
      ports:
        - containerPort: 80
```



Ahmed ElBakry

# Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

## Create Deployment

Declarative

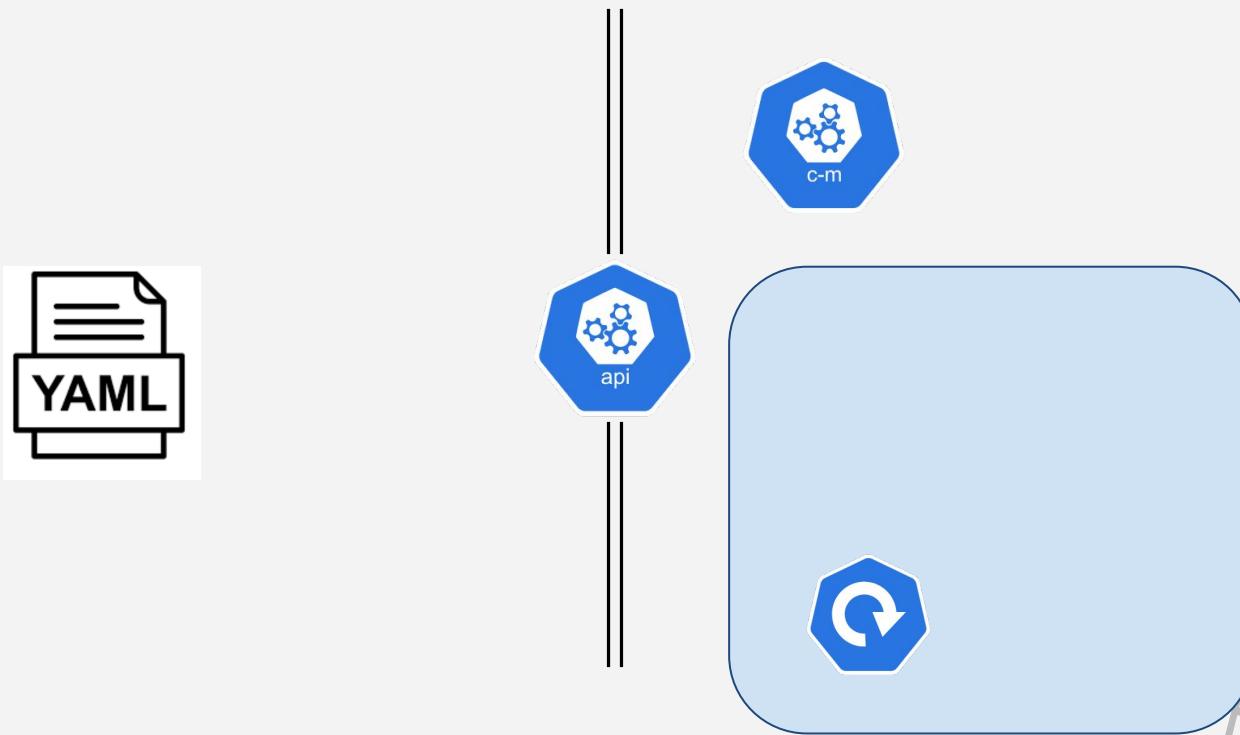
Imperative

(\* |kind-c11:default)-> kubectl  
create -f deploy.yaml



# Deployments

## Deployment Object vs Deployment Controller



# Deployments

## Demo



- ❑ Deployment YAML File
- ❑ Create Deployment
  - ❑ Declarative
  - ❑ Imperative
- ❑ Update Deployment
- ❑ Rollback Deployment
- ❑ Scale Deployment
- ❑ Revision History Limit

### Reference:

<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>



**DEMO**  
AhmedElBakry

# Deployment Strategies

Ahmed ElBakry

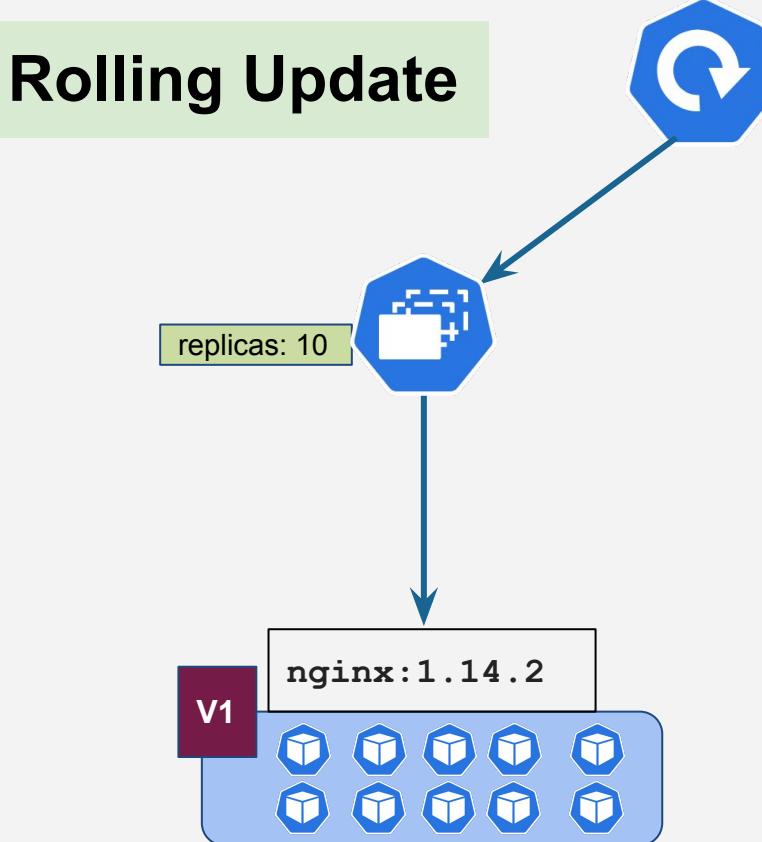
# Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

## Rolling Update

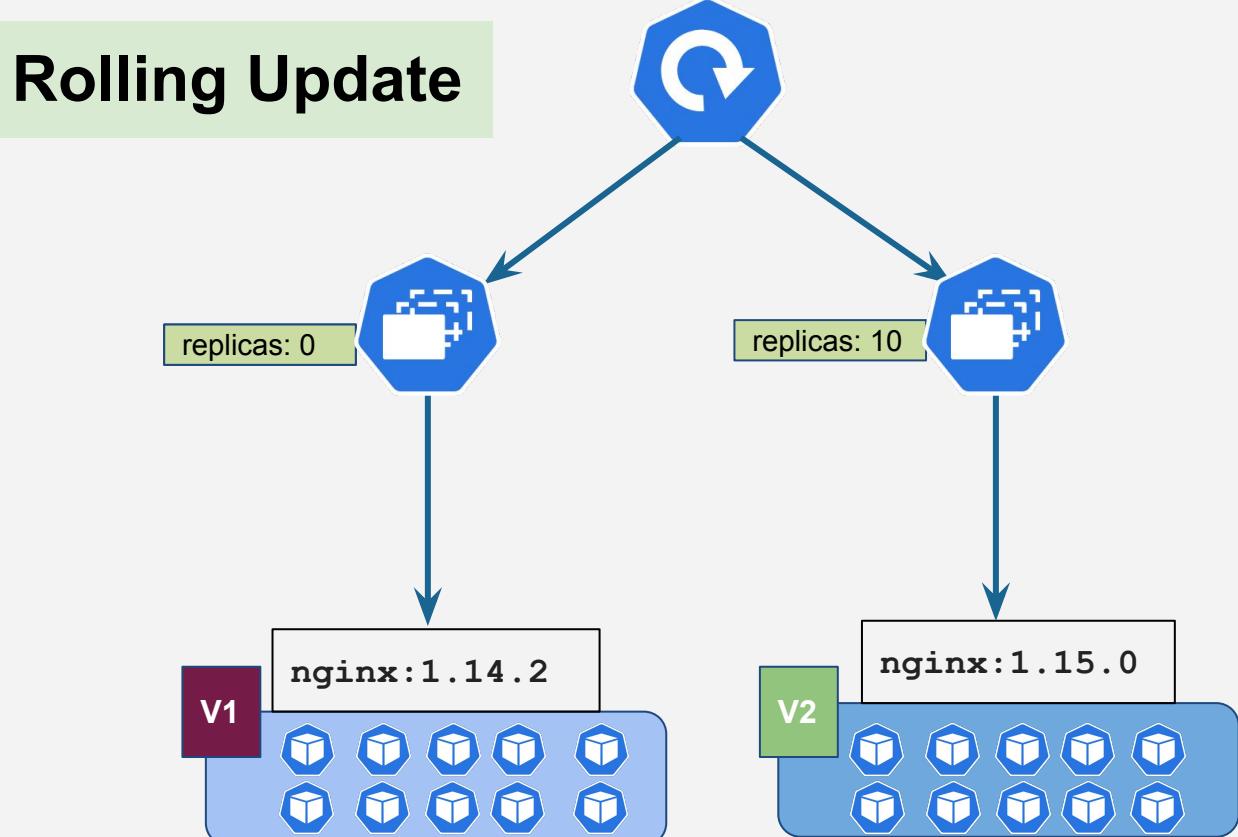
# Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 30%
      maxUnavailable: 30%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



# Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 30%
      maxUnavailable: 30%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```

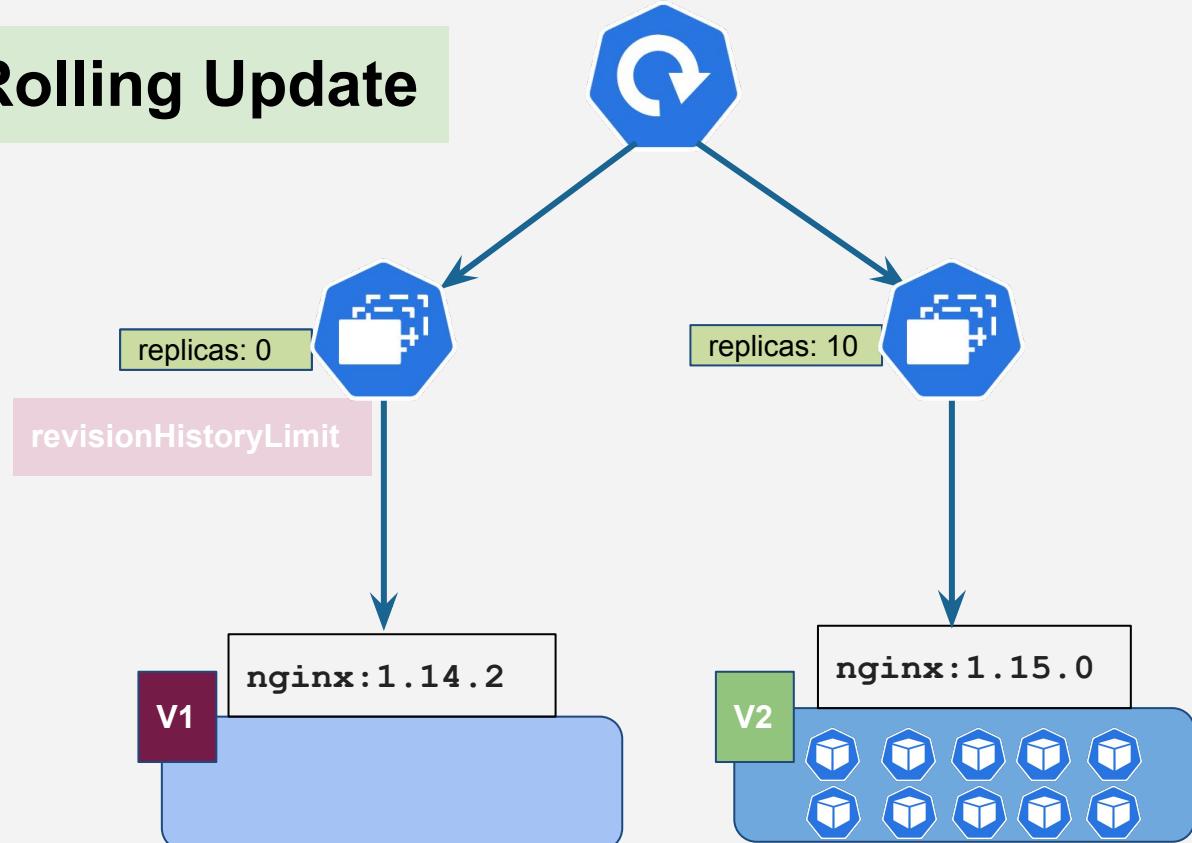


Ahmed ElBakry

# Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 30%
      maxUnavailable: 30%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```

## Rolling Update

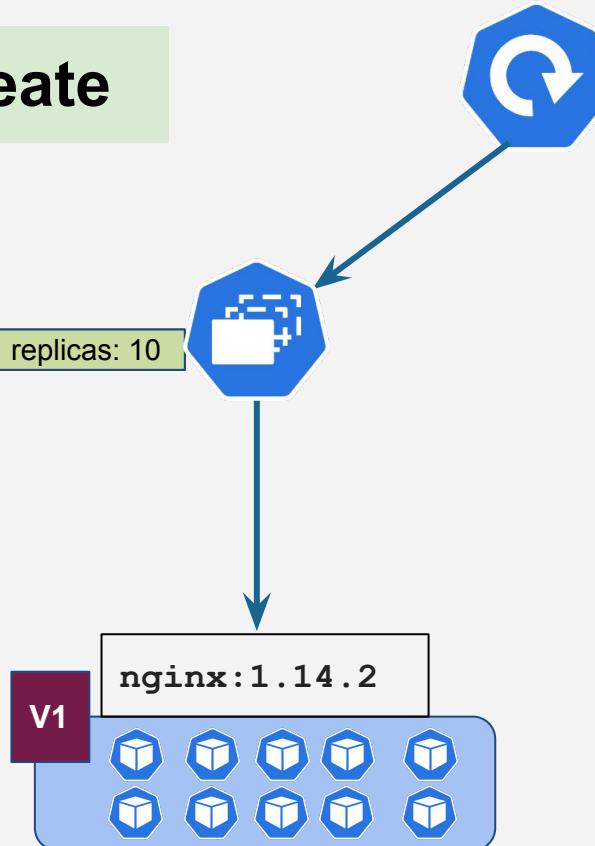


Ahmed ElBakry

# Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

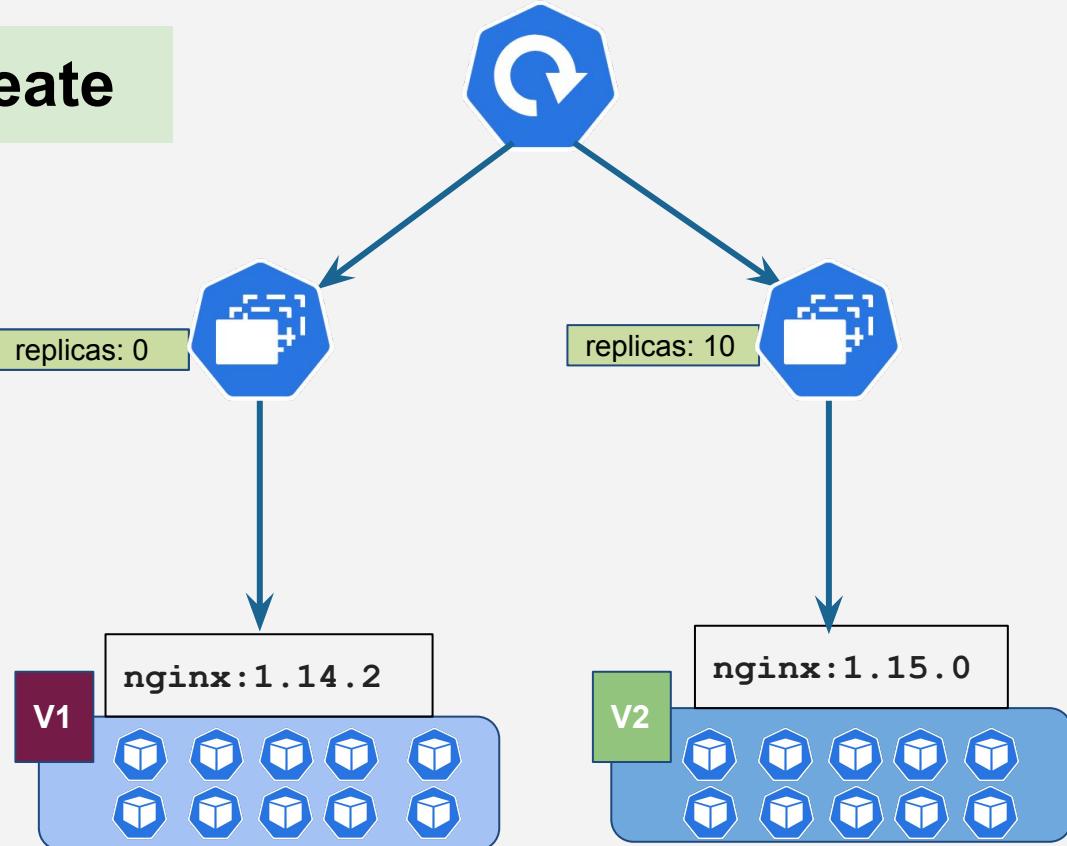
## Recreate



# Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```

## Recreate



Ahmed ElBakry

# Deployment Strategies

## Traffic Pattern





# Module 7

## Services

Ahm

# Services

01

## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

# Services

01

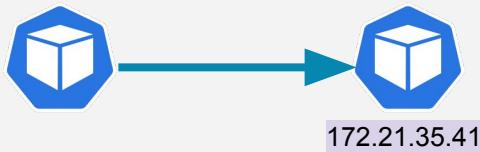
## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.



# Services

01

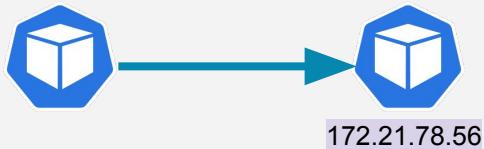
## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.



# Services

01

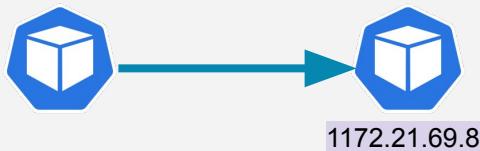
## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.



# Services

01

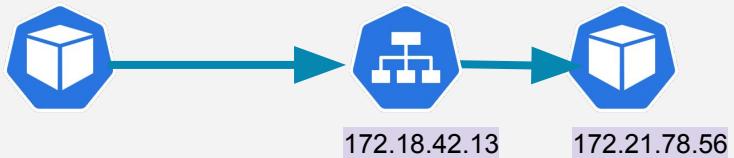
## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.



# Services

01

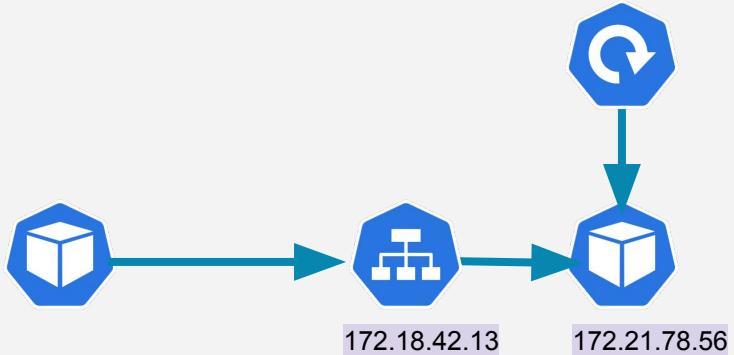
## Service Discovery

A way to connect different services or components with a single stable endpoint

02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

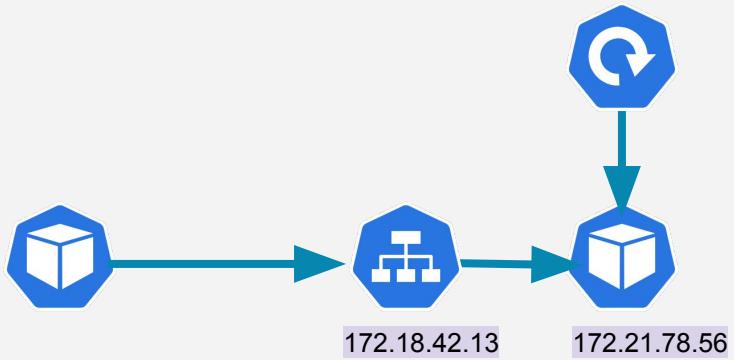


# Services

01

## Service Discovery

A way to connect different services or components with a single stable endpoint



02

## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

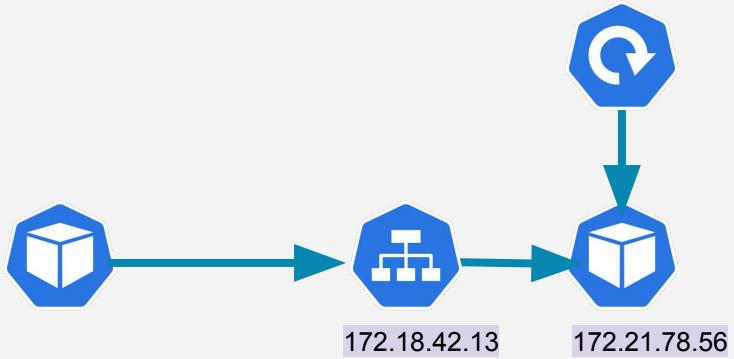


# Services

01

## Service Discovery

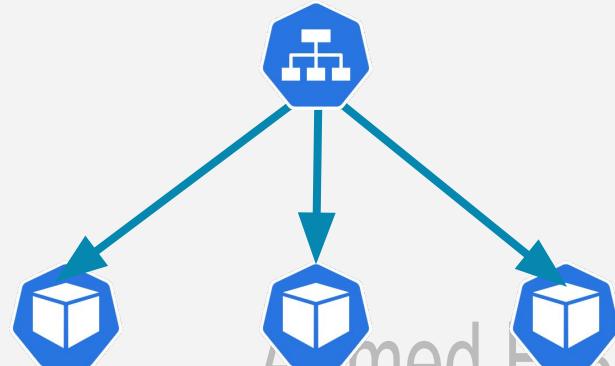
A way to connect different services or components with a single stable endpoint



02

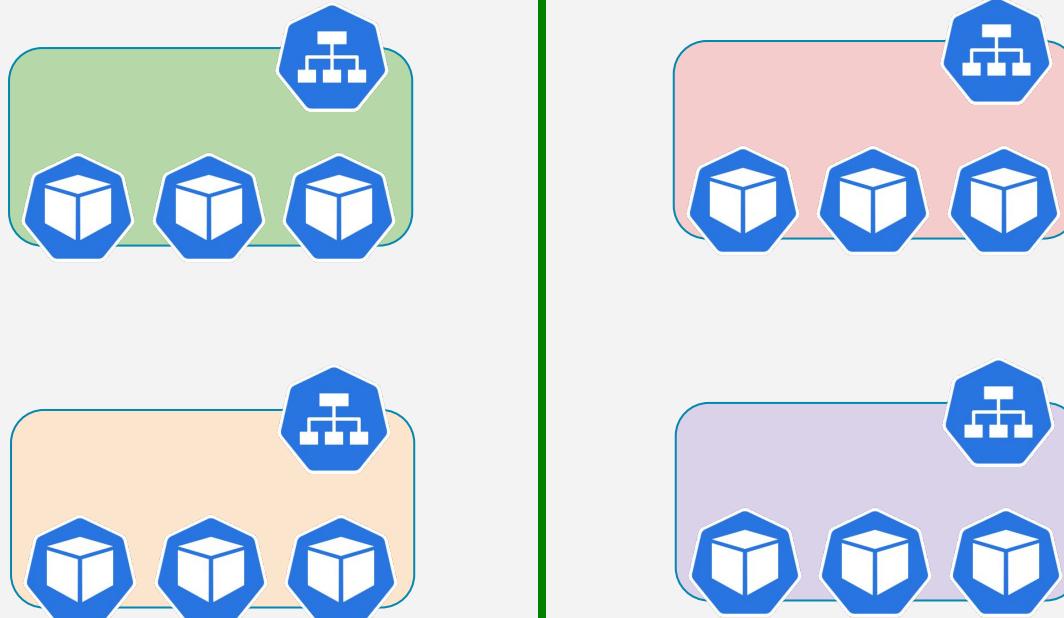
## Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

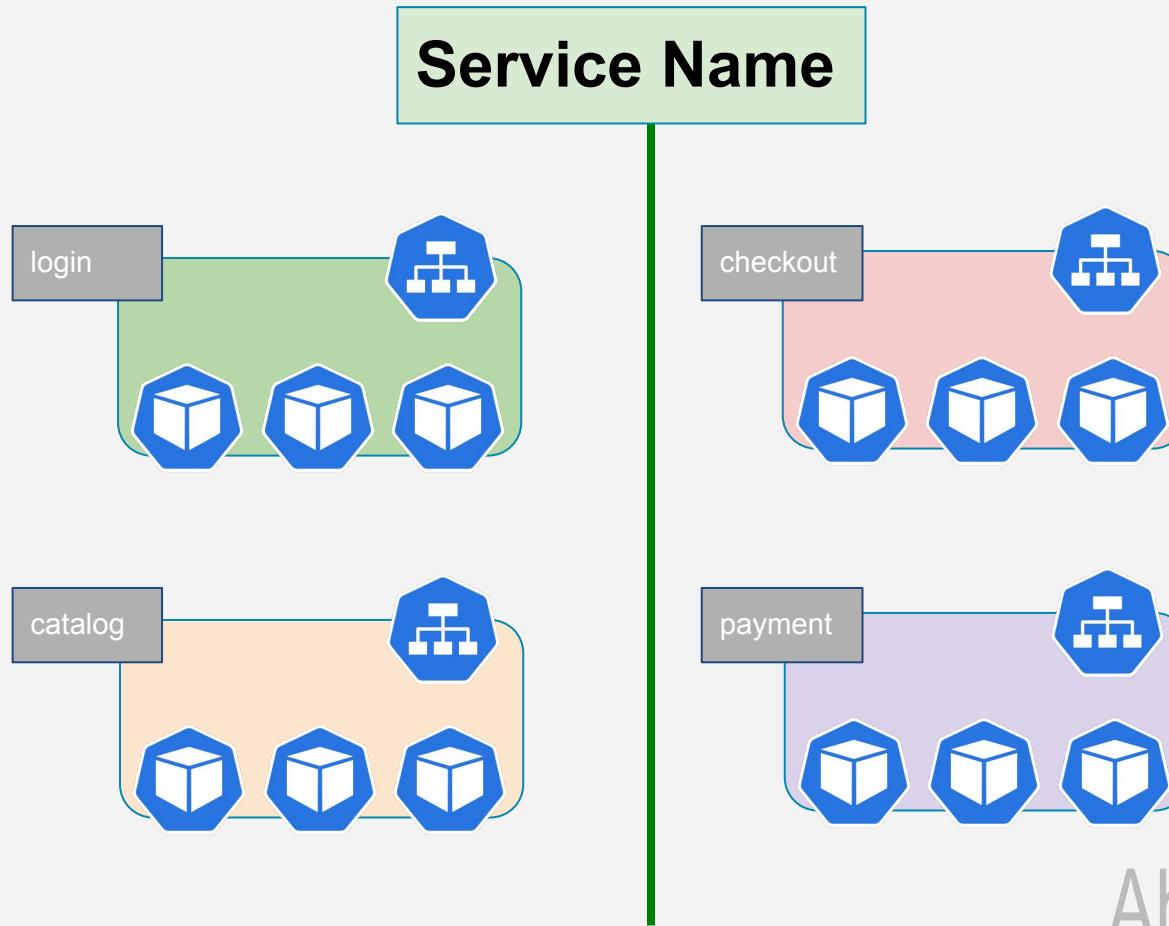


# Services

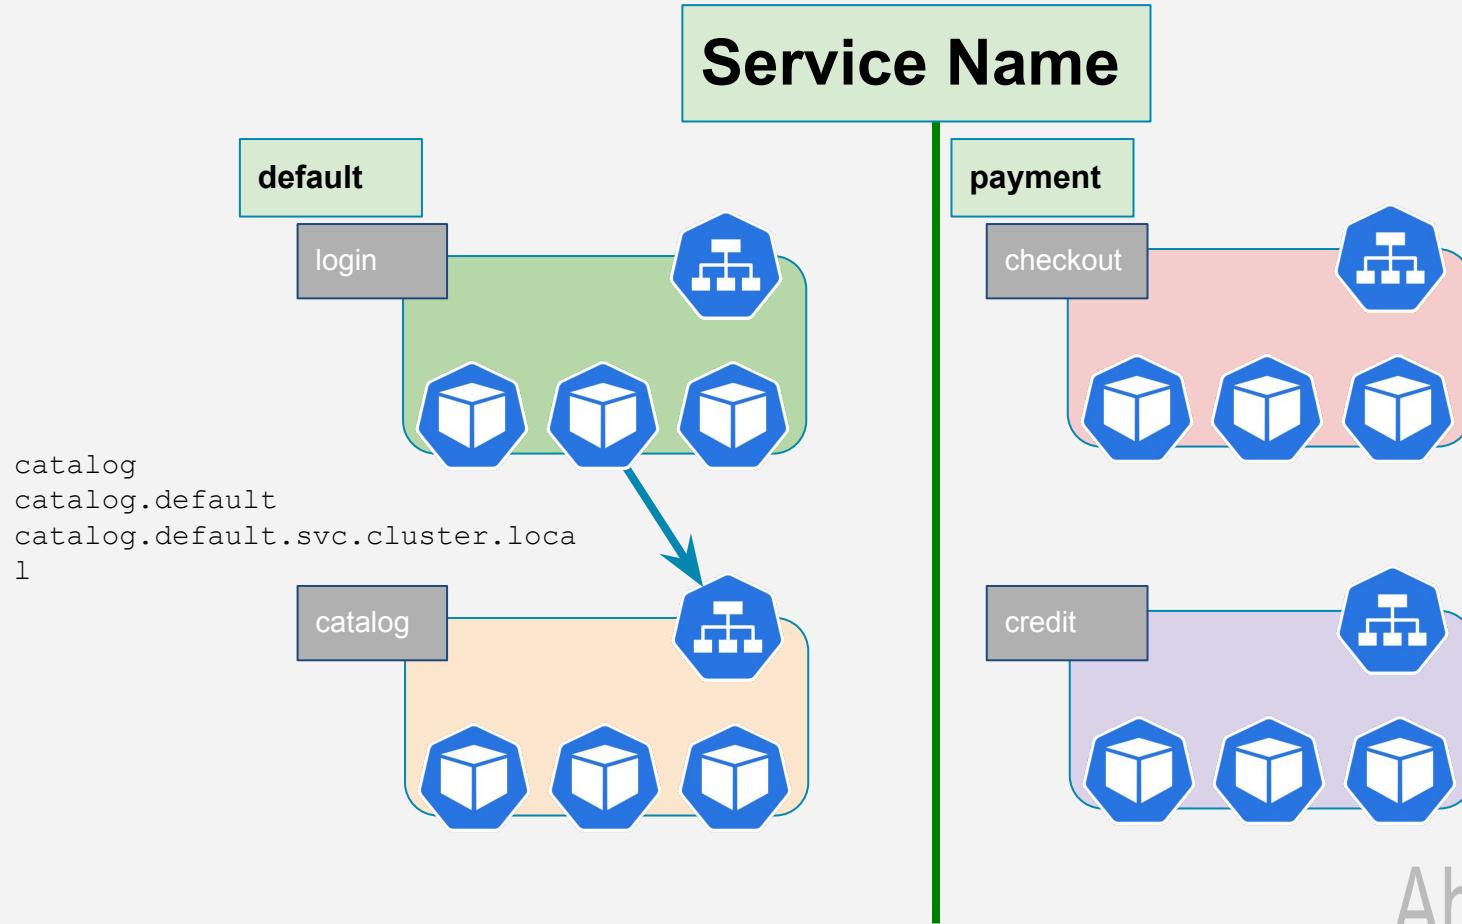
Service Name



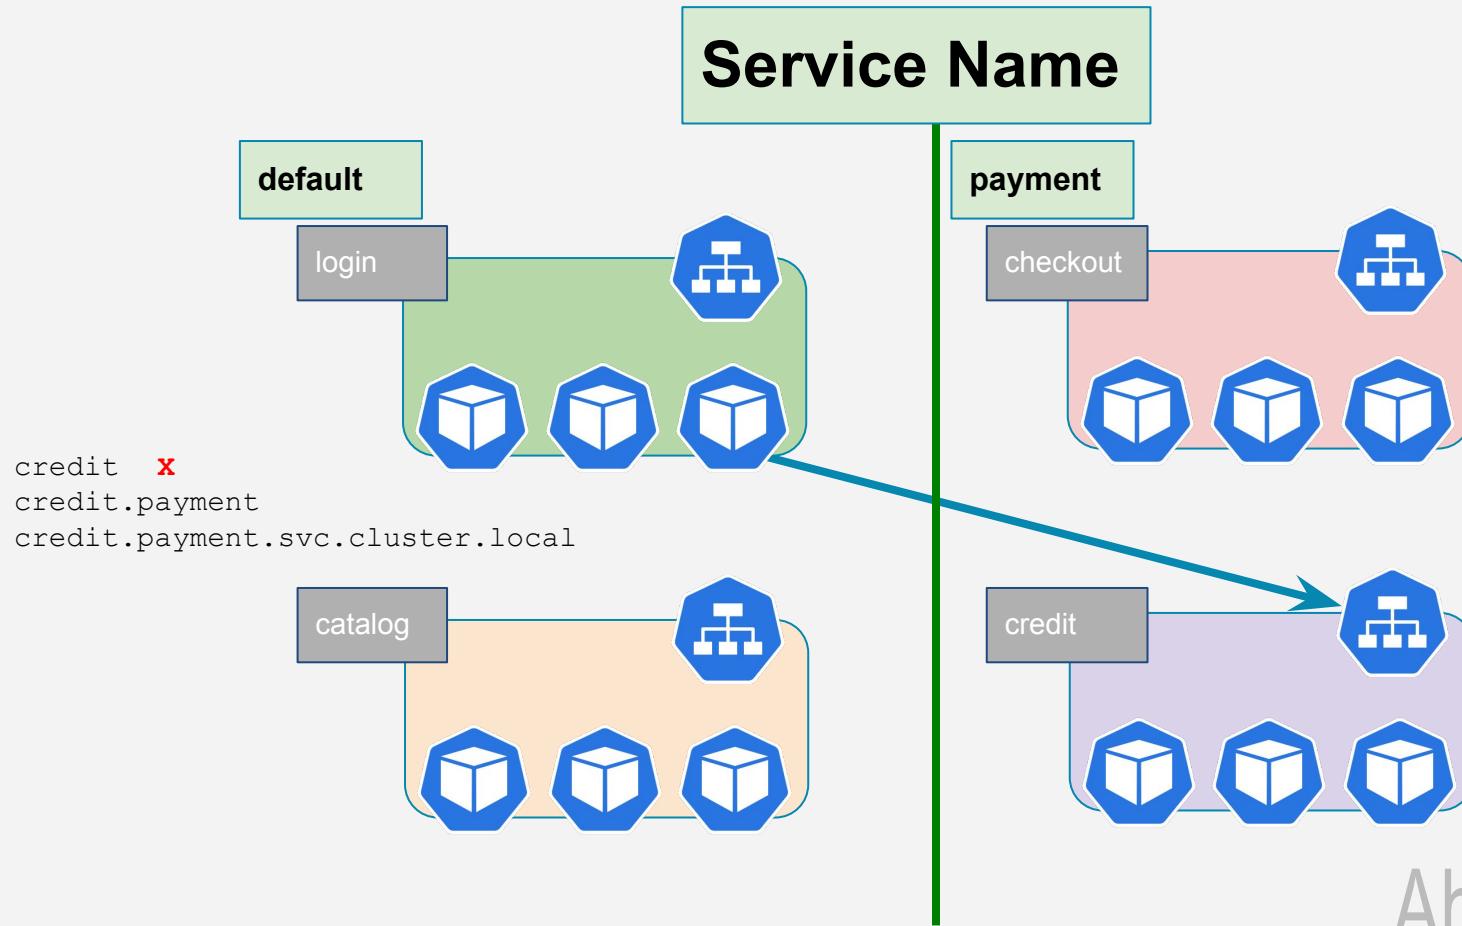
# Services



# Services



# Services



# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
  labels:
    app: nginx
spec:
  containers:
    - name: web
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
  labels:
    app: nginx
spec:
  containers:
    - name: web
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```



# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
  labels:
    app: nginx
spec:
  containers:
    - name: web
      image: nginx:1.14.2
      ports:
        containerPort: 80
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: ClusterIP
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  clusterIP: None
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: test-svc
spec:
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: test-svc
spec:
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 9376
```

```
apiVersion: v1
kind: Endpoints
metadata:
  name: test-svc
subsets:
  - addresses:
    - ip: 10.96.87.20
      ports:
        - port: 9376
```

# Services

```
apiVersion: v1
kind: Service
metadata:
  name: test-svc
spec:
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 9376
```

```
apiVersion: v1
kind: Endpoints
metadata:
  name: test-svc
subsets:
  - addresses:
    - ip: 10.96.87.20
      ports:
        - port: 9376
```

# Services

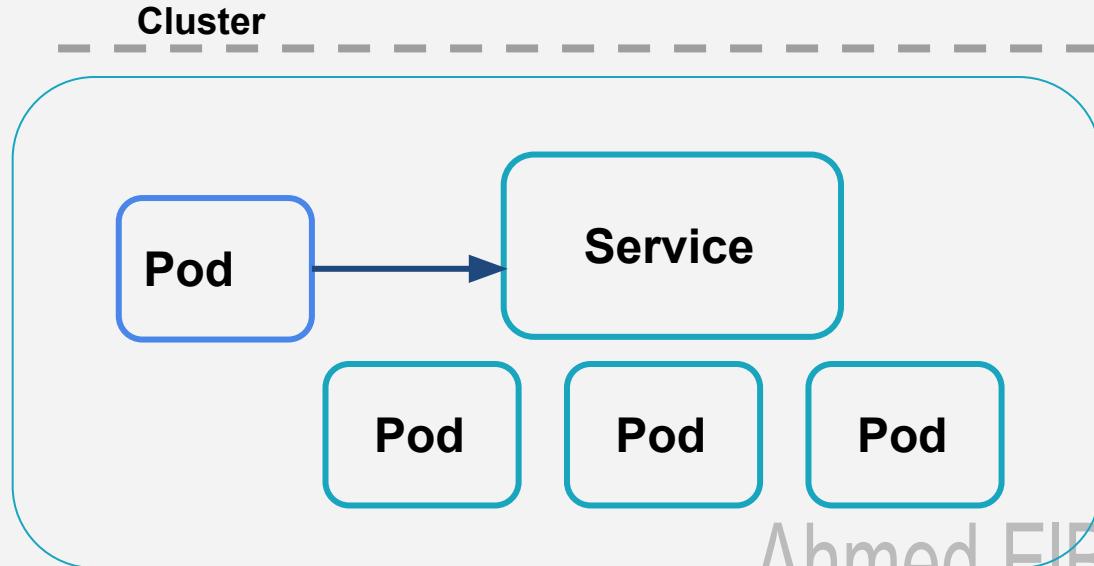
```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      name: http
      port: 80
      targetPort: 9376
    - name: https
      protocol: TCP
      port: 443
      targetPort: 9377
```

# Service Types

- ClusterIP
- NodePort
- LoadBalancer
- Headless
- ExternalName
- ExternalIP

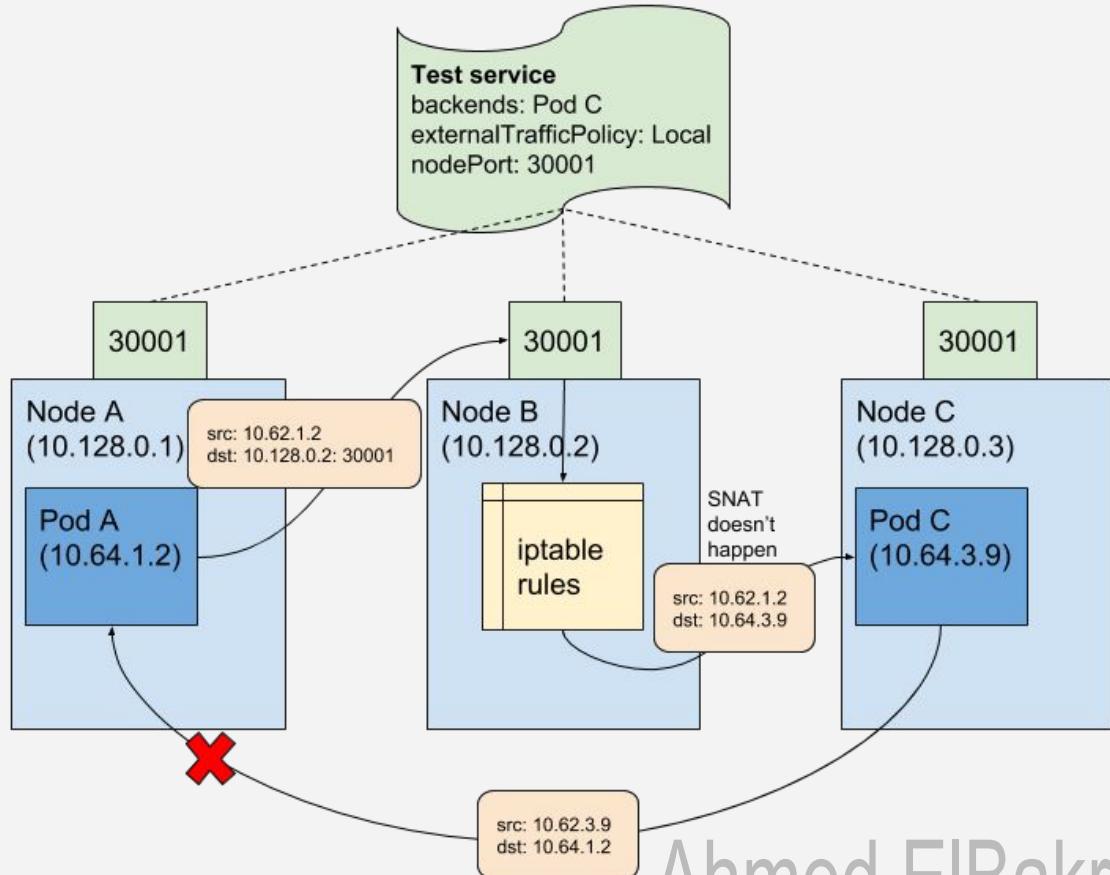
# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: ClusterIP
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```



# Services

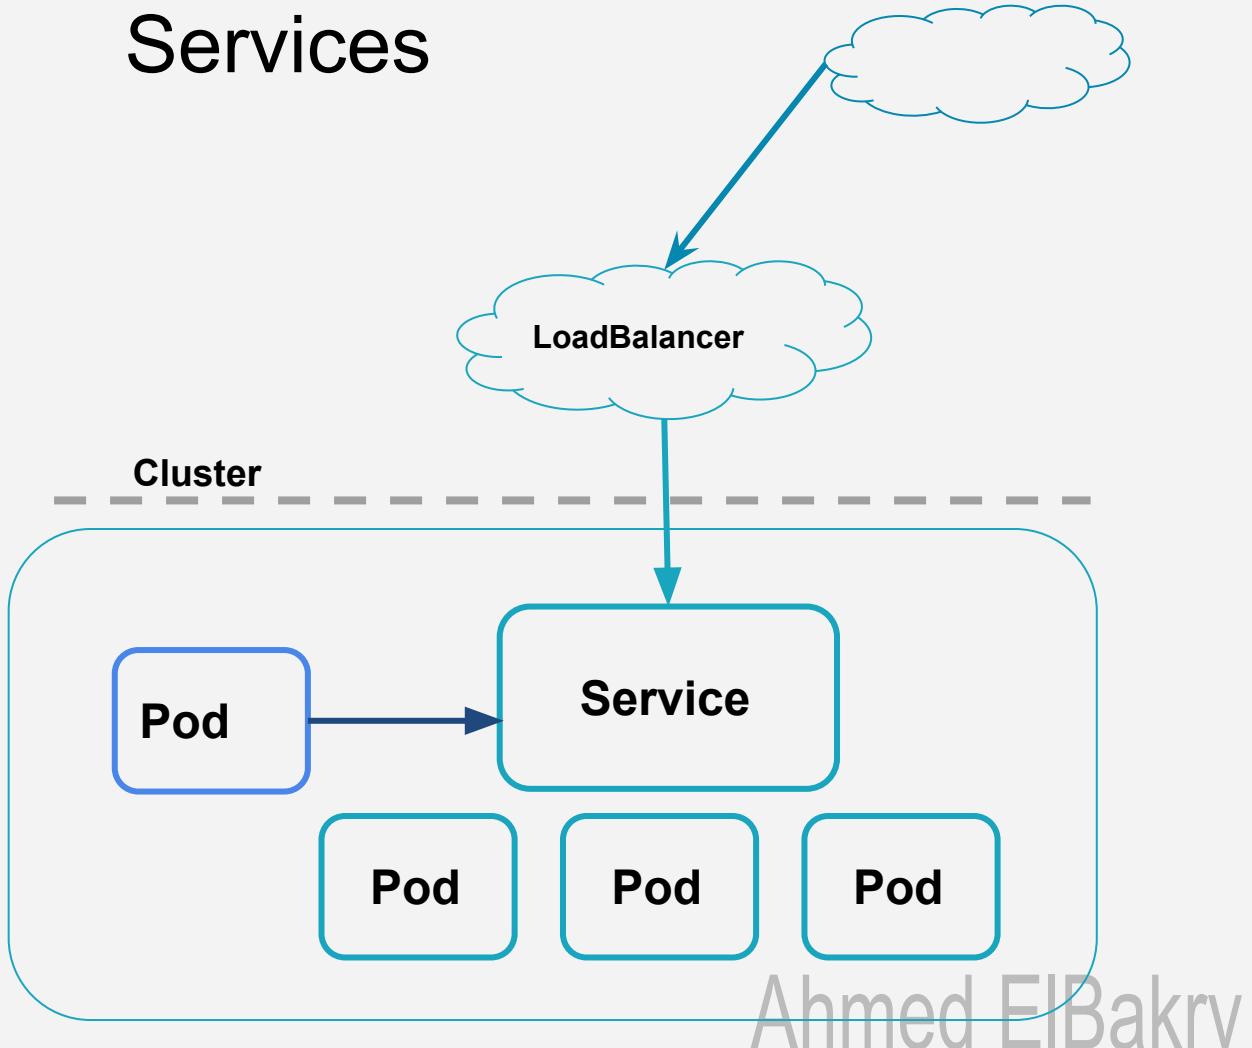
```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
      nodePort: 30305
```



Ahmed ElBakry

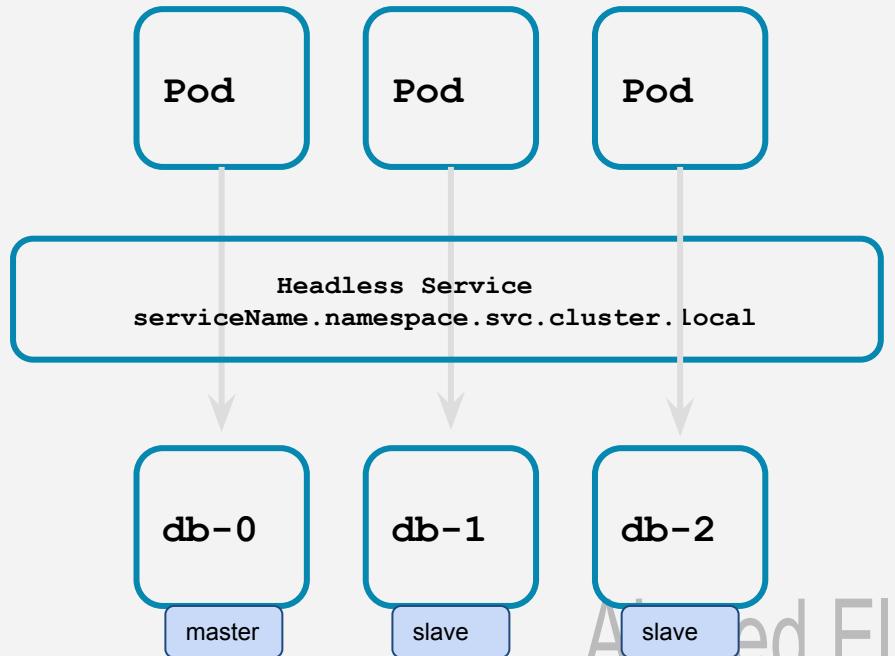
# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```



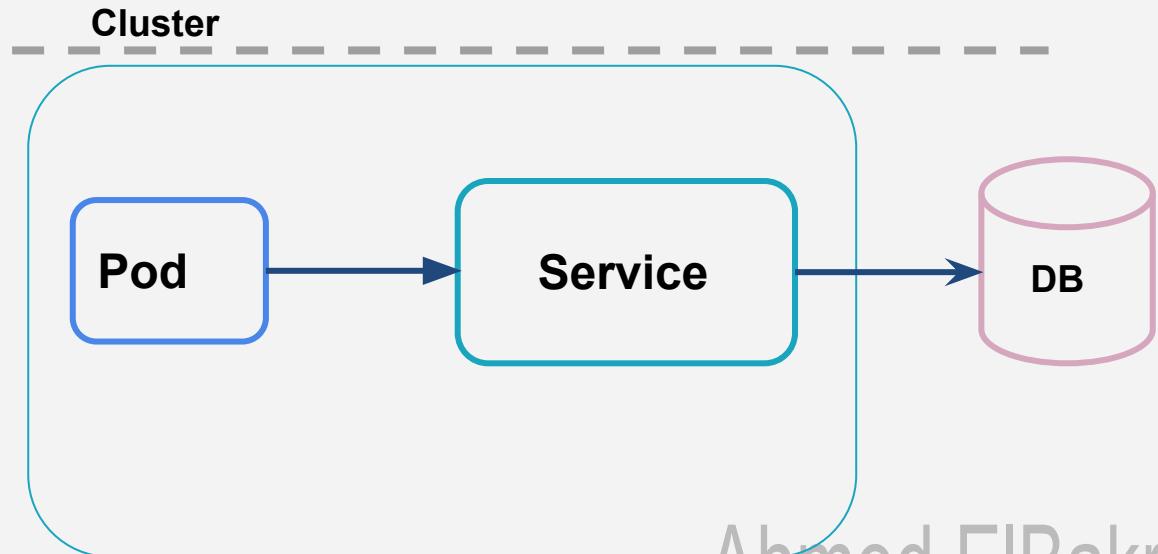
# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  clusterIP: None
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 3306
      targetPort: 3306
```



# Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: ExternalName
  externalName:
    mysql.external.domain.com
  ports:
    - protocol: TCP
      port: 3306
```





**DEMO**

Amr ElBakry

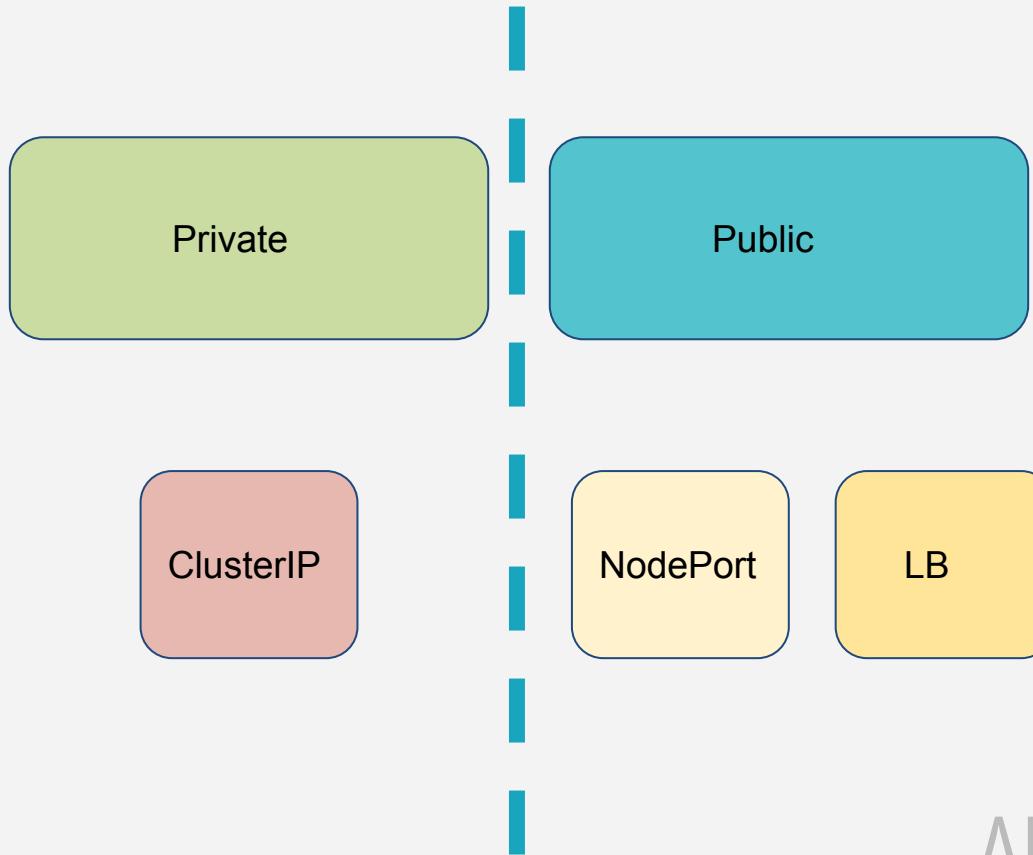


# **Module 8**

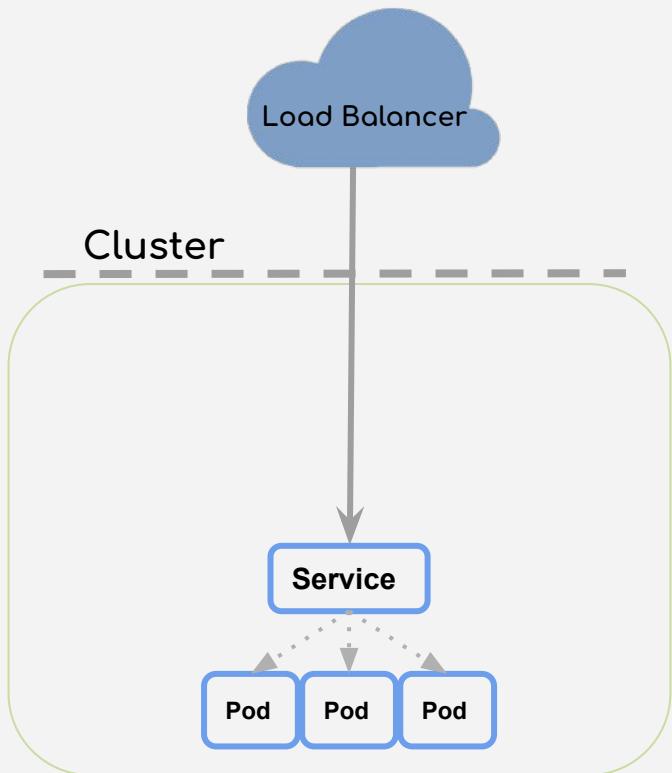
## Ingress

Ahm

# Service Types

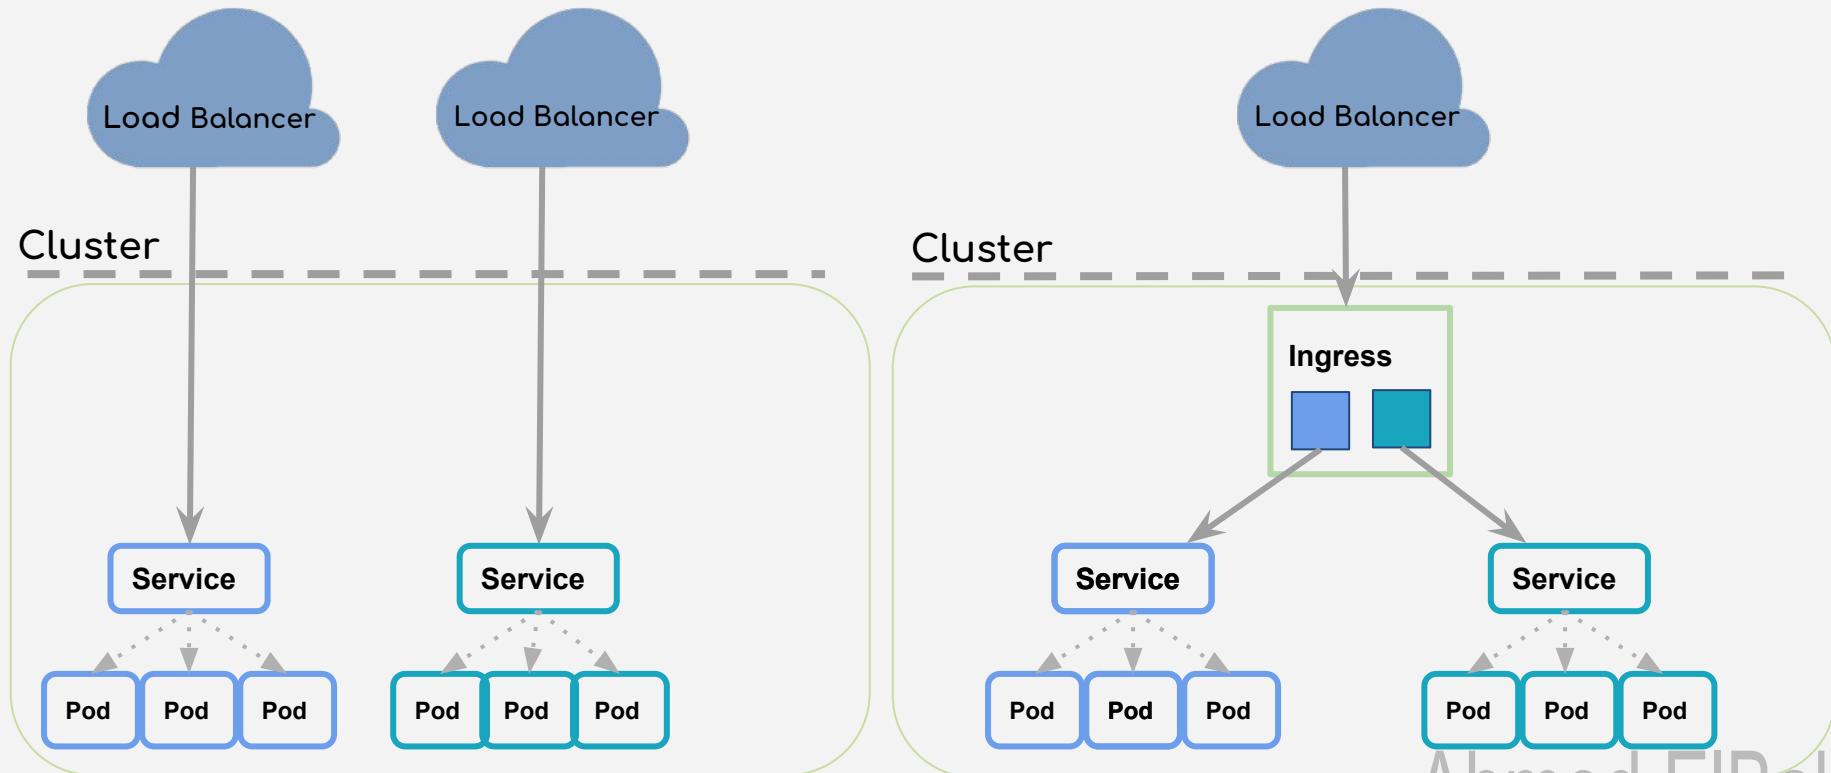


# Ingress

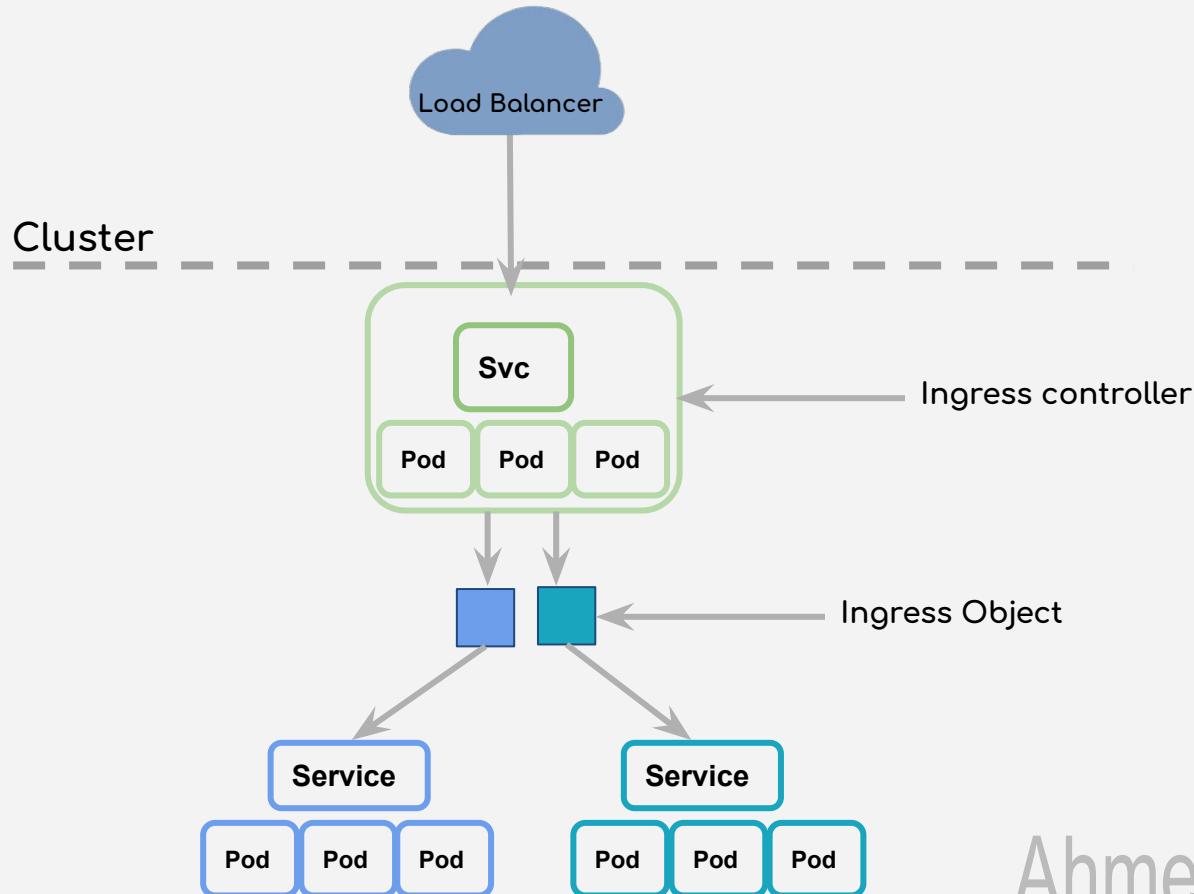


Ahmed ElBakry

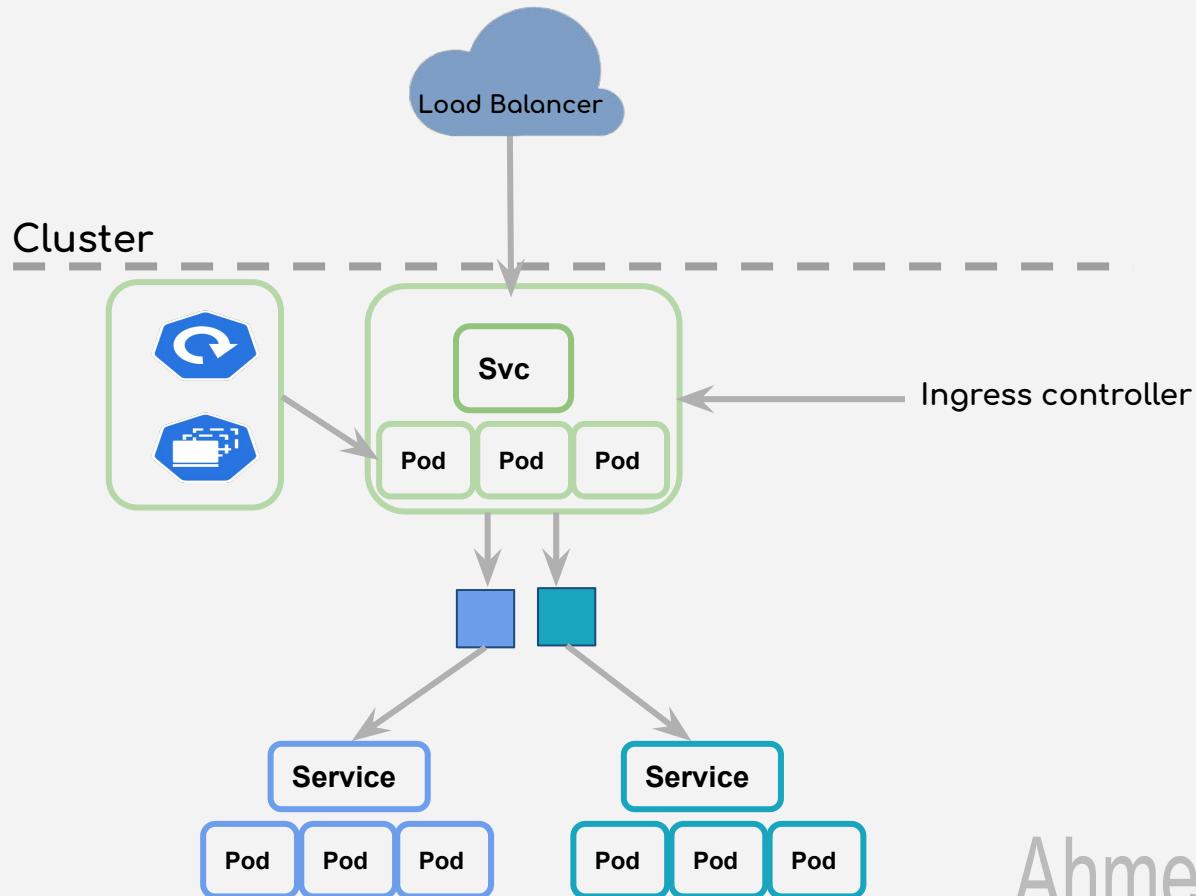
# Ingress



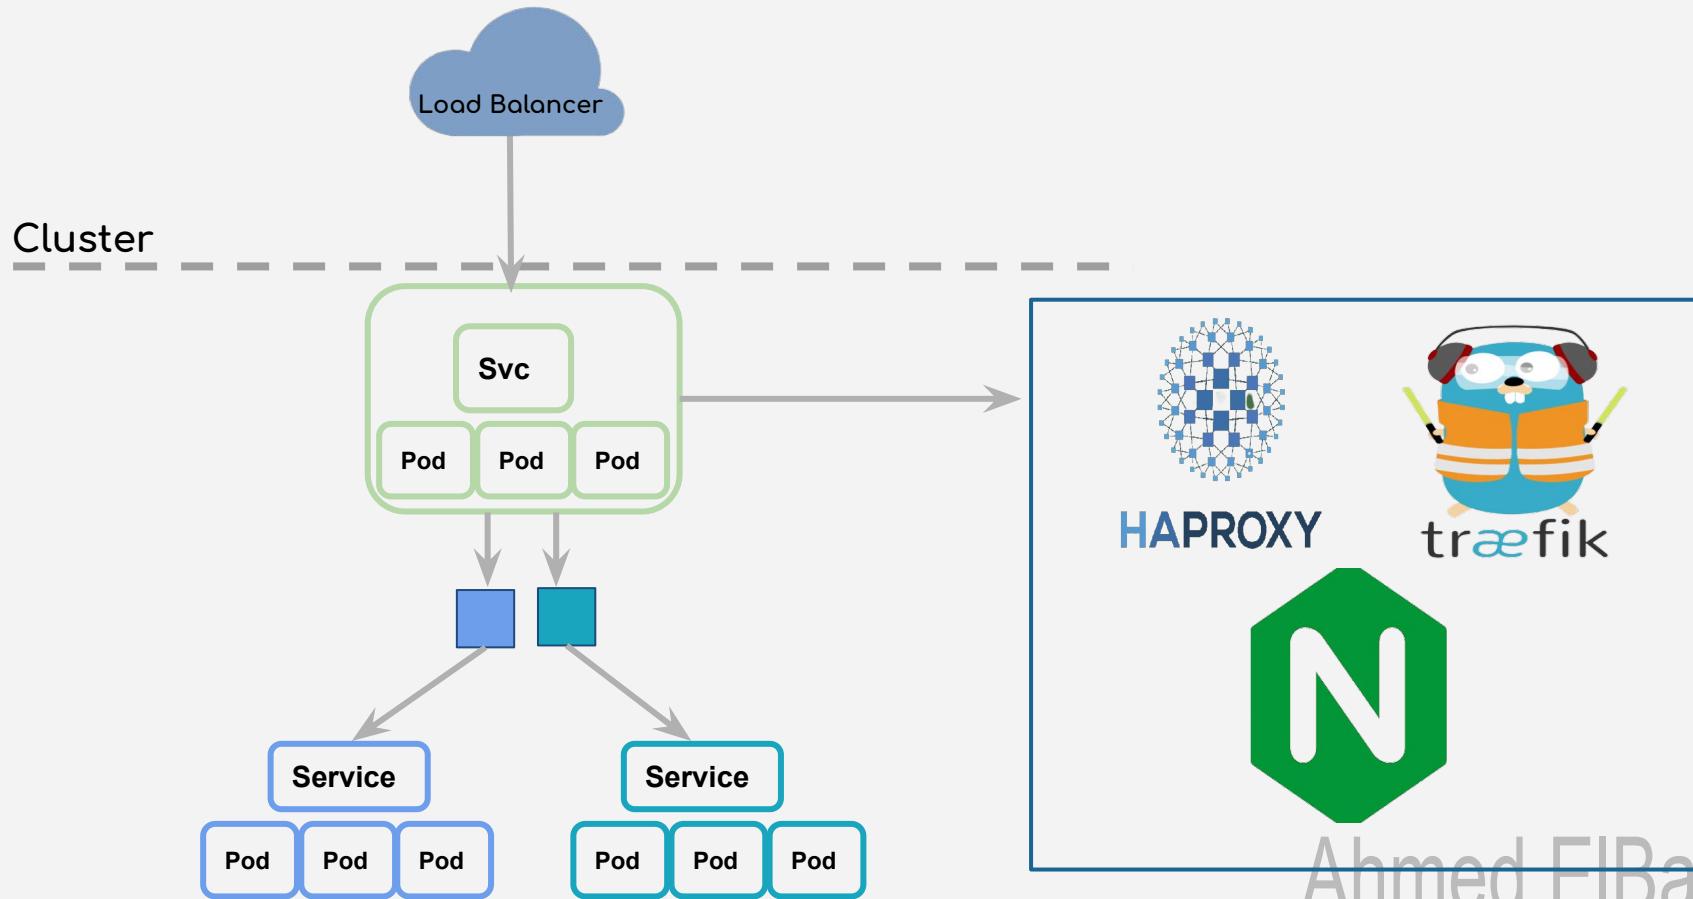
# Ingress



# Ingress

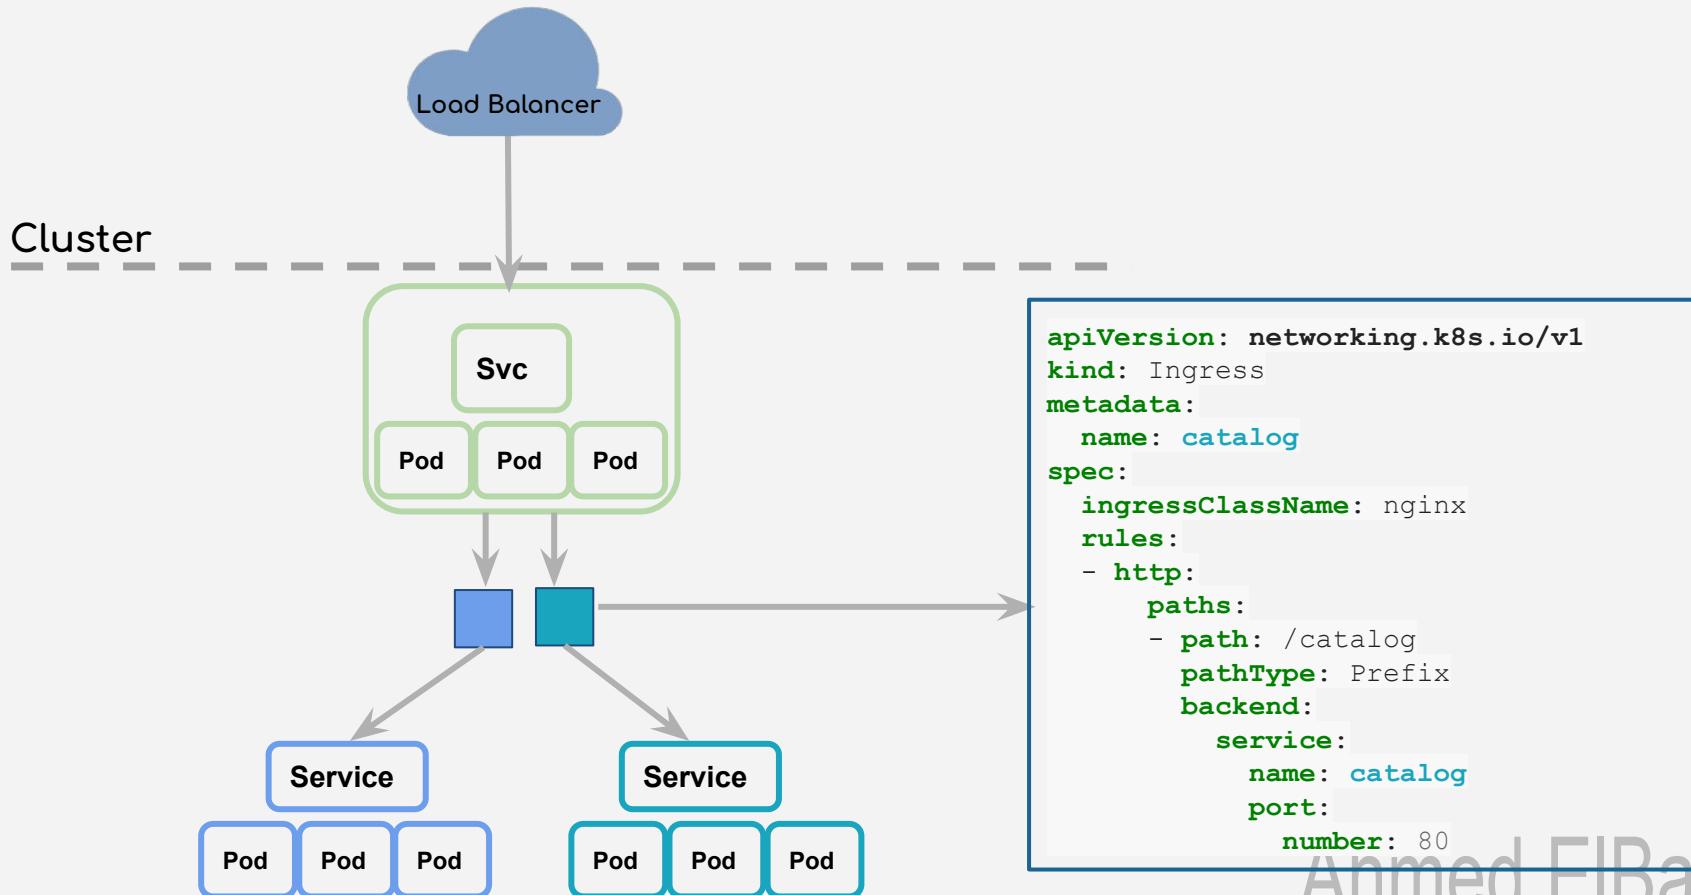


# Ingress



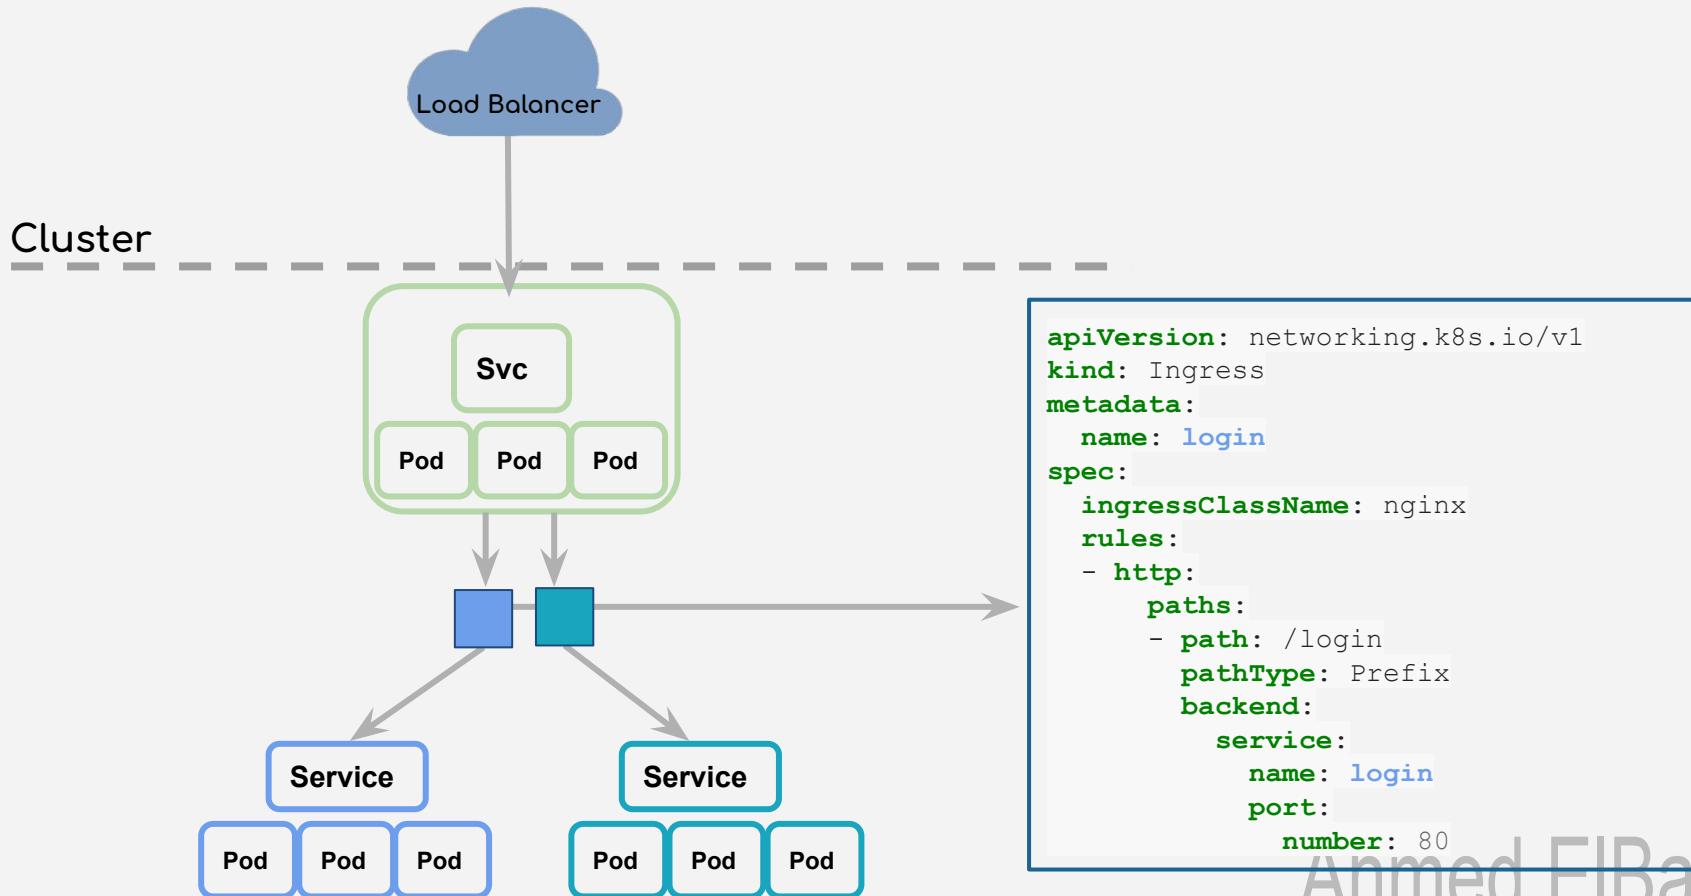
Ahmed EIBakry

# Ingress



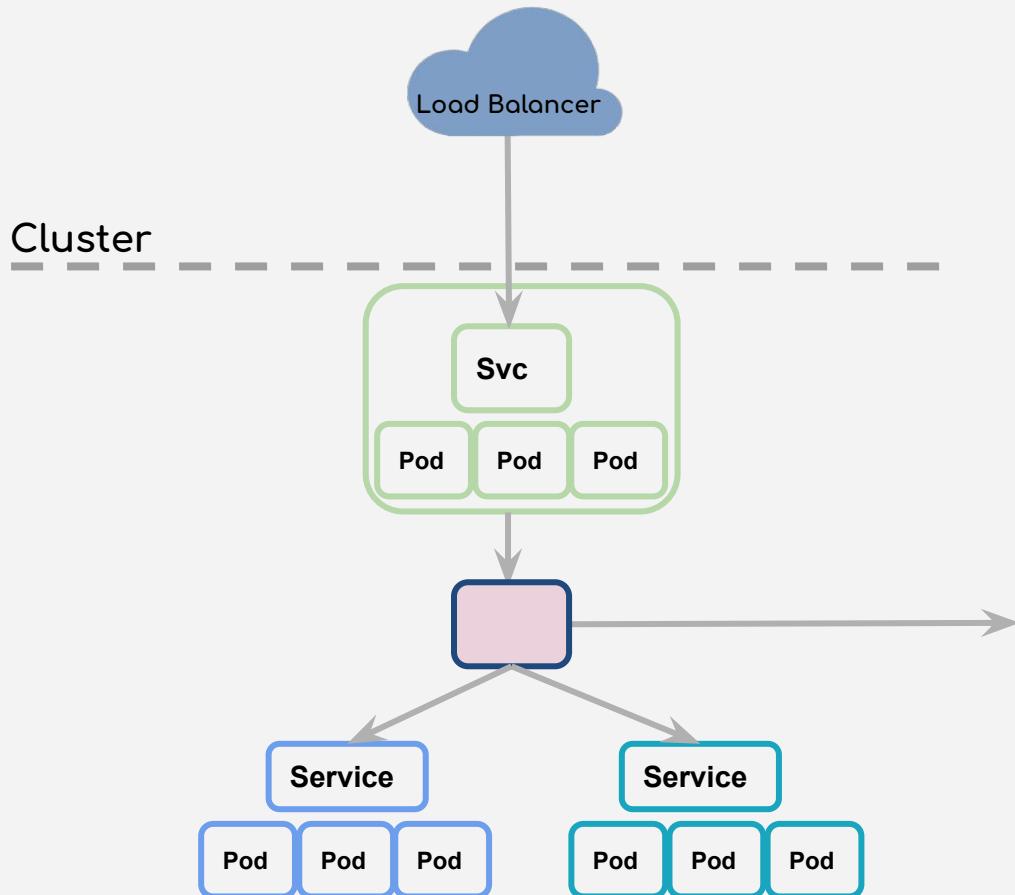
Anmed ElBakry

# Ingress



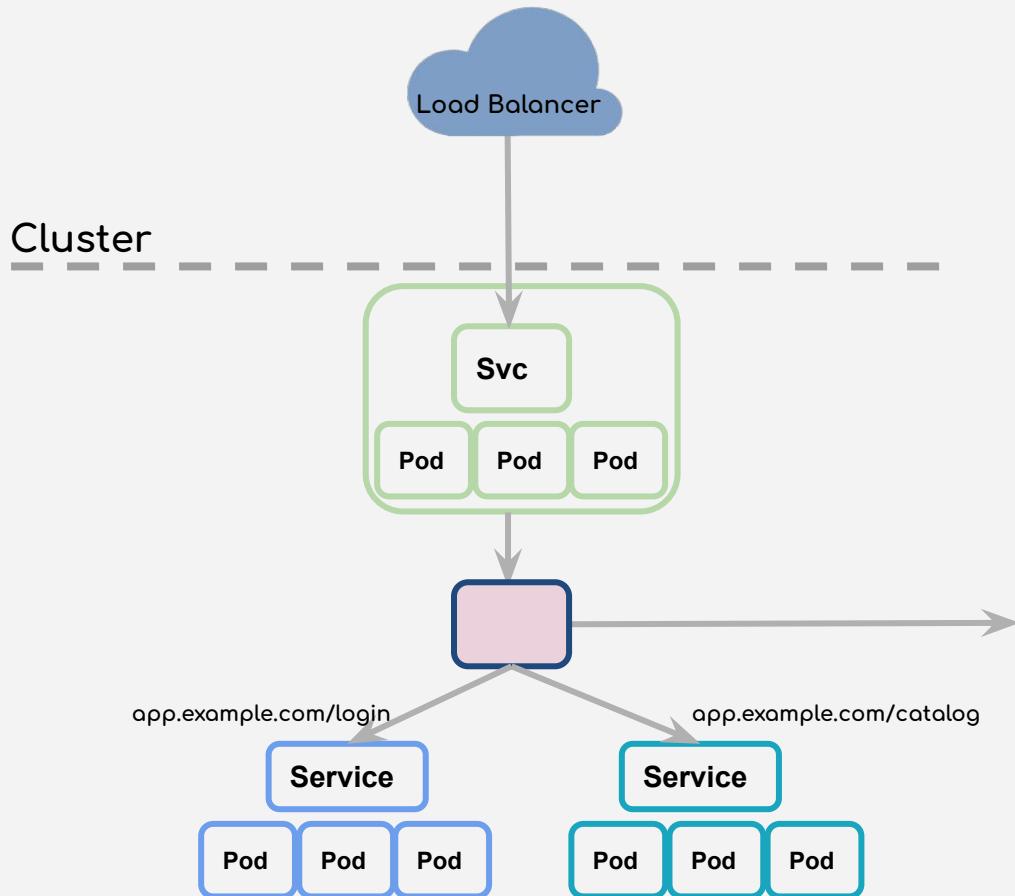
Ahmed ElBakry

# Ingress



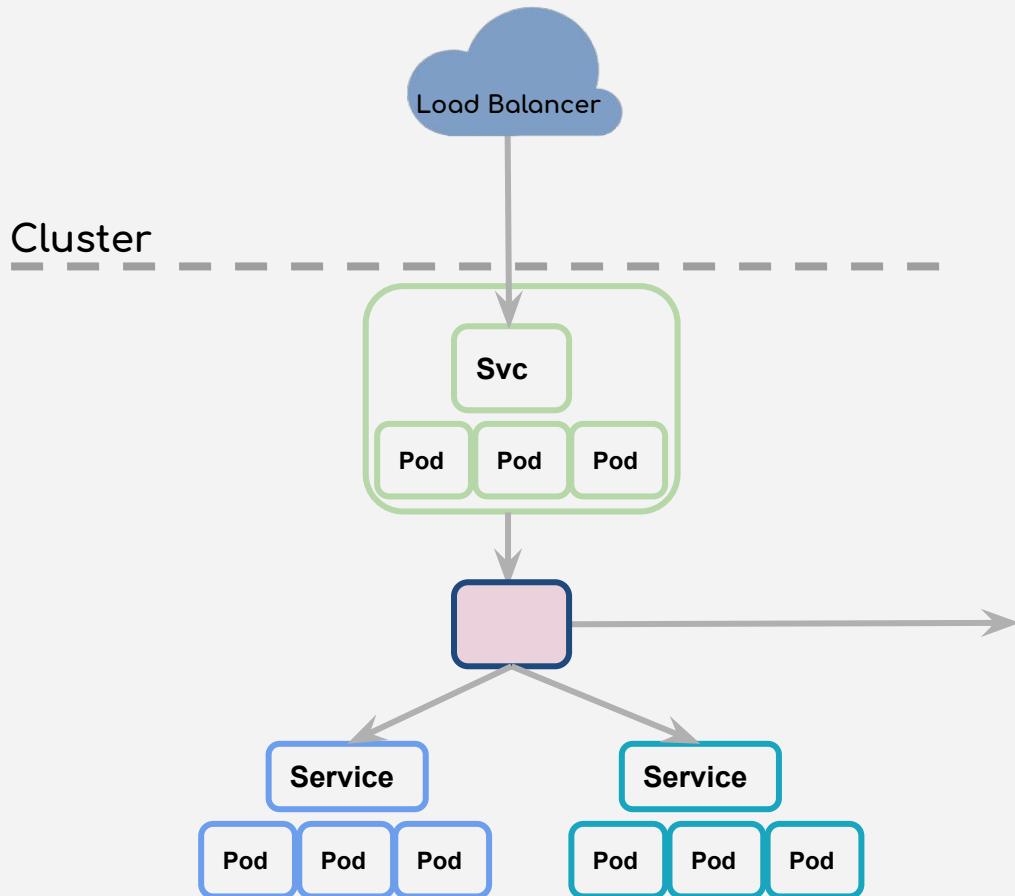
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
spec:
  rules:
  - host: app.example.com
    http:
      paths:
      - path: /login
        pathType: Prefix
        backend:
          service:
            name: login
            port:
              number: 4200
      - path: /catalog
        pathType: Prefix
        backend:
          service:
            name: catalog
            port:
              number: 8080
```

# Ingress



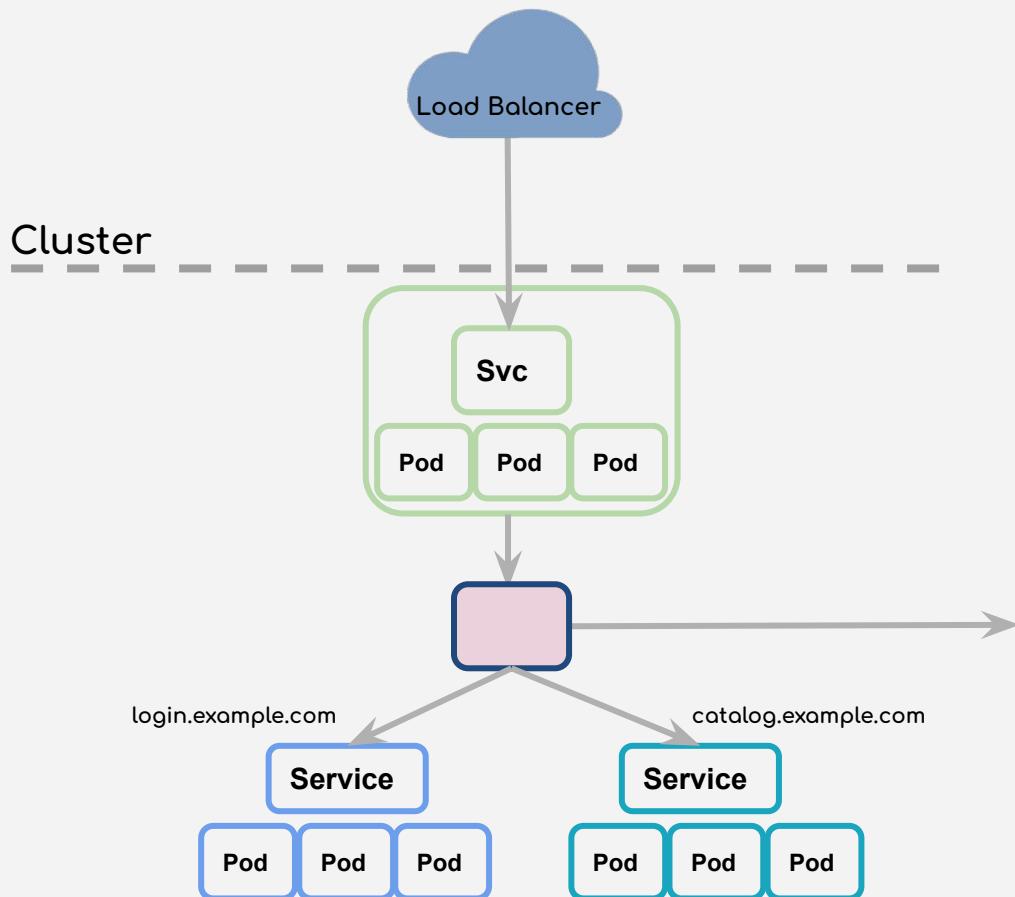
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
spec:
  rules:
  - host: app.example.com
    http:
      paths:
      - path: /login
        pathType: Prefix
        backend:
          service:
            name: login
            port:
              number: 4200
      - path: /catalog
        pathType: Prefix
        backend:
          service:
            name: catalog
            port:
              number: 8080
```

# Ingress



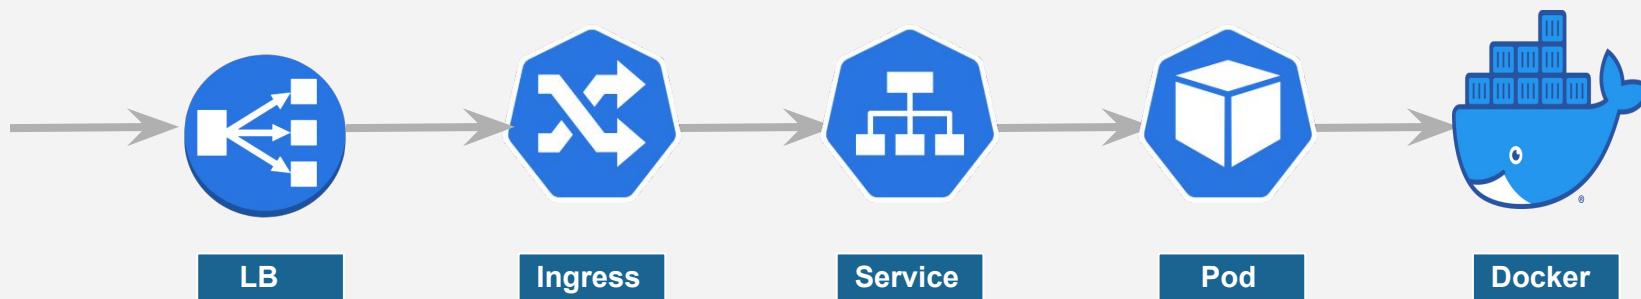
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
spec:
  rules:
  - host: login.example.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: login
            port:
              number: 80
  - host: catalog.example.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: catalog
            port:
              number: 80
```

# Ingress



```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
spec:
  rules:
  - host: login.example.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: login
            port:
              number: 80
  - host: catalog.example.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: catalog
            port:
              number: 80
```

# Ingress





**DEMO**

Amr ElBakry